

بسم الله الرحمن الرحيم

«سیستم عامل»

۱

جلسه ۳:

OS-Related Hardware & Software

Complications in real systems

Brief introduction to

- memory protection and relocation
- virtual memory & MMUs
- I/O & Interrupts

The “process” abstraction

Process scheduling

Process states


Process hierarchies

Process system calls in Unix

Why its not quite that simple ...

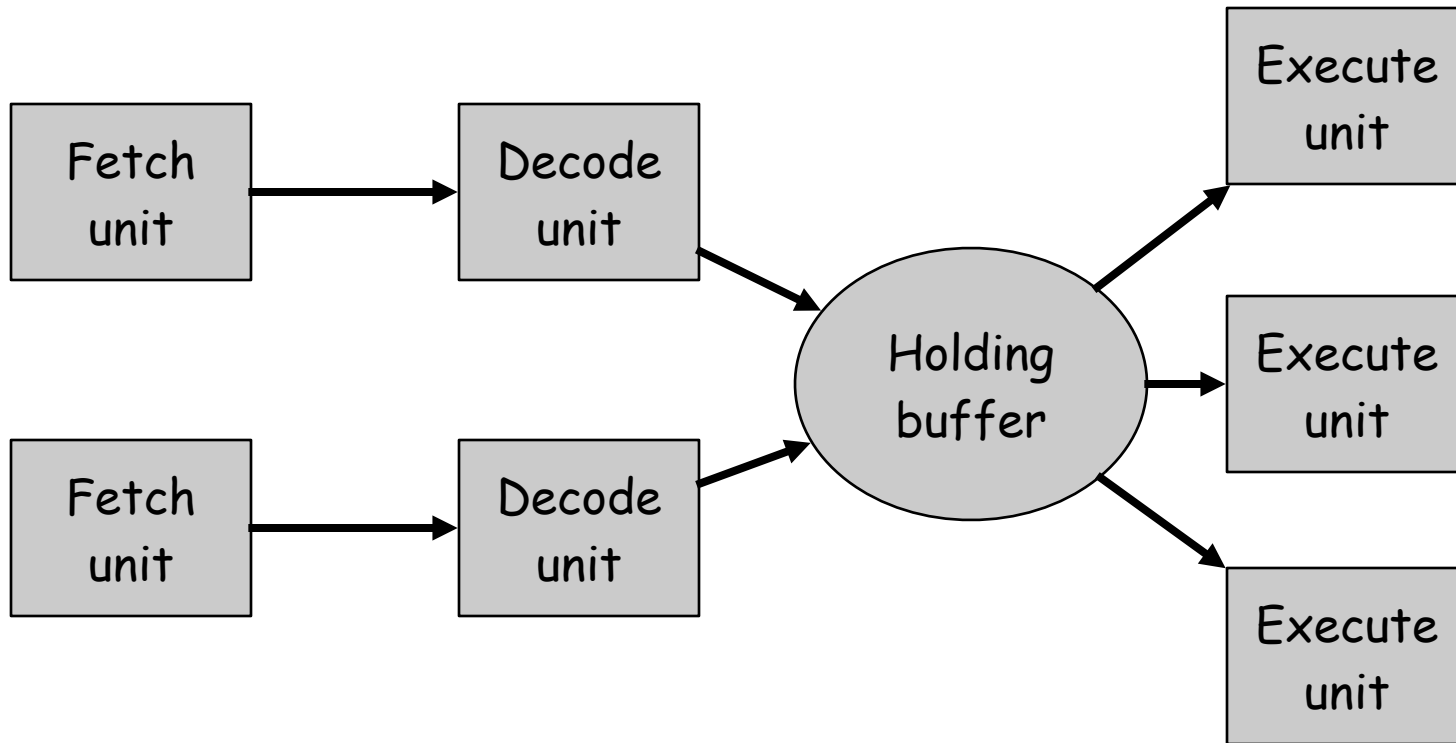
- ❑ The basic model introduced in lecture 1 still applies, but the following issues tend to complicate implementation in real systems:
 - ❖ Pipelined CPUs
 - ❖ Superscalar CPUs
 - ❖ Multi-level memory hierarchies
 - ❖ Virtual memory
 - ❖ Complexity of devices and buses

Pipelined CPUs



Execution of current instruction performed in parallel with decode of next instruction and fetch of the one after that

Superscalar CPUs



What does this mean for the OS?

What does this mean for the OS?

- ▣ **Pipelined CPUs**
 - ❖ more complexity in taking a snapshot of the state of a running application
 - ❖ more expensive to suspend and resume applications

What does this mean for the OS?

- ❑ **Pipelined CPUs**
 - ❖ more complexity in taking a snapshot of the state of a running application
 - ❖ more expensive to suspend and resume applications
- ❑ **Superscalar CPUs**
 - ❖ even more complexity in capturing state of a running application
 - ❖ even more expensive to suspend and resume applications
 - ❖ support from hardware is useful ie. precise interrupts

What does this mean for the OS?

- ❑ **Pipelined CPUs**
 - ❖ more complexity in taking a snapshot of the state of a running application
 - ❖ more expensive to suspend and resume applications
- ❑ **Superscalar CPUs**
 - ❖ even more complexity in capturing state of a running application
 - ❖ even more expensive to suspend and resume applications
 - ❖ support from hardware is useful ie. precise interrupts
- ❑ **More details, but fundamentally the same task**

What does this mean for the OS?

- ❑ **Pipelined CPUs**
 - ❖ more complexity in taking a snapshot of the state of a running application
 - ❖ more expensive to suspend and resume applications
- ❑ **Superscalar CPUs**
 - ❖ even more complexity in capturing state of a running application
 - ❖ even more expensive to suspend and resume applications
 - ❖ support from hardware is useful ie. precise interrupts
- ❑ **More details, but fundamentally the same task**
- ❑ **The BLITZ CPU is not pipelined or superscalar**
 - ❖ BLITZ has precise interrupts

The memory hierarchy

- ❑ **2GHz processor → 0.5 ns clock cycle**
- ❑ **Data/instruction cache access time → 0.5ns - 10 ns**
 - ❖ This is where the CPU looks first!
 - ❖ Memory this fast is very expensive !
 - ❖ Size ~64 kB- 1MB (too small for whole program)
- ❑ **Main memory access time → 60 ns**
 - ❖ Slow, but cheap
 - ❖ Size 1GB+
- ❑ **Magnetic disk → 10 ms, 200+ Gbytes**

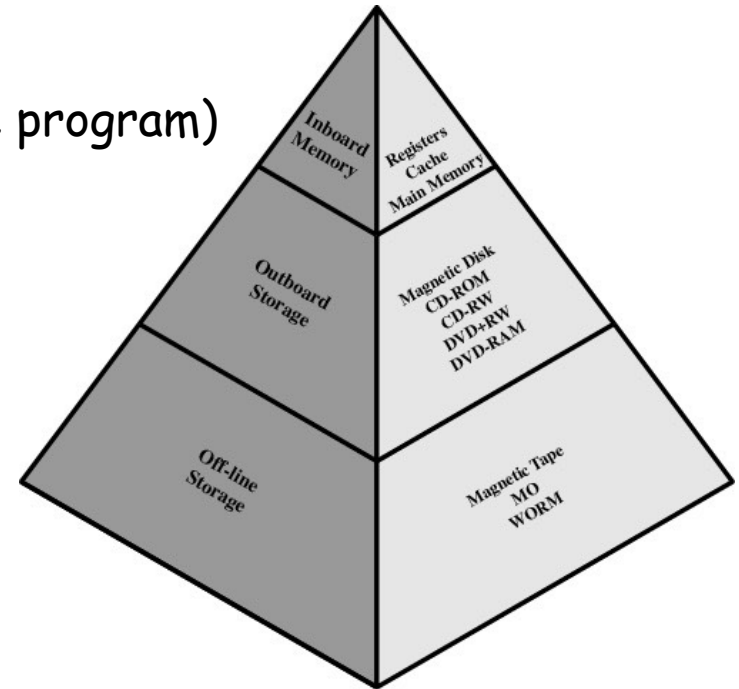


Figure 1.14 The Memory Hierarchy

Who manages the memory hierarchy?

Who manages the memory hierarchy?

- ▣ Movement of data from main memory to cache is under hardware control
 - ❖ cache lines loaded on demand automatically
 - ❖ Placement and replacement policy fixed by hardware

Who manages the memory hierarchy?

- ❑ **Movement of data from main memory to cache is under hardware control**
 - ❖ cache lines loaded on demand automatically
 - ❖ Placement and replacement policy fixed by hardware
- ❑ **Movement of data from cache to main memory can be affected by OS**
 - ❖ instructions for “flushing” the cache
 - ❖ can be used to maintain consistency of main memory

Who manages the memory hierarchy?

- ❑ **Movement of data from main memory to cache is under hardware control**
 - ❖ cache lines loaded on demand automatically
 - ❖ Placement and replacement policy fixed by hardware
- ❑ **Movement of data from cache to main memory can be affected by OS**
 - ❖ instructions for “flushing” the cache
 - ❖ can be used to maintain consistency of main memory
- ❑ **Movement of data among lower levels of the memory hierarchy is under direct control of the OS**
 - ❖ virtual memory page faults
 - ❖ file system calls

OS implications of a memory hierarchy?

OS implications of a memory hierarchy?

- How do you keep the contents of memory **consistent** across layers of the hierarchy?
- How do you **allocate** space at layers of the memory hierarchy “fairly” across different applications?

OS implications of a memory hierarchy?

- ❑ How do you keep the contents of memory **consistent** across layers of the hierarchy?
- ❑ How do you **allocate** space at layers of the memory hierarchy “fairly” across different applications?
- ❑ How do you **hide the latency** of the slower subsystems?
 - Main memory?
 - Disk

Other memory-related issues

Other memory-related issues

- How do you **protect** one application's area of memory from other applications?

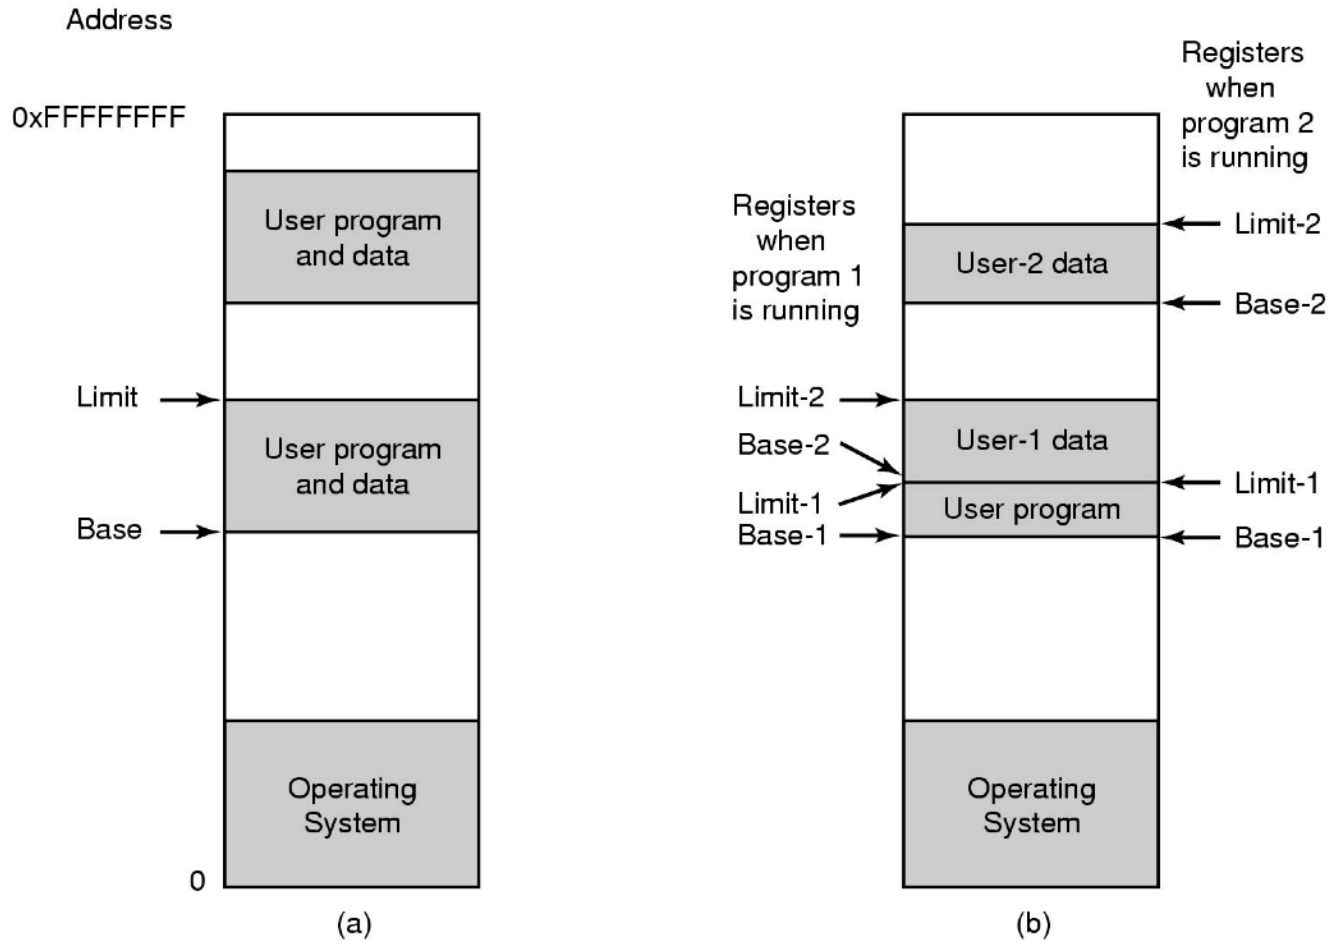
Other memory-related issues

- How do you **protect** one application's area of memory from other applications?
- How do you **relocate** an application in memory?
 - ❖ How does the programmer know where the program will ultimately reside in memory?

Memory protection and relocation ...

- **Memory protection – the basic ideas**
 - ❖ **virtual vs physical addresses**
 - address range in each application starts at 0
 - ❖ Possible solution with **base** and **limit registers**
 - Get CPU to interpret address indirectly, via registers
 - base register holds starting address
 - Limit register holds ending address
 - Add base value to address to get a real address before main memory is accessed
 - Compare address to “limit register” to keep memory references within bounds
 - ❖ Relocation
 - by changing the base register value

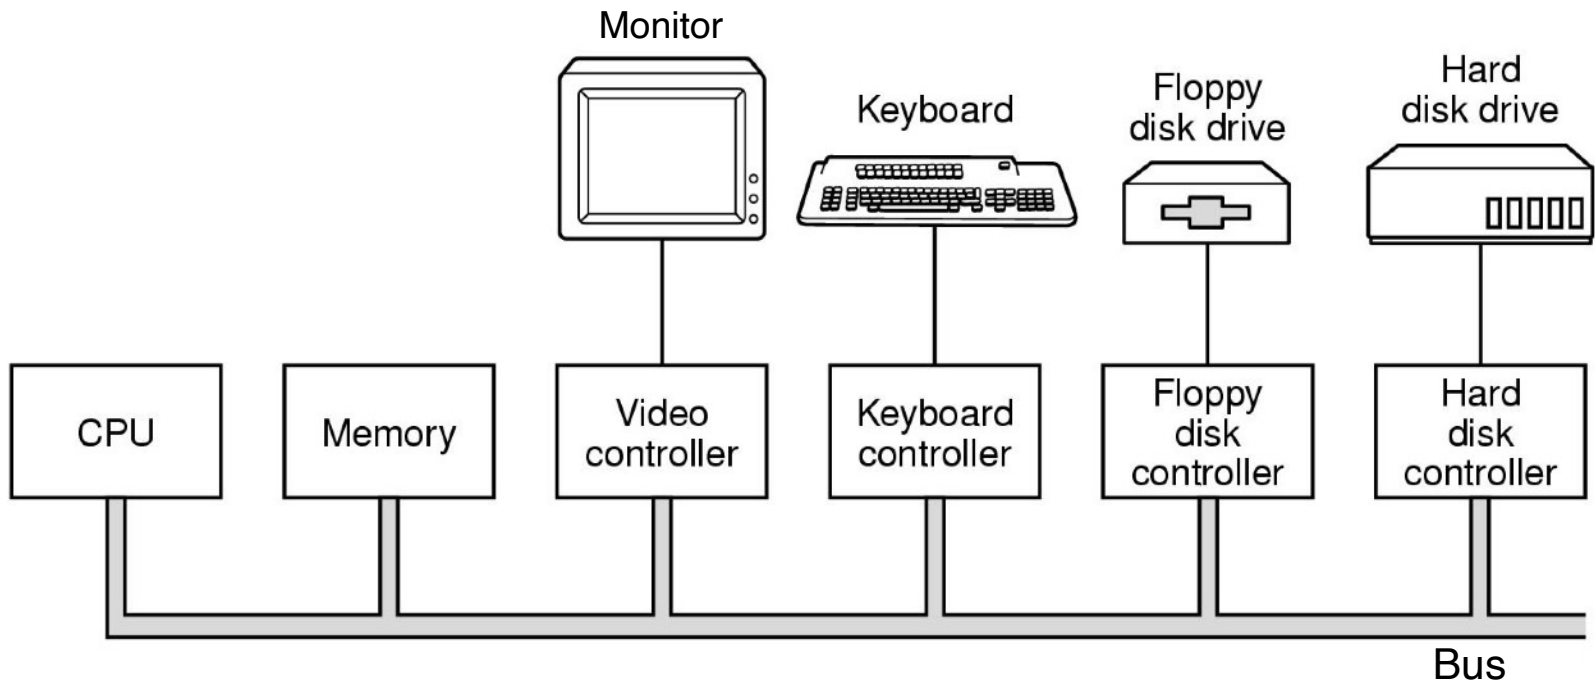
Base & Limit Registers (single & multiple)



Paged virtual memory

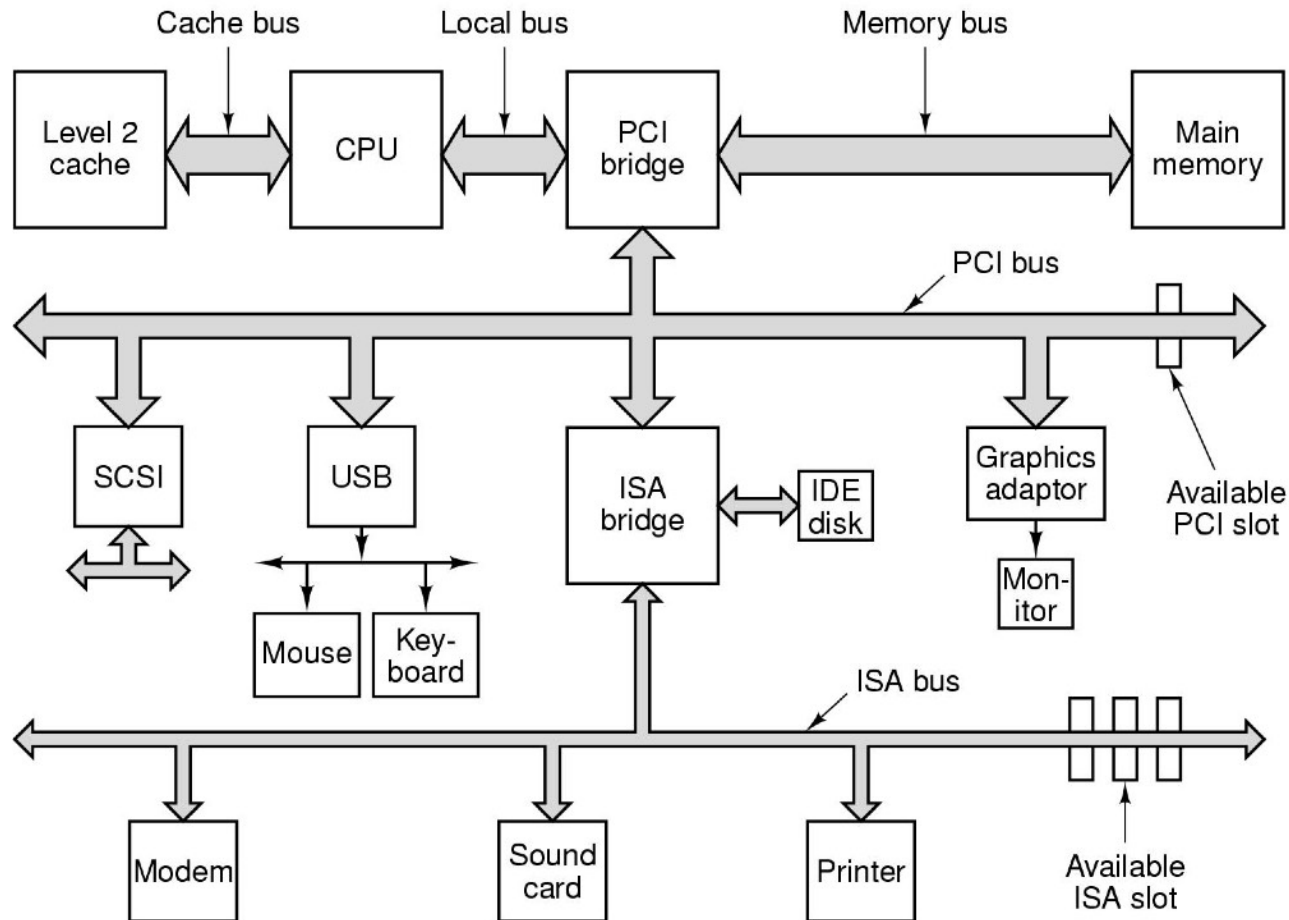
- **The same basic concept, but ...**
 - ❖ Supports non-contiguous allocation of memory
 - ❖ Allows processes to grow and shrink dynamically
 - ❖ Requires hardware support for page-based address translation
 - Sometimes referred to as a memory management unit (MMU) or a translation lookaside buffer (TLB)
 - ❖ More later ...

What about I/O devices?



A simplified view of a computer system

Structure of a Pentium system



How do programs interact with devices?

How do programs interact with devices?

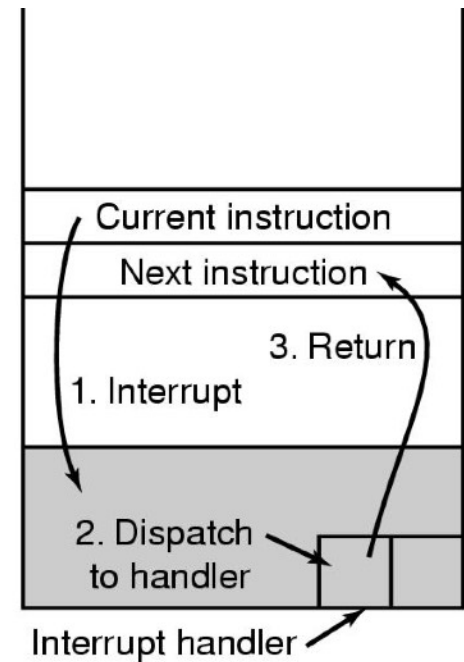
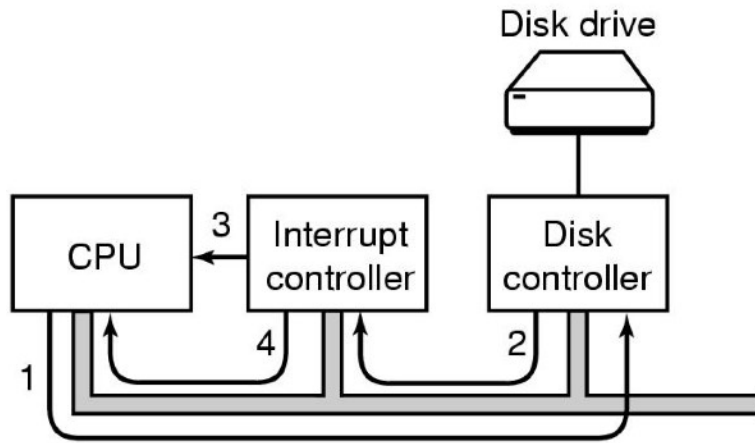
- Why protect access to devices by accessing them indirectly via the OS?
- **Devices vs device controllers vs device drivers**
 - ❖ device drivers are part of the OS (ie. Software)
 - ❖ programs call the OS which calls the device driver

How do programs interact with devices?

- ❑ Why protect access to devices by accessing them indirectly via the OS?
- ❑ **Devices vs device controllers vs device drivers**
 - ❖ device drivers are part of the OS (ie. Software)
 - ❖ programs call the OS which calls the device driver
- ❑ **Device drivers interact with device controllers**
 - ❖ either using special IO instructions
 - ❖ or by reading/writing controller registers that appear as memory locations
 - ❖ Device controllers are hardware
 - ❖ They communicate with device drivers via interrupts

How do devices interact with programs?

■ Interrupts



Different types of interrupts

- **Timer interrupts**

- ❖ Allows OS to regain control of the CPU
- ❖ One way to keep track of time

- **I/O interrupts**

- ❖ Keyboard, mouse, disks, network, etc...

- **Program generated (traps & faults)**

- ❖ Address translation faults (page fault, TLB miss)
- ❖ Programming errors: seg. faults, divide by zero, etc.
- ❖ System calls like `read()`, `write()`, `gettimeofday()`

System calls

- ❑ System calls are the mechanism by which programs communicate with the O.S.
- ❑ Implemented via a TRAP instruction
- ❑ Example UNIX system calls:
 `open(), read(), write(), close()`
 `kill(), signal()`
 `fork(), wait(), exec(), getpid()`
 `link(), unlink(), mount(), chdir()`
 `setuid(), getuid(), chown()`

The inner workings of a system call

User-level code

Process usercode

```
{  
  ...  
  read (file, buffer, n);  
  ...  
}
```

Library code

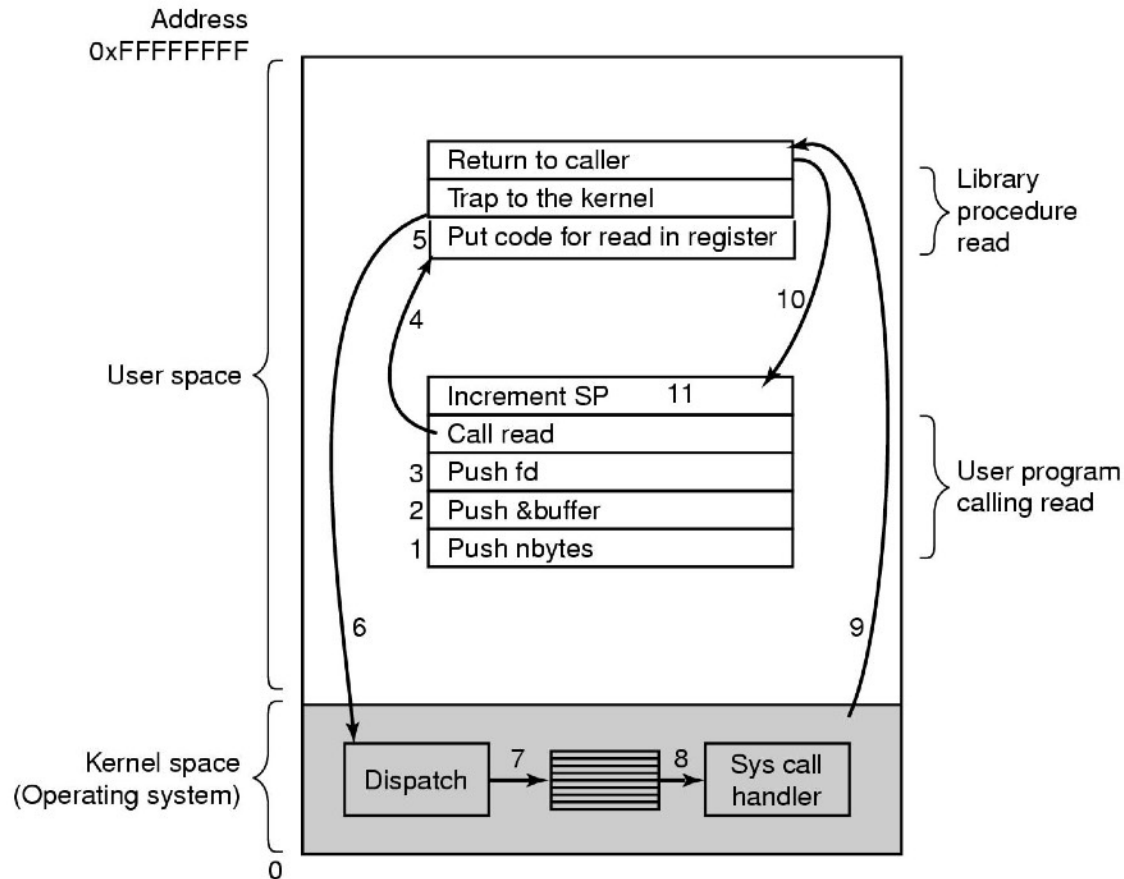
Procedure read(file, buff, n)

```
{  
  ...  
  read(file, buff, n)  
  ...  
}
```


_read:

```
LOAD r1, @SP+2  
LOAD r2, @SP+4  
LOAD r3, @SP+6  
TRAP Read_Call
```

Steps in making a read() system call



What about disks and file storage?

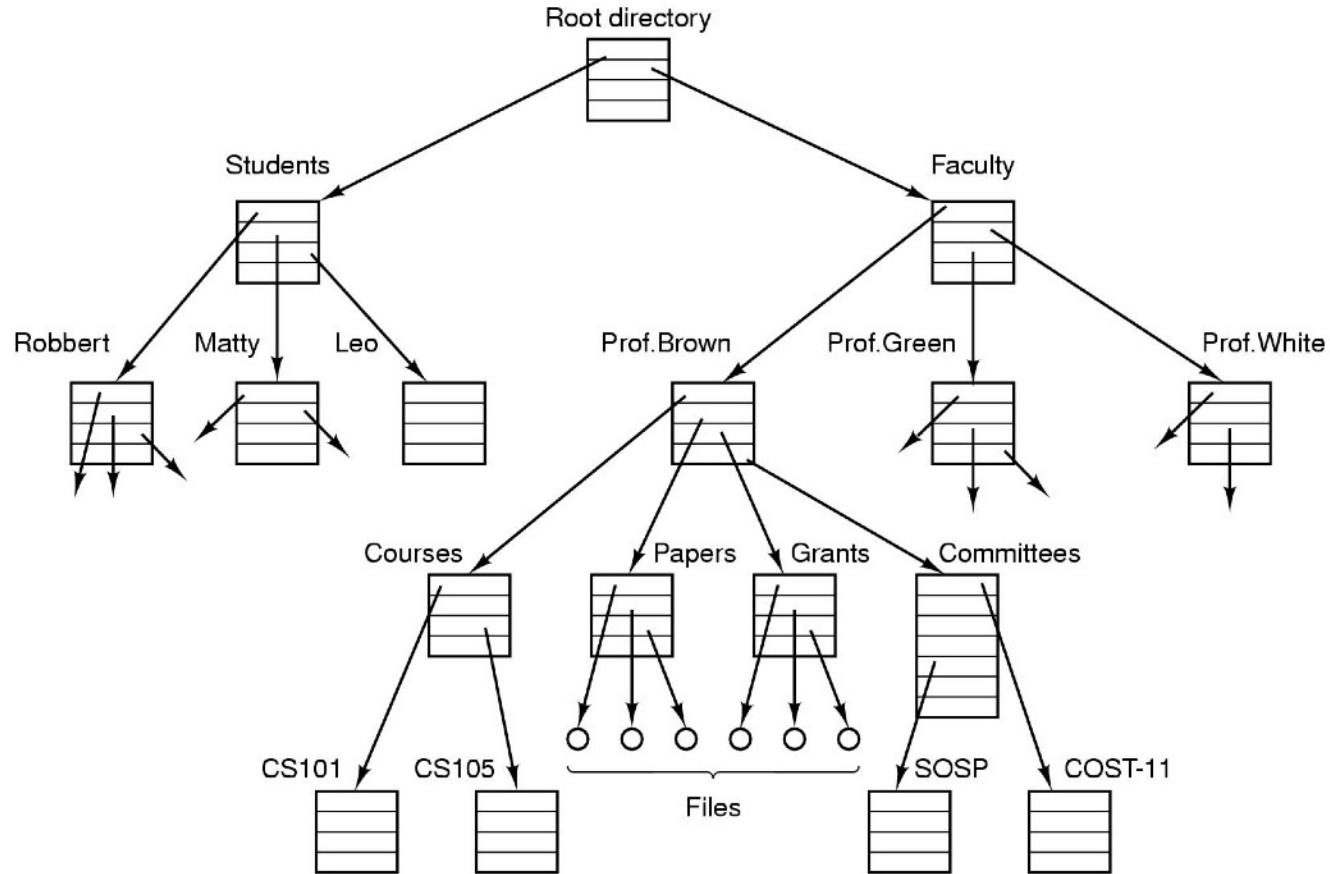


Structure of a disk drive

Disks and file storage

- ❑ **Manipulating the disk device is complicated**
 - ❖ hide some of the complexity behind disk controller, disk device driver
- ❑ **Disk blocks are not a very user-friendly abstraction for storage**
 - ❖ contiguous allocation may be difficult for large data items
 - ❖ how do you manage administrative information?
- ❑ **One application should not (automatically) be able to access another application's storage**
 - ❖ OS needs to provide a "file system"

File systems



File system - an abstraction above disk blocks

What about networks?

- ❑ **Network interfaces are just another kind of shared device/resource**
- ❑ **Need to hide complexity**
 - ❖ send and receive primitives, packets, interrupts etc
 - ❖ protocol layers
- ❑ **Need to protect the device**
 - ❖ access via the OS
- ❑ **Need to allocate resources fairly**
 - ❖ packet scheduling

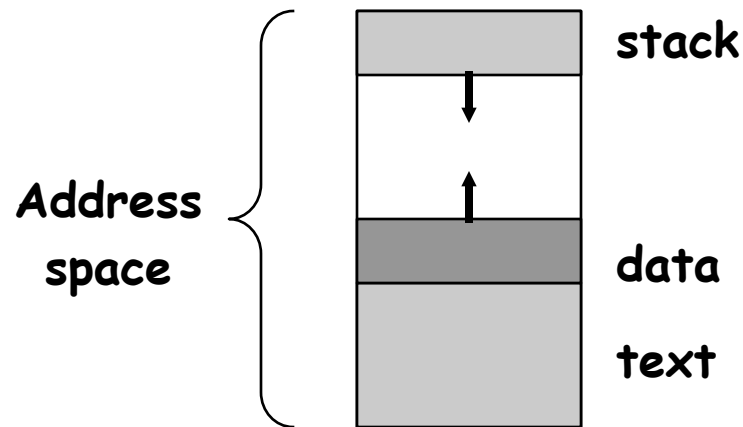
The Process Concept

The Process Concept

- **Process - a program in execution**
 - ❖ **Program**
 - description of how to perform an activity
 - instructions and static data values
 - ❖ **Process**
 - a snapshot of a program in execution
 - memory (program instructions, static and dynamic data values)
 - CPU state (registers, PC, SP, etc)
 - operating system state (open files, accounting statistics etc)

Process address space

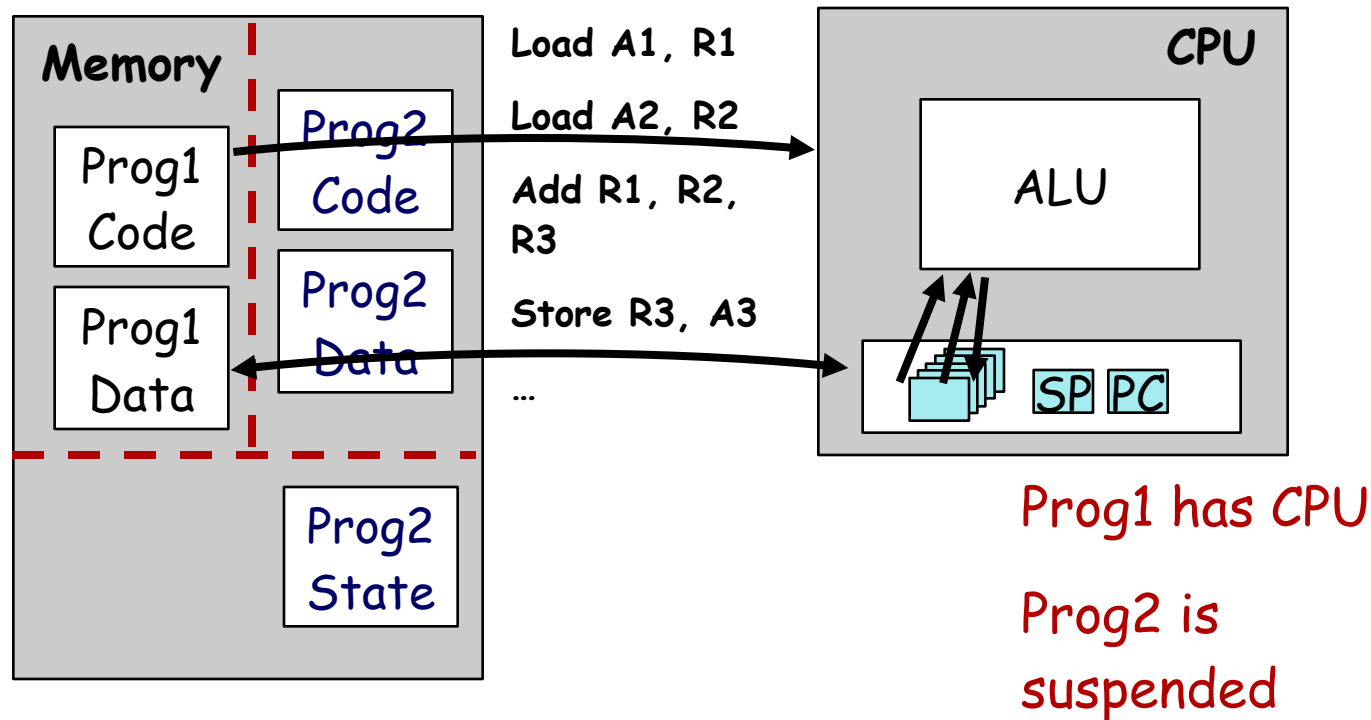
- Each process runs in its own virtual memory address space that consists of:
 - Stack space - used for function and system calls
 - Data space - variables (both static and dynamic allocation)
 - Text - the program code (usually read only)



- Invoking the same program multiple times results in the creation of multiple distinct address spaces

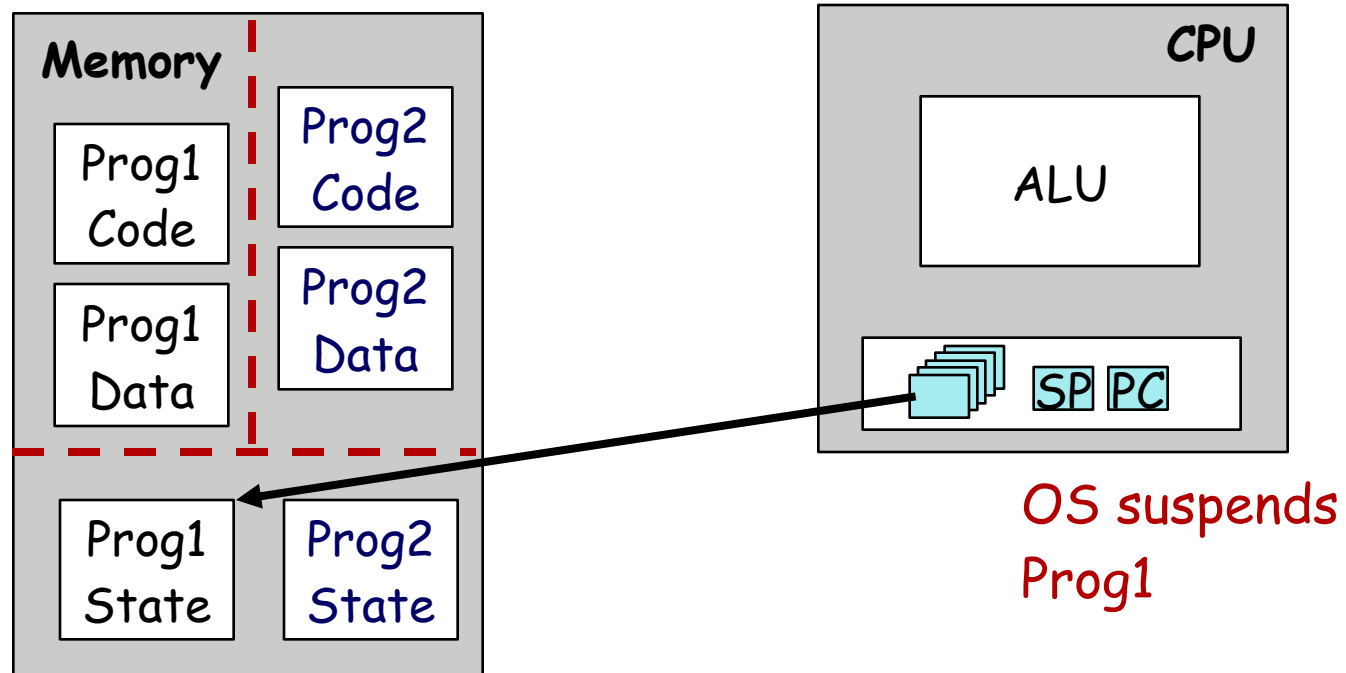
Switching among multiple processes

- Program instructions operate on operands in memory and (temporarily) in registers



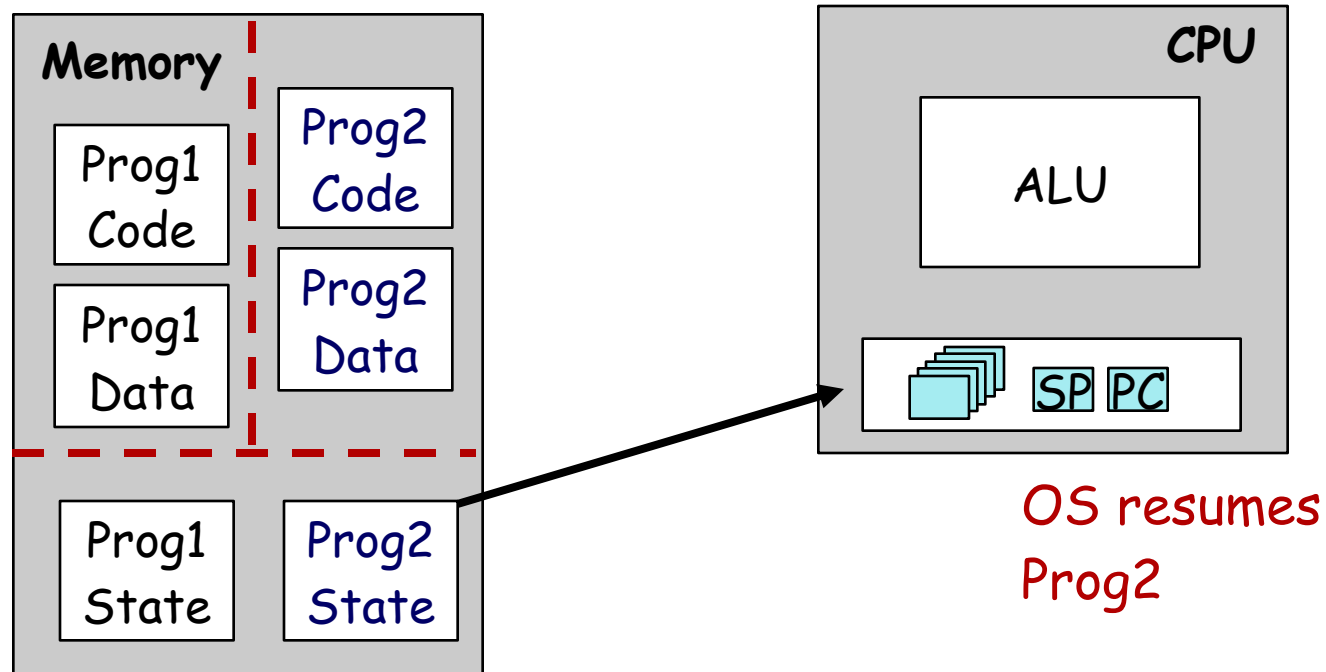
Switching among multiple processes

- Saving all the information about a process allows a process to be temporarily suspended and later resumed from the same point



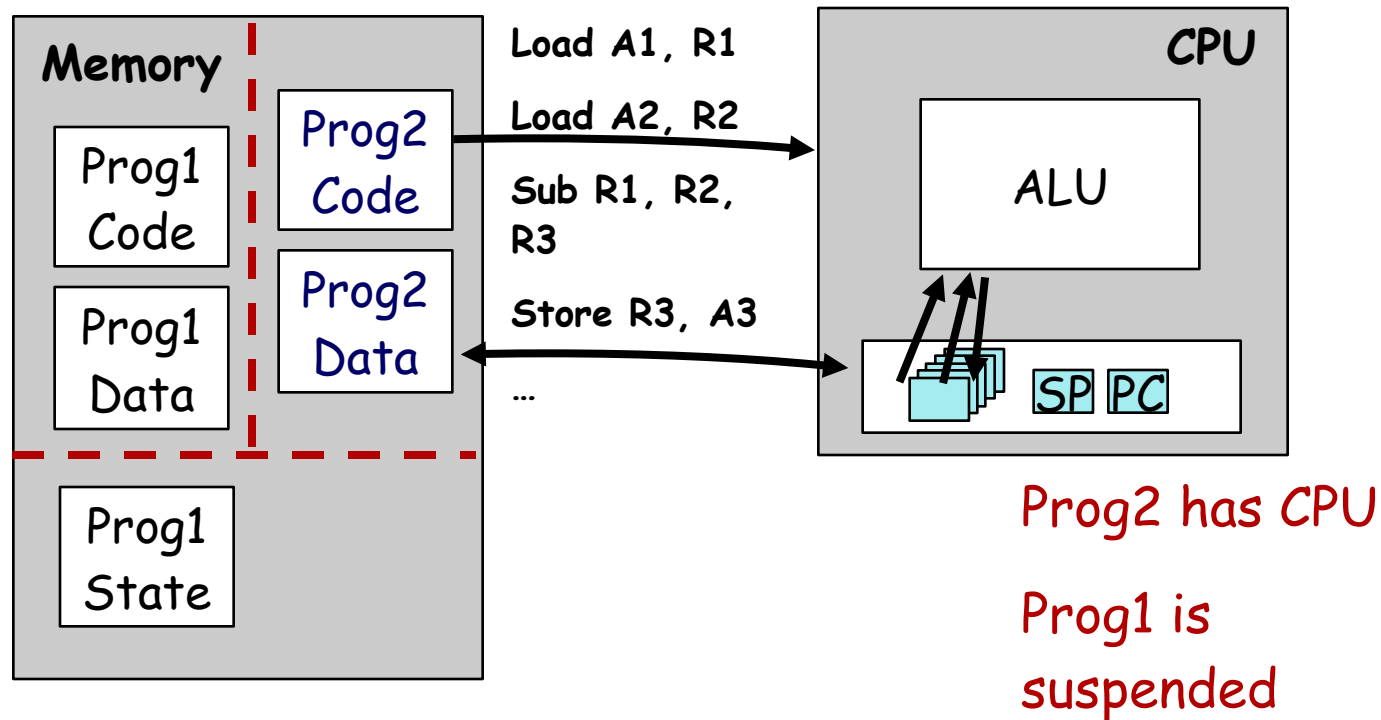
Switching among multiple processes

- Saving all the information about a process allows a process to be temporarily suspended and later resumed




Switching among multiple processes

- Program instructions operate on operands in memory and in registers




Why use the process abstraction?

One program counter




(a)

Four program counters



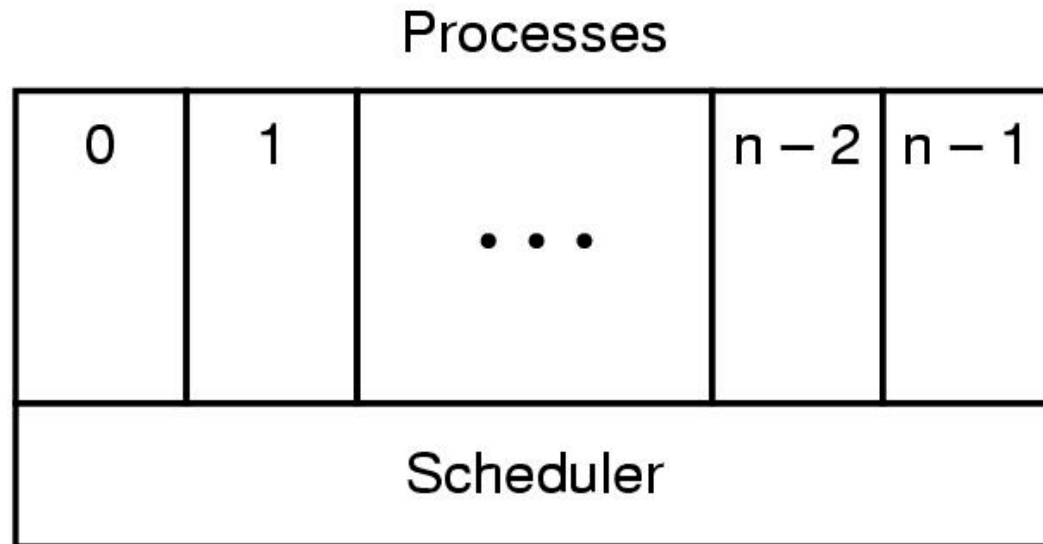
(b)



(c)

- ❑ **Multiprogramming of four programs in the same address space**
- ❑ **Conceptual model of 4 independent, sequential processes**
- ❑ **Only one program active at any instant**

The role of the scheduler



- ❑ **Lowest layer of process-structured OS**
 - ❖ handles interrupts & scheduling of processes
- ❑ **Sequential processes only exist above that layer**