

بسم الله الرحمن الرحيم

# «سیستم عامل»

۱


جلسه ۶: هم‌روندی (Concurrency)

# Interprocess communication

---

- ❑ Shared memory
- ❑ Message Passing
- ❑ Client/Server: Remote Procedure Call

# Thread



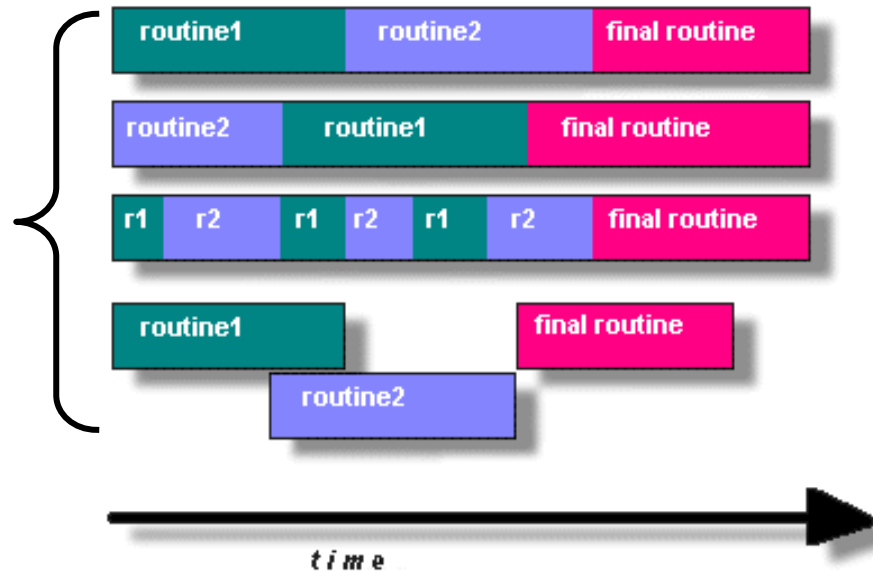
A word processor with three threads

# Concurrency

همروندی

## Assumptions:

Alternative strategies for executing multiple routines



## Example:

- ❖ One thread writes a variable
- ❖ The other thread reads from the same variable
- ❖ Problem - non-determinism:
  - The relative order of one thread's reads and the other thread's writes determines the end result!

وضعت  
رقابتي

Race  
conditions

# Race conditions

---

- ❖ A simple multithreaded program with a race:

```
i++;
```

# Race conditions

---

- ❖ A simple multithreaded program with a race:

ریسمان ۱ (i++)

...

load i to register;

increment register;

store register to i;

...

ریسمان ۲ (i--)

...

load i to register;

decrement register;

store register to i;

...

$T_0$ :	producer	execute	$register_1 = count$	{ $register_1 = 5$ }
$T_1$ :	producer	execute	$register_1 = register_1 + 1$	{ $register_1 = 6$ }
$T_2$ :	consumer	execute	$register_2 = count$	{ $register_2 = 5$ }
$T_3$ :	consumer	execute	$register_2 = register_2 - 1$	{ $register_2 = 4$ }
$T_4$ :	producer	execute	$count = register_1$	{ $count = 6$ }
$T_5$ :	consumer	execute	$count = register_2$	{ $count = 4$ }

# Race conditions

---

- ❑ **Why did this race condition occur?**
  - ❖ two or more threads have an inconsistent view of a shared memory region (I.e., a variable)
  - ❖ values of memory locations replicated in registers during execution
  - ❖ context switches at arbitrary times during execution
  - ❖ threads can see "stale" memory values in registers



# Race Conditions

---

- ❑ Race condition: whenever the output depends on the precise execution order of the processes!
- ❑ What solutions can we apply?
  - ❖ prevent context switches by preventing interrupts
  - ❖ make threads coordinate with each other to ensure mutual exclusion in accessing critical sections of code

**mutex —**

وضعت  
رقابتي

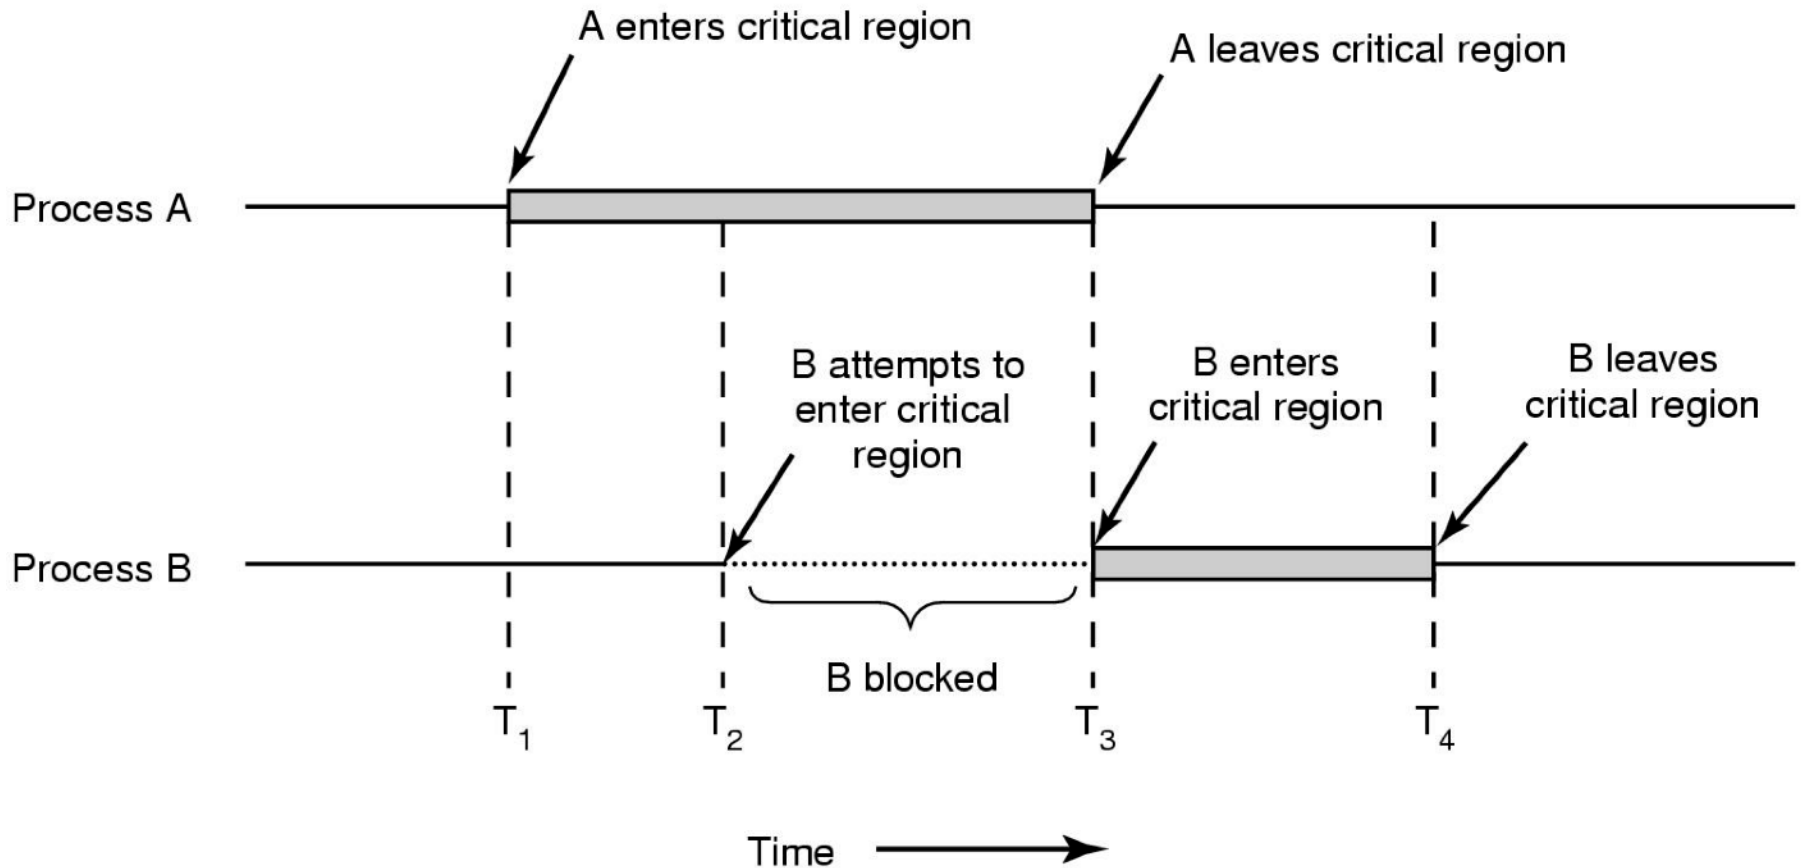
**Race  
conditions**

# Mutual exclusion conditions

---

- ❑ **Mutual exclusion.** No two processes simultaneously in critical section
  - ❑ No assumptions made about speeds or numbers of CPUs
- ❑ **Progress.** No process running outside its critical section may block another process
- ❑ **Bounded waiting.** No process must wait forever to enter its critical section

# Using mutual exclusion for critical sections



# How can we enforce mutual exclusion?

---

- ❑ What about using locks ?
- ❑ Locks solve the problem of exclusive access to shared data.
  - ❖ Acquiring a lock prevents concurrent access
  - ❖ Expresses intention to enter critical section
- ❑ **Assumption:**
  - ❖ Each shared data item has an associated lock
  - ❖ All threads set the lock before accessing the shared data
  - ❖ Every thread releases the lock after it is done

# Acquiring and releasing locks

---

Thread B

Thread C

Thread A

Thread D

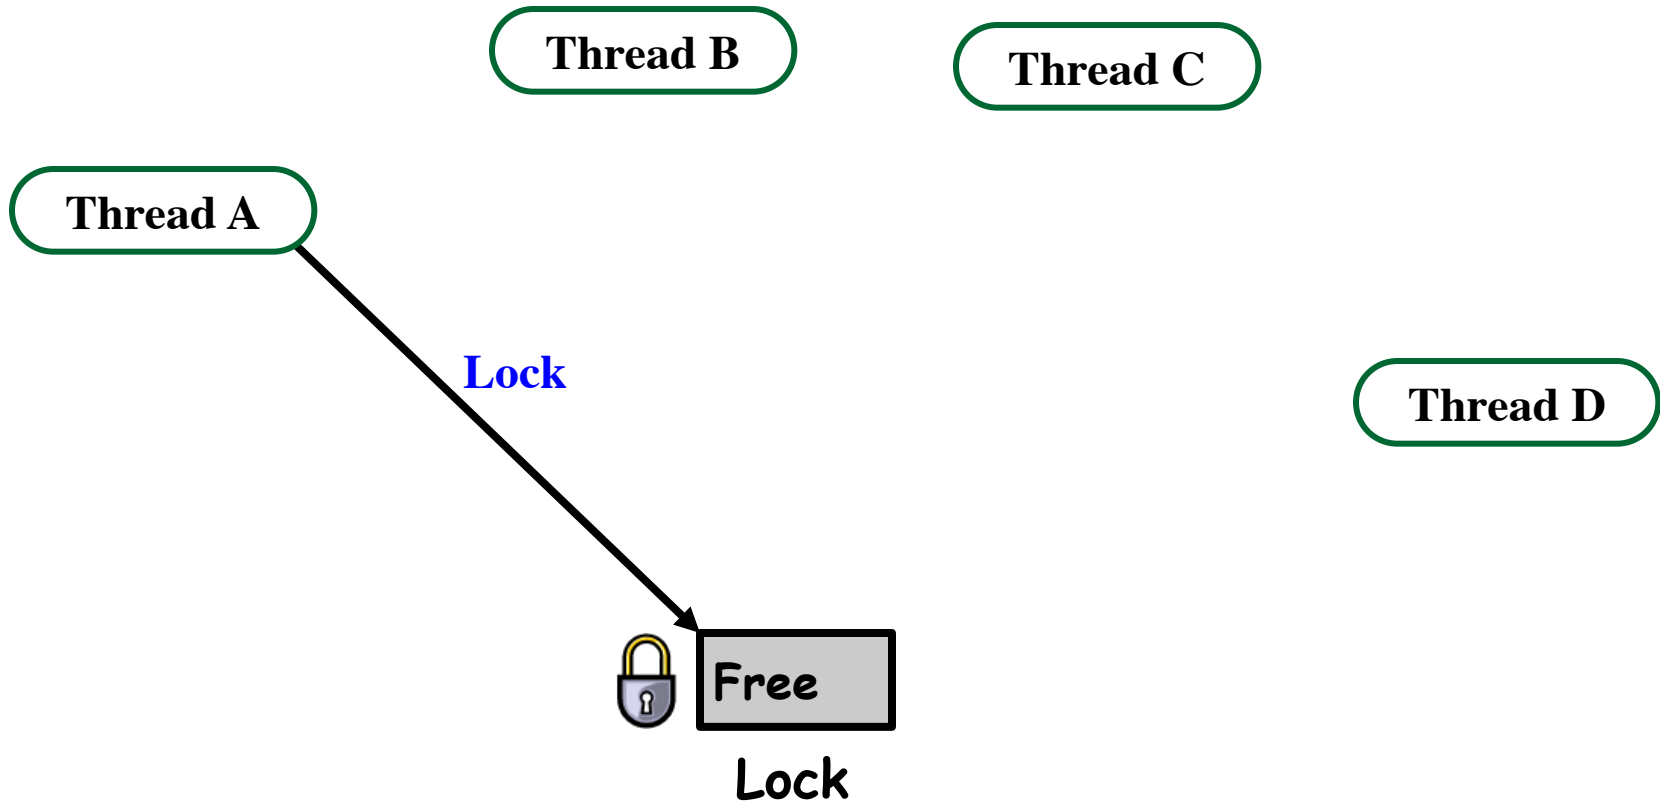


Free

Lock

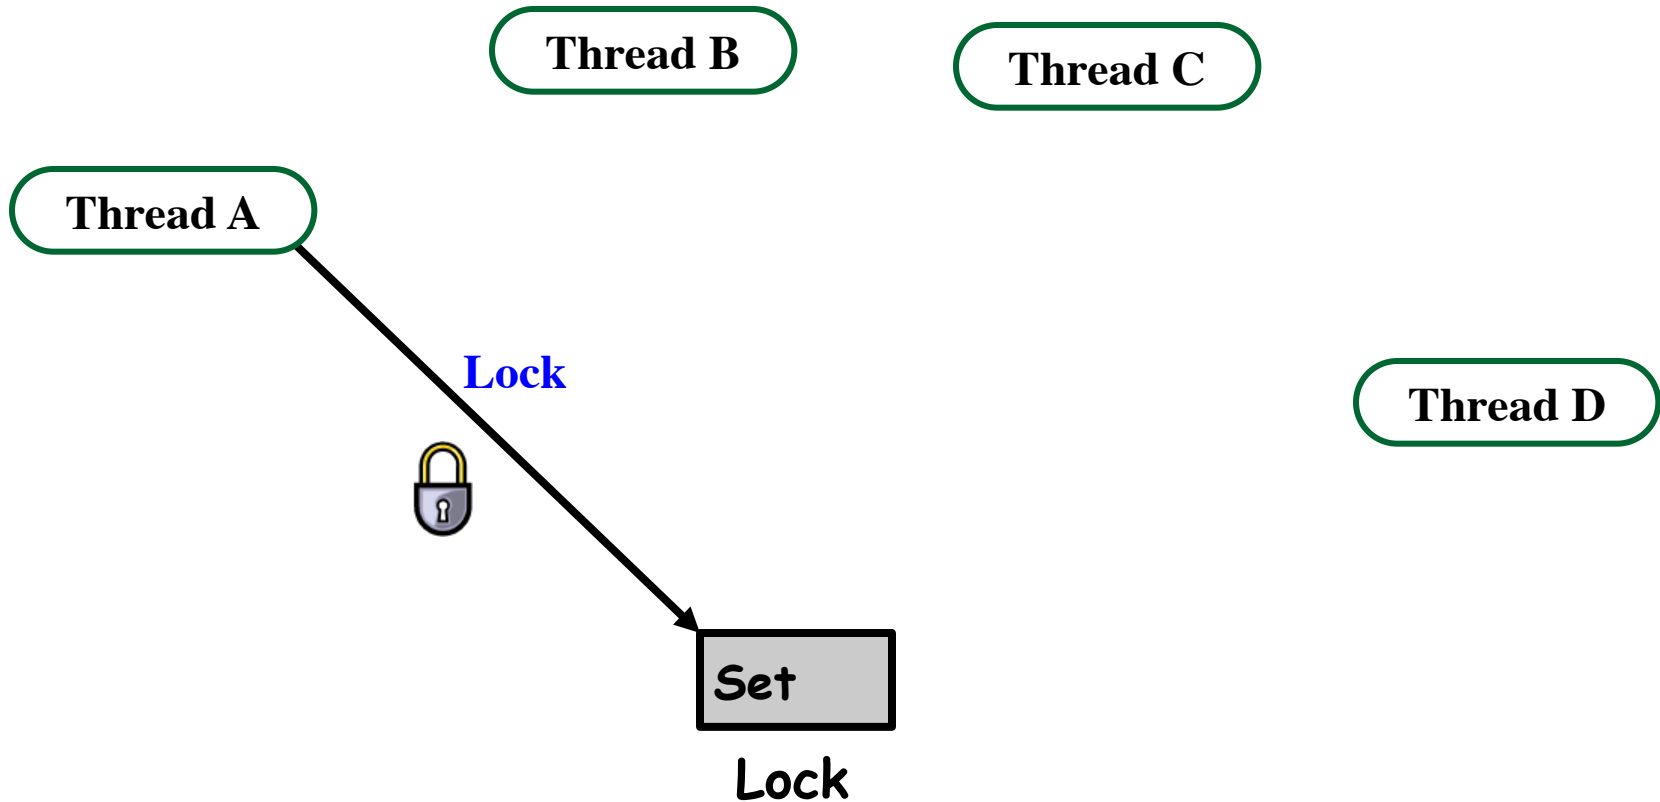
# Acquiring and releasing locks

---



# Acquiring and releasing locks

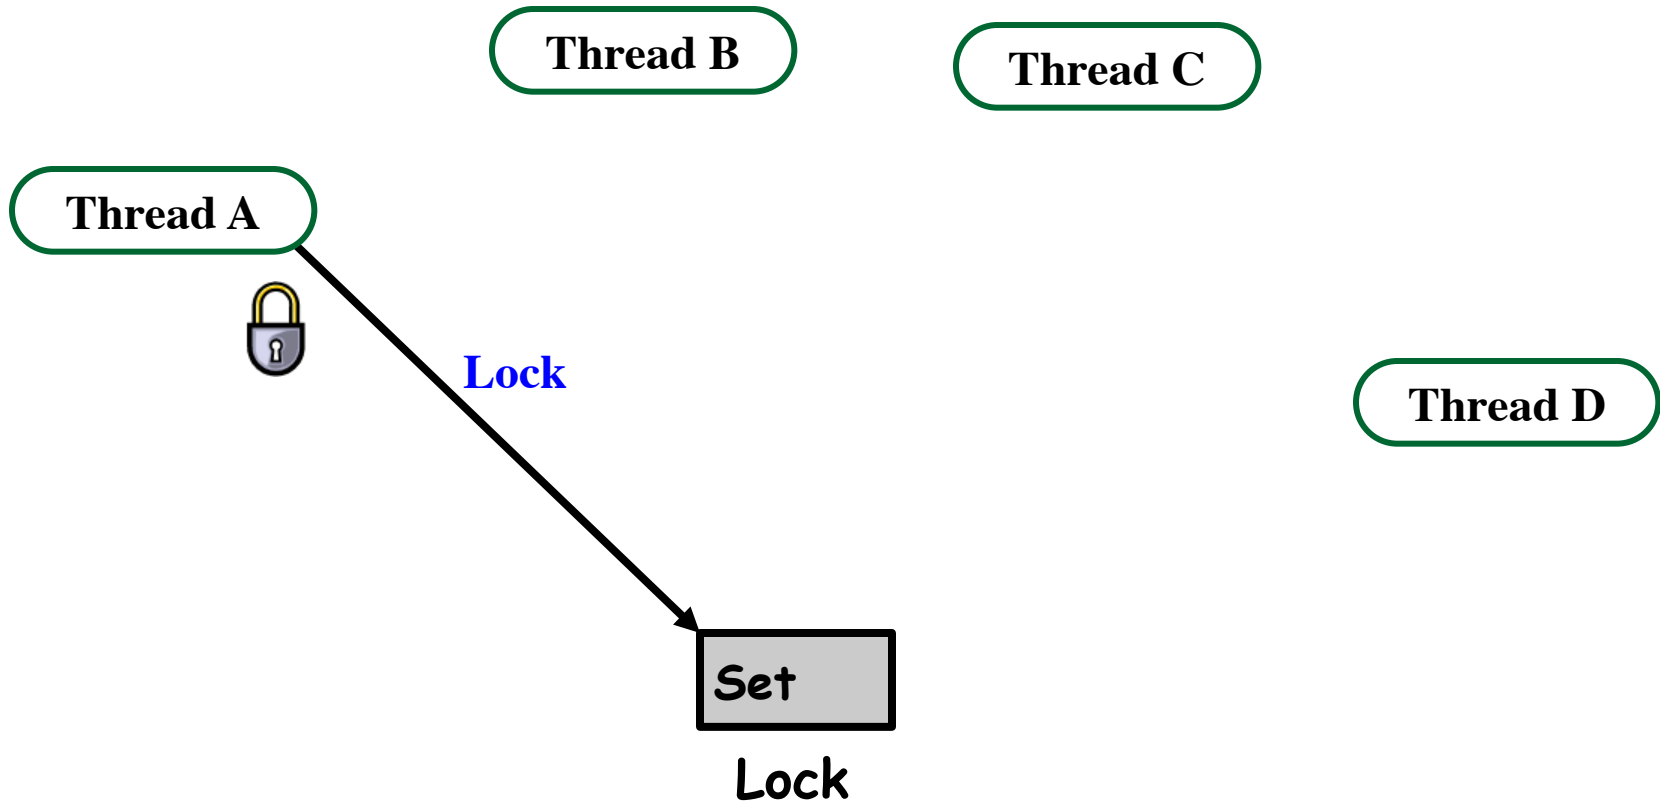
---





# Acquiring and releasing locks

---



# Acquiring and releasing locks

---

Thread B

Thread C

Thread A



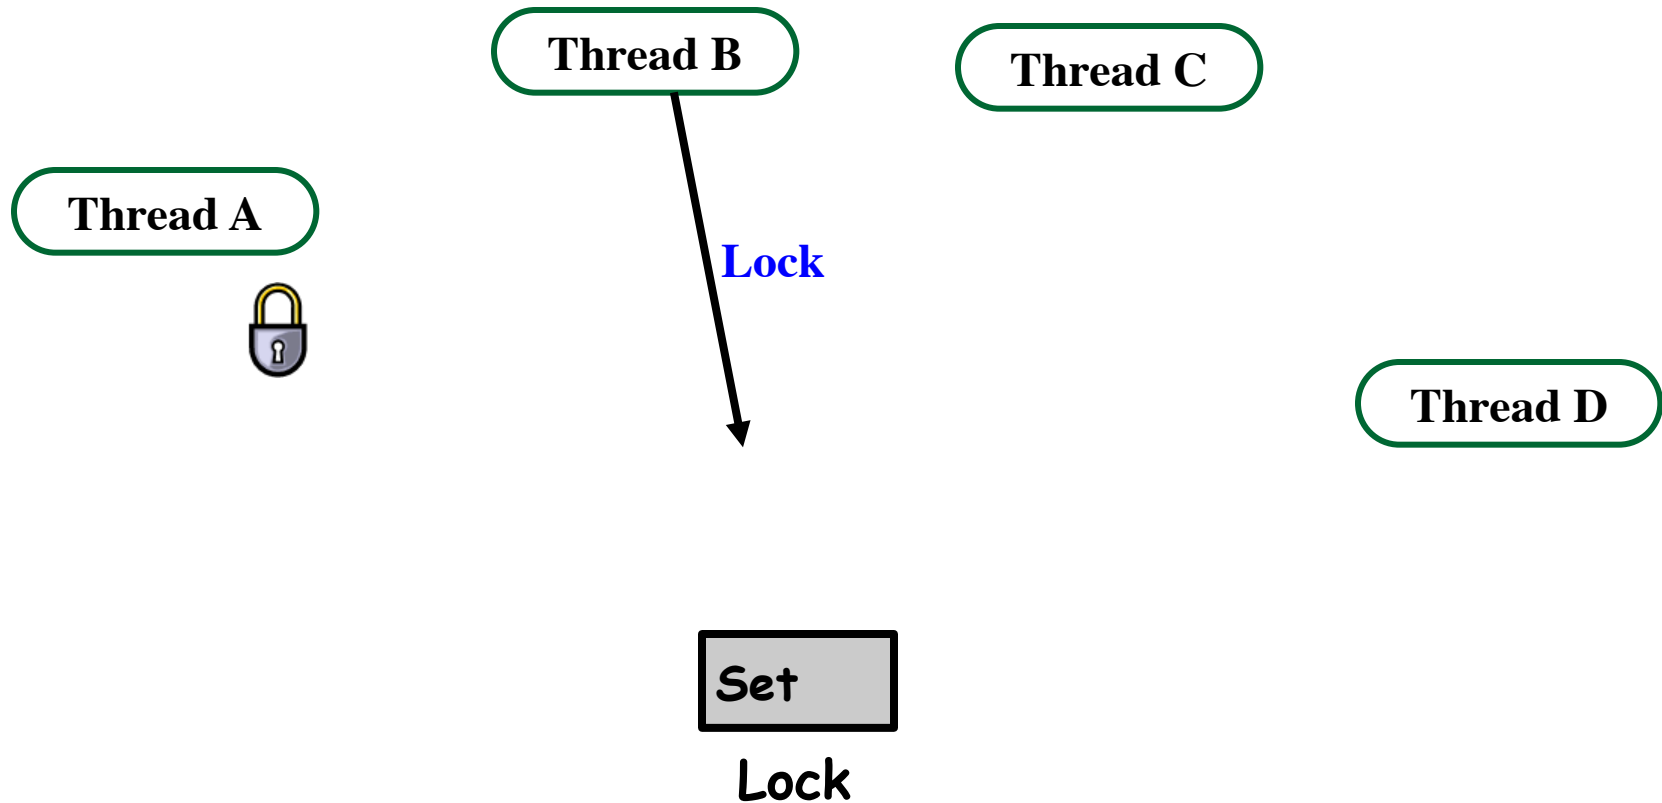
Thread D

Set

Lock

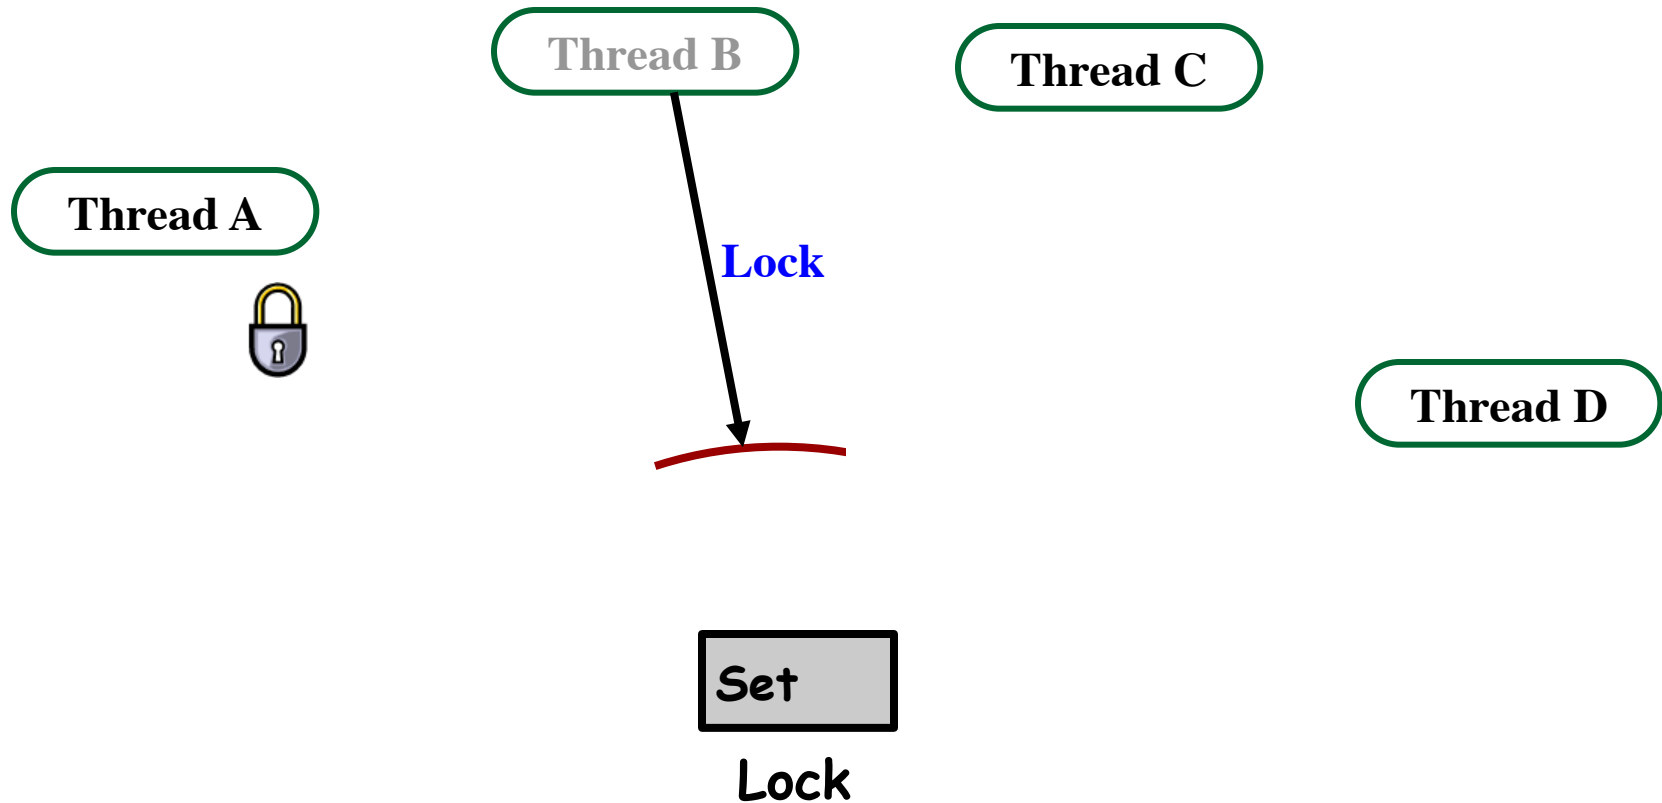
# Acquiring and releasing locks

---



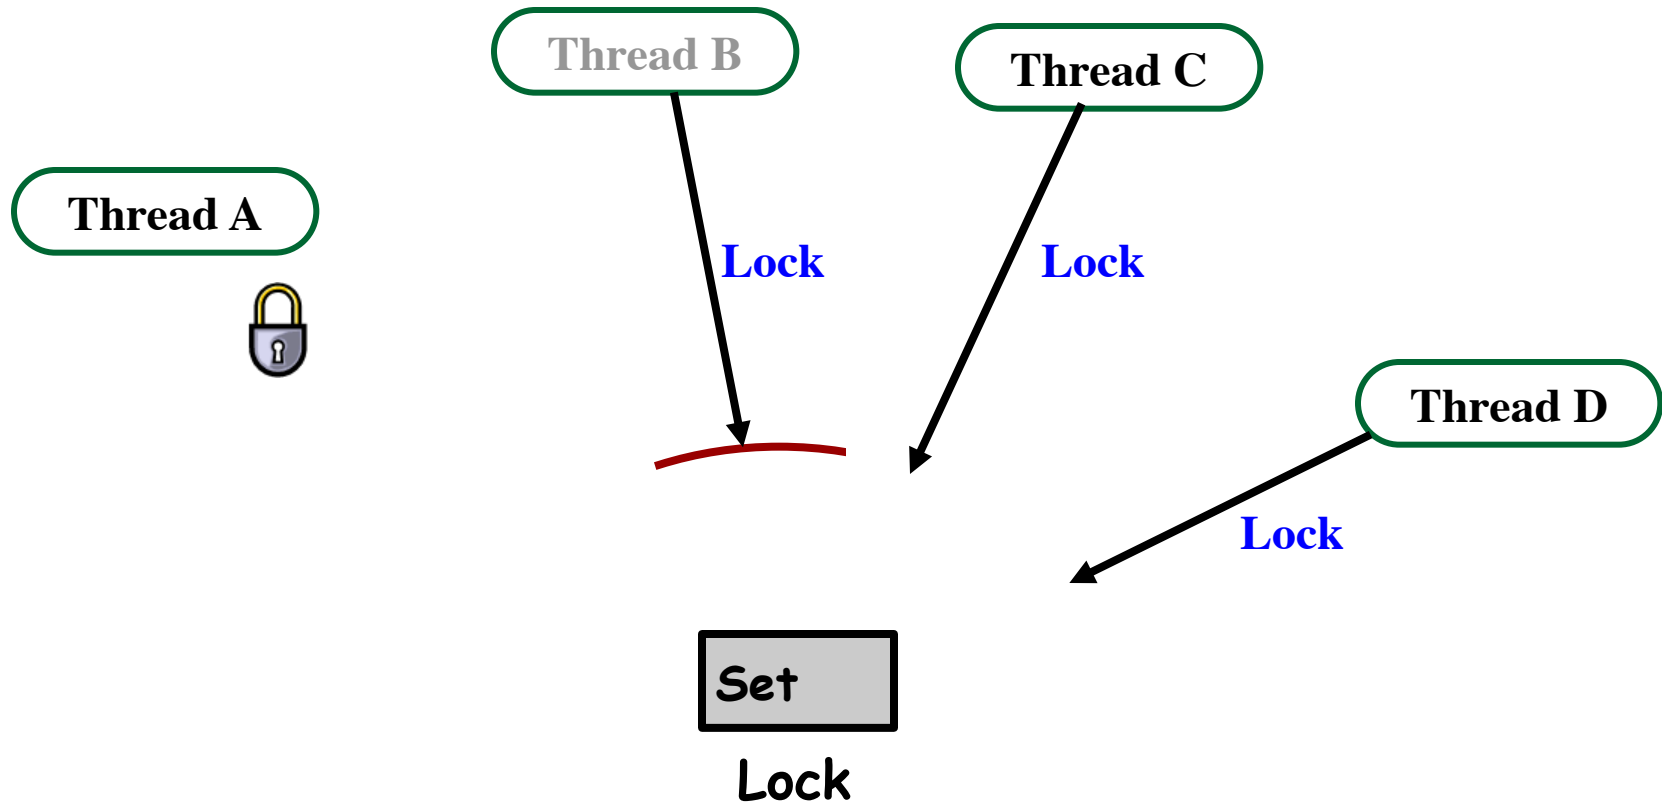
# Acquiring and releasing locks

---




# Acquiring and releasing locks

---




# Acquiring and releasing locks

---




# Acquiring and releasing locks

---



# Acquiring and releasing locks


---





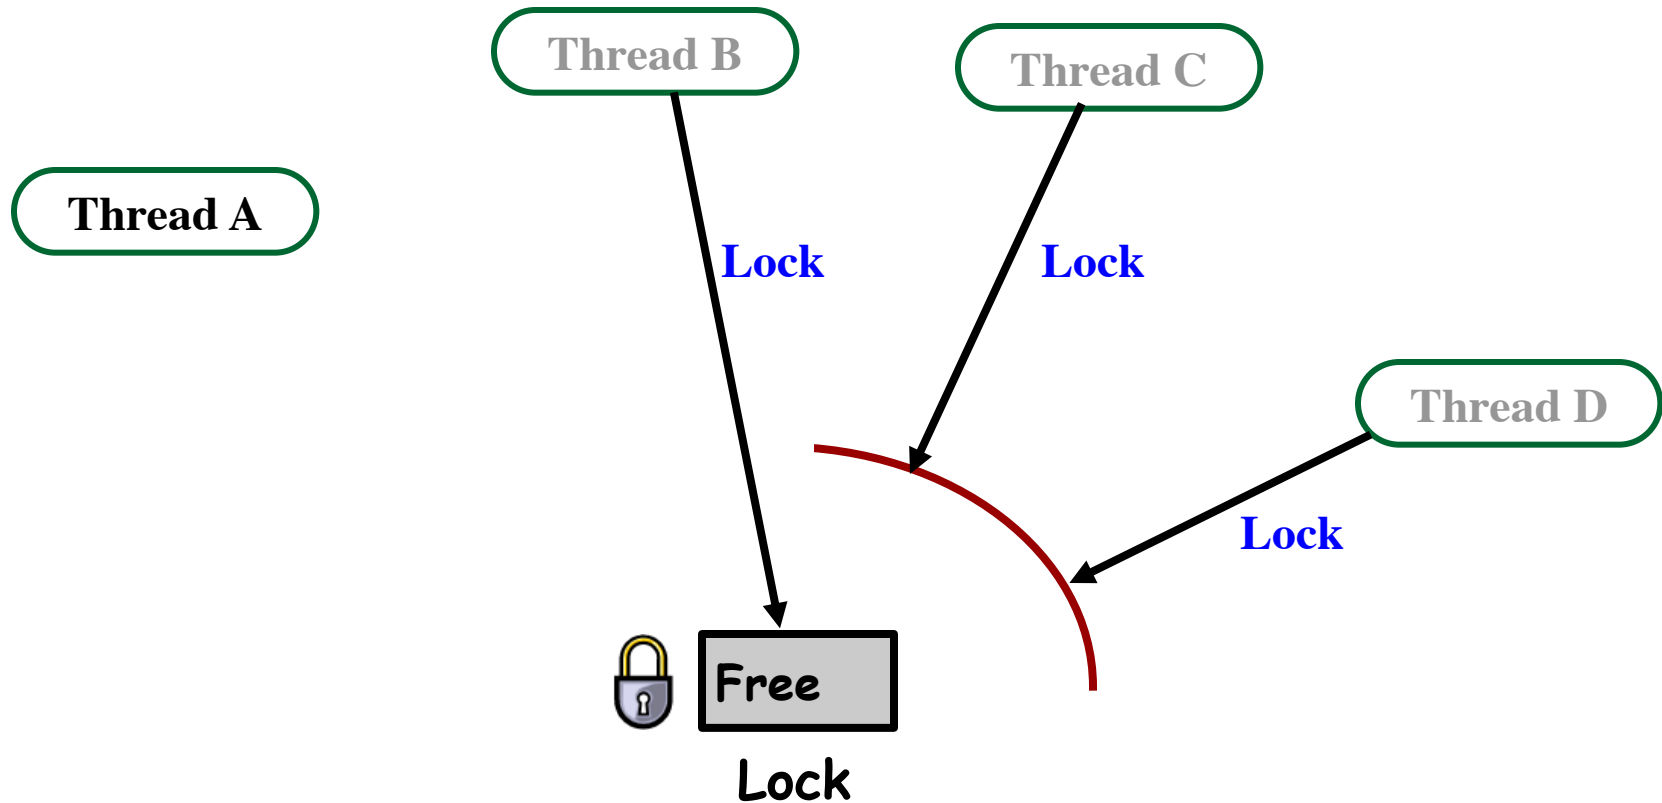
# Acquiring and releasing locks

---



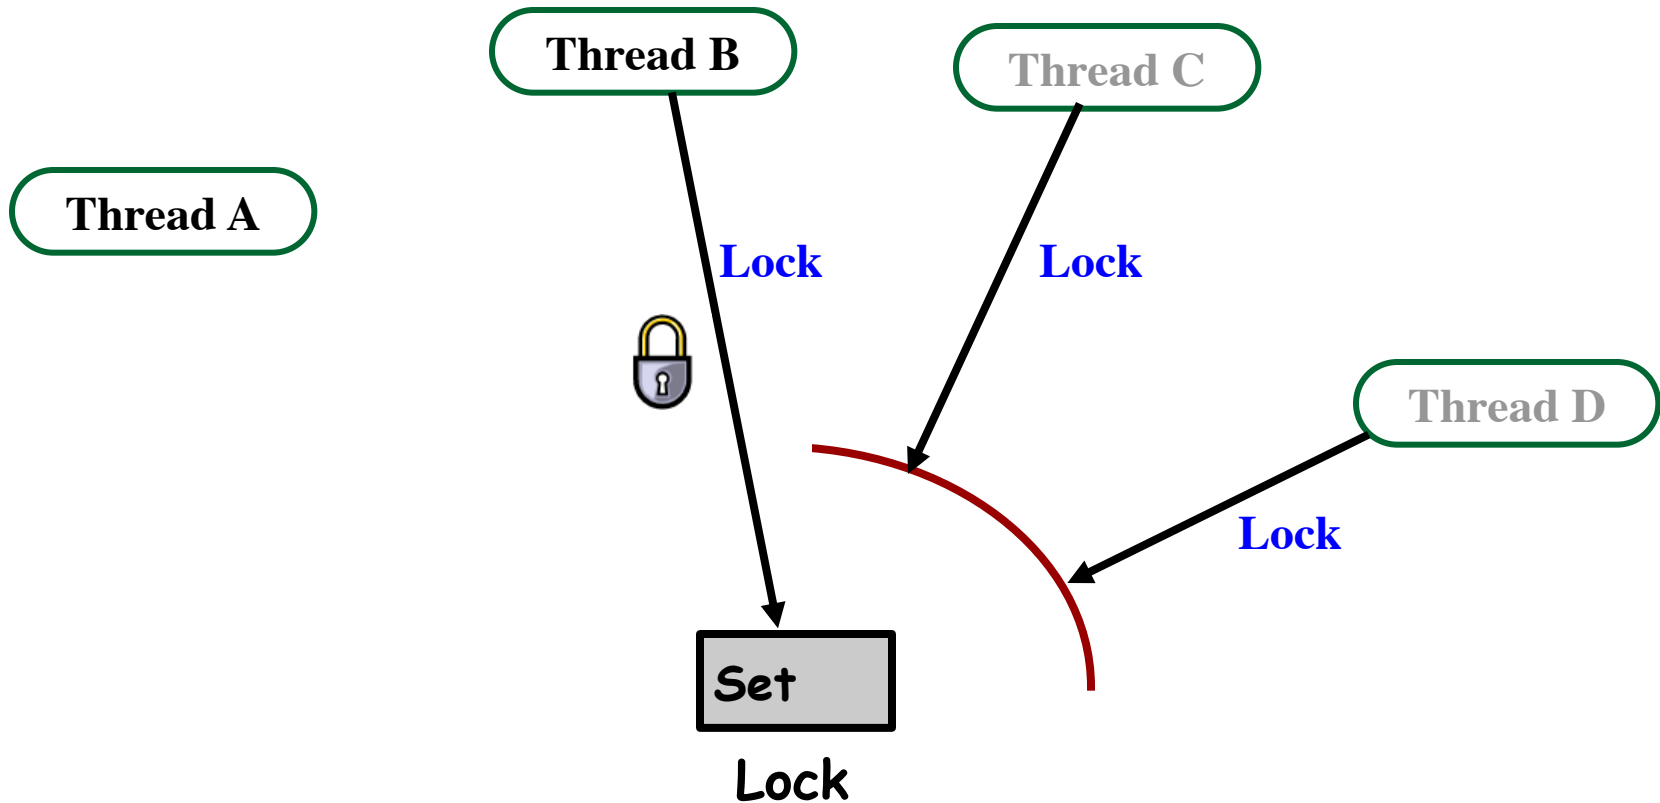
# Acquiring and releasing locks

---



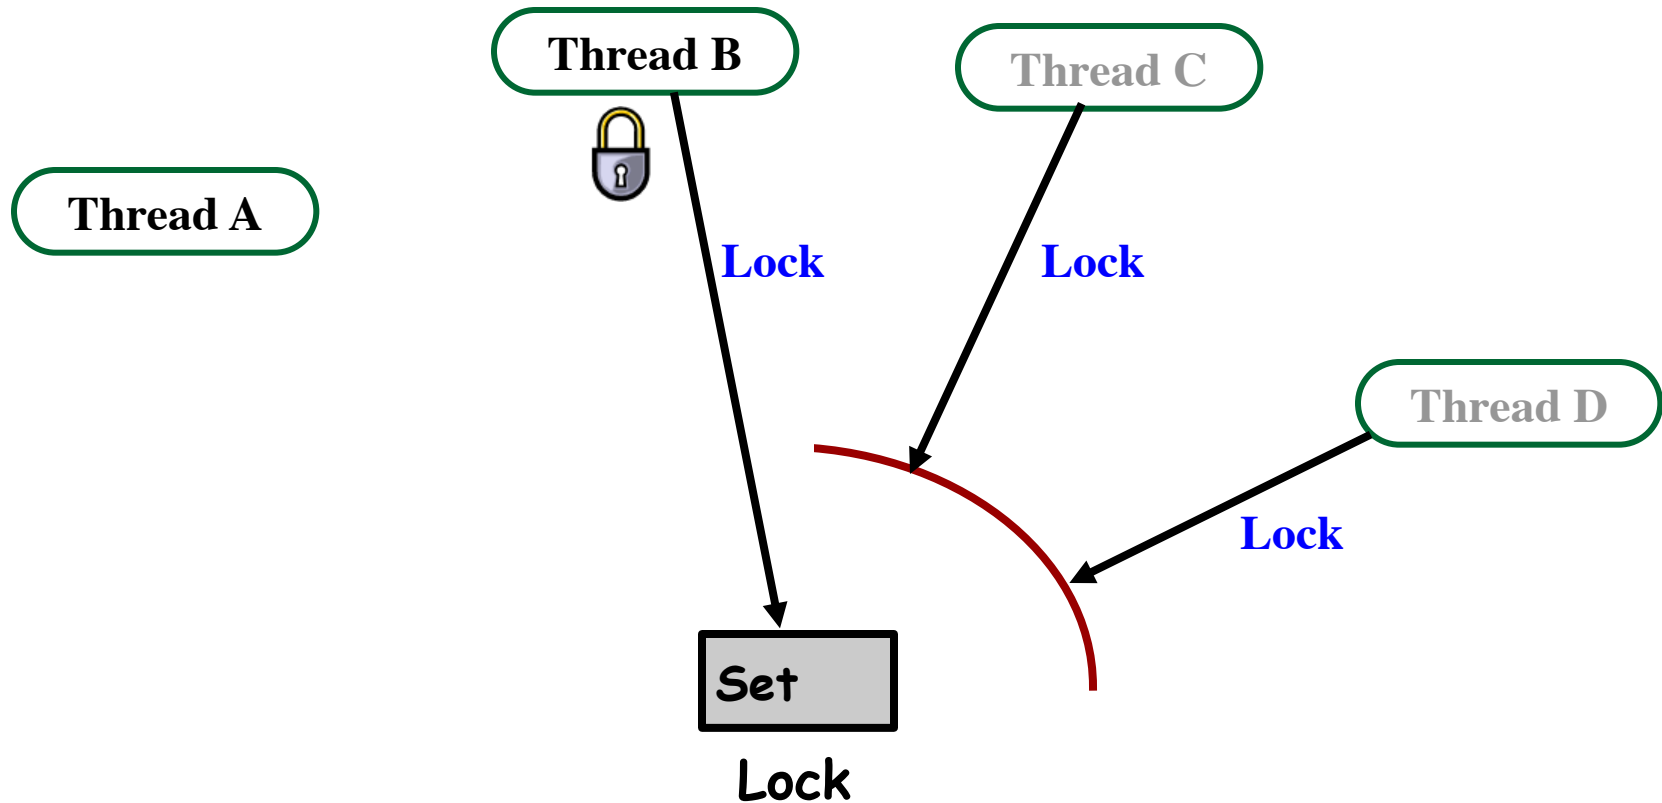
# Acquiring and releasing locks

---



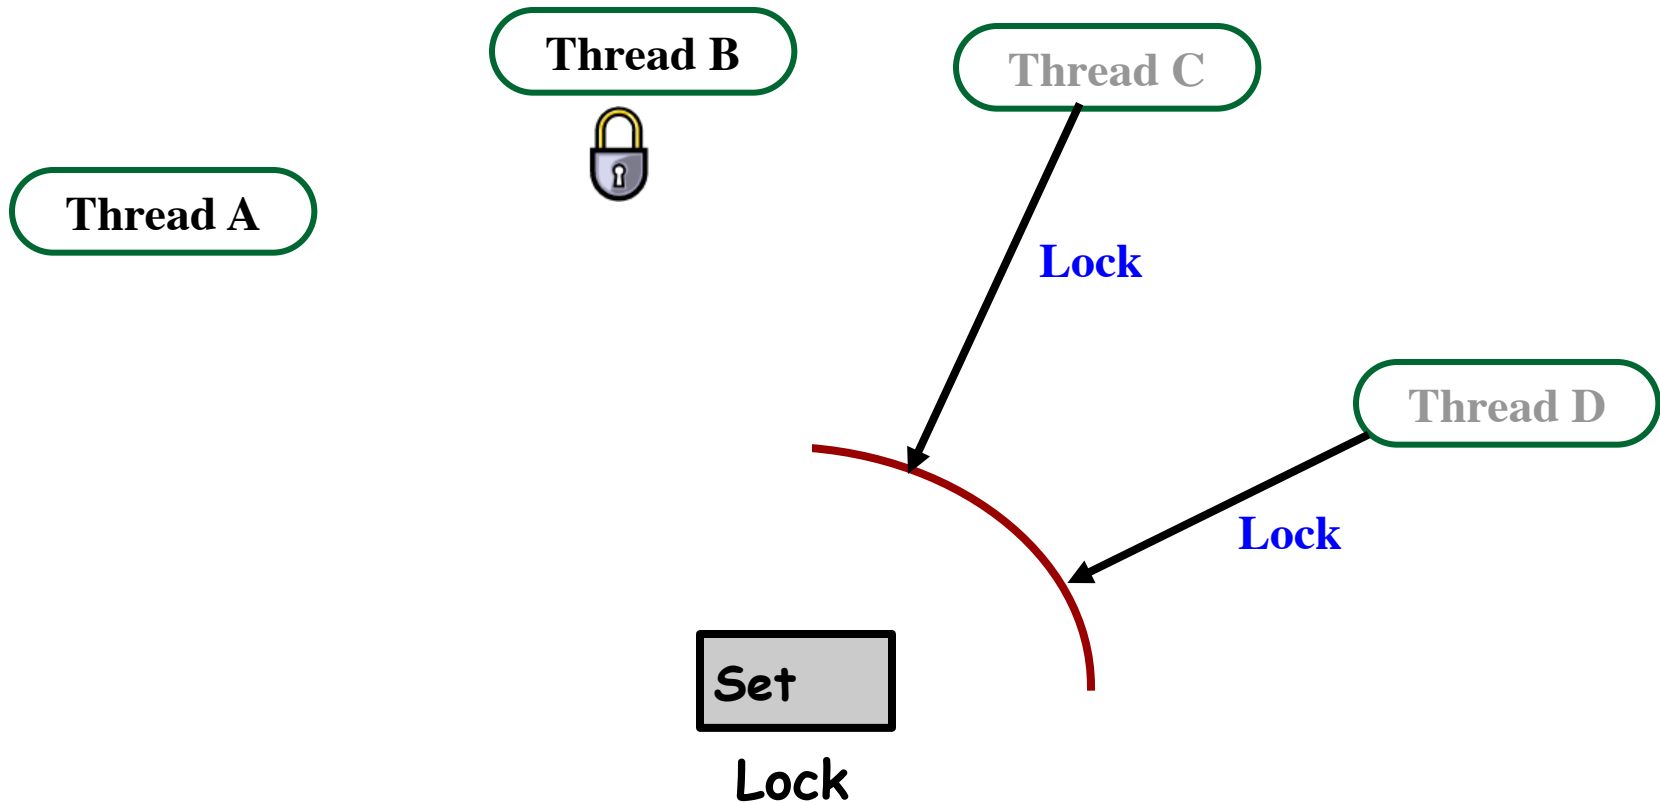
# Acquiring and releasing locks

---



# Acquiring and releasing locks

---



# Mutual exclusion (mutex) locks

---

- ❑ An abstract data type
- ❑ Used for synchronization
- ❑ The mutex is either:
  - ❖ Locked ("the lock is held")
  - ❖ Unlocked ("the lock is free")

# Mutex lock operations

---

- ❑ **Lock (mutex)**
  - ❖ Acquire the lock if it is free ... and continue
  - ❖ Otherwise wait until it can be acquired
- ❑ **Unlock (mutex)**
  - ❖ Release the lock
  - ❖ If there are waiting threads wake up one of them

# How to use a mutex?

---

## Shared data:

Mutex myLock;

```
1 repeat
2   Lock(myLock) ;
3   critical section
4   Unlock(myLock) ;
5   remainder section
6 until FALSE
```

```
1 repeat
2   Lock(myLock) ;
3   critical section
4   Unlock(myLock) ;
5   remainder section
6 until FALSE
```



# But how can we implement a mutex?

---

- ❑ What if the lock is a binary variable
- ❑ How would we implement the lock and unlock procedures?

# But how can we implement a mutex?

---

- ❑ **Lock** and **Unlock** operations must be atomic !
- ❑ Many computers have some limited hardware support for setting locks
  - ❖ Atomic Test and Set Lock instruction
  - ❖ Atomic compare and swap operation
- ❑ These can be used to implement mutex locks

# Test-and-set-lock instruction (TSL, tset)

---

- A lock is a single word variable with two values
  - ❖ 0 = FALSE = not locked
  - ❖ 1 = TRUE = locked
- Test-and-set does the following atomically:
  - ❖ Get the (old) value
  - ❖ Set the lock to TRUE
  - ❖ Return the old value

**If** the returned value was FALSE...


**Then** you got the lock!!!

**If** the returned value was TRUE...

**Then** someone else has the lock  
    (so try again later)

# Test and set lock

---




FALSE

Lock


# Test and set lock

---




# Test and set lock

---




# Test and set lock

---



# Test and set lock


---






# Test and set lock

---




# Test and set lock

---




# Test and set lock

---




# Test and set lock

---



# Test and set lock

---



# Using TSL directly for critical sections

---

```
1 repeat I  
2   while(TSL(lock))  
3     no-op;  
4   critical section  
5   Lock = FALSE;  
6   remainder section  
7 until FALSE
```

```
1 repeat J  
2. while(TSL(lock))  
3   no-op;  
4   critical section  
5   Lock = FALSE;  
6   remainder section  
7 until FALSE
```

- Guarantees that only one thread at a time will enter its critical section

# Implementing a mutex with TSL

---

```
1 repeat
2   while(TSL(mylock)) }      Lock (mylock)
3   no-op;
4   critical section
5   mylock = FALSE; }      Unlock (mylock)
6   remainder section
7 until FALSE
```

- Note that processes are **busy** while waiting
  - ❖ this kind of mutex is called a **spin lock**

# Busy waiting

---

- ❑ Also called polling or spinning
  - ❖ The thread consumes CPU cycles to evaluate when the lock becomes free !
- ❑ Problem on a single CPU system...
  - ❖ A busy-waiting thread can prevent the lock holder from running & completing its critical section & releasing the lock!
    - time spent spinning is wasted on a single CPU system
  - ❖ Why not block instead of busy wait ?



# Blocking synchronization primitives

---

- ❑ **Sleep**
  - ❖ Put a thread to sleep
  - ❖ Thread becomes BLOCKED
- ❑ **Wakeup**
  - ❖ Move a BLOCKED thread back onto "Ready List"
  - ❖ Thread becomes READY (or RUNNING)
- ❑ **Yield**
  - ❖ Put calling thread on ready list and schedule next thread
  - ❖ Does not BLOCK the calling thread!
    - Just gives up the current time-slice

کنترل  
هم‌روندی در  
هسته

# But how can these be implemented?

---

- ❑ **In User Programs:**
  - ❖ System calls to the kernel
- ❑ **In Kernel:**
  - ❖ Calls to the thread **scheduler** routines

# Concurrency control in user programs

---

- ❑ User threads call sleep and wakeup system calls
- ❑ Scheduler routines in the kernel implement sleep and wakeup
  - ❖ they manipulate the “ready list”
  - ❖ but the ready list is **shared data**
  - ❖ the code that manipulates it is a **critical section**
    - What if a timer interrupt occurs during a sleep or wakeup call?
- ❑ **Problem:**
  - ❖ How can scheduler routines be programmed to execute correctly in the face of concurrency?

# Concurrency in the kernel

---

## **Solution 1: Disable interrupts during critical sections**

- ❖ Ensures that interrupt handling code will not run
- ❖ ... but what if there are multiple CPUs?

## **Solution 2: Use mutex locks based on TSL for critical sections**

- ❖ Ensures mutual exclusion for all code that follows that convention
- ❖ ... but what if your hardware doesn't have TSL?

# Disabling interrupts

---

- ❑ Disabling interrupts in the OS vs disabling interrupts in user processes
  - ❖ why not allow user processes to disable interrupts?
  - ❖ is it ok to disable interrupts in the OS?
  - ❖ what precautions should you take?

# Disabling interrupts in the kernel

---

Scenario 1:

A thread is running; wants to access shared data

Disable interrupts

Access shared data ("critical section")

Enable interrupts

# Disabling interrupts in the kernel

---

**Problem:**

**Interrupts are already disabled and a thread wants to access the critical section  
...using the above sequence...**

- ▣ **Ie. One critical section gets nested inside another**



# Disabling interrupts in the kernel

---

**Problem: Interrupts are already disabled.**

- ❖ Thread wants to access critical section using the previous sequence...

Save previous interrupt status (enabled/disabled)

Disable interrupts

Access shared data ("critical section")

Restore interrupt status to what it was before

# Disabling interrupts is not enough on MPs...

---

- ❑ **Disabling interrupts during critical sections**
  - ❖ Ensures that interrupt handling code will not run
  - ❖ But what if there are multiple CPUs?
  - ❖ A thread on a different CPU might make a system call which invokes code that manipulates the ready queue
- ❑ **Using a mutex lock (based on TSL) for critical sections**
  - ❖ Ensures mutual exclusion for all code that follows that convention

# Mutex is not enough

---

- ❑ Interrupt inside interrupt handler
  - ❖

# Some tricky issues ...

---

- ❑ **The interrupt handling code that saves interrupted state is a critical section**
  - ❖ It could be executed concurrently if multiple almost simultaneous interrupts happen
  - ❖ Interrupts must be disabled during this (short) time period to ensure critical state is not lost
- ❑ **What if this interrupt handling code attempts to lock a mutex that is held?**
  - ❖ What happens if we sleep with interrupts disabled?
  - ❖ What happens if we busy wait (spin) with interrupts disabled?

# Implementing mutex locks without TSL

---

- If your CPU did not have TSL, how would you implement blocking mutex lock and unlock calls using interrupt disabling?
  - ❖ ... this is your next Blitz project !