

بسم الله الرحمن الرحيم

# «سیستم عامل»

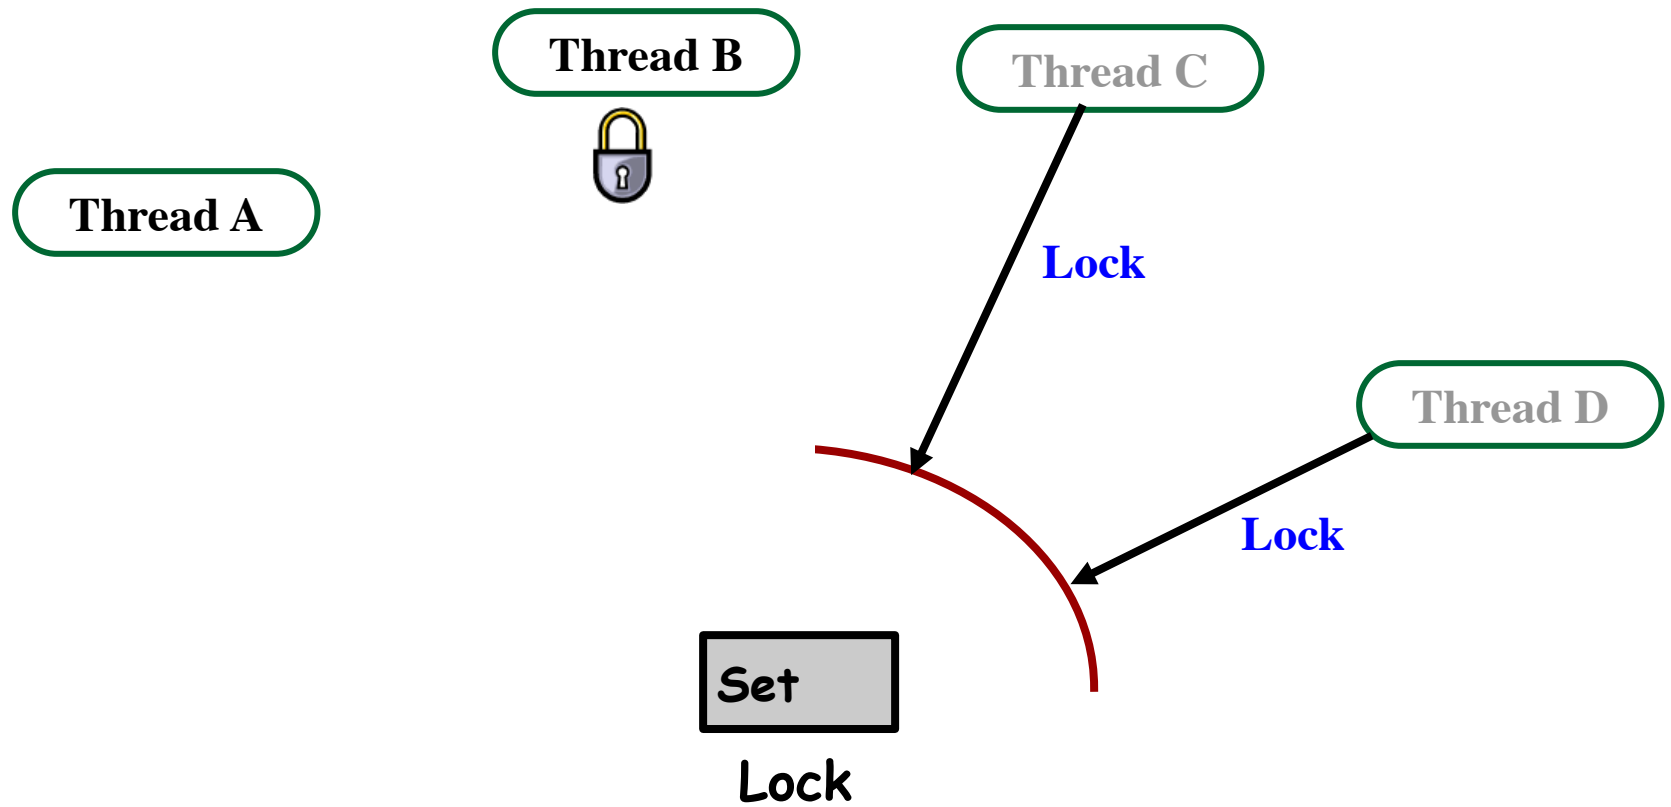
۱

جلسه ۷: کنترل هم‌روندی

# یادآوری

- وضعیت رقابتی:
- حافظه مشترک، اجرای موازی ریسمان‌ها،
- نتایج متفاوت به ازای ترتیب متفاوت اجرا

# Mutex



# How to use a mutex?

---

## Shared data:

Mutex myLock;

```
1 repeat
2   Lock(myLock) ;
3   critical section
4   Unlock(myLock) ;
5   remainder section
6 until FALSE
```

```
1 repeat
2   Lock(myLock) ;
3   critical section
4   Unlock(myLock) ;
5   remainder section
6 until FALSE
```

# ابزار Test-and-set

- اول مقداردهی و برگرداندن مقدار قبلی
- به صورت (atomic)
- `<==` می توانیم mutex را پیاده سازی کنیم.

# کنترل همروندی در هسته

- خاموش کردن interrupt

- صبر تا وقتی mutex آزاد شود

- برای پردازنده‌ها و core‌های دیگر

- شرط: وقتی توسط یک پردازنده mutex گرفته شده، وقفه روی همان پردازنده خاموش است.

- پردازنده‌ای که mutex را دارد، به زودی آزادش می‌کند.

تولیدکننده و  
مصرفکننده

مثالی از  
کنترل  
همروندی:

# The Producer-Consumer Problem

---


- ❑ An example of the **pipelined model**
  - ❖ One thread produces data items
  - ❖ Another thread consumes them
- ❑ Use a bounded buffer between the threads
- ❑ The buffer is a shared resource
  - ❖ Code that manipulates it is a **critical section**
- ❑ Must suspend the producer thread if the buffer is full
- ❑ Must suspend the consumer thread if the buffer is empty



# Is this busy-waiting solution correct?

```
thread producer {
    while(1){
        // Produce char c
        while (count==n) {
            no_op
        }
        buf[InP] = c
        InP = InP + 1 mod n
        count++
    }
}
```

```
thread consumer {
    while(1){
        while (count==0) {
            no_op
        }
        c = buf[OutP]
        OutP = OutP + 1 mod n
        count--
        // Consume char
    }
}
```



Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

# This code is incorrect!

---

- ❑ The “count” variable can be corrupted:
  - ❖ Increments or decrements may be lost!
  - ❖ Possible Consequences:
    - Both threads may spin forever
    - Buffer contents may be over-written
- ❑ What is this problem called?

# This code is incorrect!

---

- ❑ The “count” variable can be corrupted:
  - ❖ Increments or decrements may be lost!
  - ❖ Possible Consequences:
    - Both threads may sleep forever
    - Buffer contents may be over-written
- ❑ What is this problem called? **Race Condition**
- ❑ Code that manipulates count must be made into a **???** and protected using **???**

# This code is incorrect!

---

- ❑ The “count” variable can be corrupted:
  - ❖ Increments or decrements may be lost!
  - ❖ Possible Consequences:
    - Both threads may sleep forever
    - Buffer contents may be over-written
- ❑ What is this problem called? **Race Condition**
- ❑ Code that manipulates count must be made into a **critical section** and protected using **mutual exclusion!**

# Some more problems with this code


---

- ❑ **What if buffer is full?**
  - ❖ Producer will busy-wait
  - ❖ On a single CPU system the consumer will not be able to empty the buffer
- ❑ **What if buffer is empty?**
  - ❖ Consumer will busy-wait
  - ❖ On a single CPU system the producer will not be able to fill the buffer
- ❑ **We need a solution based on blocking!**

# Producer/Consumer with Blocking - 1<sup>st</sup> attempt

```
0  thread producer {
1.  while(1) {
2.      // Produce char c
3.      if (count==n) {
4.          sleep(full)
5.      }
6.      buf[InP] = c;
7.      InP = InP + 1 mod n
8.      count++
9.      if (count == 1)
10.         wakeup(empty)
11.  }
12 }
```

```
0  thread consumer {
1.  while(1) {
2.      if (count==0) {
3.          sleep(empty)
4.      }
5.      c = buf[OutP]
6.      OutP = OutP + 1 mod n
7.      count--;
8.      if (count == n-1)
9.          wakeup(full)
10.     // Consume char
11. }
12 }
```




Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

# Use a mutex to fix the race condition in this code

```
0  thread producer {
1.  while(1) {
2.      // Produce char c
3.      if (count==n) {
4.          sleep(full)
5.      }
6.      buf[InP] = c;
7.      InP = InP + 1 mod n
8.      count++
9.      if (count == 1)
10.         wakeup(empty)
11.  }
12 }
```

```
0  thread consumer {
1.  while(1) {
2.      if (count==0) {
3.          sleep(empty)
4.      }
5.      c = buf[OutP]
6.      OutP = OutP + 1 mod n
7.      count--;
8.      if (count == n-1)
9.          wakeup(full)
10.     // Consume char
11. }
12 }
```



Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

# Problems

---

- ❑ Sleeping while holding the mutex causes deadlock !
- ❑ Releasing the mutex then sleeping opens up a window during which a context switch might occur ... again risking deadlock
- ❑ How can we release the mutex and sleep in a single atomic operation?
- ❑ We need a more powerful synchronization primitive



# Semaphores

---

- An abstract data type that can be used for condition synchronization and mutual exclusion

What is the difference between mutual exclusion and condition synchronization?

# Semaphores

---

- ❑ An abstract data type that can be used for condition synchronization and mutual exclusion
- ❑ **Condition synchronization**
  - ❖ **wait** until condition holds before proceeding
  - ❖ **signal** when condition holds so others may proceed
- ❑ **Mutual exclusion**
  - ❖ only one at a time in a critical section

# Semaphores

---

- **An abstract data type**
  - ❖ containing an integer variable ( $S$ )
  - ❖ Two operations: Wait ( $S$ ) and Signal ( $S$ )
- **Alternative names for the two operations**
  - ❖  $\text{Wait}(S) = \text{Down}(S) = P(S)$
  - ❖  $\text{Signal}(S) = \text{Up}(S) = V(S)$
- **Blitz names its semaphore operations Down and Up**

# Classical Definition of Wait and Signal

---

```
Wait(S)
{
    while S <= 0 do noop;    /* busy wait! */
    S = S - 1;               /* S >= 0 */
}

Signal (S)
{
    S = S + 1;
}
```

# Problems with classical definition

---

- **Waiting threads hold the CPU**
  - ❖ Waste of time in single CPU systems
  - ❖ Required preemption to avoid deadlock

# Blocking implementation of semaphores

---

Semaphore  $S$  has a value,  $S.val$ , and a thread list,  $S.list$ .

## Wait ( $S$ )

```
S.val = S.val - 1
```

```
If S.val < 0
```

```
/* negative value of S.val */
```

```
{ add calling thread to S.list; /* is # waiting threads */  
  block; /* sleep */
```

```
}
```

## Signal ( $S$ )

```
S.val = S.val + 1
```

```
If S.val <= 0
```

```
{ remove a thread T from S.list;  
  wakeup (T);
```

```
}
```

# Implementing semaphores

---

- Wait () and Signal () are assumed to be **atomic**

How can we ensure that they are atomic?

# Implementing semaphores

---

- Wait () and Signal () are assumed to be **atomic**

How can we ensure that they are atomic?

- Implement Wait() and Signal() as system calls?
  - ❖ how can the kernel ensure Wait() and Signal() are completed atomically?
  - ❖ Same solutions as before
    - Disable interrupts, or
    - Use TSL-based mutex



# Semaphores with interrupt disabling

---

```
struct semaphore {  
    int val;  
    list L;  
}
```

```
Wait(semaphore sem)  
    DISABLE_INTS  
    sem.val--  
    if (sem.val < 0) {  
        add thread to sem.L  
        sleep(thread)  
    }  
    ENABLE_INTS
```

```
Signal(semaphore sem)  
    DISABLE_INTS  
    sem.val++  
    if (sem.val <= 0) {  
        th = remove next  
        thread from sem.L  
        wakeup(th)  
    }  
    ENABLE_INTS
```

# Semaphores with interrupt disabling

---

```
struct semaphore {  
    int val;  
    list L;  
}
```

```
Wait(semaphore sem)  
    DISABLE_INTS  
    sem.val--  
    if (sem.val < 0) {  
        add thread to sem.L  
        sleep(thread)  
    }  
    ENABLE_INTS
```

```
Signal(semaphore sem)  
    DISABLE_INTS  
    sem.val++  
    if (sem.val <= 0) {  
        th = remove next  
        thread from sem.L  
        wakeup(th)  
    }  
    ENABLE_INTS
```

# Blitz code for Semaphore.wait

---

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

# Blitz code for Semaphore.wait

---

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

# Blitz code for Semaphore.wait

---

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

# Blitz code for Semaphore.wait

---

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

# But what is `currentThread.Sleep ()` ?

---

- ❑ If `sleep` stops a thread from executing, how, where, and when does it return?
  - ❖ which thread enables interrupts following `sleep`?
  - ❖ the thread that called `sleep` shouldn't return until another thread has called `signal` !
  - ❖ ... but how does that other thread get to run?
  - ❖ ... where exactly does the `thread switch` occur?
- ❑ Trace down through the Blitz code until you find a call to `switch()`
  - ❖ Switch is called in one thread but returns in another!
  - ❖ See where registers are saved and restored

# Look at the following Blitz source code

---

- ❑ **Thread.c**
  - ❖ Thread.Sleep ()
  - ❖ Run (nextThread)
  
- ❑ **Switch.s**
  - ❖ Switch (prevThread, nextThread)



# Blitz code for Semaphore.signal

---

```
method Signal ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during
            'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)

endMethod
```

# Blitz code for Semaphore.signal

---

```
method Signal ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during
            'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)

endMethod
```

# Blitz code for Semaphore.signal

---

```
method Signal ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during
            'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)

endMethod
```

# Blitz code for Semaphore.signal

---

```
method Signal ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during
            'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)

endMethod
```

# Semaphores using atomic instructions

---

- ❑ **Implementing semaphores with interrupt disabling only works on uni-processors**
  - ❖ What should we do on a multiprocessor?
- ❑ **As we saw earlier, hardware provides special atomic instructions for synchronization**
  - ❖ test and set lock (TSL)
  - ❖ compare and swap (CAS)
  - ❖ etc
- ❑ **Semaphore can be built using atomic instructions**
  1. build mutex locks from atomic instructions
  2. build semaphores from mutex locks

# Building spinning mutex locks using TSL

---

Mutex\_lock:

TSL REGISTER,MUTEX	copy mutex to register and set mutex to 1
CMP REGISTER,#0	was mutex zero?
JZE ok	if it was zero, mutex is unlocked, so return
JMP mutex_lock	try again

Ok: RET | return to caller; enter critical section

Mutex\_unlock:

MOVE MUTEX,#0	store a 0 in mutex
RET	return to caller

# Using Mutex Locks to Build Semaphores

---

- How would you modify the Blitz code to do this?

# What if you had a blocking mutex lock?

---

Problem: Implement a counting semaphore

Up ()

Down ()

...using just Mutex locks

- Goal: Make use of the mutex lock's blocking behavior rather than reimplementing it for the semaphore operations



# How about this solution?

---

```
var cnt: int = 0           -- Signal count
var m1: Mutex = unlocked   -- Protects access to "cnt"
    m2: Mutex = locked     -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Lock (m2)
    Unlock (m1)
else
    Unlock (m1)
endIf
```

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

# How about this solution?

```
var cnt: int = 0           -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

Lock (m1)

cnt = cnt - 1

if cnt < 0

Lock (m2)

Unlock (m1)

else

Unlock (m1)

endIf

Up () :

Lock (m1)

cnt = cnt + 1

if cnt <= 0

Unlock (m2)

endIf

Unlock (m1)

Contains a  
Deadlock!

# How about this solution then?

---

```
var cnt: int = 0           -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endIf
```

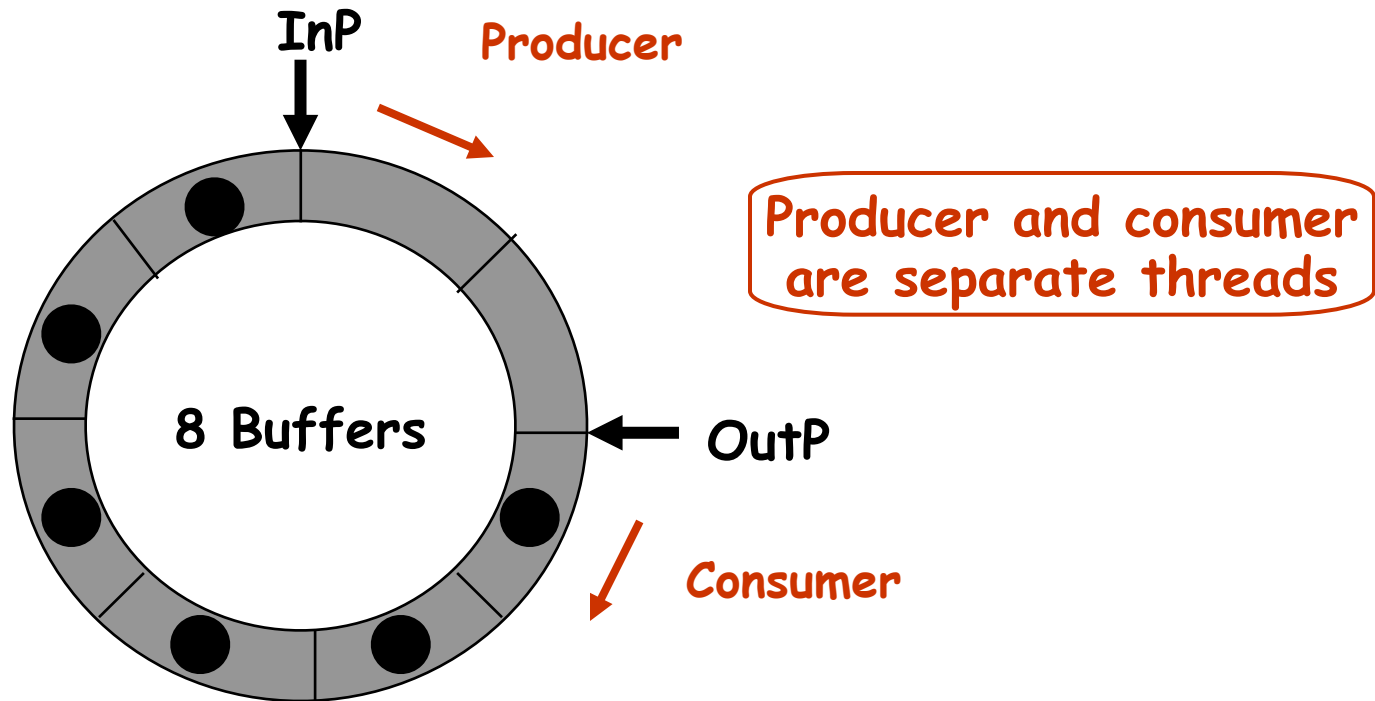
Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

# Producer consumer problem

---


- Also known as the bounded buffer problem



# Is this a valid solution?

```
thread producer {  
    while(1){  
        // Produce char c  
        while (count==n) {  
            no_op  
        }  
        buf[InP] = c  
        InP = InP + 1 mod n  
        count++  
    }  
}
```

```
thread consumer {  
    while(1){  
        while (count==0) {  
            no_op  
        }  
        c = buf[OutP]  
        OutP = OutP + 1 mod n  
        count--  
        // Consume char  
    }  
}
```



Global variables:

```
char buf[n]  
int InP = 0    // place to add  
int OutP = 0   // place to get  
int count
```

# Does this solution work?

---

Global variables

```
semaphore full_buffs = 0;  
semaphore empty_buffs = n;  
char buff[n];  
int InP, OutP;
```

```
0 thread producer {  
1.   while(1){  
2.       // Produce char c...  
3.       down(empty_buffs)  
4.       buf[InP] = c  
5.       InP = InP + 1 mod n  
6.       up(full_buffs)  
7.   }  
8 }
```

```
0 thread consumer {  
1.   while(1){  
2.       down(full_buffs)  
3.       c = buf[OutP]  
4.       OutP = OutP + 1 mod n  
5.       up(empty_buffs)  
6.       // Consume char...  
7.   }  
8 }
```

درست است؟


# Use a mutex to fix the race condition in this code

```
0 thread producer {  
1. while(1) {  
2.     // Produce char c  
3.     if (count==n) {  
4.         sleep(full)  
5.     }  
6.     buf[InP] = c;  
7.     InP = InP + 1 mod n  
8.     count++  
9.     if (count == 1)  
10.        wakeup(empty)  
11. }  
12 }
```

wait(m)  
signal(m)  
wait(m)  
signal(m)

```
0 thread consumer {  
1. while(1) {  
2.     if (count==0) {  
3.         sleep(empty)  
4.     }  
5.     c = buf[OutP]  
6.     OutP = OutP + 1 mod n  
7.     count--;  
8.     if (count == n-1)  
9.         wakeup(full)  
10.    // Consume char  
11. }  
12 }
```

wait(m)  
signal(m)  
wait(m)  
signal(m)



Global variables:

```
char buf[n]  
int InP = 0    // place to add  
int OutP = 0   // place to get  
int count
```

# How about this solution then?

---

```
var cnt: int = n          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endIf
```

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```



# How about this solution then?

---

```
var cnt: int = n          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endIf
```

.. ابتدا n = ..

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

# How about this solution then?

```
var cnt: int = n          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endIf
```

.. ابتدا n = ..

۱- چند  
ریمان اینجا

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

# How about this solution then?

```
var cnt: int = n          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

.. ابتدا n = ..

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endIf
```

۱- چند  
ریسمان اینجا

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

۲- چندین بار  
این صدا شود

# How about this solution then?

```
var cnt: int = n          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock (m1)
cnt = cnt - 1
if cnt < 0
    Unlock (m1)
    Lock (m2)
else
    Unlock (m1)
endIf
```

.. ابتدا n = ۰

۱- چند  
ریسمان اینجا

Up () :

```
Lock (m1)
cnt = cnt + 1
if cnt <= 0
    Unlock (m2)
endIf
Unlock (m1)
```

۲- چندین بار  
این صدا شود

چون هنوز  
Lock(m2) صدا  
نشده، نتیجه چندین  
بار باز کردن قفل باز  
نامعلوم است

# پیاده‌سازی Semaphore با Mutex (روش ۱)

---

**var**

```
mutex=1: binary-semaphore;  
delay=0: binary-semaphore;  
C={initvalue}: integer;
```

**Procedure** Wait()

```
begin  
wait(mutex);  
C:=C-1;  
if C < 0 then begin  
    signal(mutex);  
    wait(delay);  
end  
signal(mutex);  
end
```

**Procedure** Signal()

```
begin  
wait(mutex);  
C:=C+1;  
if C <= 0 then  
    signal(delay)  
else  
    signal(mutex)  
end
```

## پیاده‌سازی Semaphore با Mutex (روش ۲)

---

**var**

```
mutex=1: binary-semaphore;  
delay=0: binary-semaphore;  
barrier=1: binary-semaphore;  
C={initvalue}: integer;
```

**Procedure** Wait()

**begin**

wait(barrier);

wait(mutex);

C:=C-1;

**if** C < 0 **then begin**

    signal(mutex);

    wait(delay);

**end**

**else**

    signal(mutex);

signal(barrier);

**end**

**Procedure** Signal()

**begin**

wait(mutex);

C:=C+1;

**if** C = 1 **then**

    signal(delay)

signal(mutex)

**end**

## پیاده‌سازی Semaphore با Mutex (روش ۳)

---

**var**

```
mutex=1: binary-semaphore;  
delay={min(1,initvalue)}: binary-semaphore;  
C={initvalue}: integer;
```

**Procedure** Wait()

```
begin  
wait(delay);  
wait(mutex);  
C:=C-1;  
if C > 0 then  
    signal(delay);  
signal(mutex);  
end
```

**Procedure** Signal()

```
begin  
wait(mutex);  
C:=C+1;  
if C = 1 then  
    signal(delay)  
signal(mutex)  
end
```