

بسم الله الرحمن الرحيم

«سیستم عامل»

۱۹۱

جلسه ۲۰: مدیریت حافظه (۸)

The optimal page replacement algorithm

- Idea:
 - ❖ Given all the data, how to find the optimal page replacement?
 - ❖ **Longest Forward Distance (LFD):** Select the page that will not be needed for the longest time

$$\text{LFD} = \text{OPT}$$

The optimal page replacement algorithm

- ❑ Idea:
 - ❖ Select the page that will not be needed for the longest time
- ❑ Problem?

The optimal page replacement algorithm

- ❑ Idea:
 - ❖ Select the page that will not be needed for the longest time
- ❑ Problem:
 - ❖ Can't know the future of a program
 - ❖ Can't know when a given page will be needed next
 - ❖ The optimal algorithm is unrealizable

The optimal page replacement algorithm

- However:
 - ❖ We can use it as a control case for simulation studies
 - Run the program once
 - Generate a log of all memory references
 - Do we need all of them?
 - Use the log to simulate various page replacement algorithms
 - Can compare others to “optimal” algorithm

FIFO page replacement algorithm

- Always replace the oldest page ...
 - ❖ "Replace the page that has been in memory for the longest time."

FIFO page replacement algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a

Page	0	a									
Frames	1	b		a	a		a				
	2	c					b				
	3	d	c	c	c	c					
				d	d						

Page faults

X

FIFO page replacement algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a

Page	0	a									
Frames	1	b	a	a	a	a	a	a	a		
	2	c			b	b	b	b	b		
	3	d	c	c	c	e	e	e	e		
				d	d	d	d	d	d		

Page faults

X

X

FIFO page replacement algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time	0	1	2	3	4	5	6	7	8	9	10
Requests		c	a	d	b	e	b	a	b	c	a

Page	0	a									
Frames	1	b	a	a	a	a	a	a	a	c	
	2	c			b	b	b	b	b	b	
	3	d	c	c	c	e	e	e	e	e	
				d	d	d	d	d	d	d	

Page faults

X

X

X

FIFO page replacement algorithm

- Replace the page that was first brought into memory
- Example: Memory system with 4 frames:

Time		0	1	2	3	4	5	6	7	8	9	10
Requests			c	a	d	b	e	b	a	b	c	a
<hr/>												
Page	0	a										
Frames	1	b		a	a	a	a	a	a	a	c	c
	2	c				b	b	b	b	b	b	b
	3	d	c	c	c	c	e	e	e	e	e	e
					d	d	d	d	d	d	d	a
Page faults							X				X	X

FIFO page replacement algorithm

- ❑ Always replace the oldest page.
 - ❖ “Replace the page that has been in memory for the longest time.”
- ❑ Implementation
 - ❖ Maintain a linked list of all pages in memory
 - ❖ Keep it in order of when they came into memory
 - ❖ The page at the tail of the list is oldest
 - ❖ Add new page to head of list

FIFO page replacement algorithm

- ❑ Disadvantage?

FIFO page replacement algorithm

- ❑ Disadvantage:
 - ❖ The oldest page may be needed again soon
 - ❖ Some page may be important throughout execution
 - ❖ It will get old, but replacing it will cause an immediate page fault

How can we do better?

- ❑ Need an approximation of how likely each frame is to be accessed in the future
 - ❖ If we base this on past behavior we need a way to track past behavior
 - ❖ Tracking memory accesses requires hardware support to be efficient

Page table: referenced and dirty bits

- Each page table entry (and TLB entry!) has a
 - ❖ Referenced bit - set by TLB when page read / written
 - ❖ Dirty / modified bit - set when page is written
 - ❖ If TLB entry for this page is valid, it has the most up to date version of these bits for the page
 - OS must copy them into the page table entry during fault handling
- Idea: use the information contained in these bits to drive the page replacement algorithm

Page table: referenced and dirty bits

- ❑ Some hardware does not have support for the dirty bit
- ❑ Instead, memory protection can be used to emulate it
- ❑ Idea:
 - ❖ Software sets the protection bits for all pages to “read only”
 - ❖ When program tries to update the page...
 - A trap occurs
 - Software sets the Dirty Bit in the page table and clears the ReadOnly bit
 - Resumes execution of the program

Not recently used page replacement alg.

- ▣ Uses the Referenced Bit and the Dirty Bit
- ▣ Initially, all pages have
 - ❖ Referenced Bit = 0
 - ❖ Dirty Bit = 0
- ▣ Periodically... (e.g. whenever a timer interrupt occurs)
 - ❖ Clear the Referenced Bit
 - ❖ Referenced bit now indicates "recent" access

Not recently used page replacement alg.

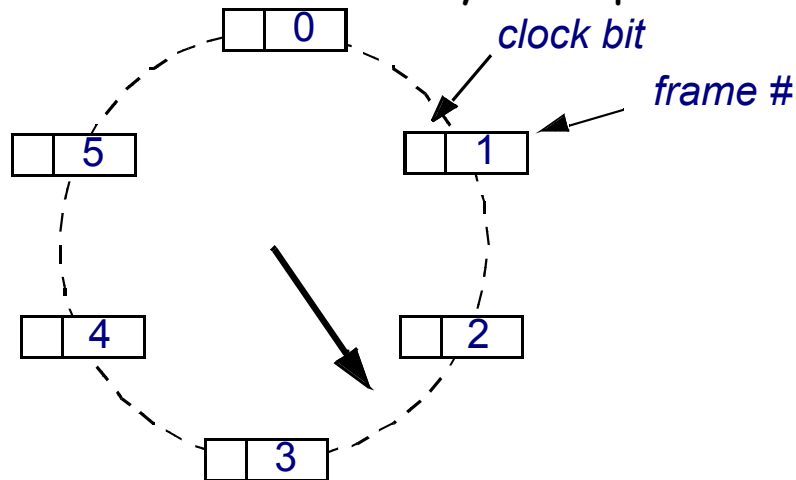
- ❑ When a page fault occurs...
- ❑ Categorize each page...
 - ❖ Class 1: Referenced = 0 Dirty = 0
 - ❖ Class 2: Referenced = 0 Dirty = 1
 - ❖ Class 3: Referenced = 1 Dirty = 0
 - ❖ Class 4: Referenced = 1 Dirty = 1
- ❑ Choose a victim page from class 1 ... why?
- ❑ If none, choose a page from class 2 ... why?
- ❑ If none, choose a page from class 3 ... why?
- ❑ If none, choose a page from class 4 ... why?

Second chance page replacement alg.

- ❑ An implementation of NRU based on FIFO
- ❑ Pages kept in a linked list
 - ❖ Oldest is at the front of the list
- ❑ Look at the oldest page
 - ❖ If its "referenced bit" is 0...
 - Select it for replacement
 - ❖ Else
 - It was used recently; don't want to replace it
 - Clear its "referenced bit"
 - Move it to the end of the list
 - ❖ Repeat
- ❑ What if every page was used in last clock tick?
 - ❖ Select a page at random
- ❑

Clock algorithm (an implementation of NRU)

- ❑ Maintain a circular list of pages in memory
- ❑ Set a bit for the page when a page is referenced
- ❑ Clock sweeps over memory looking for a victim page that does not have the referenced bit set
 - ❖ If the bit is set, clear it and move on to the next page
 - ❖ Replaces pages that haven't been referenced for one complete clock revolution - essentially an implementation of NRU



Least recently used algorithm (LRU)

- A refinement of NRU that orders how recently a page was used
 - ❖ Keep track of when a page is used
 - ❖ Replace the page that has been used least recently

LRU page replacement

- Replace the page that hasn't been referenced in the longest time

Time		0	1	2	3	4	5	6	7	8	9	10
Requests			c	a	d	b	e	b	a	b	c	d
<hr/>												
Page	0	a	a	a	a	a	a	a	a	a	a	a
Frames	1	b	b	b	b	b	b	b	b	b	b	b
	2	c	c	c	c	e	e	e	e	e	e	d
	3	d	d	d	d	d	d	d	d	d	c	c
Page faults							X				X	X

Least recently used algorithm (LRU)

- But how can we implement this?

Least recently used algorithm (LRU)

- ❑ But how can we implement this?
- ❑ Implementation #1:
 - ❖ Keep a linked list of all pages
 - ❖ On every memory reference,
 - **Move that page to the front of the list**
 - ❖ The page at the tail of the list is replaced

Least recently used algorithm (LRU)

- ❑ But how can we implement this?
 - ❖ ... without requiring “every access” to be recorded?
- ❑ Implementation #2:
 - ❖ MMU (hardware) maintains a counter
 - ❖ Incremented on every clock cycle
 - ❖ Every time a page table entry is used
 - MMU writes the value to the page table entry
 - This timestamp value is the time-of-last-use
 - ❖ When a page fault occurs
 - Software looks through the page table
 - Identifies the entry with the oldest timestamp

Least recently used algorithm (LRU)

- ❑ What if we don't have hardware support for a counter?
- ❑ Implementation #3:
 - ❖ Maintain a counter in software
 - ❖ On every timer interrupt...
 - Increment counter
 - Run through the page table
 - For every entry that has "ReferencedBit" = 1
 - Update its timestamp
 - Clear the ReferencedBit
 - ❖ Approximates LRU
 - ❖ If several have oldest time, choose one arbitrarily