

بسم الله الرحمن الرحيم

«سیستم عامل»

۲۲۲

جلسه ۲۲: مدیریت حافظه (۹)

Working set page replacement

- ❑ Demand paging
 - ❖ Pages are only loaded when accessed
 - ❖ When process begins, all pages marked INVALID

Working set page replacement

- ❑ **Demand paging**
 - ❖ Pages are only loaded when accessed
 - ❖ When process begins, all pages marked INVALID
- ❑ **Locality of reference**
 - ❖ Processes tend to use only a small fraction of their pages

Working set page replacement

- ❑ **Demand paging**
 - ❖ Pages are only loaded when accessed
 - ❖ When process begins, all pages marked INVALID
- ❑ **Locality of Reference**
 - ❖ Processes tend to use only a small fraction of their pages
- ❑ **Working Set**
 - ❖ The set of pages a process needs
 - ❖ If working set is in memory, no page faults
 - ❖ What if you can't get working set into memory?

Working set page replacement


- ❑ **Thrashing**
 - ❖ If you can't get working set into memory page faults occur every few instructions
 - ❖ Little work gets done
 - ❖ Most of the CPU's time is going on overhead

Working set page replacement

- ❑ **Based on prepaging (prefetching)**
 - ❖ Load pages before they are needed
- ❑ **Main idea:**
 - ❖ Try to identify the process's "working set"
- ❑ **How big is the working set?**
 - ❖ Look at the last K memory references
 - ❖ As K gets bigger, more pages needed.
 - ❖ In the limit, all pages are needed.

Working set page replacement

- The size of the working set for last k time intervals at time t :

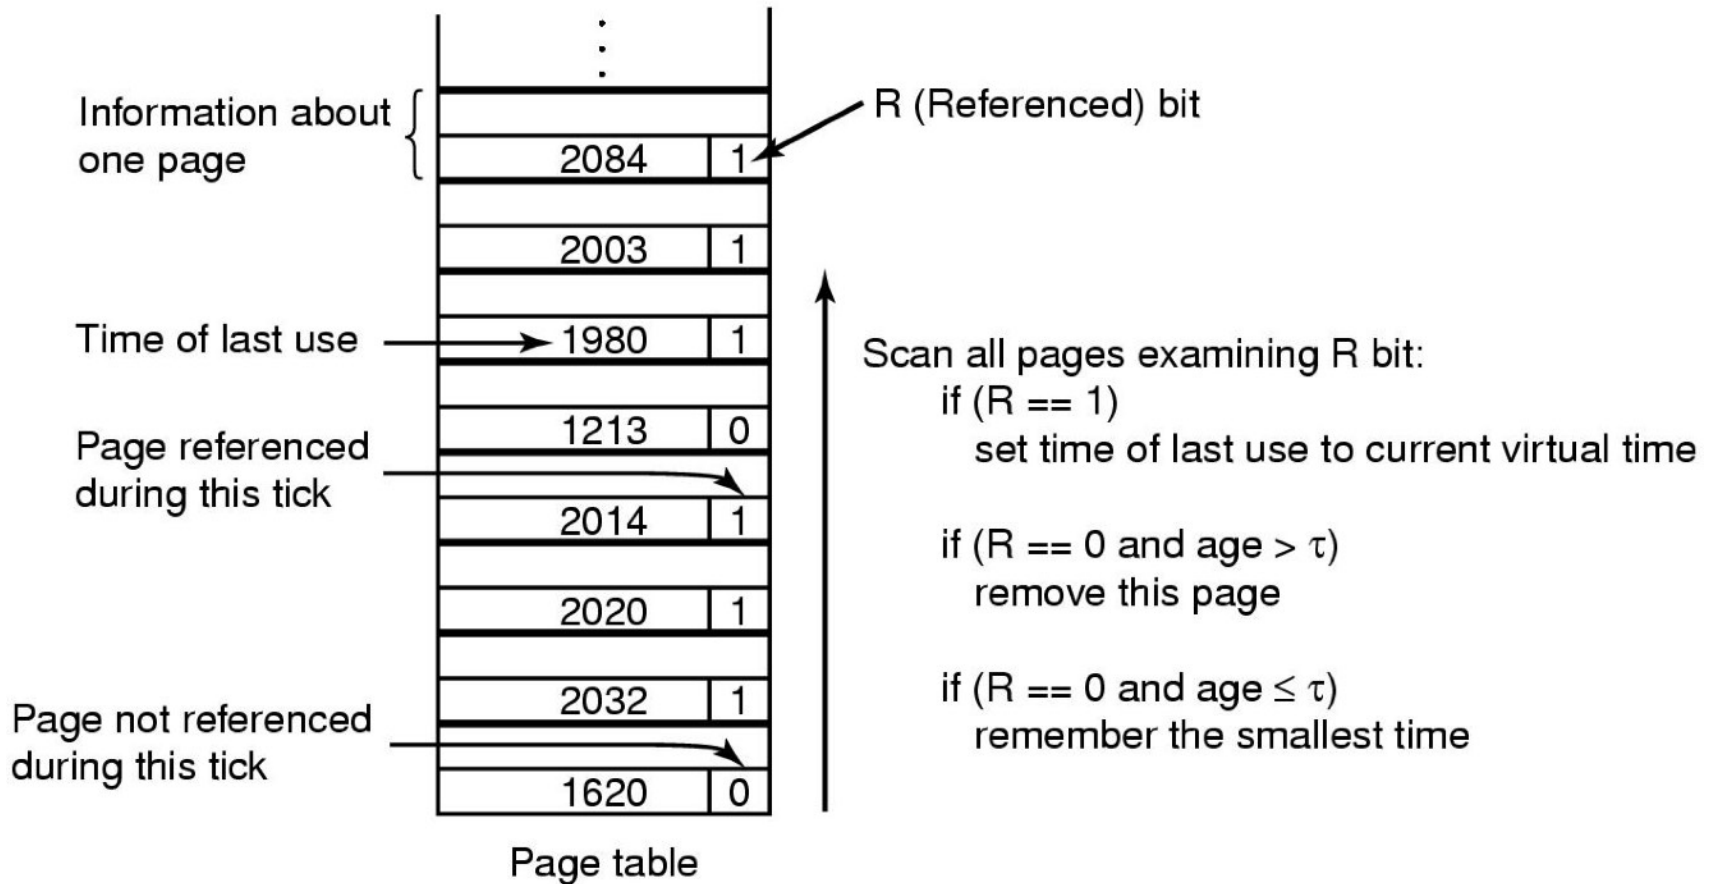


Working set page replacement

- ❑ Idea:
 - ❖ Look back over the last T msec of time
 - ❖ Which pages were referenced?
 - This is the working set.
- ❑ **Current Virtual Time**
 - ❖ Only consider how much CPU time this process has seen.
- ❑ **Implementation**
 - ❖ On each clock tick, look at each page
 - ❖ Was it referenced?
 - **Yes: Make a note of Current Virtual Time**
 - ❖ If a page has not been used in the last T msec,
 - **It is not in the working set!**
 - **Evict it; write it out if it is dirty.**

Working set page replacement

2204 Current virtual time



WSClock page replacement algorithm

- ❑ An implementation of the working set algorithm
- ❑ All pages are kept in a circular list (ring)
- ❑ As pages are added, they go into the ring
- ❑ The “clock hand” advances around the ring
- ❑ Each entry contains “time of last use”
- ❑ Upon a page fault...
 - ❖ If Reference Bit = 1...
 - Page is in use now. Do not evict.
 - Clear the Referenced Bit.
 - Update the “time of last use” field.


WSClock page replacement algorithm

- **If Reference Bit = 0**
 - ❖ If the age of the page is less than T...
 - This page is in the working set.
 - Advance the hand and keep looking
 - ❖ If the age of the page is greater than T...
 - If page is clean
 - Reclaim the frame and we are done!
 - If page is dirty
 - Schedule a write for the page
 - Advance the hand and keep looking

Proactive use of replacement algorithm

- ❑ Replacing victim frame on each page fault typically requires two disk accesses per page fault
- ❑ Alternative → the O.S. can keep several pages free in anticipation of upcoming page faults.

In Unix: low and high water marks



low water mark

high water mark

$\text{low} < \# \text{ free pages} < \text{high}$

Free pages and the clock algorithm

- ❑ The rate at which the clock sweeps through memory determines the number of pages that are kept free:
 - ❖ Too high a rate --> Too many free pages marked
 - ❖ Too low a rate --> Not enough (or no) free pages marked
- ❑ Large memory system considerations
 - ❖ As memory systems grow, it takes longer and longer for the hand to sweep through memory
 - ❖ This washes out the effect of the clock somewhat

The UNIX memory model

- ❑ UNIX page replacement
 - ❖ clock algorithm for page replacement
 - If page has not been accessed move it to the free list for use as allocatable page
 - If modified/dirty → write to disk (still keep stuff in memory though)
 - If unmodified → just move to free list
 - ❖ High and low water marks for free pages
 - ❖ Pages on the free-list can be re-allocated if they are accessed again before being overwritten