

بسم الله الرحمن الرحيم

«سیستم عامل»

۲۳۶

جلسه ۲۳: مدیریت حافظه (۱۰)

But which algorithm is best?

Comparing algorithms through modeling

- **Run a program**
 - ❖ Look at all memory references
 - ❖ Don't need all this data
 - ❖ Look at which pages are accessed
0000001222333300114444001123444
 - ❖ Eliminate duplicates
012301401234
- **This defines the Reference String**
 - ❖ Use this to evaluate different page replacement algorithms
 - ❖ Count page faults given the same reference string

Summary

Algorithm	Comment
Optimal	Not implementable, but useful as a benchmark
NRU (Not Recently Used)	Very crude
FIFO (First-In, First-Out)	Might throw out important pages
Second chance	Big improvement over FIFO
Clock	Realistic
LRU (Least Recently Used)	Excellent, but difficult to implement exactly
NFU (Not Frequently Used)	Fairly crude approximation to LRU
Aging	Efficient algorithm that approximates LRU well
Working set	Somewhat expensive to implement
WSClock	Good efficient algorithm

Local vs. global page replacement

- ❑ Assume several processes: A, B, C, \dots
- ❑ Some process gets a page fault (say, process A)
- ❑ Choose a page to replace.
- ❑ **Local page replacement**
 - ❖ Only choose one of A 's pages
- ❑ **Global page replacement**
 - ❖ Choose any page

Local vs. global page replacement

- Example: Process has a page fault...

	Age
A0	10
A1	7
A2	5
A3	4
A4	6
A5	3
B0	9
B1	4
B2	6
B3	2
B4	5
B5	6
B6	12
C1	3
C2	5
C3	6

Original

A0
A1
A2
A3
A4
A6
B0
B1
B2
B3
B4
B5
B6
C1
C2
C3

Local

A0
A1
A2
A3
A4
A5
B0
B1
B2
A6
B4
B5
B6
C1
C2
C3


Global

Local vs. global page replacement

- ❑ **Assume we have**
 - ❖ 5,000 frames in memory
 - ❖ 10 processes
- ❑ **Idea: Give each process 500 frames**
- ❑ **Fairness?**
 - ❖ Small processes: do not need all those pages
 - ❖ Large processes: may benefit from even more frames
- ❑ **Idea:**
 - ❖ Look at the size of each process (... but how?)
 - ❖ Give them a pro-rated number of frames
 - ❖ With a minimum of (say) 10 frames per process


Page fault frequency

- “If you give a process more pages,
❖ its page fault frequency will decline.”



Page fault frequency

- “If you give a process more pages,
❖ its page fault frequency will decline.”



Page fault frequency

- ❑ Measure the page fault frequency of each process.
- ❑ Count the number of faults every second.
- ❑ May want to consider the past few seconds as well.

Page fault frequency

- ❑ Measure the page fault frequency of each process.
- ❑ Count the number of faults every second.
- ❑ May want to consider the past few seconds as well.
- ❑ Aging:
 - ❖ Keep a running value.
 - ❖ Every second
 - Count number of page faults
 - Divide running value by 2
 - Add in the count for this second

Load control

- Assume:
 - ❖ The best page replacement algorithm
 - ❖ Optimal global allocation of page frames

Load control

- ❑ Assume:
 - ❖ The best page replacement algorithm
 - ❖ Optimal global allocation of page frames
- ❑ Thrashing is still possible!

Load control

- ❑ Assume:
 - ❖ The best page replacement algorithm
 - ❖ Optimal global allocation of page frames
- ❑ Thrashing is still possible!
 - ❖ Too many page faults!
 - ❖ No useful work is getting done!
 - ❖ Demand for frames is too great!

Load Control

- ❑ Assume:
 - ❖ The best page replacement algorithm
 - ❖ Optimal global allocation of page frames
- ❑ Thrashing is still possible!
 - ❖ Too many page faults!
 - ❖ No useful work is getting done!
 - ❖ Demand for frames is too great!
- ❑ Solution:
 - ❖ Get rid of some processes (temporarily).
 - ❖ Swap them out.
 - ❖ "Two-level scheduling"

Spare slides

Belady's anomaly

- If you have more page frames (i.e., more memory)...
 - ❖ You will have fewer page faults, right???

Belady's anomaly

- ❑ If you have more page frames (i.e., more memory)...
 - ❖ You will have fewer page faults, right???
- ❑ Not always!

Belady's anomaly

- ❑ If you have more page frames (i.e., more memory)...
 - ❖ You will have fewer page faults, right???
- ❑ Not always!
- ❑ Consider FIFO page replacement
- ❑ Look at this reference string:
012301401234

Belady's anomaly

- ❑ If you have more page frames (i.e., more memory)...
 - ❖ You will have fewer page faults, right???
- ❑ Not always!
- ❑ Consider FIFO page replacement
- ❑ Look at this reference string
012301401234
- ❑ **Case 1:**
 - ❖ 3 frames available --> 9 page faults
- ❑ **Case 2:**
 - ❖ 4 frames available --> 10 page faults

Belady's anomaly

All pages frames initially empty

	0	1	2	3	0	1	4	0	1	2	3	4
Youngest page	0	1	2	3	0	1	4	4	4	2	3	3
		0	1	2	3	0	1	1	1	4	2	2
Oldest page			0	1	2	3	0	0	0	1	4	4
	P	P	P	P	P	P	P			P	P	

9 Page faults

FIFO with 3 page frames

Belady's anomaly

All pages frames initially empty

		0	1	2	3	0	1	4	0	1	2	3	4
Youngest page		0	1	2	3	0	1	4	4	4	2	3	3
			0	1	2	3	0	1	1	1	4	2	2
Oldest page				0	1	2	3	0	0	0	1	4	4
		P	P	P	P	P	P				P	P	

9 Page faults

FIFO with 3 page frames

		0	1	2	3	0	1	4	0	1	2	3	4
Youngest page		0	1	2	3	3	3	4	0	1	2	3	4
			0	1	2	2	2	3	4	0	1	2	3
Oldest page				0	1	1	1	2	3	4	0	1	2
					0	0	0	1	2	3	4	0	1
		P	P	P	P			P	P	P	P	P	P

10 Page faults

FIFO with 4 page frames

Belady's anomaly

LRU ???

Which page size is best?

- Smaller page sizes...

- ❖ Advantages

- Less internal fragmentation
 - On average: half of the last page is wasted
 - Working set takes less memory
 - Less unused program in memory

- ❖ Disadvantages

- Page tables are larger
 - Disk-seek time dominates transfer time (It takes the same time to read a large page as a small page)

Which page size is best?

Let

s = size of average process

e = bytes required for each page table entry

p = size of page, in bytes

s/p = Number of pages per process

es/p = Size of page table

$p/2$ = space wasted due to internal fragmentation

overhead = $se/p + p/2$

Which page size is best?

- $\text{Overhead} = se/p + p/2$
- Want to choose p to minimize overhead.
- Take derivative w.r.t. p and set to zero
$$-se/p^2 + 1/2 = 0$$
- Solving for p ...
$$p = \sqrt{2se}$$

Which page size is best?

s = size of average process

e = bytes required for each page table entry

- ▣ Solving for p ...

$$p = \sqrt{2se}$$

Example:

$$p = \sqrt{2 * 1\text{MB} * 8} = 4\text{K}$$

Which page size is best?

s = size of average process

e = bytes required for each page table entry

- ▣ Solving for p ...

$$p = \sqrt{2se}$$

Example:

$$p = \sqrt{2 * 8\text{MB} * 4} = 8\text{K}$$