

**CS 333**  
**Introduction to Operating Systems**  
**Class 1 - Introduction to OS-related**  
**Hardware and Software**

**Jonathan Walpole**  
**Computer Science**  
**Portland State University**

# About the instructor

---

- **Instructor - Jonathan Walpole**
  - ❖ Professor at PSU since 2004, OGI 1989 - 2004
  - ❖ Research Interests: *Operating System Design, Parallel and Distributed Computing Systems*
  - ❖ <http://www.cs.pdx.edu/~walpole>

# About CS 333

---

## □ **Goals of the class**

- ❖ understand the basic concepts of operating systems
  - *designing & building* operating systems, not *using* them!
- ❖ gain some practical experience so its not just words!

## □ **Expectations**

- ❖ reading assignments should be read *before* class
- ❖ active participation in class discussions
- ❖ no cheating!

# Grading

---

- **Exams**
  - ❖ Mid-term - 25%
  - ❖ Final - 25%
  
- **Coursework**
  - ❖ Project - 50%

# Text books

---

"Operating Systems Concepts: Essentials" by Silberschatz,  
Galvin, Gagne

"The BLITZ System" by Harry Porter

# The project

---

- ❑ You will read, understand and write real operating system code!
- ❑ We will be using the BLITZ system, written by Harry Porter
- ❑ About BLITZ
  - ❖ CPU emulator, assembler, high-level language, operating system, and debugging environment
  - ❖ Simple enough to understand *in detail* how everything works!
  - ❖ Realistic enough to understand *in detail* how everything works!
  - ❖ Runs on the departmental Sun machines (cs.pdx.edu), plus Macs and x86/Linux

# Administrative

---

- **Class web site**
  - ❖ [www.cs.pdx.edu/~walpole/class/cs333/spring2012/home.html](http://www.cs.pdx.edu/~walpole/class/cs333/spring2012/home.html)
  - ❖ Find my website from the faculty listing on the department website. Follow *teaching* link to current classes
- **Class mailing list**
  - ❖ <https://mailhost.cecs.pdx.edu/cgi-bin/mailman/listinfo/cs333>
- **Project 0**
  - ❖ read the class web site
  - ❖ join the class mailing list
- **Project 1**
  - ❖ due next week!
  - ❖ see class web site for project assignments

# Class 1 - Introduction to OS-related Hardware and Software



# Overview

---

- ❑ **What is an Operating System?**
- ❑ **A review of OS-related hardware**

# What is an operating system?

---

- Operating system -- “a *program* that controls the execution of application programs and implements an interface between the user of a computer and the computer hardware”
  - ❖ Narrow view of a computer and OS
    - Traditional computer with applications running on it (e.g. PCs, Workstations, Servers)
  - ❖ Broad view of a computer and OS
    - Anything that needs to manage resources (e.g. router OS, embedded system, cell phone OS ...)

# Two key OS functions

---

## □ Abstract Machine

- ❖ Hides complex details of the underlying hardware
- ❖ Provides common API to applications and services
- ❖ Simplifies application writing

## □ Resource Manager

- ❖ Controls accesses to *shared* resources
  - CPU, memory, disks, network, ...
- ❖ Allows for global policies to be implemented

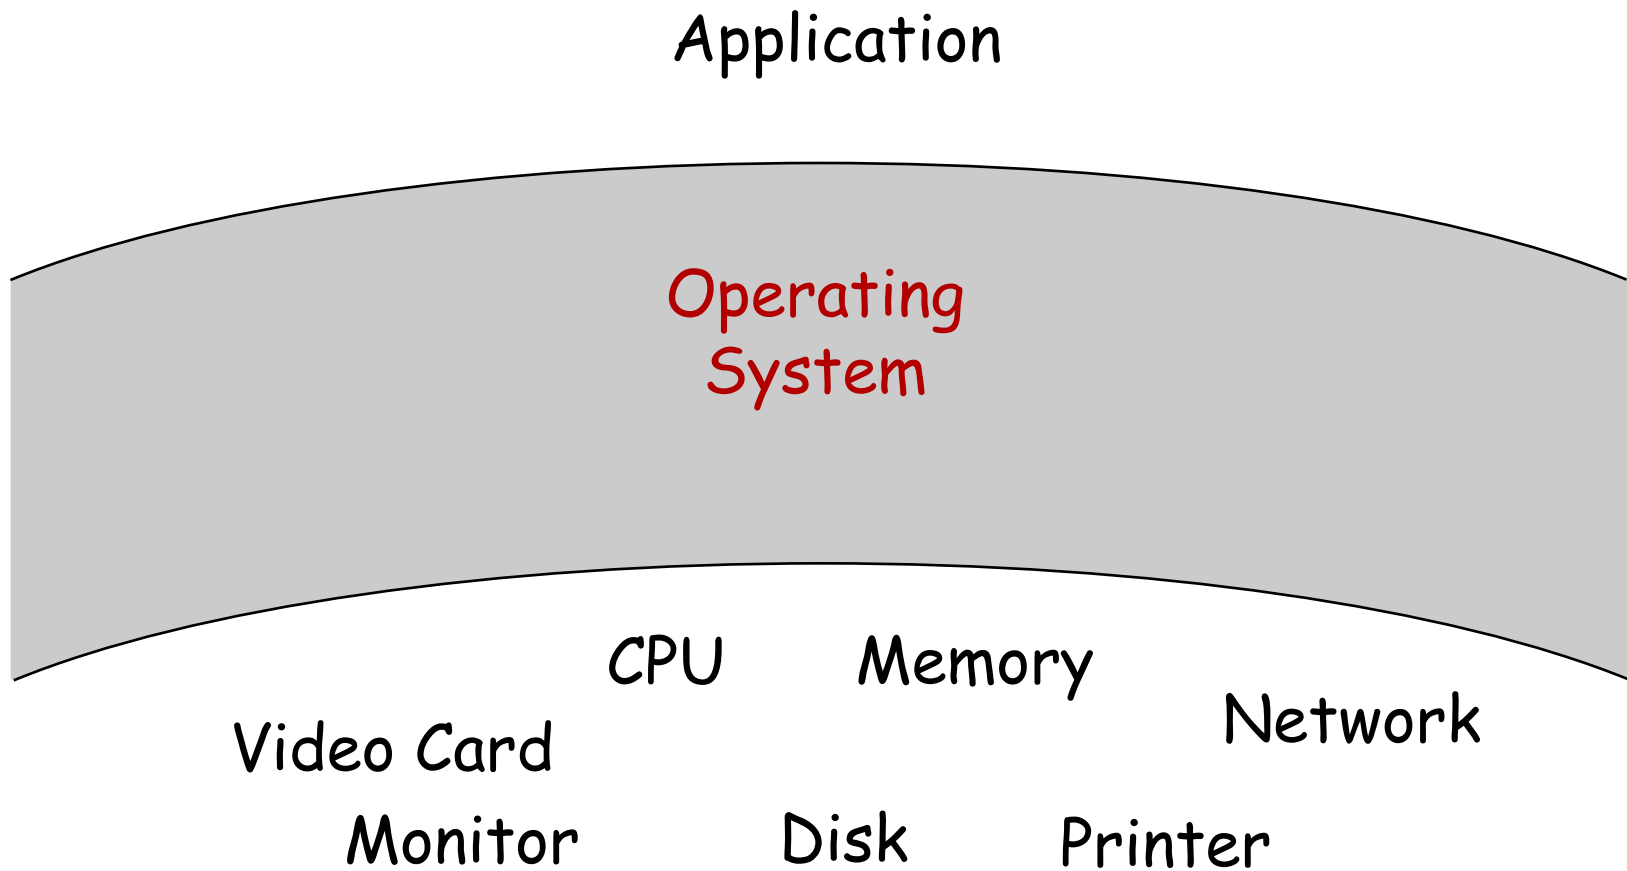
# Why is abstraction important?

---

- **Without OSs and abstract interfaces, application writers must program all device access directly**
  - ❖ load device command codes into device registers
  - ❖ handle initialization, recalibration, sensing, timing etc for physical devices
  - ❖ understand physical characteristics and layout
  - ❖ control motors
  - ❖ interpret return codes ... etc
- **Applications suffer severe code bloat!**
  - ❖ very complicated maintenance and upgrading
  - ❖ no portability
  - ❖ writing this code once, and sharing it, is how OS began!

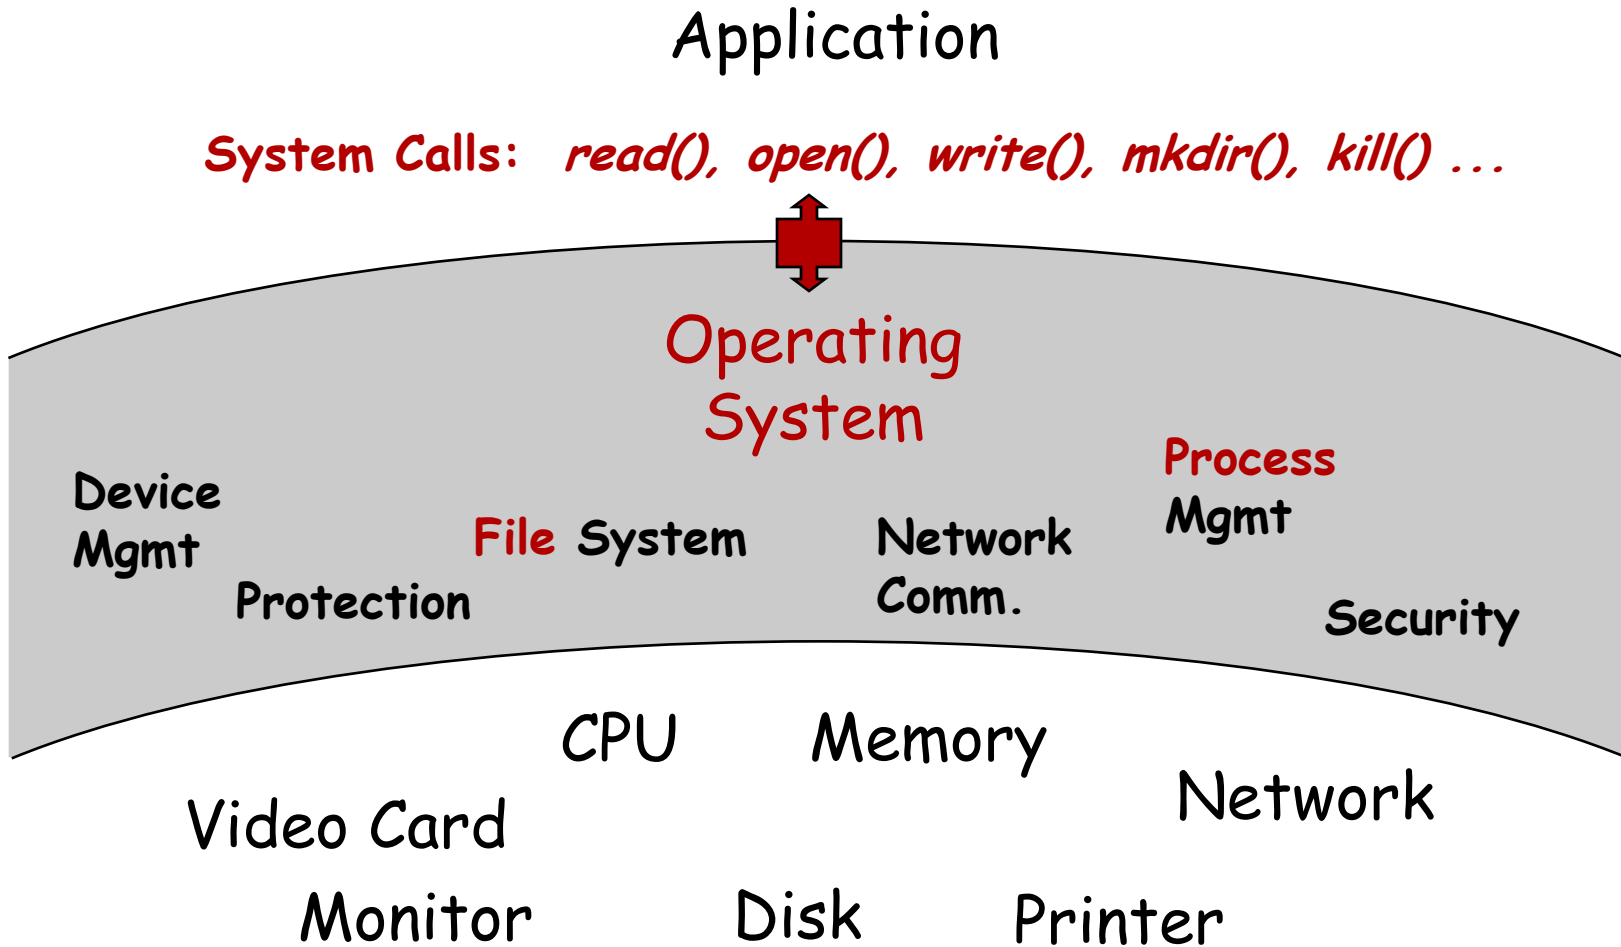
# Providing abstraction via system calls

---



# Providing abstraction via system calls

---



# OS as a resource manager

---

- ❑ **Allocating resources to applications across space and time**
  - ❖ time sharing a resource (scheduling)
  - ❖ space sharing a resource (allocation)
- ❑ **Making efficient use of limited resources**
  - ❖ improving utilization
  - ❖ minimizing overhead
  - ❖ improving throughput/good put
- ❑ **Protecting applications from each other**
  - ❖ enforcement of boundaries

# Problems an OS must solve

---

- ❑ Time sharing the CPU among applications
- ❑ Space sharing the memory among applications
- ❑ Space sharing the disk among users
- ❑ Time sharing access to the disk
- ❑ Time sharing access to the network



# More problems an OS must solve

---

- ❑ **Protection**
  - ❖ of applications from each other
  - ❖ of user data from other users
  - ❖ of hardware/devices
  - ❖ of the OS itself!
  
- ❑ **The OS is just a program! It needs help from the hardware to accomplish these tasks!**
  - ❖ When an application is running, the OS is not running!
  - ❖ When the OS is not running, it can't do anything!

# Overview

---

- What is an Operating System?
- A review of OS-related hardware

# Instruction sets

---

- A CPU's instruction set defines what it can do
  - ❖ different for different CPU architectures
  - ❖ all have **load** and **store** instructions for moving items between memory and registers
    - *Load* a word located at an address in memory into a register
    - *Store* the contents of a register to a word located at an address in memory
  - ❖ many instructions for comparing and combining values in registers and putting result into a register
- Look at the Blitz instruction set which is similar to a SUN SPARC instruction set

# Basic anatomy on a CPU

---

- Program Counter (PC)

# Basic anatomy on a CPU

---

- **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction

# Basic anatomy on a CPU

---

- ❑ **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction
- ❑ **Instruction Register (IR)**

# Basic anatomy on a CPU

---

- ❑ **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction
- ❑ **Instruction Register (IR)**
  - ❖ Holds the instruction currently being executed

# Basic anatomy on a CPU

---

- ❑ **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction
- ❑ **Instruction Register (IR)**
  - ❖ holds the instruction currently being executed
- ❑ **General Registers (Reg. 1..n)**



# Basic anatomy on a CPU

---

- ❑ **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction
- ❑ **Instruction Register (IR)**
  - ❖ holds the instruction currently being executed
- ❑ **General Registers (Reg. 1..n)**
  - ❖ hold variables and temporary results

# Basic anatomy on a CPU

---

- ❑ **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction
- ❑ **Instruction Register (IR)**
  - ❖ holds the instruction currently being executed
- ❑ **General Registers (Reg. 1..n)**
  - ❖ hold variables and temporary results
- ❑ **Arithmetic and Logic Unit (ALU)**

# Basic anatomy on a CPU

---

- ❑ **Program Counter (PC)**
  - ❖ Holds the memory address of the next instruction
- ❑ **Instruction Register (IR)**
  - ❖ holds the instruction currently being executed
- ❑ **General Registers (Reg. 1..n)**
  - ❖ hold variables and temporary results
- ❑ **Arithmetic and Logic Unit (ALU)**
  - ❖ performs arithmetic functions and logic operations

# Basic anatomy on a CPU

---

- Stack Pointer (SP)

# Basic anatomy on a CPU

---

- **Stack Pointer (SP)**

- ❖ holds memory address of a stack with a frame for each active procedure's parameters & local variables

# Basic anatomy on a CPU

---

- **Stack Pointer (SP)**
  - ❖ holds memory address of a stack with a frame for each active procedure's parameters & local variables
- **Processor Status Word (PSW)**

# Basic anatomy on a CPU

---

- ❑ **Stack Pointer (SP)**
  - ❖ holds memory address of a stack with a frame for each active procedure's parameters & local variables
- ❑ **Processor Status Word (PSW)**
  - ❖ contains various control bits including the **mode bit** which determines whether privileged instructions can be executed at this time

# Basic anatomy on a CPU

---

- ❑ **Stack Pointer (SP)**
  - ❖ holds memory address of a stack with a frame for each active procedure's parameters & local variables
- ❑ **Processor Status Word (PSW)**
  - ❖ contains various control bits including the mode bit which determines whether privileged instructions can be executed
- ❑ **Memory Address Register (MAR)**
  - ❖ contains address of memory to be loaded from/stored to
- ❑ **Memory Data Register (MDR)**
  - ❖ contains memory data loaded or to be stored



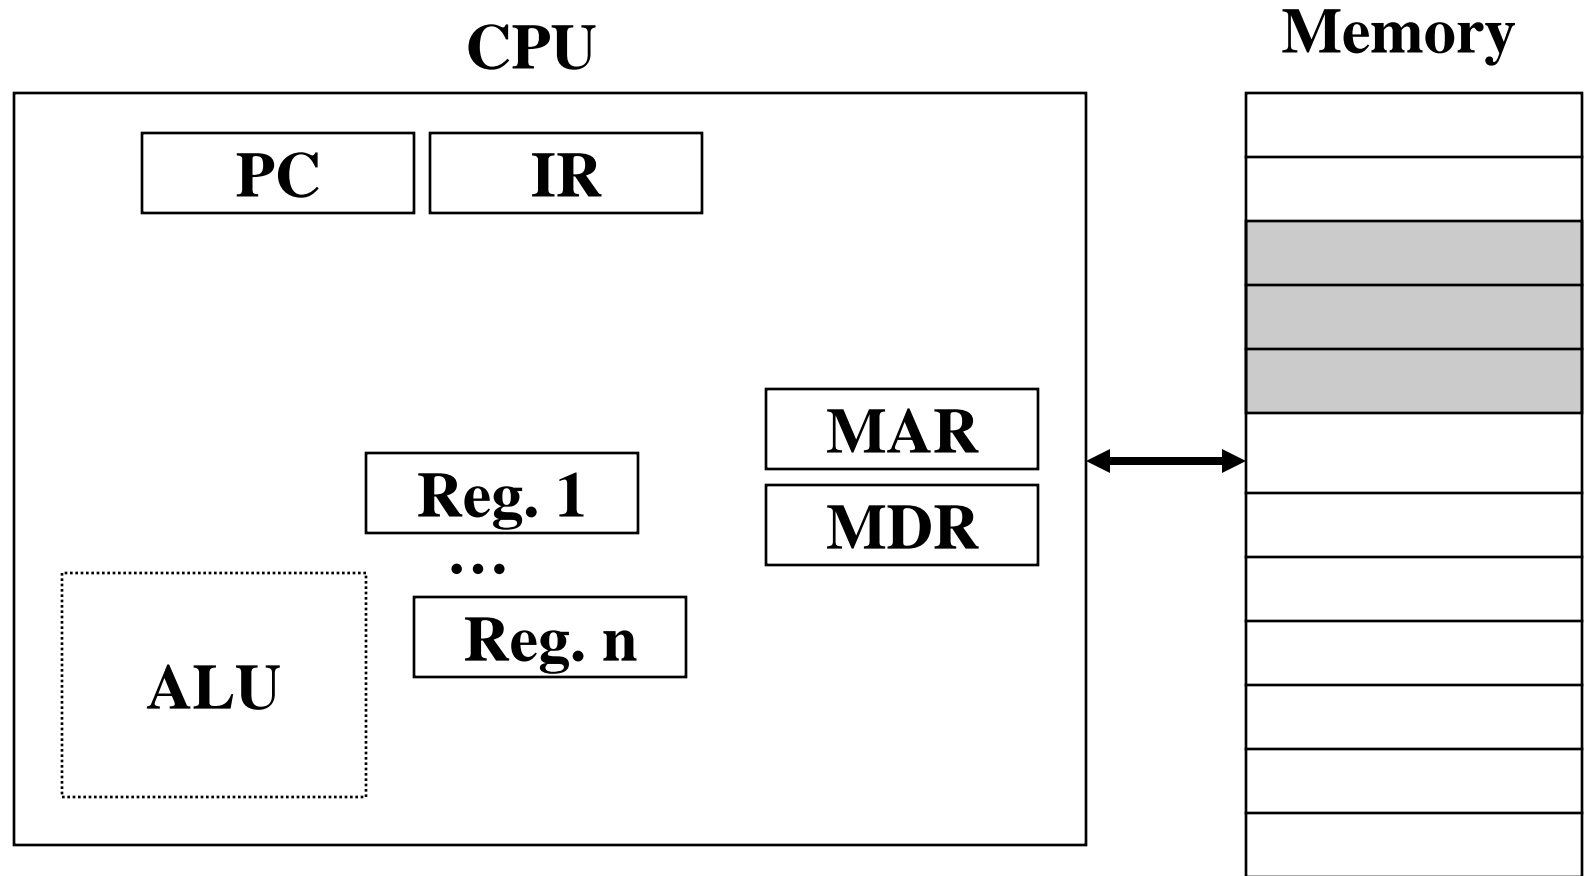
# Program execution

---

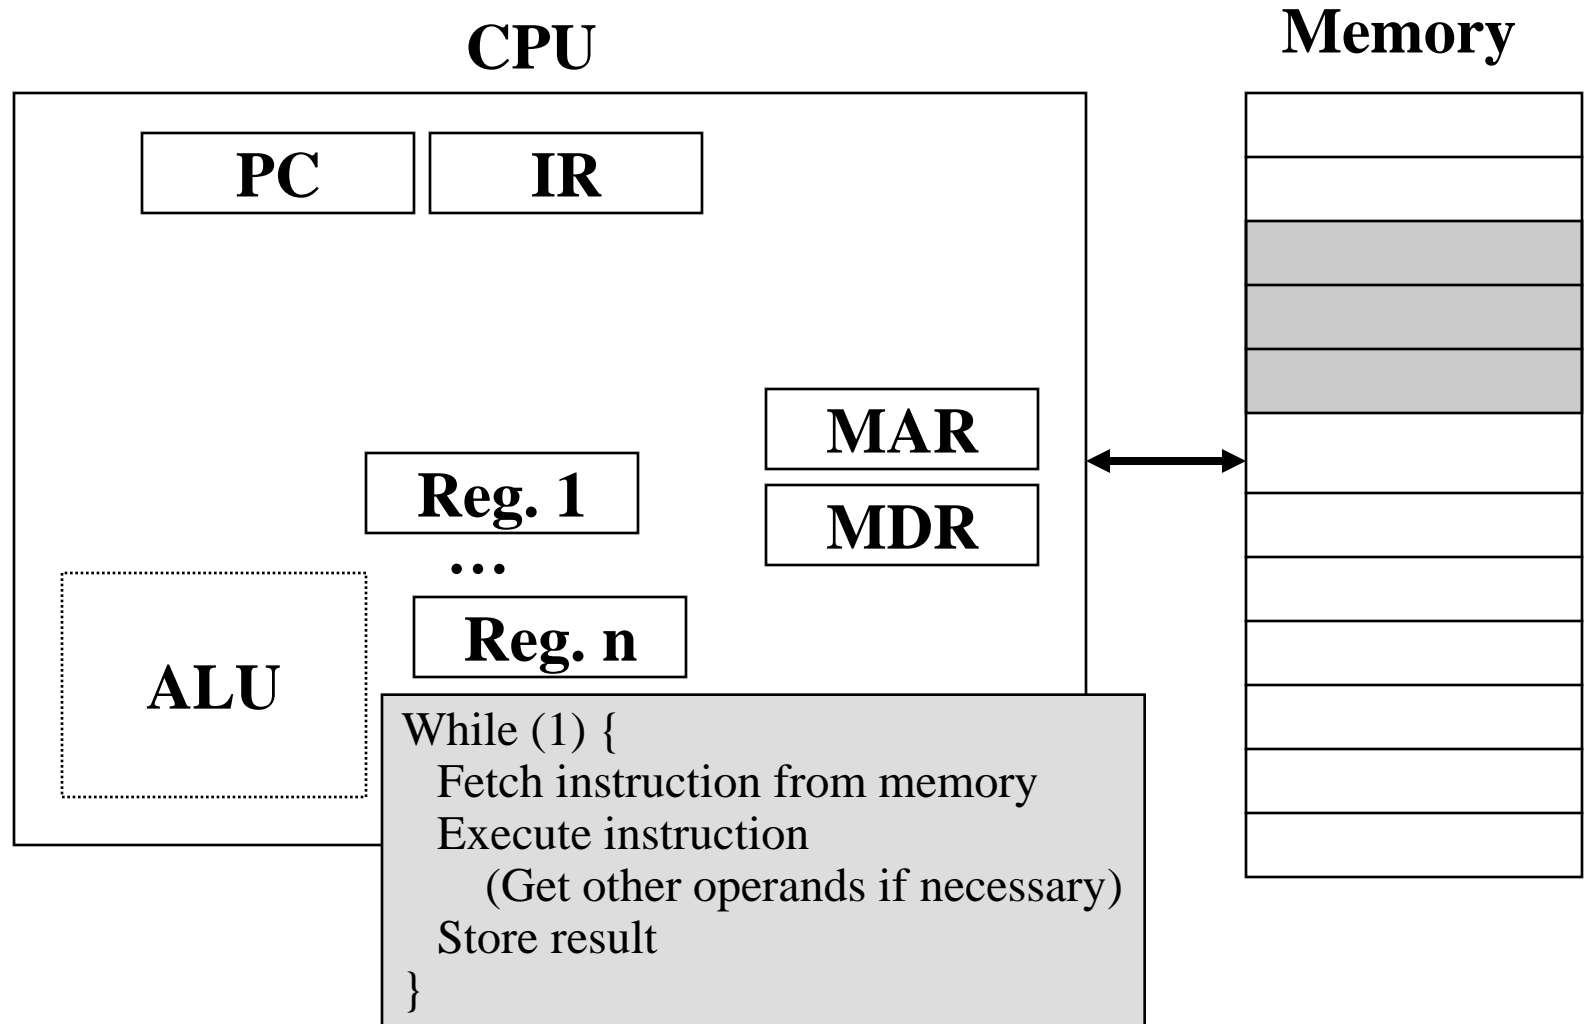
- **The Fetch/Decode/Execute cycle**
  - ❖ fetch next instruction pointed to by PC
  - ❖ decode it to find its type and operands
  - ❖ execute it
  - ❖ repeat
  
- **At a fundamental level, fetch/decode/execute is all a CPU does, regardless of which program it is executing**

# Fetch/decode/execute cycle

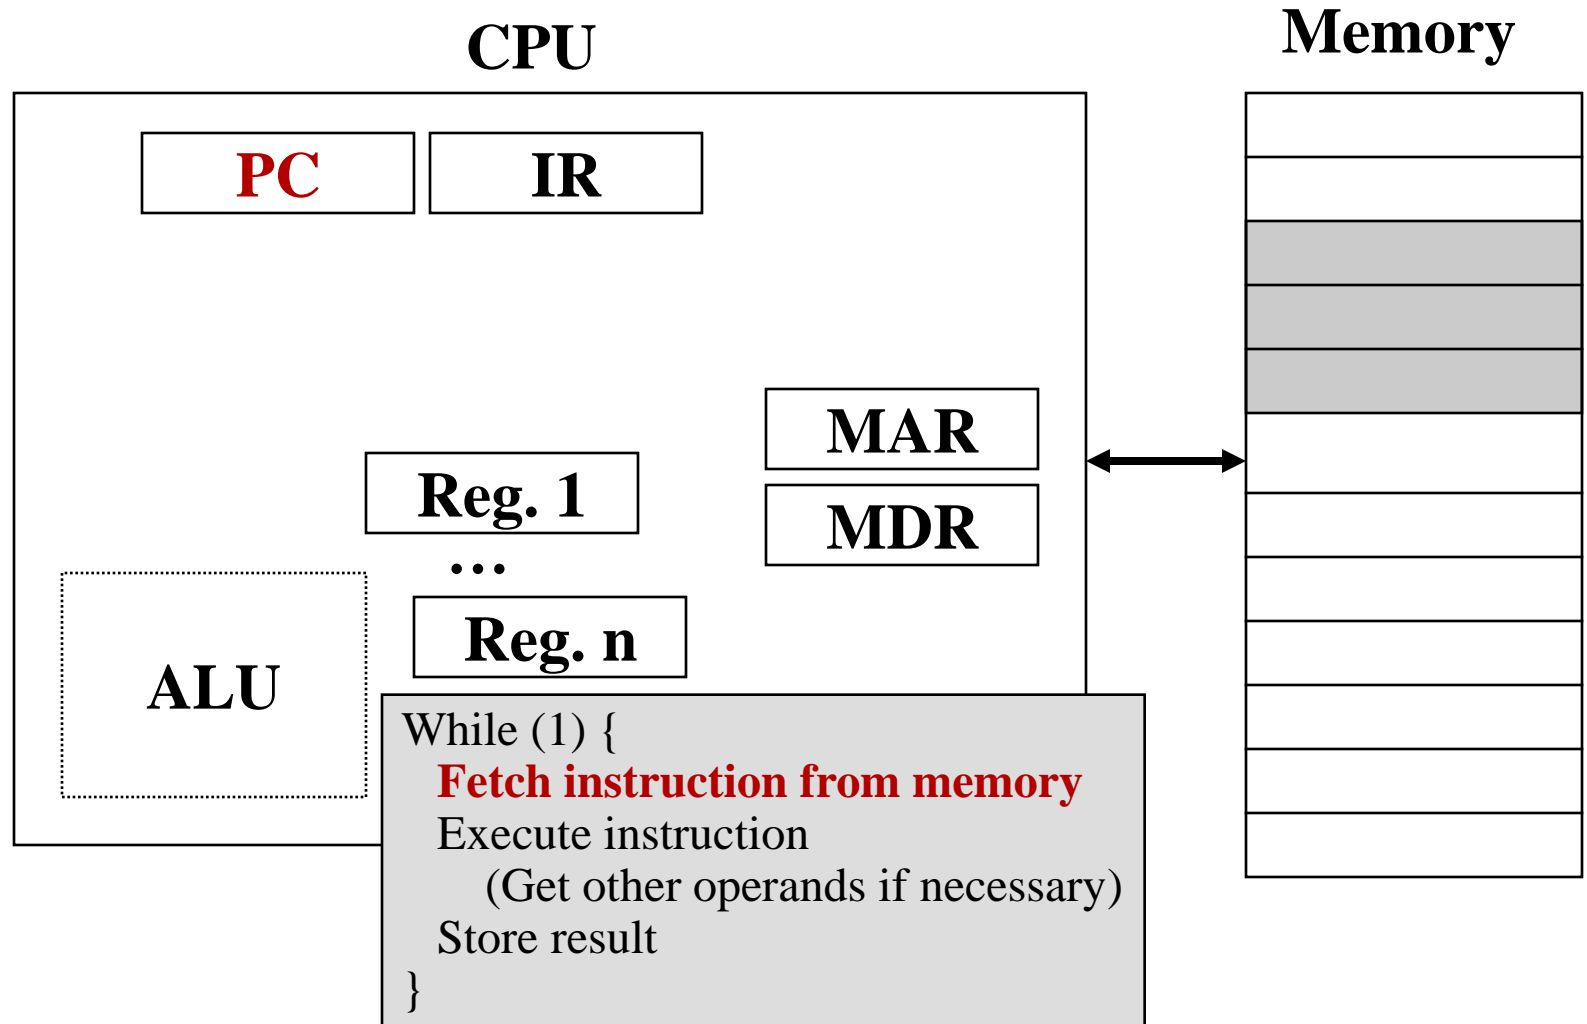
---



# Fetch/decode/execute cycle

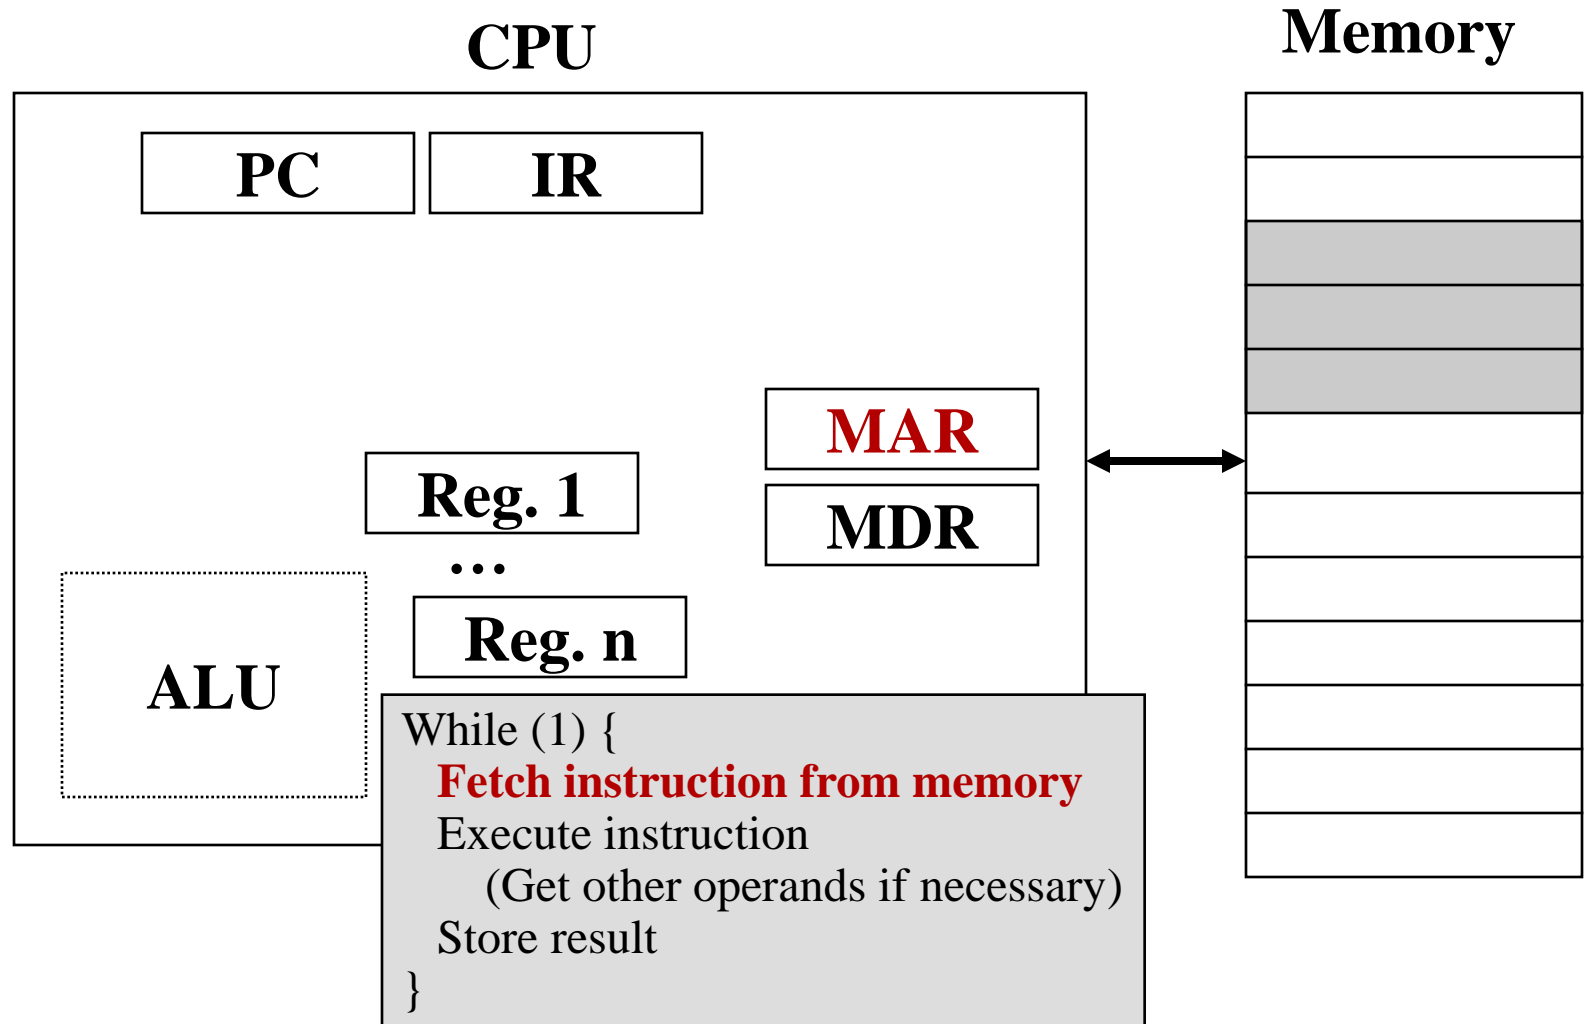


# Fetch/decode/execute cycle



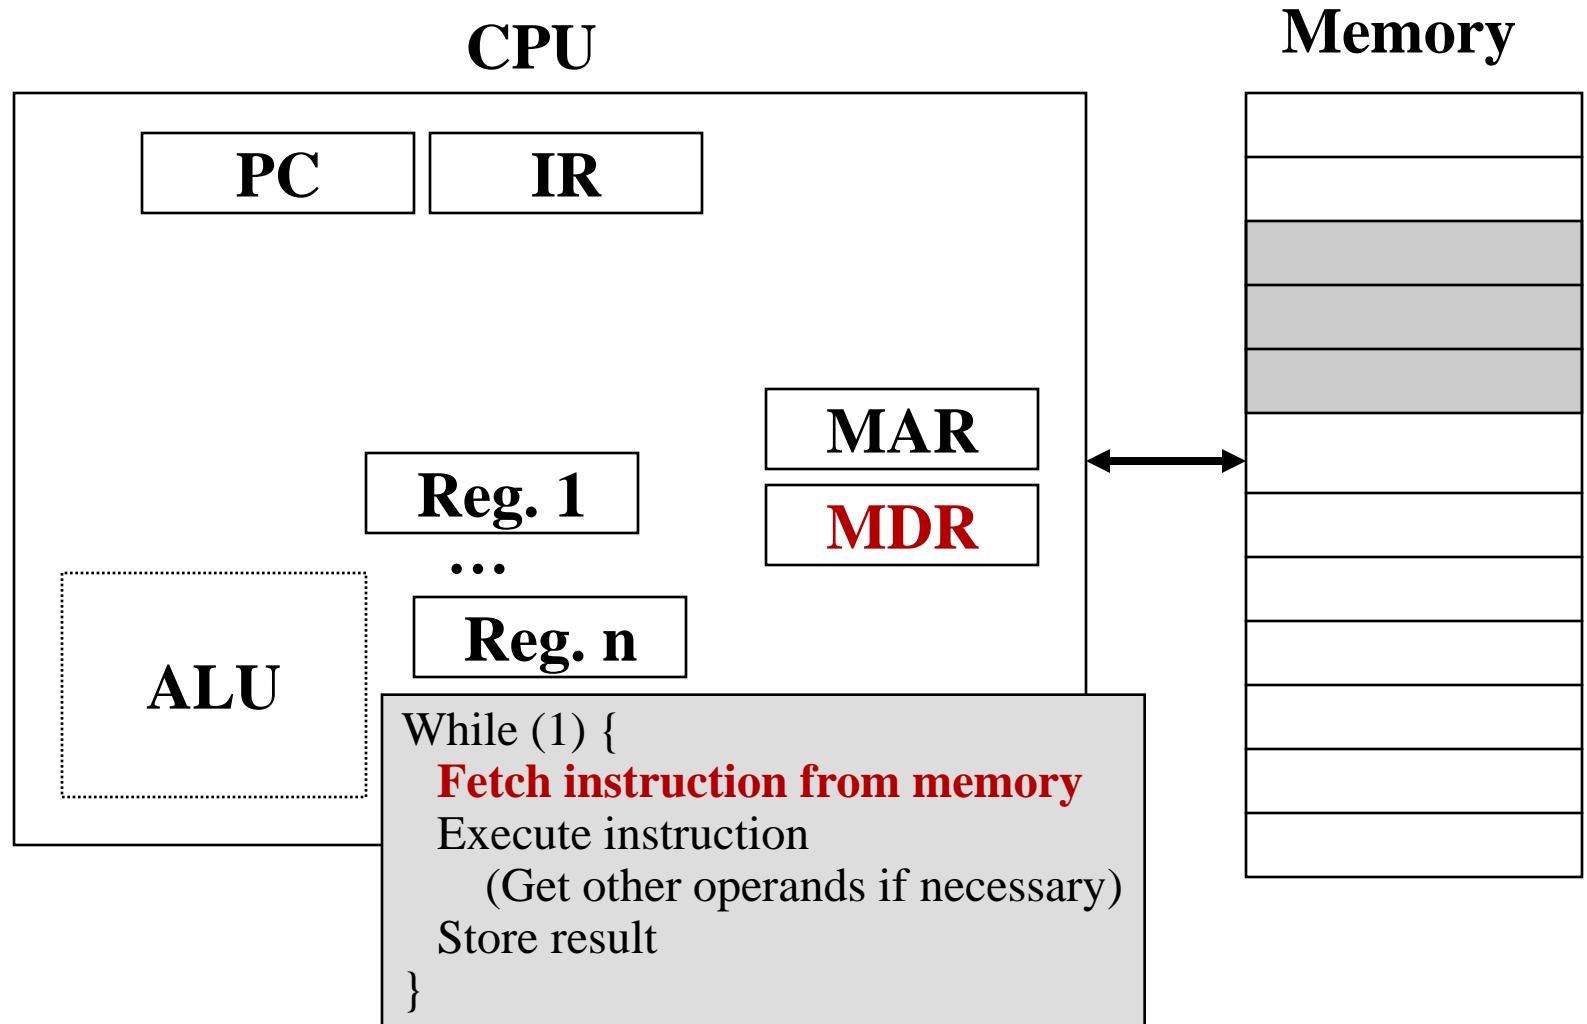
# Fetch/decode/execute cycle

---



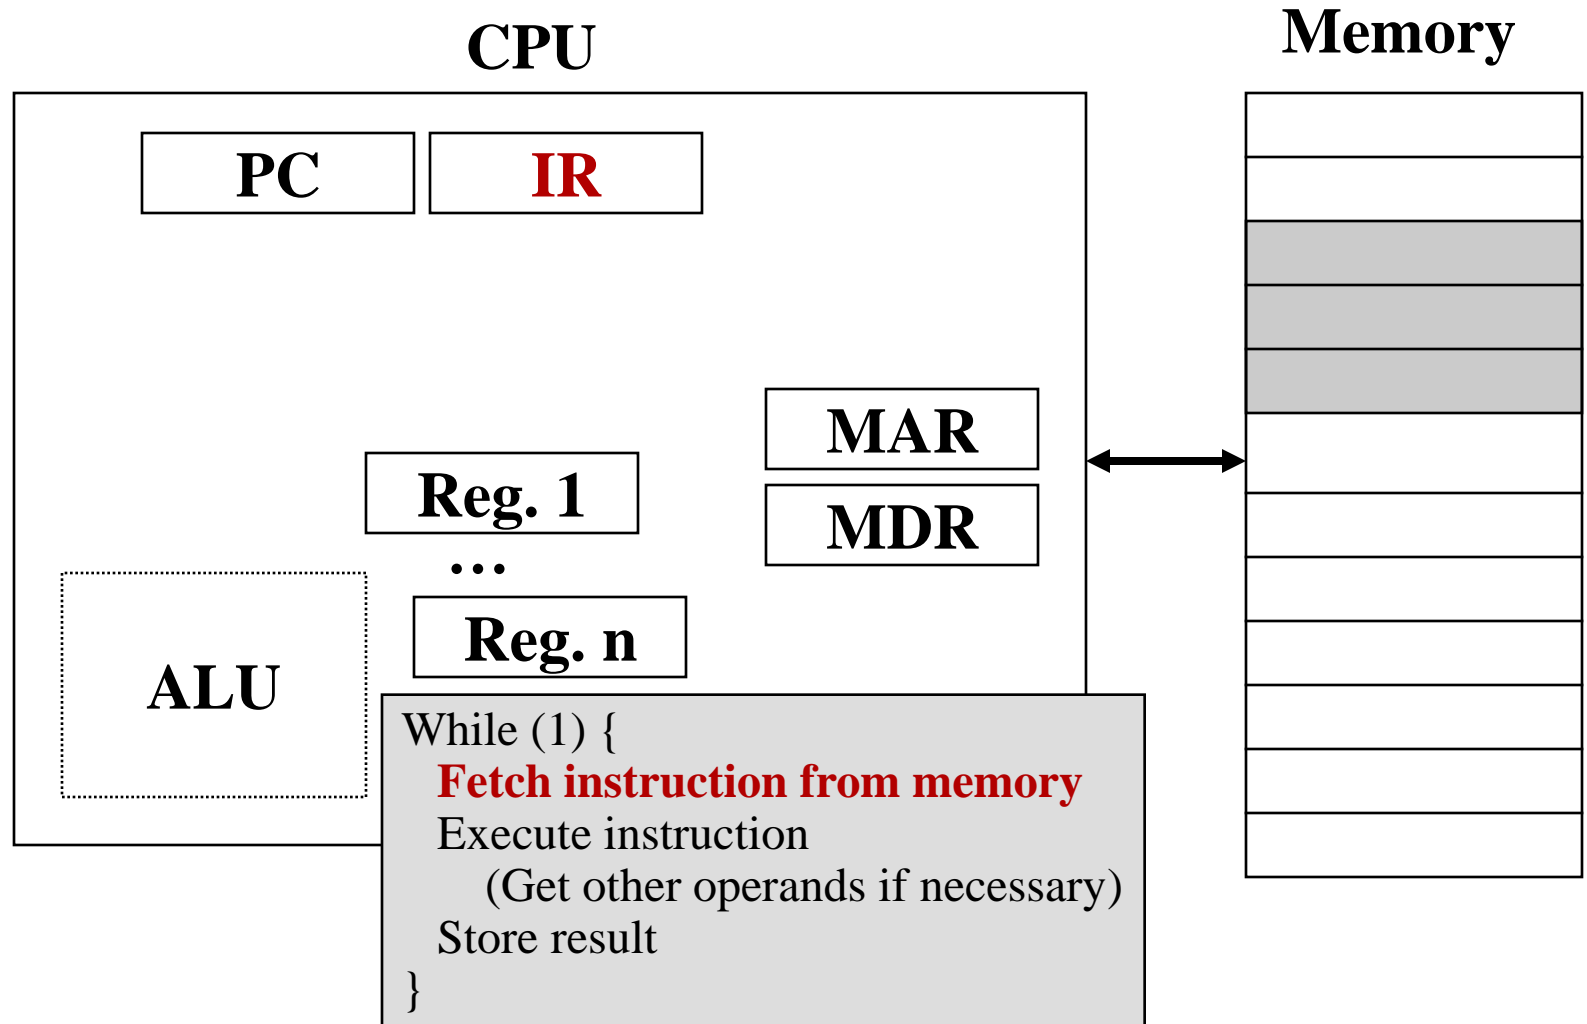
# Fetch/decode/execute cycle

---

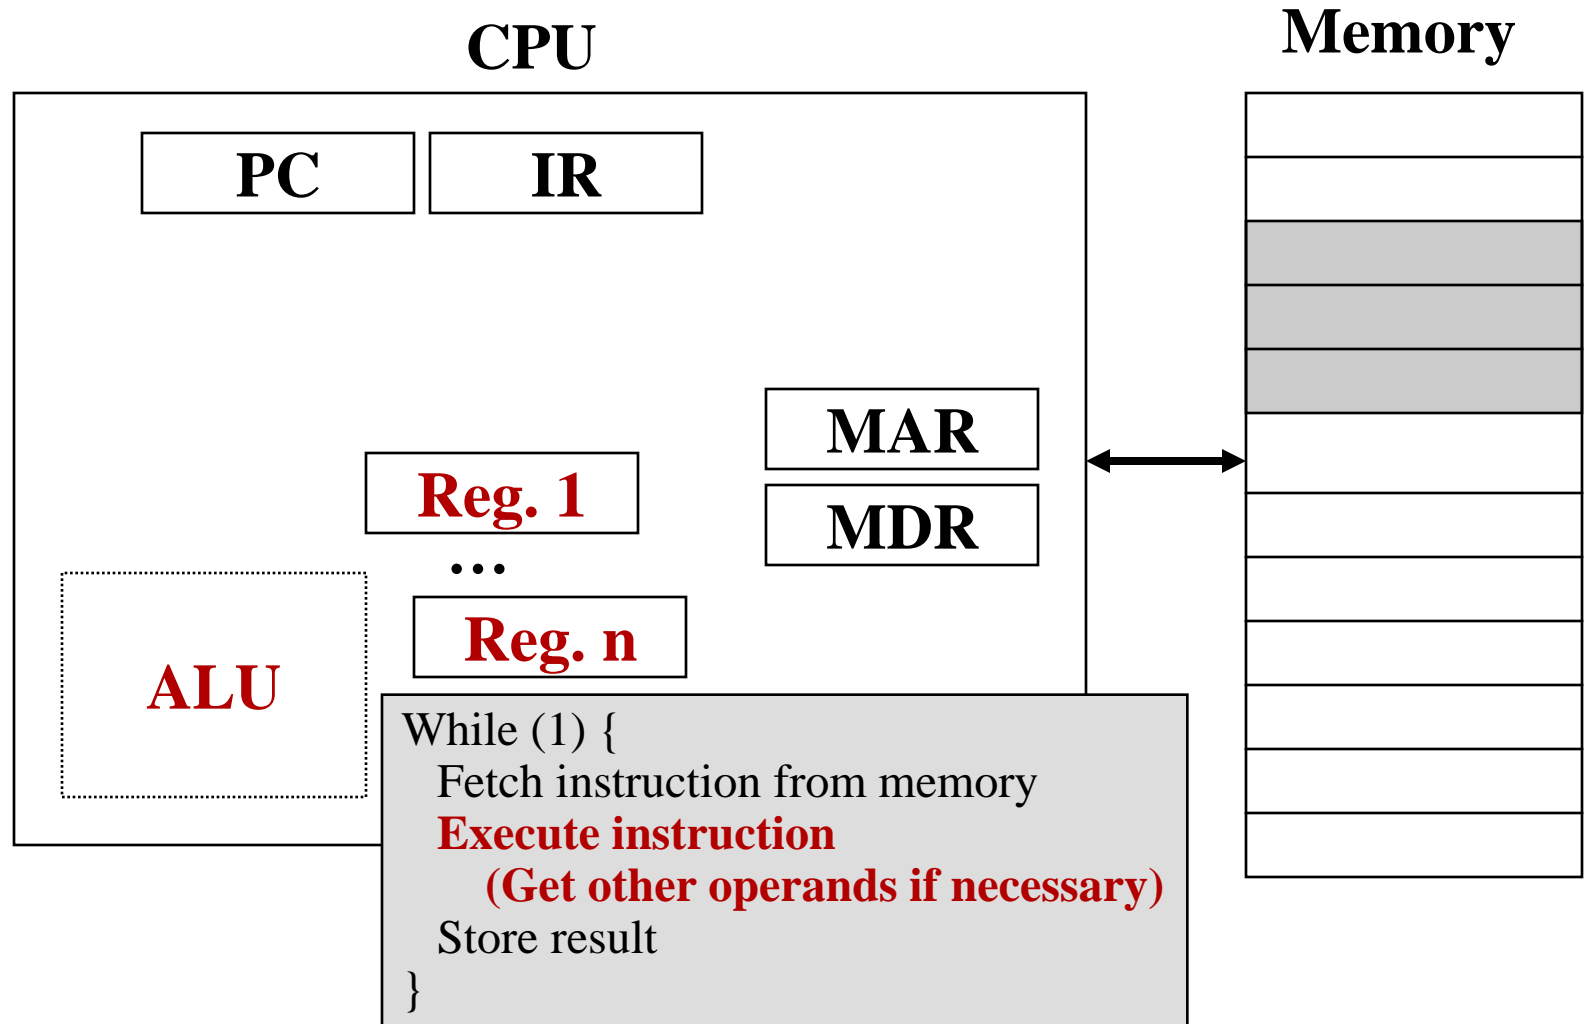


# Fetch/decode/execute cycle

---

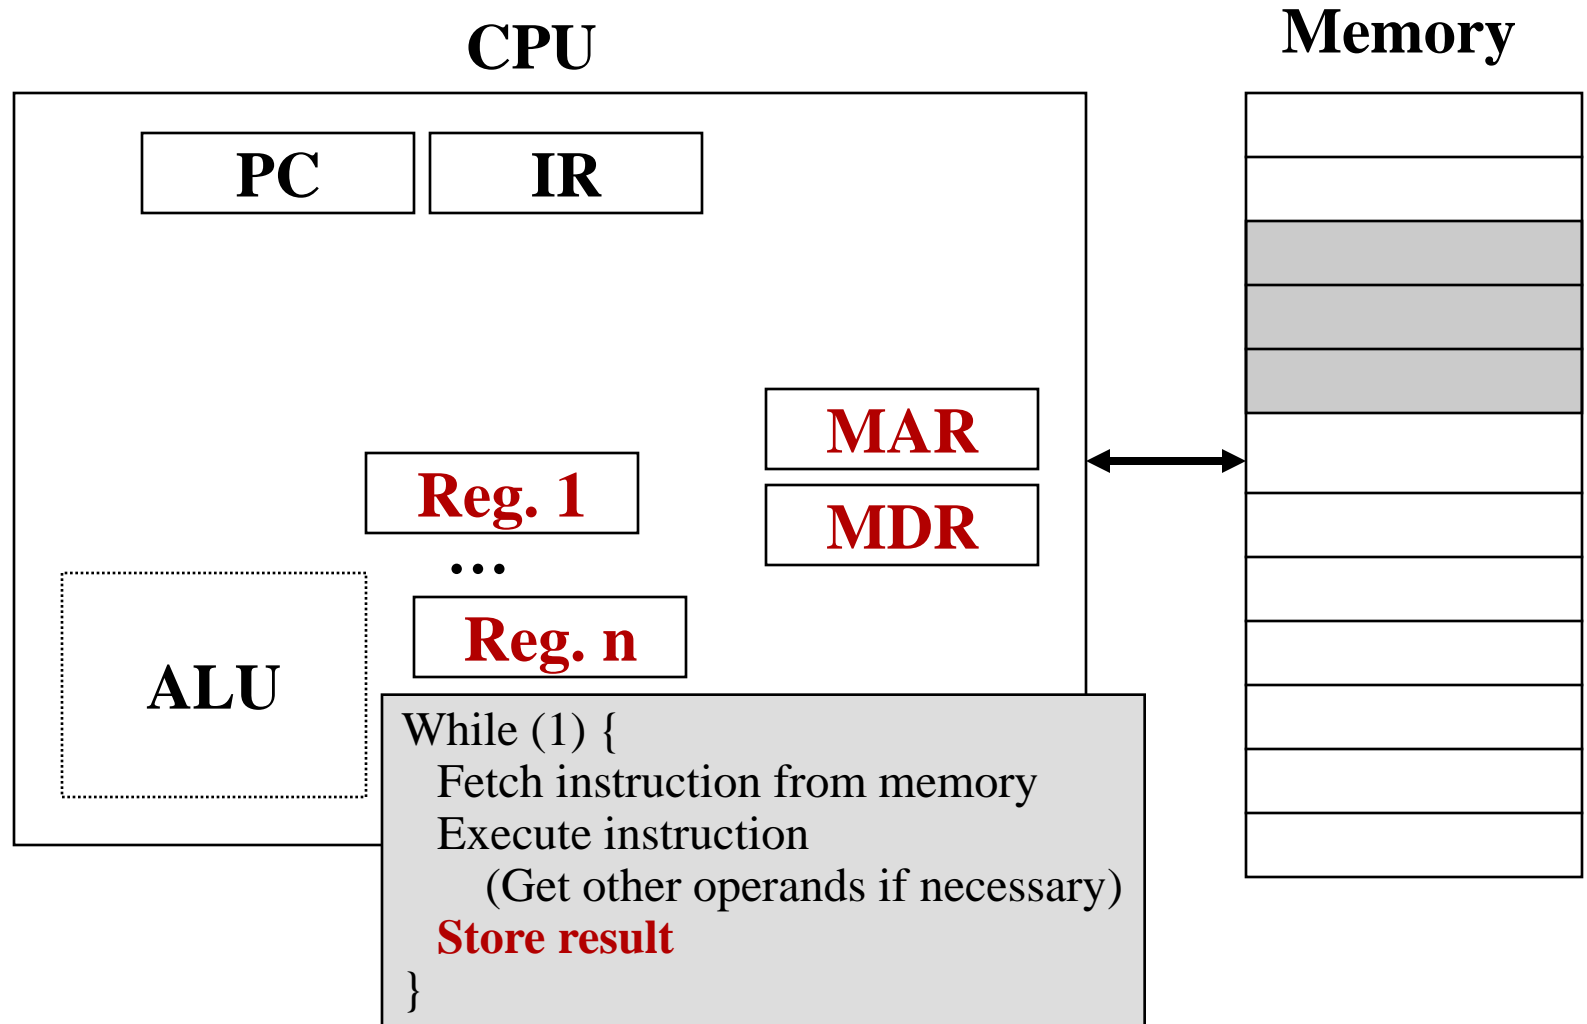


# Fetch/decode/execute cycle



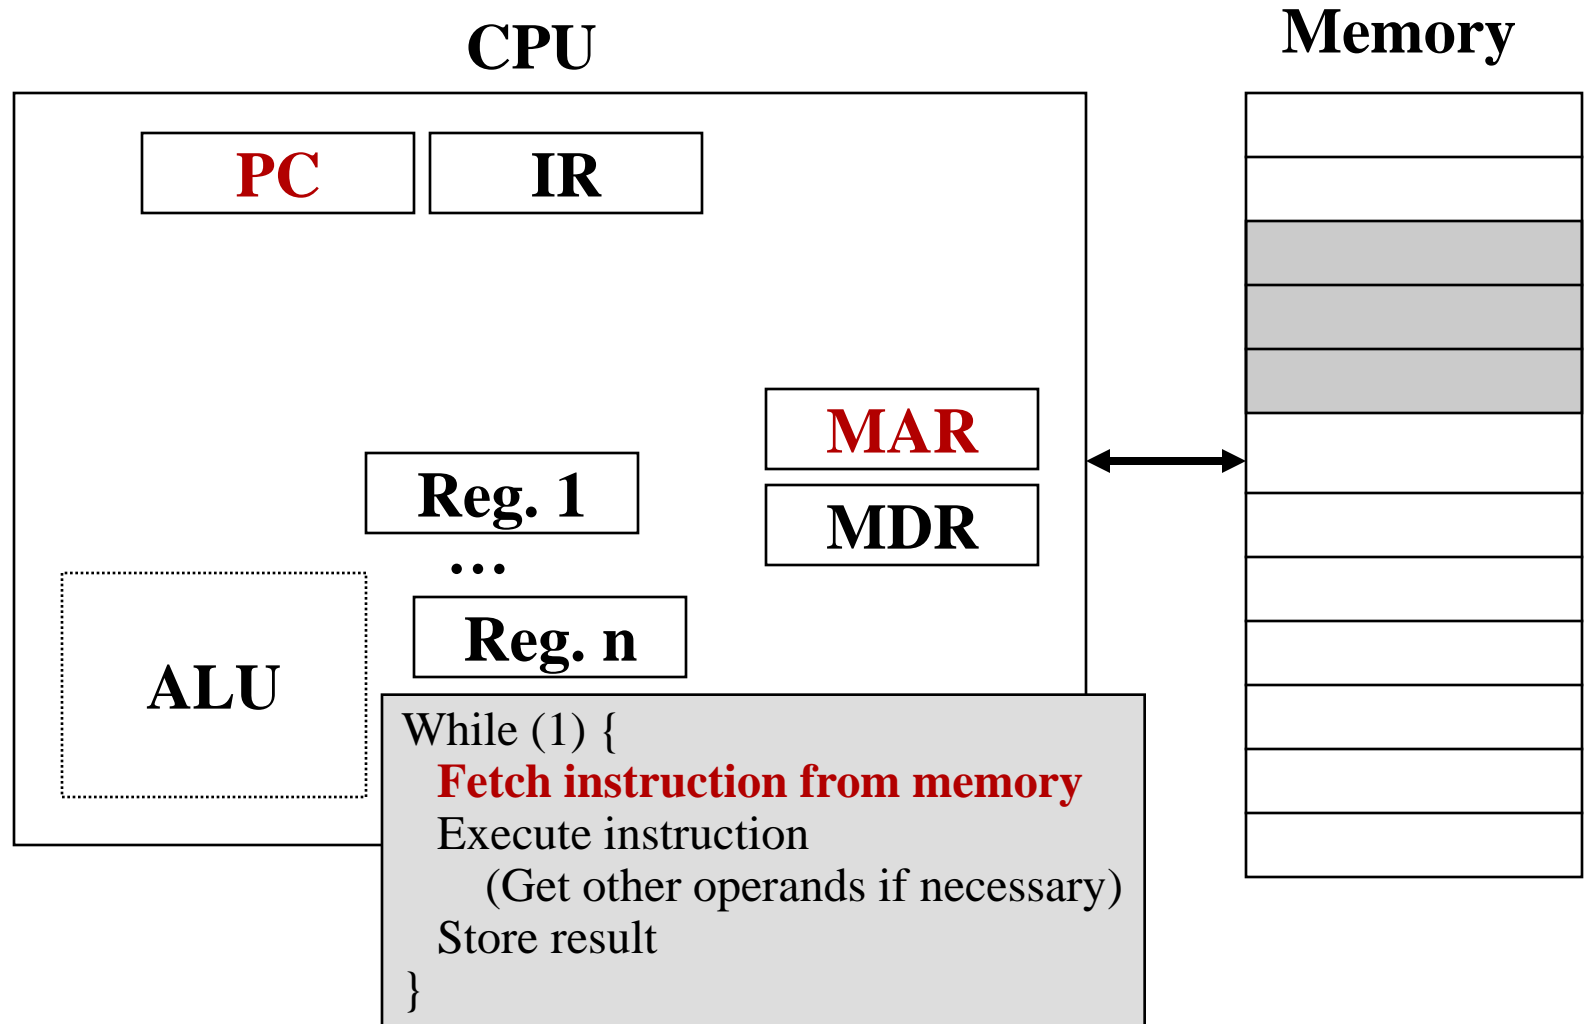


# Fetch/decode/execute cycle



# Fetch/decode/execute cycle

---



# The OS is just a program!

---

- The OS is a sequence of instructions that the CPU will fetch/decode/execute
  - ❖ How can the OS cause application programs to run?
  - ❖ How can the OS switch the CPU to run a different application and later resume the first one?
  - ❖ How can the OS maintain control?
  - ❖ In what ways can application code try to seize control indefinitely (ie. cheat)?
  - ❖ And how can the OS prevent such cheating?
  - ❖ How can applications programs cause the OS to run?

# How can the OS invoke an application?

---

# How can the OS invoke an application?

---

- Somehow, the OS must load the address of the application's starting instruction into the PC
  - ❖ The computer boots and begins running the OS
    - OS code must be loaded into memory somehow
    - fetch/decode/execute OS instructions
    - OS requests user input to identify application "file"
    - OS loads application file (executable) into memory
    - OS loads the memory address of the application's starting instruction into the PC
    - CPU fetches/decodes/executes the application's instructions

# How can OS guarantee to regain control?

---

- ❑ What if a running application doesn't make a system call and hence hogs the CPU?
  - ❖ OS needs interrupts from a timer device!
  - ❖ OS must register a future timer interrupt before it hands control of the CPU over to an application
  - ❖ When the timer interrupt goes off the interrupt hardware jumps control back into the OS at a pre-specified location called an interrupt handler
  - ❖ The interrupt handler is just a program (part of the OS)
  - ❖ The address of the interrupt handler's first instruction is placed in the PC by the interrupt h/w

# What if the application tries to cheat?

---

- What stops the running application from disabling the future timer interrupt so that the OS can not take control back from it?
  - ❖ Disabling interrupts must be a privileged instruction which is not executable by applications
  - ❖ The CPU knows whether or not to execute privileged instructions based on the value of the mode bit in the PSW!
- Privileged instructions can only be executed when the mode bit is set
  - ❖ disabling interrupts
  - ❖ setting the mode bit!
  - ❖ Attempted execution in non-privileged mode generally causes an interrupt (trap) to occur

# What other ways are there to cheat?

---

- What stops the running application from modifying the OS?
  - ❖ Specifically, what stops it from modifying the timer interrupt handler to jump control back to the application?



# What other ways are there to cheat?

---

- ❑ What stops the running application from modifying the OS?
  - ❖ Memory protection!
  - ❖ Memory protection instructions must be privileged
  - ❖ They can only be executed with the mode bit set ...
- ❑ Why must the OS clear the mode bit before it hands control to an application?

# How can applications invoke the OS?

---

- Why not just set PC to an OS instruction address and transfer control that way?

# How can applications invoke the OS?

---

- ❑ Special instruction causes a trap / interrupt
- ❑ Trap instruction changes PC to point to a predetermined OS entry point instruction and simultaneously sets the mode bit
  - ❖ application calls a library procedure that includes the appropriate trap instruction
  - ❖ fetch/decode/execute cycle begins at a pre-specified OS entry point called a **system call**
  - ❖ CPU is now running in privileged mode
- ❑ Traps, like interrupts, are hardware events, but they are caused by the executing program rather than a device external to the CPU

# How can the OS switch to a new application?

---

- To suspend execution of an application simply capture its memory state and processor state
  - ❖ preserve all the memory values of this application
  - ❖ copy values of all CPU registers into a data structure which is saved in memory
  - ❖ restarting the application from the same point just requires reloading the register values

# Recap

---

- ❖ Why do we need a timer device?
- ❖ Why do we need an interrupt mechanism?
- ❖ Why are system calls different to procedure calls?
- ❖ How are system calls different to interrupts?
- ❖ Why is memory protection necessary?

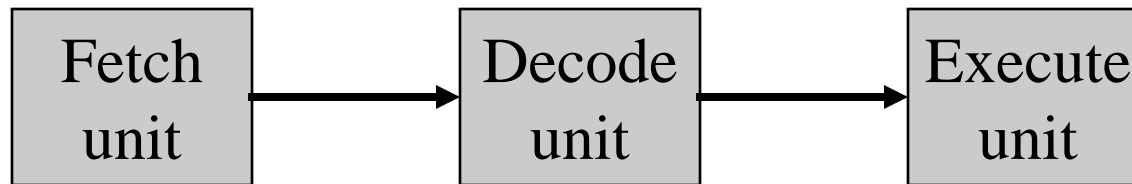
# Why its not quite that simple ...

---

- ❑ Pipelined CPUs
- ❑ Superscalar CPUs
- ❑ Multi-level memory hierarchies
- ❑ Virtual memory
- ❑ Complexity of devices and buses
- ❑ Heterogeneity of hardware

# Pipelined CPUs

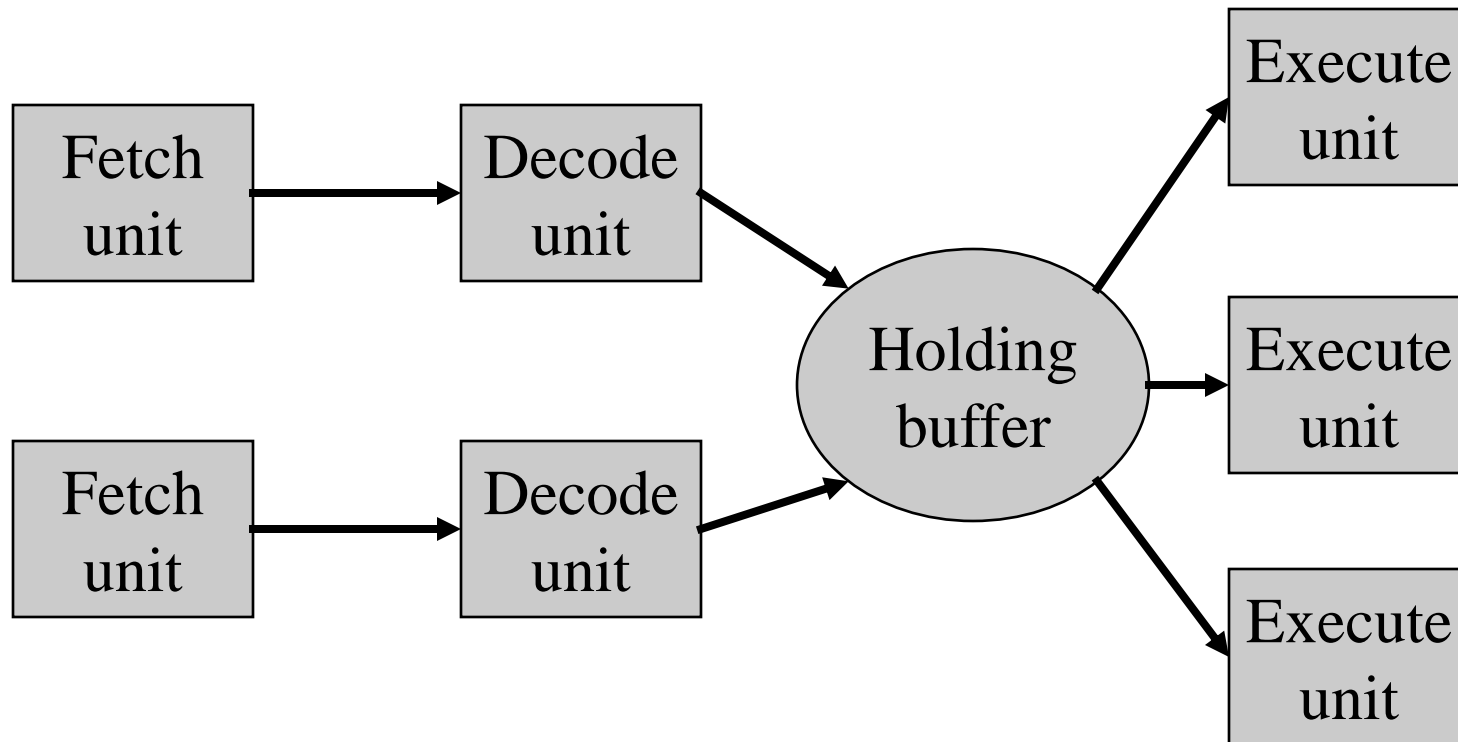
---



Execution of current instruction performed in parallel with decode of next instruction and fetch of the one after that

# Superscalar CPUs

---





# What does this mean for the OS?

---

- ❑ **Pipelined CPUs**
  - ❖ more complexity in capturing state of a running application
  - ❖ more expensive to suspend and resume applications
- ❑ **Superscalar CPUs**
  - ❖ even more complexity in capturing state of a running application
  - ❖ even more expensive to suspend and resume applications
- ❑ **More details, but fundamentally the same task**
- ❑ **The BLITZ CPU is not pipelined or superscalar**

# The memory hierarchy

- ❑ **2GHz processor → 0.5 ns**
- ❑ **Data/inst. cache → 0.5ns - 10 ns, 64 kB- 1MB**  
*(this is where the CPU looks first!)*
- ❑ **Main memory → 60 ns, 512 MB - 1GB**
- ❑ **Magnetic disk → 10 ms, 160 Gbytes**

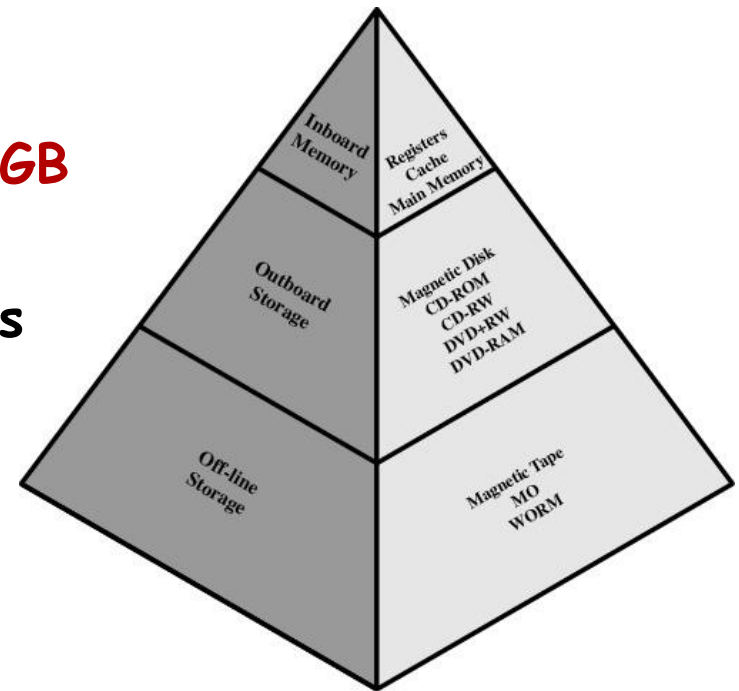


Figure 1.14 The Memory Hierarchy

# Who manages the memory hierarchy?

---

- **Movement of data from main memory to cache is under hardware control**
  - ❖ cache lines loaded on demand automatically
  - ❖ replacement policy fixed by hardware
- **Movement of data from cache to main memory can be affected by OS**
  - ❖ instructions for “flushing” the cache
  - ❖ can be used to maintain consistency of main memory
- **Movement of data among lower levels of the memory hierarchy is under direct control of the OS**
  - ❖ virtual memory page faults
  - ❖ file system calls

# OS implications of a memory hierarchy?

---

- ❑ How do you keep the contents of memory consistent across layers of the hierarchy?
- ❑ How do you allocate space at layers of the memory hierarchy “fairly” across different applications?
- ❑ How do you hide the latency of the slower subsystems?
  - These include main memory as well as disk!
- ❑ How do you protect one application's area of memory from other applications?
- ❑ How do you relocate an application in memory?

# Quiz

---

**How does the OS solve these problems:**

- ❖ Time sharing the CPU among applications?
- ❖ Space sharing the memory among applications?
- ❖ Protection of applications from each other?
- ❖ Protection of hardware/devices?
- ❖ Protection of the OS itself?

# What to do before next class

---

- ❑ Reading for today's class
- ❑ Reading for Wednesday's class
- ❑ Assignment 0 - read class web page and join class email list
- ❑ Start project 1 - Introduction to BLITZ