

CS 333
Introduction to Operating Systems
Class 5 - Semaphores and Classical
Synchronization Problems

Jonathan Walpole
Computer Science
Portland State University

An Example Synchronization Problem


The Producer-Consumer Problem

- ❑ An example of the **pipelined model**
 - ❖ One thread produces data items
 - ❖ Another thread consumes them
- ❑ Use a bounded buffer between the threads
- ❑ The buffer is a shared resource
 - ❖ Code that manipulates it is a **critical section**
- ❑ Must suspend the producer thread if the buffer is full
- ❑ Must suspend the consumer thread if the buffer is empty

Is this busy-waiting solution correct?

```
thread producer {
    while(1){
        // Produce char c
        while (count==n) {
            no_op
        }
        buf[InP] = c
        InP = InP + 1 mod n
        count++
    }
}
```

```
thread consumer {
    while(1){
        while (count==0) {
            no_op
        }
        c = buf[OutP]
        OutP = OutP + 1 mod n
        count--
        // Consume char
    }
}
```



Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

This code is incorrect!

- ❑ The “count” variable can be corrupted:
 - ❖ Increments or decrements may be lost!
 - ❖ Possible Consequences:
 - Both threads may spin forever
 - Buffer contents may be over-written
- ❑ *What is this problem called?*

This code is incorrect!

- ❑ The “count” variable can be corrupted:
 - ❖ Increments or decrements may be lost!
 - ❖ Possible Consequences:
 - Both threads may sleep forever
 - Buffer contents may be over-written
- ❑ *What is this problem called?* **Race Condition**
- ❑ Code that manipulates count must be made into a **???** and protected using **???**

This code is incorrect!

- ❑ The “count” variable can be corrupted:
 - ❖ Increments or decrements may be lost!
 - ❖ Possible Consequences:
 - Both threads may sleep forever
 - Buffer contents may be over-written
- ❑ *What is this problem called? **Race Condition***
- ❑ Code that manipulates count must be made into a **critical section** and protected using **mutual exclusion!**


Some more problems with this code

- ❑ **What if buffer is full?**
 - ❖ Producer will busy-wait
 - ❖ On a single CPU system the consumer will not be able to empty the buffer
- ❑ **What if buffer is empty?**
 - ❖ Consumer will busy-wait
 - ❖ On a single CPU system the producer will not be able to fill the buffer
- ❑ **We need a solution based on blocking!**

Producer/Consumer with Blocking - 1st attempt

```
0  thread producer {
1    while(1) {
2      // Produce char c
3      if (count==n) {
4        sleep(full)
5      }
6      buf[InP] = c;
7      InP = InP + 1 mod n
8      count++
9      if (count == 1)
10         wakeup(empty)
11    }
12 }
```

```
0  thread consumer {
1    while(1) {
2      if(count==0) {
3        sleep(empty)
4      }
5      c = buf[OutP]
6      OutP = OutP + 1 mod n
7      count--;
8      if (count == n-1)
9        wakeup(full)
10     // Consume char
11   }
12 }
```




Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

Use a mutex to fix the race condition in this code

```
0  thread producer {
1    while(1) {
2      // Produce char c
3      if (count==n) {
4        sleep(full)
5      }
6      buf[InP] = c;
7      InP = InP + 1 mod n
8      count++
9      if (count == 1)
10         wakeup(empty)
11    }
12 }
```

```
0  thread consumer {
1    while(1) {
2      if(count==0) {
3        sleep(empty)
4      }
5      c = buf[OutP]
6      OutP = OutP + 1 mod n
7      count--;
8      if (count == n-1)
9        wakeup(full)
10     // Consume char
11   }
12 }
```



Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

Problems

- ❑ Sleeping while holding the mutex causes deadlock !
- ❑ Releasing the mutex then sleeping opens up a window during which a context switch might occur ... again risking deadlock
- ❑ How can we release the mutex and sleep in a single atomic operation?
- ❑ We need a more powerful synchronization primitive

Semaphores

- An abstract data type that can be used for condition synchronization and mutual exclusion

What is the difference between mutual exclusion and condition synchronization?

Semaphores

- ❑ An abstract data type that can be used for condition synchronization and mutual exclusion
- ❑ Condition synchronization
 - ❖ **wait** until condition holds before proceeding
 - ❖ **signal** when condition holds so others may proceed
- ❑ Mutual exclusion
 - ❖ only one at a time in a critical section

Semaphores

- **An abstract data type**
 - ❖ containing an integer variable (S)
 - ❖ Two operations: Wait (S) and Signal (S)
- **Alternative names for the two operations**
 - ❖ $Wait(S) = Down(S) = P(S)$
 - ❖ $Signal(S) = Up(S) = V(S)$
- **Blitz names its semaphore operations Down and Up**

Classical Definition of Wait and Signal

Wait(S)

```
{  
    while S <= 0 do noop;    /* busy wait! */  
    S = S - 1;               /* S >= 0 */  
}
```

Signal (S)

```
{  
    S = S + 1;  
}
```

Problems with classical definition

- **Waiting threads hold the CPU**
 - ❖ Waste of time in single CPU systems
 - ❖ Required preemption to avoid deadlock

Blocking implementation of semaphores

Semaphore S has a value, $S.val$, and a thread list, $S.list$.

Wait (S)

```
S.val = S.val - 1
```

```
If S.val < 0
```

```
{ add calling thread to S.list;
  block;
}
```

```
/* negative value of S.val */
```

```
/* is # waiting threads */
```

```
/* sleep */
```

Signal (S)

```
S.val = S.val + 1
```

```
If S.val <= 0
```

```
{ remove a thread T from S.list;
  wakeup (T);
}
```

Implementing semaphores

- Wait () and Signal () are assumed to be **atomic**

How can we ensure that they are atomic?

Implementing semaphores

- Wait () and Signal () are assumed to be **atomic**

How can we ensure that they are atomic?

- Implement Wait() and Signal() as system calls?
 - ❖ how can the kernel ensure Wait() and Signal() are completed atomically?
 - ❖ Same solutions as before
 - Disable interrupts, or
 - Use TSL-based mutex

Semaphores with interrupt disabling

```
struct semaphore {  
    int val;  
    list L;  
}
```

```
Wait(semaphore sem)
```

```
    DISABLE_INTS
```

```
    sem.val--
```

```
    if (sem.val < 0){
```

```
        add thread to sem.L
```

```
        sleep(thread)
```

```
    }
```

```
    ENABLE_INTS
```

```
Signal(semaphore sem)
```

```
    DISABLE_INTS
```

```
    sem.val++
```

```
    if (sem.val <= 0) {
```

```
        th = remove next
```

```
        thread from sem.L
```

```
        wakeup(th)
```

```
    }
```

```
    ENABLE_INTS
```

Semaphores with interrupt disabling

```
struct semaphore {  
    int val;  
    list L;  
}
```

```
Wait(semaphore sem)  
    DISABLE_INTS  
    sem.val--  
    if (sem.val < 0){  
        add thread to sem.L  
        sleep(thread)  
    }  
    ENABLE_INTS
```

```
Signal(semaphore sem)  
    DISABLE_INTS  
    sem.val++  
    if (sem.val <= 0) {  
        th = remove next  
            thread from sem.L  
        wakeup(th)  
    }  
    ENABLE_INTS
```

Blitz code for Semaphore.wait

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait`
        operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Blitz code for Semaphore.wait

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
            operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Blitz code for Semaphore.wait

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
            operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```


Blitz code for Semaphore.wait

```
method Wait ()
    var oldIntStat: int
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x80000000
        FatalError ("Semaphore count underflowed during 'Wait'
            operation")
    EndIf
    count = count - 1
    if count < 0 waitingThreads.AddToEnd (currentThread)
        currentThread.Sleep ()
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

But what is `currentThread.Sleep ()`?

- ❑ If `sleep` stops a thread from executing, how, where, and when does it return?
 - ❖ which thread enables interrupts following `sleep`?
 - ❖ the thread that called `sleep` shouldn't return until another thread has called `signal` !
 - ❖ ... but how does that other thread get to run?
 - ❖ ... where exactly does the `thread switch` occur?
- ❑ Trace down through the Blitz code until you find a call to `switch()`
 - ❖ Switch is called in one thread but returns in another!
 - ❖ See where registers are saved and restored

Look at the following Blitz source code

- ❑ **Thread.c**
 - ❖ Thread.Sleep ()
 - ❖ Run (nextThread)

- ❑ **Switch.s**
 - ❖ Switch (prevThread, nextThread)

Blitz code for Semaphore.signal

```
method Signal ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during
            'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Blitz code for Semaphore.signal

```
method Signal ()
  var oldIntStat: int
  t: ptr to Thread
  oldIntStat = SetInterruptsTo (DISABLED)
  if count == 0x7fffffff
    FatalError ("Semaphore count overflowed during
      'Signal' operation")
  endIf
  count = count + 1
  if count <= 0
    t = waitingThreads.Remove ()
    t.status = READY
    readyList.AddToEnd (t)
  endIf
  oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Blitz code for Semaphore.signal

```
method Signal ()
  var oldIntStat: int
  t: ptr to Thread
  oldIntStat = SetInterruptsTo (DISABLED)
  if count == 0x7fffffff
    FatalError ("Semaphore count overflowed during
      'Signal' operation")
  endIf
  count = count + 1
  if count <= 0
    t = waitingThreads.Remove ()
    t.status = READY
    readyList.AddToEnd (t)
  endIf
  oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Blitz code for Semaphore.signal

```
method Signal ()
    var oldIntStat: int
    t: ptr to Thread
    oldIntStat = SetInterruptsTo (DISABLED)
    if count == 0x7fffffff
        FatalError ("Semaphore count overflowed during
            'Signal' operation")
    endIf
    count = count + 1
    if count <= 0
        t = waitingThreads.Remove ()
        t.status = READY
        readyList.AddToEnd (t)
    endIf
    oldIntStat = SetInterruptsTo (oldIntStat)
endMethod
```

Semaphores using atomic instructions

- **Implementing semaphores with interrupt disabling only works on uni-processors**
 - ❖ What should we do on a multiprocessor?
- **As we saw earlier, hardware provides special atomic instructions for synchronization**
 - ❖ test and set lock (TSL)
 - ❖ compare and swap (CAS)
 - ❖ etc
- **Semaphore can be built using atomic instructions**
 1. build mutex locks from atomic instructions
 2. build semaphores from mutex locks

Building *spinning* mutex locks using TSL

Mutex_lock:

TSL REGISTER,MUTEX	copy mutex to register and set mutex to 1
CMP REGISTER,#0	was mutex zero?
JZE ok	if it was zero, mutex is unlocked, so return
JMP mutex_lock	try again

Ok: RET	return to caller; enter critical section
---------	--

Mutex_unlock:

MOVE MUTEX,#0	store a 0 in mutex
RET	return to caller

Using Mutex Locks to Build Semaphores

- How would you modify the Blitz code to do this?

What if you had a blocking mutex lock?

Problem: Implement a counting semaphore

Up ()

Down ()

...using just Mutex locks

- Goal: Make use of the mutex lock's blocking behavior rather than reimplementing it for the semaphore operations

How about this solution?

```
var cnt: int = 0          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock(m1)
cnt = cnt - 1
if cnt < 0
    Lock(m2)
    Unlock(m1)
else
    Unlock(m1)
endIf
```

Up () :

```
Lock(m1)
cnt = cnt + 1
if cnt <= 0
    Unlock(m2)
endIf
Unlock(m1)
```

How about this solution?

```
var cnt: int = 0          -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock(m1)
cnt = cnt - 1
if cnt < 0
    Lock(m2)
    Unlock(m1)
else
    Unlock(m1)
endIf
```

Up () :

```
Lock(m2)
cnt = cnt + 1
if cnt <= 0
    Unlock(m2)
endIf
Unlock(m1)
```

Contains a
Deadlock!

How about this solution then?

```
var cnt: int = 0           -- Signal count
var m1: Mutex = unlocked  -- Protects access to "cnt"
    m2: Mutex = locked    -- Locked when waiting
```

Down () :

```
Lock(m1)
cnt = cnt - 1
if cnt < 0
    Unlock(m1)
    Lock(m2)
else
    Unlock(m1)
endIf
```

Up () :


```
Lock(m1)
cnt = cnt + 1
if cnt <= 0
    Unlock(m2)
endIf
Unlock(m1)
```

Classical Synchronization problems

- ❑ Producer Consumer (bounded buffer)
- ❑ Dining philosophers
- ❑ Sleeping barber
- ❑ Readers and writers

Producer consumer problem


- Also known as the bounded buffer problem



Is this a valid solution?

```
thread producer {
    while(1){
        // Produce char c
        while (count==n) {
            no_op
        }
        buf[InP] = c
        InP = InP + 1 mod n
        count++
    }
}
```

```
thread consumer {
    while(1){
        while (count==0) {
            no_op
        }
        c = buf[OutP]
        OutP = OutP + 1 mod n
        count--
        // Consume char
    }
}
```



Global variables:

```
char buf[n]
int InP = 0    // place to add
int OutP = 0   // place to get
int count
```

Does this solution work?

Global variables

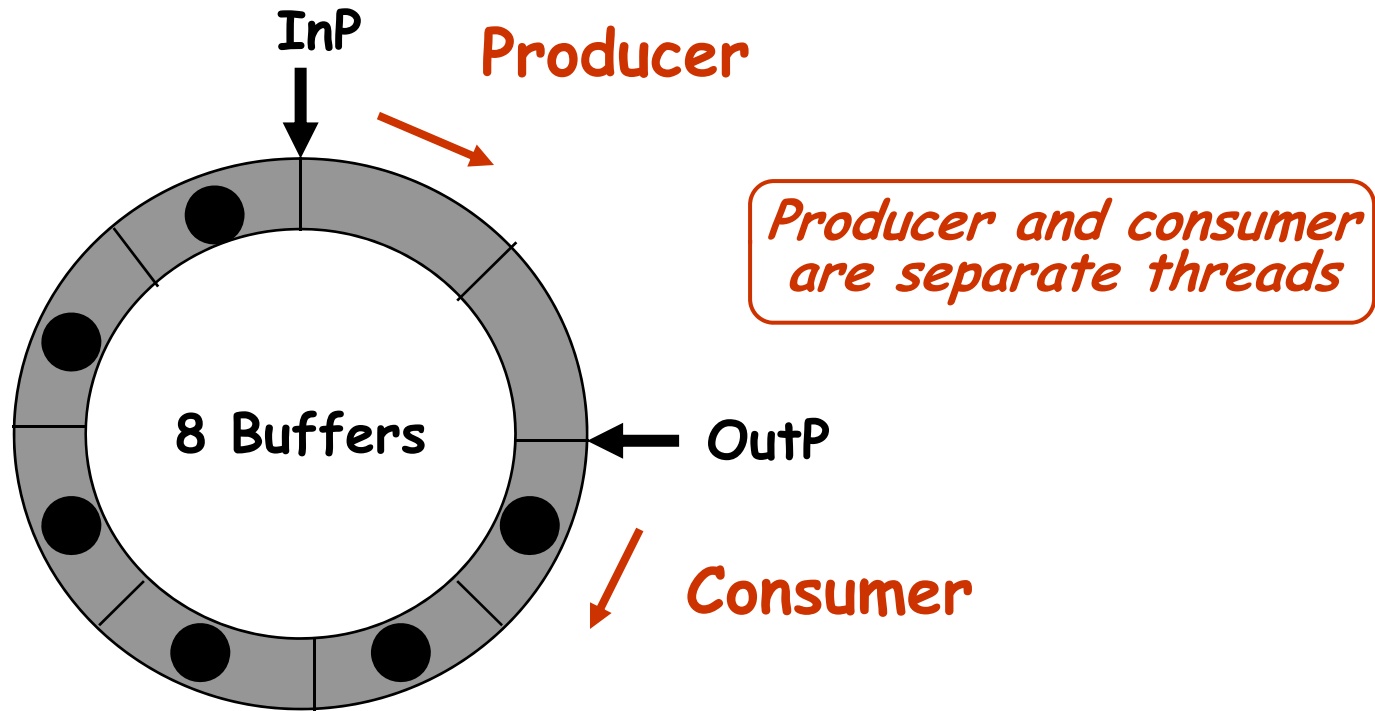
```
semaphore full_buffs = 0;  
semaphore empty_buffs = n;  
char buff[n];  
int InP, OutP;
```

```
0 thread producer {  
1   while(1){  
2       // Produce char c...  
3       down(empty_buffs)  
4       buf[InP] = c  
5       InP = InP + 1 mod n  
6       up(full_buffs)  
7   }  
8 }
```

```
0 thread consumer {  
1   while(1){  
2       down(full_buffs)  
3       c = buf[OutP]  
4       OutP = OutP + 1 mod n  
5       up(empty_buffs)  
6       // Consume char...  
7   }  
8 }
```

Producer consumer problem

- ❑ What is the shared state in the last solution?
- ❑ Does it apply mutual exclusion? If so, how?




Problems with solution

- What if we have multiple producers and multiple consumers?
 - ❖ Producer-specific and consumer-specific data becomes shared
 - ❖ We need to define and protect critical sections

Dining philosophers problem

- ❑ Five philosophers sit at a table
- ❑ One fork between each philosopher



Each philosopher is modeled with a thread

```
while(TRUE) {  
    Think();  
    Grab first fork;  
    Grab second fork;  
    Eat();  
    Put down first fork;  
    Put down second fork;  
}
```

- ❑ *Why do they need to synchronize?*
- ❑ *How should they do it?*

Is this a valid solution?

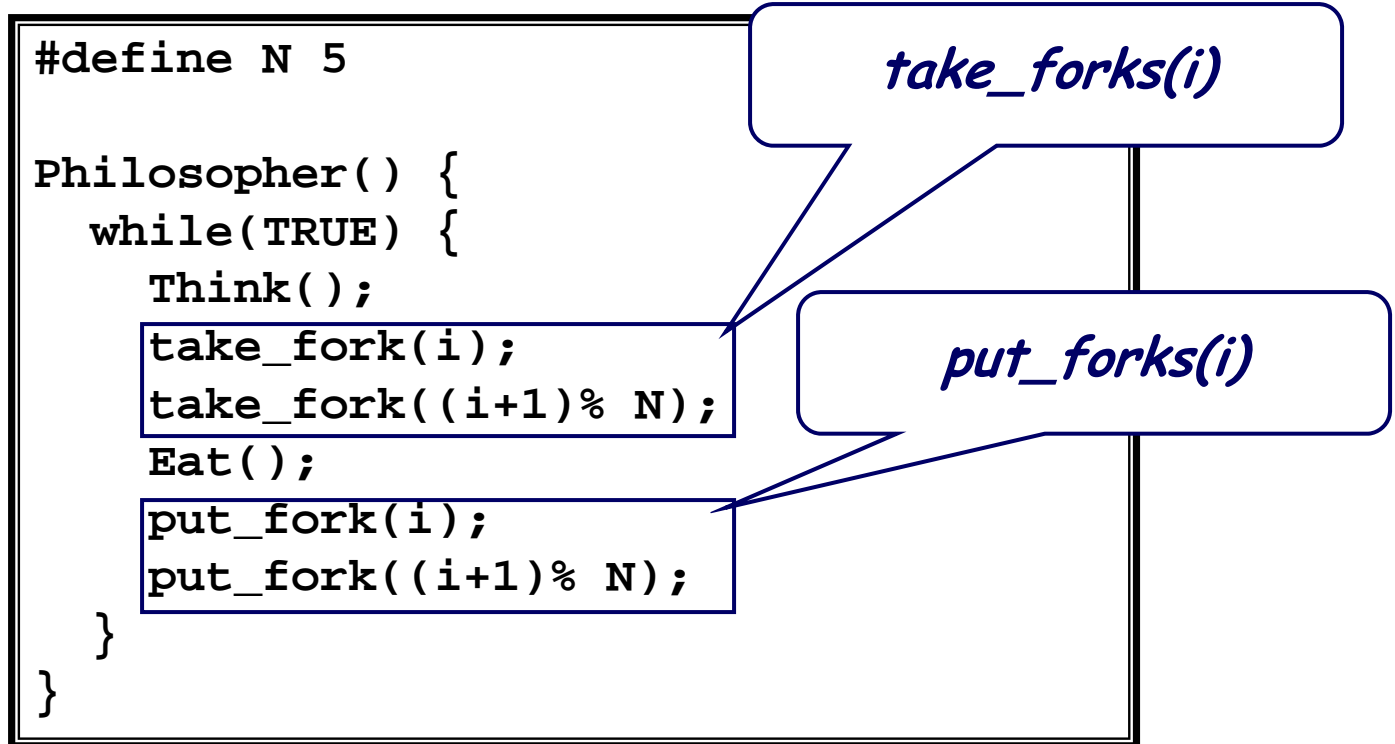
```
#define N 5

Philosopher() {
    while(TRUE) {
        Think();
        take_fork(i);
        take_fork((i+1)% N);
        Eat();
        put_fork(i);
        put_fork((i+1)% N);
    }
}
```

Problems

- Potential for deadlock !

Working towards a solution ...



Working towards a solution ...

```
#define N 5

Philosopher() {
    while(TRUE) {
        Think();
        take_forks(i);
        Eat();
        put_forks(i);
    }
}
```

Picking up forks

```
int state[N]
semaphore mutex = 1
semaphore sem[i]
```

```
take_forks(int i) {
    wait(mutex);
    state[i] = HUNGRY;
    test(i);
    signal(mutex);
    wait(sem[i]);
}
```

```
// only called with mutex set!

test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        signal(sem[i]);
    }
}
```

Putting down forks

```
int state[N]
semaphore mutex = 1
semaphore sem[i]
```

```
put_forks(int i) {
    wait(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    signal(mutex);
}
```


```
// only called with mutex set!

test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        signal(sem[i]);
    }
}
```

Dining philosophers

- ❑ Is the previous solution correct?
- ❑ What does it mean for it to be correct?
- ❑ Is there an easier way?

The sleeping barber problem



The sleeping barber problem

□ *Barber:*

- ❖ While there are people waiting for a hair cut, put one in the barber chair, and cut their hair
- ❖ When done, move to the next customer
- ❖ Else go to sleep, until someone comes in

□ *Customer:*

- ❖ If barber is asleep wake him up for a haircut
- ❖ If someone is getting a haircut wait for the barber to become free by sitting in a chair
- ❖ If all chairs are all full, leave the barbershop

Designing a solution

- ❑ How will we model the barber and customers?
- ❑ What state variables do we need?
 - ❖ .. and which ones are shared?
 - ❖ and how will we protect them?
- ❑ How will the barber sleep?
- ❑ How will the barber wake up?
- ❑ How will customers wait?
- ❑ What problems do we need to look out for?

Is this a good solution?

```
const CHAIRS = 5
var customers: Semaphore
    barbers: Semaphore
    lock: Mutex
    numWaiting: int = 0
```

Barber Thread:

```
while true
    Wait(customers)
    Lock(lock)
    numWaiting = numWaiting-1
    Signal(barbers)
    Unlock(lock)
    CutHair()
endWhile
```

Customer Thread:

```
Lock(lock)
if numWaiting < CHAIRS
    numWaiting = numWaiting+1
    Signal(customers)
    Unlock(lock)
    Wait(barbers)
    GetHaircut()
else -- give up & go home
    Unlock(lock)
endIf
```


The readers and writers problem

- ❑ Multiple readers and writers want to access a database (each one is a thread)
- ❑ Multiple readers can proceed concurrently
- ❑ Writers must synchronize with readers and other writers
 - ❖ *only one writer at a time !*
 - ❖ *when someone is writing, there must be no readers !*

Goals:

- ❖ Maximize concurrency.
- ❖ Prevent starvation.

Designing a solution

- ❑ How will we model the readers and writers?
- ❑ What state variables do we need?
 - ❖ .. and which ones are shared?
 - ❖ and how will we protect them?
- ❑ How will the writers wait?
- ❑ How will the writers wake up?
- ❑ How will readers wait?
- ❑ How will the readers wake up?
- ❑ What problems do we need to look out for?

Is this a valid solution to readers & writers?

```
var mut: Mutex = unlocked
    db: Semaphore = 1
    rc: int = 0
```

Writer Thread:

```
while true
    ...Remainder Section...
    Wait(db)
    ...Write shared data...
    Signal(db)
endWhile
```

Reader Thread:

```
while true
    Lock(mut)
    rc = rc + 1
    if rc == 1
        Wait(db)
    endIf
    Unlock(mut)
    ... Read shared data...
    Lock(mut)
    rc = rc - 1
    if rc == 0
        Signal(db)
    endIf
    Unlock(mut)
    ... Remainder Section...
endWhile
```

Readers and writers solution

- Does the previous solution have any problems?
 - ❖ is it "fair"?
 - ❖ can any threads be starved? If so, how could this be fixed?
 - ❖ ... and how much confidence would you have in your solution?

Quiz

- ❑ What is a race condition?
- ❑ How can we protect against race conditions?
- ❑ Can locks be implemented simply by reading and writing to a binary variable in memory?
- ❑ How can a kernel make synchronization-related system calls atomic on a uniprocessor?
 - ❖ Why wouldn't this work on a multiprocessor?
- ❑ Why is it better to block rather than spin on a uniprocessor?
- ❑ Why is it sometimes better to spin rather than block on a multiprocessor?

Quiz

- ❑ When faced with a concurrent programming problem, what strategy would you follow in designing a solution?
- ❑ What does all of this have to do with Operating Systems?