

CS 333
Introduction to Operating Systems
Class 6 - Monitors and Message Passing

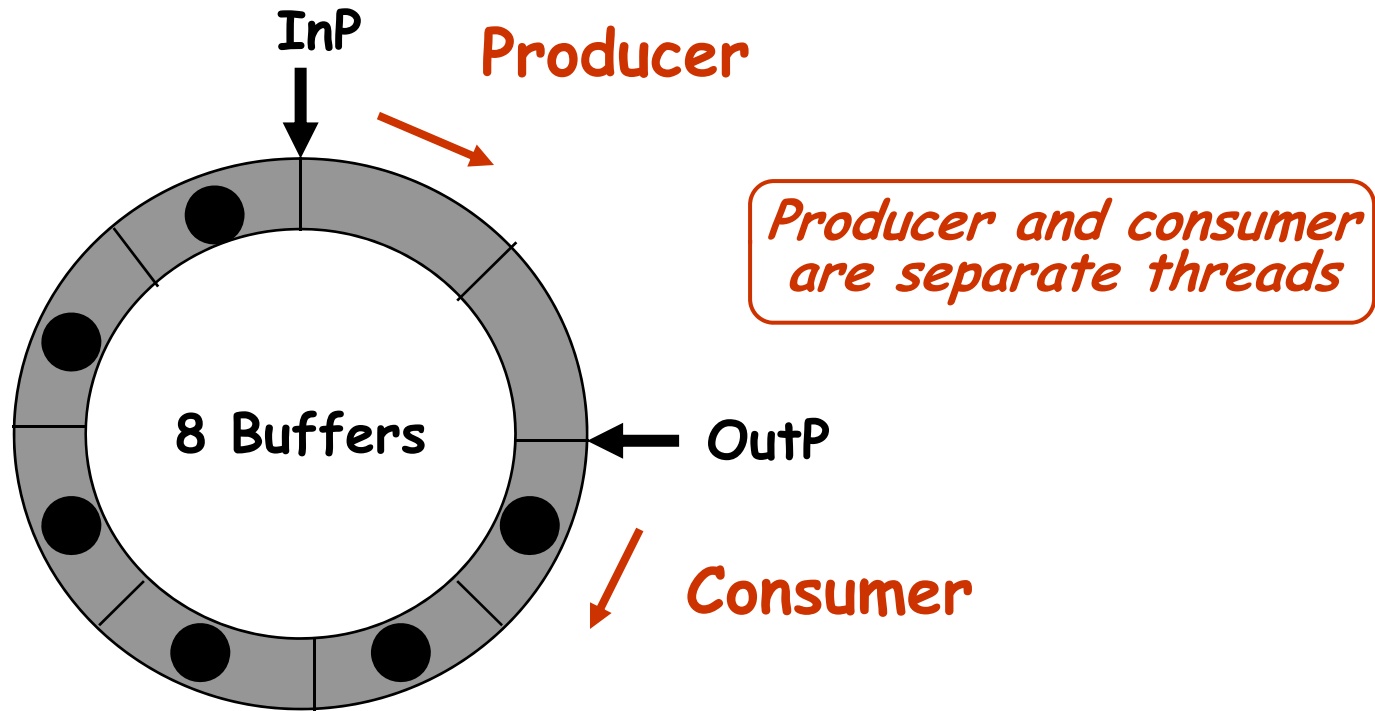
Jonathan Walpole
Computer Science
Portland State University

But first ...

- Continuation of Class 5 - Classical Synchronization Problems

Producer consumer problem

- Also known as the bounded buffer problem



Does this solution work?

Global variables


```
semaphore full_buffs = 0;
semaphore empty_buffs = n;
char buff[n];
int InP, OutP;
```

```
0 thread producer {
1   while(1){
2     // Produce char c...
3     down(empty_buffs)
4     buf[InP] = c
5     InP = InP + 1 mod n
6     up(full_buffs)
7   }
8 }
```

```
0 thread consumer {
1   while(1){
2     down(full_buffs)
3     c = buf[OutP]
4     OutP = OutP + 1 mod n
5     up(empty_buffs)
6     // Consume char...
7   }
8 }
```

Producer consumer problem

- ❑ What is the shared state in the last solution?
- ❑ Does it apply mutual exclusion? If so, how?




Problems with solution

- What if we have multiple producers and multiple consumers?
 - ❖ Producer-specific and consumer-specific data becomes shared
 - ❖ We need to define and protect critical sections

Dining philosophers problem

- ❑ Five philosophers sit at a table
- ❑ One fork/chopstick between each philosopher - need two to eat



Each philosopher is modeled with a thread

```
while(TRUE) {  
    Think();  
    Grab first fork;  
    Grab second fork;  
    Eat();  
    Put down first fork;  
    Put down second fork;  
}
```

- ❑ *Why do they need to synchronize?*
- ❑ *How should they do it?*

Is this a valid solution?


```
#define N 5

Philosopher() {
    while(TRUE) {
        Think();
        take_fork(i);
        take_fork((i+1)% N);
        Eat();
        put_fork(i);
        put_fork((i+1)% N);
    }
}
```


Problems

- Holding one fork while you wait for the other can lead to deadlock!
 - ❖ You should not hold on to a fork unless you can get both
 - ❖ Is there a deterministic, deadlock-free, starvation-free solution to doing this?

Working towards a solution ...



Working towards a solution ...

```
#define N 5

Philosopher() {
    while(TRUE) {
        Think();
        take_forks(i);
        Eat();
        put_forks(i);
    }
}
```

Picking up forks

```
int state[N]
semaphore mutex = 1
semaphore sem[i]
```

```
take_forks(int i) {
    wait(mutex);
    state[i] = HUNGRY;
    test(i);
    signal(mutex);
    wait(sem[i]);
}
```

```
// only called with mutex set!

test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        signal(sem[i]);
    }
}
```

Putting down forks

```
int state[N]
semaphore mutex = 1
semaphore sem[i]
```

```
put_forks(int i) {
    wait(mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    signal(mutex);
}
```


```
// only called with mutex set!

test(int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        signal(sem[i]);
    }
}
```

Dining philosophers

- ❑ Is the previous solution correct?
- ❑ What does it mean for it to be correct?
- ❑ How could you generate output to help detect common problems?
 - ❖ What would a race condition look like?
 - ❖ What would deadlock look like?
 - ❖ What would starvation look like?

The sleeping barber problem



The sleeping barber problem

□ *Barber:*

- ❖ While there are customers waiting for a hair cut put one in the barber chair and cut their hair
- ❖ When done move to the next customer else go to sleep, until a customer comes in

□ *Customer:*

- ❖ If barber is asleep wake him up for a haircut
- ❖ If someone is getting a haircut wait for the barber to become free by sitting in a chair
- ❖ If all chairs are all full, leave the barbershop

Designing a solution

- ❑ How will we model the barber(s) and customers?
- ❑ What state variables do we need?
 - ❖ .. and which ones are shared?
 - ❖ and how will we protect them?
- ❑ How will the barber sleep?
- ❑ How will the barber wake up?
- ❑ How will customers wait?
- ❑ How will they proceed?
- ❑ What problems do we need to look out for?

Is this a good solution?

```
const CHAIRS = 5
var customers: Semaphore
    barbers: Semaphore
    lock: Mutex
    numWaiting: int = 0
```

Barber Thread:

```
while true
    Wait(customers)
    Lock(lock)
    numWaiting = numWaiting-1
    Signal(barbers)
    Unlock(lock)
    CutHair()
endWhile
```

Customer Thread:

```
Lock(lock)
if numWaiting < CHAIRS
    numWaiting = numWaiting+1
    Signal(customers)
    Unlock(lock)
    Wait(barbers)
    GetHaircut()
else -- give up & go home
    Unlock(lock)
endIf
```

The readers and writers problem

- ❑ Multiple readers and writers want to access a database (each one is a thread)
- ❑ Multiple readers can proceed concurrently
 - ❖ *No race condition if nobody is modifying data*
- ❑ Writers must synchronize with readers and other writers
 - ❖ *only one writer at a time !*
 - ❖ *when someone is writing, there must be no readers !*

Goals:

- ❖ Maximize concurrency.
- ❖ Prevent starvation.

Designing a solution

- ❑ How will we model the readers and writers?
- ❑ What state variables do we need?
 - ❖ .. and which ones are shared?
 - ❖ and how will we protect them?
- ❑ How will the writers wait?
- ❑ How will the writers wake up?
- ❑ How will readers wait?
- ❑ How will the readers wake up?
- ❑ What problems do we need to look out for?

Is this a valid solution to readers & writers?

```
var mut: Mutex = unlocked
    db: Semaphore = 1
    rc: int = 0
```

Writer Thread:

```
while true
    ...Remainder Section...
    Wait(db)
    ...Write shared data...
    Signal(db)
endWhile
```

Reader Thread:

```
while true
    Lock(mut)
    rc = rc + 1
    if rc == 1
        Wait(db)
    endIf
    Unlock(mut)
    ... Read shared data...
    Lock(mut)
    rc = rc - 1
    if rc == 0
        Signal(db)
    endIf
    Unlock(mut)
    ... Remainder Section...
endWhile
```

Readers and writers solution

- Does the previous solution have any problems?
 - ❖ is it "fair"?
 - ❖ can any threads be starved? If so, how could this be fixed?

Monitors

Monitors

- ❑ **It is difficult to produce correct programs using semaphores**
 - ❖ correct ordering of wait and signal is tricky!
 - ❖ avoiding race conditions and deadlock is tricky!
 - ❖ boundary conditions are tricky!
- ❑ **Can we get the compiler to generate the correct semaphore code for us?**
 - ❖ what are suitable higher level abstractions for synchronization?

Monitors

- ❑ Related shared objects are collected together
- ❑ Compiler enforces encapsulation/mutual exclusion
 - ❖ Encapsulation:
 - Local data variables are accessible only via the monitor's entry procedures (like methods)
 - ❖ Mutual exclusion
 - A monitor has an associated mutex lock
 - Threads must acquire the monitor's mutex lock before invoking one of its procedures

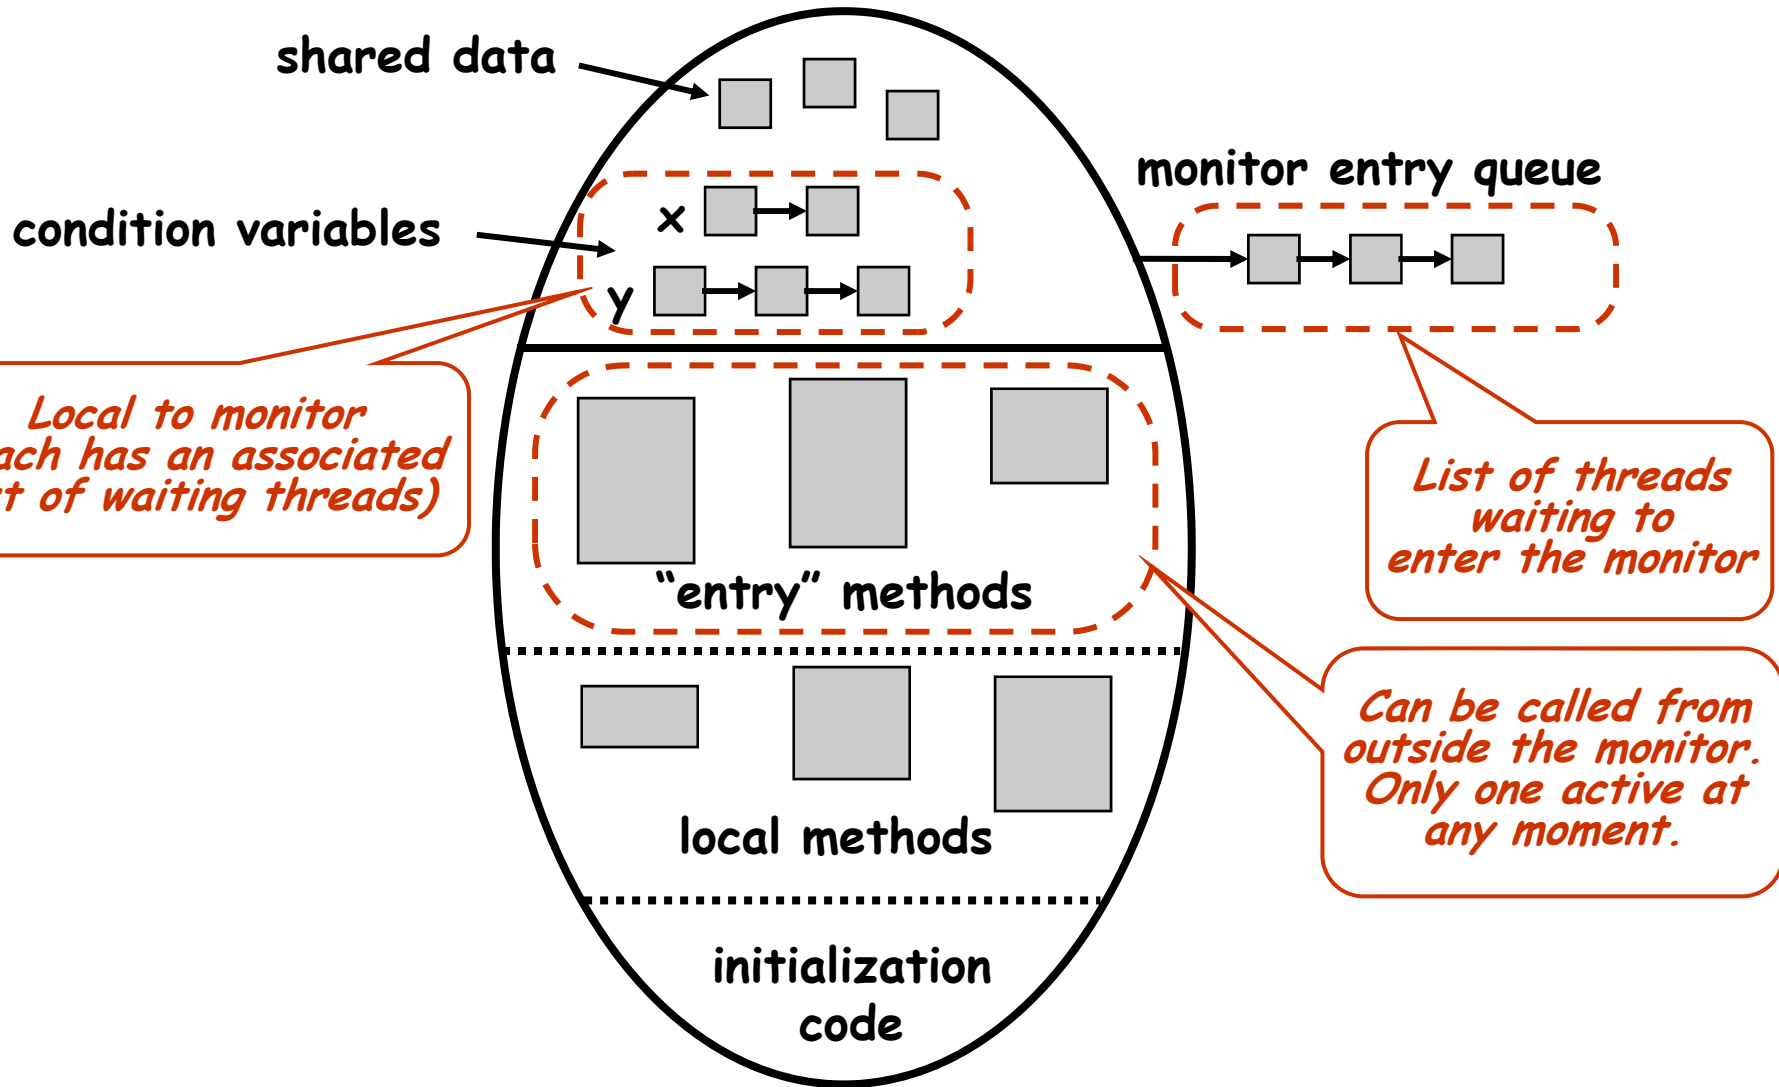
Monitors and condition variables

- But we need two flavors of synchronization
 - ❖ Mutual exclusion
 - Only one at a time in the critical section
 - Handled by the monitor's mutex
 - ❖ Condition synchronization
 - Wait until a certain condition holds
 - Signal waiting threads when the condition holds

Monitors and condition variables

- Condition variables (cv) for use within monitors
 - ❖ `cv.wait(mon-mutex)`
 - thread blocked (queued) until condition holds
 - Must not block while holding mutex!
 - monitor mutex must be released!
 - ❖ `cv.signal()`
 - signals the condition and unblocks (dequeues) a thread

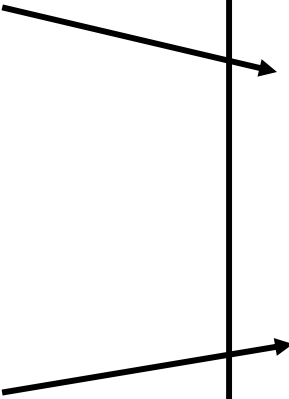
Monitor structures



Monitor example for mutual exclusion

```
process Producer
begin
  loop
    <produce char "c">
    BoundedBuffer.deposit(c)
  end loop
end ProducerL
```

```
process Consumer
begin
  loop
    BoundedBuffer.remove(c)
    <consume char "c">
  end loop
end Consumer
```



```
monitor: BoundedBuffer
var  buffer : ...;
    nextIn, nextOut :... ;

entry deposit(c: char)
begin
  ...
end

entry remove(var c: char)
begin
  ...
end

end BoundedBuffer
```

Observations

- ❑ That's much simpler than the semaphore-based solution to producer/consumer (bounded buffer)!
- ❑ ... but where is the mutex?
- ❑ ... and what do the bodies of the monitor procedures look like?

Monitor example with condition variables

```
monitor : BoundedBuffer
var buffer          : array[0..n-1] of char
    nextIn,nextOut   : 0..n-1 := 0
    fullCount        : 0..n    := 0
    notEmpty, notFull : condition
```

```
entry deposit(c:char)
begin
    if (fullCount = n) then
        wait(notFull)
    end if

    buffer[nextIn] := c
    nextIn := nextIn+1 mod n
    fullCount := fullCount+1

    signal(notEmpty)
end deposit
```

```
entry remove(var c: char)
begin
    if (fullCount = n) then
        wait(notEmpty)
    end if

    c := buffer[nextOut]
    nextOut := nextOut+1 mod n
    fullCount := fullCount-1

    signal(notFull)
end remove
```

```
end BoundedBuffer
```

Condition variables

“Condition variables allow processes to synchronize based on some state of the monitor variables.”

Condition variables in producer/consumer

“NotFull” condition

“NotEmpty” condition

- Operations **Wait()** and **Signal()** allow synchronization within the monitor
- When a producer thread adds an element...
 - ❖ A consumer may be sleeping
 - ❖ Need to wake the consumer... **Signal**

Condition synchronization semantics

- ❑ *"Only one thread can be executing in the monitor at any one time."*
- ❑ **Scenario:**
 - ❖ Thread A is executing in the monitor
 - ❖ Thread A does a **signal** waking up thread B
 - ❖ What happens now?
 - ❖ Signaling and signaled threads can not both run!
 - ❖ ... so which one runs, which one blocks, and on what queue?

Monitor design choices

- Condition variables introduce a problem for mutual exclusion
 - ❖ only one process active in the monitor at a time, so what to do when a process is unblocked on **signal**?
 - ❖ must not block holding the mutex, so what to do when a process blocks on **wait**?

Monitor design choices

- **Choices when A signals a condition that unblocks B**
 - ❖ A waits for B to exit the monitor or block again
 - ❖ B waits for A to exit the monitor or block
 - ❖ Signal causes A to immediately exit the monitor or block (... but awaiting what condition?)
- **Choices when A signals a condition that unblocks B & C**
 - ❖ B is unblocked, but C remains blocked
 - ❖ C is unblocked, but B remains blocked
 - ❖ Both B & C are unblocked ... and compete for the mutex?
- **Choices when A calls wait and blocks**
 - ❖ a new external process is allowed to enter
 - ❖ but which one?

Option 1: Hoare semantics

- **What happens when a Signal is performed?**
 - ❖ signaling thread (A) is suspended
 - ❖ signaled thread (B) wakes up and runs immediately
- **Result:**
 - ❖ B can assume the condition is now true/satisfied
 - ❖ Hoare semantics give strong guarantees
 - ❖ Easier to prove correctness
- **When B leaves monitor, A can run.**
 - A might resume execution immediately
 - ... or maybe another thread (C) will slip in!

Option 2: MESA Semantics (Xerox PARC)

- **What happens when a Signal is performed?**
 - ❖ the signaling thread (A) continues.
 - ❖ the signaled thread (B) waits.
 - ❖ when A leaves monitor, then B runs.
- **Issue: What happens while B is waiting?**
 - ❖ can the condition that caused A to generate the signal be changed before B runs?
- **In MESA semantics a signal is more like a hint**
 - ❖ Requires B to recheck the condition on which it waited to see if it can proceed or must wait some more

Code for the “deposit” entry routine

```
monitor BoundedBuffer
  var buffer: array[n] of char
    nextIn, nextOut: int = 0
    cntFull: int = 0
    notEmpty: Condition
    notFull: Condition

  entry deposit(c: char)
    if cntFull == N
      notFull.Wait()
    endIf
    buffer[nextIn] = c
    nextIn = (nextIn+1) mod N
    cntFull = cntFull + 1
    notEmpty.Signal()
  endEntry

  entry remove()
    ...

endMonitor
```

} Hoare Semantics

Code for the “deposit” entry routine

```
monitor BoundedBuffer
  var buffer: array[n] of char
    nextIn, nextOut: int = 0
    cntFull: int = 0
    notEmpty: Condition
    notFull: Condition

  entry deposit(c: char)
    while cntFull == N
      notFull.Wait()
    endWhile
    buffer[nextIn] = c
    nextIn = (nextIn+1) mod N
    cntFull = cntFull + 1
    notEmpty.Signal()
  endEntry

  entry remove()
    ...

endMonitor
```

} MESA Semantics

Code for the "remove" entry routine

```
monitor BoundedBuffer
  var buffer: array[n] of char
    nextIn, nextOut: int = 0
    cntFull: int = 0
    notEmpty: Condition
    notFull: Condition

  entry deposit(c: char)
    ...

  entry remove()
    if cntFull == 0
      notEmpty.Wait()
    endIf
    c = buffer[nextOut]
    nextOut = (nextOut+1) mod N
    cntFull = cntFull - 1
    notFull.Signal()
  endEntry

endMonitor
```

} Hoare Semantics

Code for the "remove" entry routine

```
monitor BoundedBuffer
  var buffer: array[n] of char
    nextIn, nextOut: int = 0
    cntFull: int = 0
    notEmpty: Condition
    notFull: Condition

  entry deposit(c: char)
    ...

  entry remove()
    while cntFull == 0
      notEmpty.Wait()
    endWhile
    c = buffer[nextOut]
    nextOut = (nextOut+1) mod N
    cntFull = cntFull - 1
    notFull.Signal()
  endEntry

endMonitor
```

} MESA Semantics

"Hoare Semantics"

What happens when a Signal is performed?

The signaling thread (A) is suspended.

The signaled thread (B) wakes up and runs immediately.

B can assume the condition is now true/satisfied

From the original Hoare Paper:

"No other thread can intervene [and enter the monitor] between the signal and the continuation of exactly one waiting thread."

"If more than one thread is waiting on a condition, we postulate that the signal operation will reactivate the longest waiting thread. This gives a simple neutral queuing discipline which ensures that every waiting thread will eventually get its turn."

Implementing Hoare Semantics

- ❑ Thread A holds the monitor lock
- ❑ Thread A **signals** a condition that thread B was waiting on
- ❑ Thread B is moved back to the ready queue?
 - ❖ B should run immediately
 - ❖ Thread A must be suspended...
 - ❖ the monitor lock must be passed from A to B
- ❑ When B finishes it releases the monitor lock
- ❑ Thread A must re-acquire the lock
 - ❖ A is blocked, waiting to re-aquire the lock

Implementing Hoare Semantics

- **Problem:**
 - ❖ Possession of the monitor lock must be passed **directly** from A to B and then eventually back to A

Implementing Hoare Semantics

□ Implementation Ideas:

- ❖ Consider a signaled thread like B to be “urgent” after A releases the monitor lock
 - Thread C trying to gain initial entry to the monitor is not “urgent”
- ❖ Consider two wait lists associated with each *MonitorLock* (so now this is not exactly a mutex)
 - UrgentlyWaitingThreads
 - NonurgentlyWaitingThreads
- ❖ Want to wake up urgent threads first, if any
- ❖ Alternatively, B could be added to the front of the monitor lock queue

Implementing Hoare Semantics

- **Recommendation for Project 4 implementation:**
 - ❖ Do not modify the mutex methods provided, because future code will use them
 - ❖ Create new classes:
 - **MonitorLock** -- similar to Mutex
 - **HoareCondition** -- similar to Condition

Brinch-Hansen Semantics

□ Hoare Semantics

- ❖ On signal, allow signaled process to run
- ❖ Upon its exit from the monitor, signaling process continues.

□ Brinch-Hansen Semantics

- ❖ Signaler must immediately exit following any invocation of signal
- ❖ Restricts the kind of solutions that can be written
- ❖ ... but monitor implementation is easier

Review of a Practical Concurrent Programming Issue - Reentrant Functions

Reentrant code

- A function/method is said to be **reentrant** if...

A function that has been invoked may be invoked again before the first invocation has returned, and will still work correctly

- In the context of concurrent programming...

A reentrant function can be executed simultaneously by more than one thread, with no ill effects

Reentrant Code

- Consider this function...

```
var count: int = 0  
function GetUnique () returns int  
    count = count + 1  
    return count  
endFunction
```

- What if it is executed by different threads concurrently?

Reentrant Code

- Consider this function...

```
var count: int = 0  
function GetUnique () returns int  
    count = count + 1  
    return count  
endFunction
```

- What if it is executed by different threads concurrently?
 - ❖ The results may be incorrect!
 - ❖ This routine is not reentrant!

When is code reentrant?

- ❑ **Some variables are**
 - ❖ "local" -- to the function/method/routine
 - ❖ "global" -- sometimes called "static"
- ❑ **Access to local variables?**
 - ❖ A new stack frame is created for each invocation
 - ❖ Each thread has its own stack
- ❑ **What about access to global variables?**
 - ❖ **Must use synchronization!**

Does this work?

```
var count: int = 0
```

```
    myLock: Mutex
```

```
function GetUnique ( ) returns int
```

```
    myLock.Lock( )
```

```
    count = count + 1
```

```
    myLock.Unlock( )
```

```
    return count
```

```
endFunction
```

What about this?

```
var count: int = 0
```

```
    myLock: Mutex
```

```
function GetUnique ( ) returns int
```

```
    myLock.Lock( )
```

```
    count = count + 1
```

```
    return count
```

```
    myLock.Unlock( )
```

```
endFunction
```

Making this function reentrant

```
var count: int = 0
```

```
    myLock: Mutex
```

```
function GetUnique () returns int
```

```
    var i: int
```

```
    myLock.Lock()
```

```
    count = count + 1
```

```
    i = count
```

```
    myLock.Unlock()
```

```
    return i
```

```
endFunction
```


Message Passing

Message Passing

- ❑ **Interprocess Communication**
 - ❖ via shared memory
 - ❖ across machine boundaries
- ❑ **Message passing can be used for synchronization or general communication**
- ❑ **Processes use *send* and *receive* primitives**
 - ❖ *receive* can block (like *waiting* on a Semaphore)
 - ❖ *send* unblocks a process blocked on *receive* (just as a *signal* unblocks a *waiting* process)

Producer-consumer with message passing

- The basic idea:
 - ❖ After producing, the producer sends the data to consumer in a message
 - ❖ The system buffers messages
 - The producer can out-run the consumer
 - The messages will be kept in order
 - ❖ But how does the producer avoid overflowing the buffer?
 - After consuming the data, the consumer sends back an “empty” message
 - ❖ A fixed number of messages ($N=100$)
 - ❖ The messages circulate back and forth.

Producer-consumer with message passing

```
const N = 100                -- Size of message buffer
var em: char
for i = 1 to N                -- Get things started by
    Send (producer, &em)     --      sending N empty messages
endFor
```

```
thread consumer
    var c, em: char
    while true
        Receive(producer, &c) -- Wait for a char
        Send(producer, &em)   -- Send empty message back
        // Consume char...
    endwhile
end
```

Producer-consumer with message passing

```
thread producer
  var c, em: char
  while true
    // Produce char c...
    Receive(consumer, &em)    -- Wait for an empty msg
    Send(consumer, &c)         -- Send c to consumer
  endwhile
end
```

Design choices for message passing

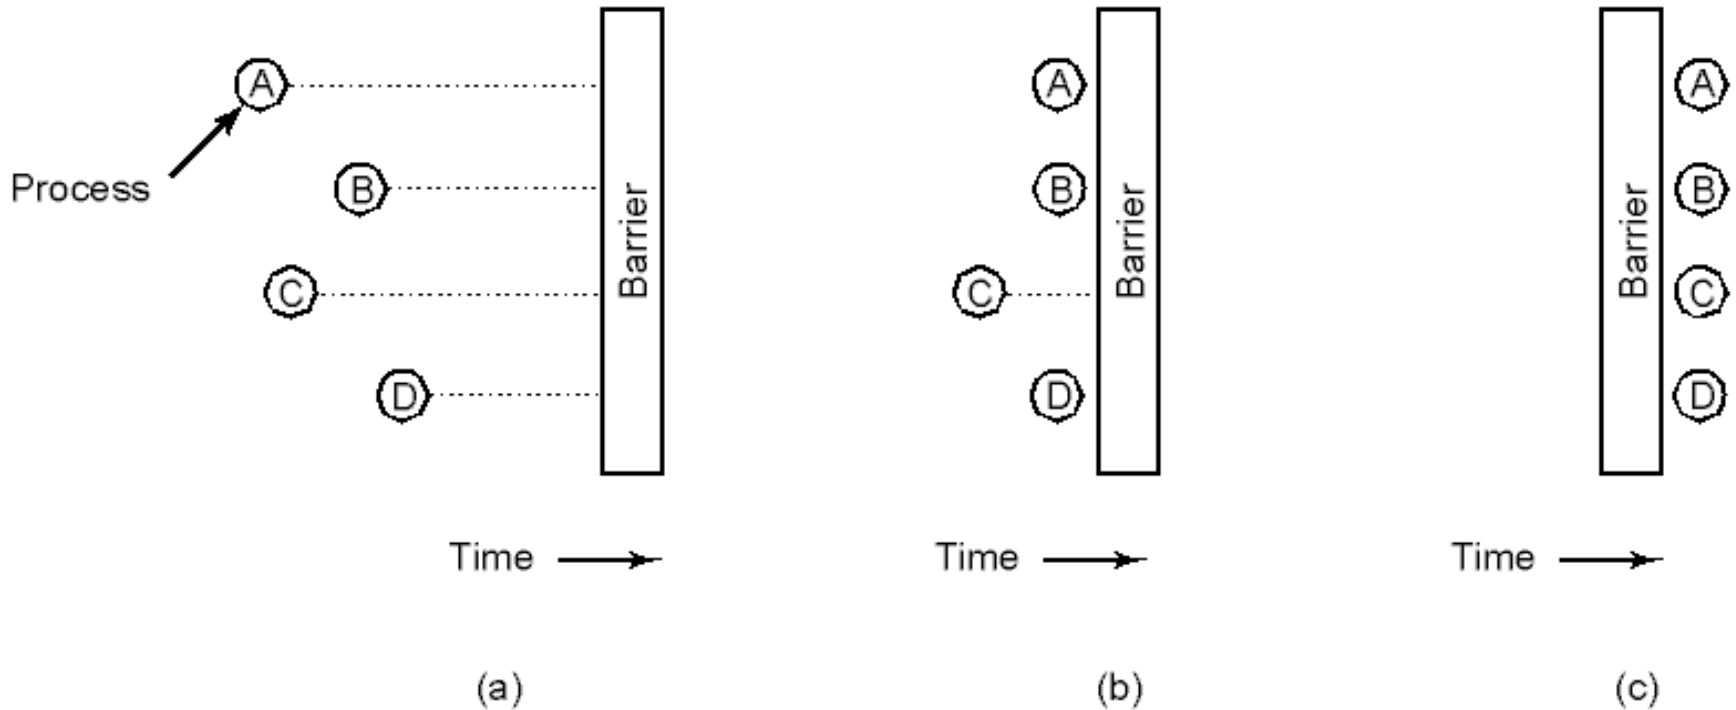
❑ Option 1: Mailboxes

- ❖ System maintains a buffer of sent, but not yet received, messages
- ❖ Must specify the size of the mailbox ahead of time
- ❖ Sender will be blocked if the buffer is full
- ❖ Receiver will be blocked if the buffer is empty

Design choices for message passing

- ❑ Option 2: No buffering
 - ❖ If Send happens first, the sending thread blocks
 - ❖ If Receiver happens first, the receiving thread blocks
 - ❖ Sender and receiver must Rendezvous (ie. meet)
 - ❖ Both threads are ready for the transfer
 - ❖ The data is copied / transmitted
 - ❖ Both threads are then allowed to proceed

Barriers



- ❖ Processes approaching a barrier
- ❖ All processes but one blocked at barrier
- ❖ Last process arrives; all are let through

Quiz

- ❑ What is the difference between a monitor and a semaphore?
 - ❖ Why might you prefer one over the other?
- ❑ How do the wait/signal methods of a condition variable differ from the wait/signal methods of a semaphore?
- ❑ What is the difference between Hoare and Mesa semantics for condition variables?
 - ❖ What implications does this difference have for code surrounding a wait() call?