

CS 333
Introduction to Operating Systems
Class 9 - Memory Management

Jonathan Walpole
Computer Science
Portland State University

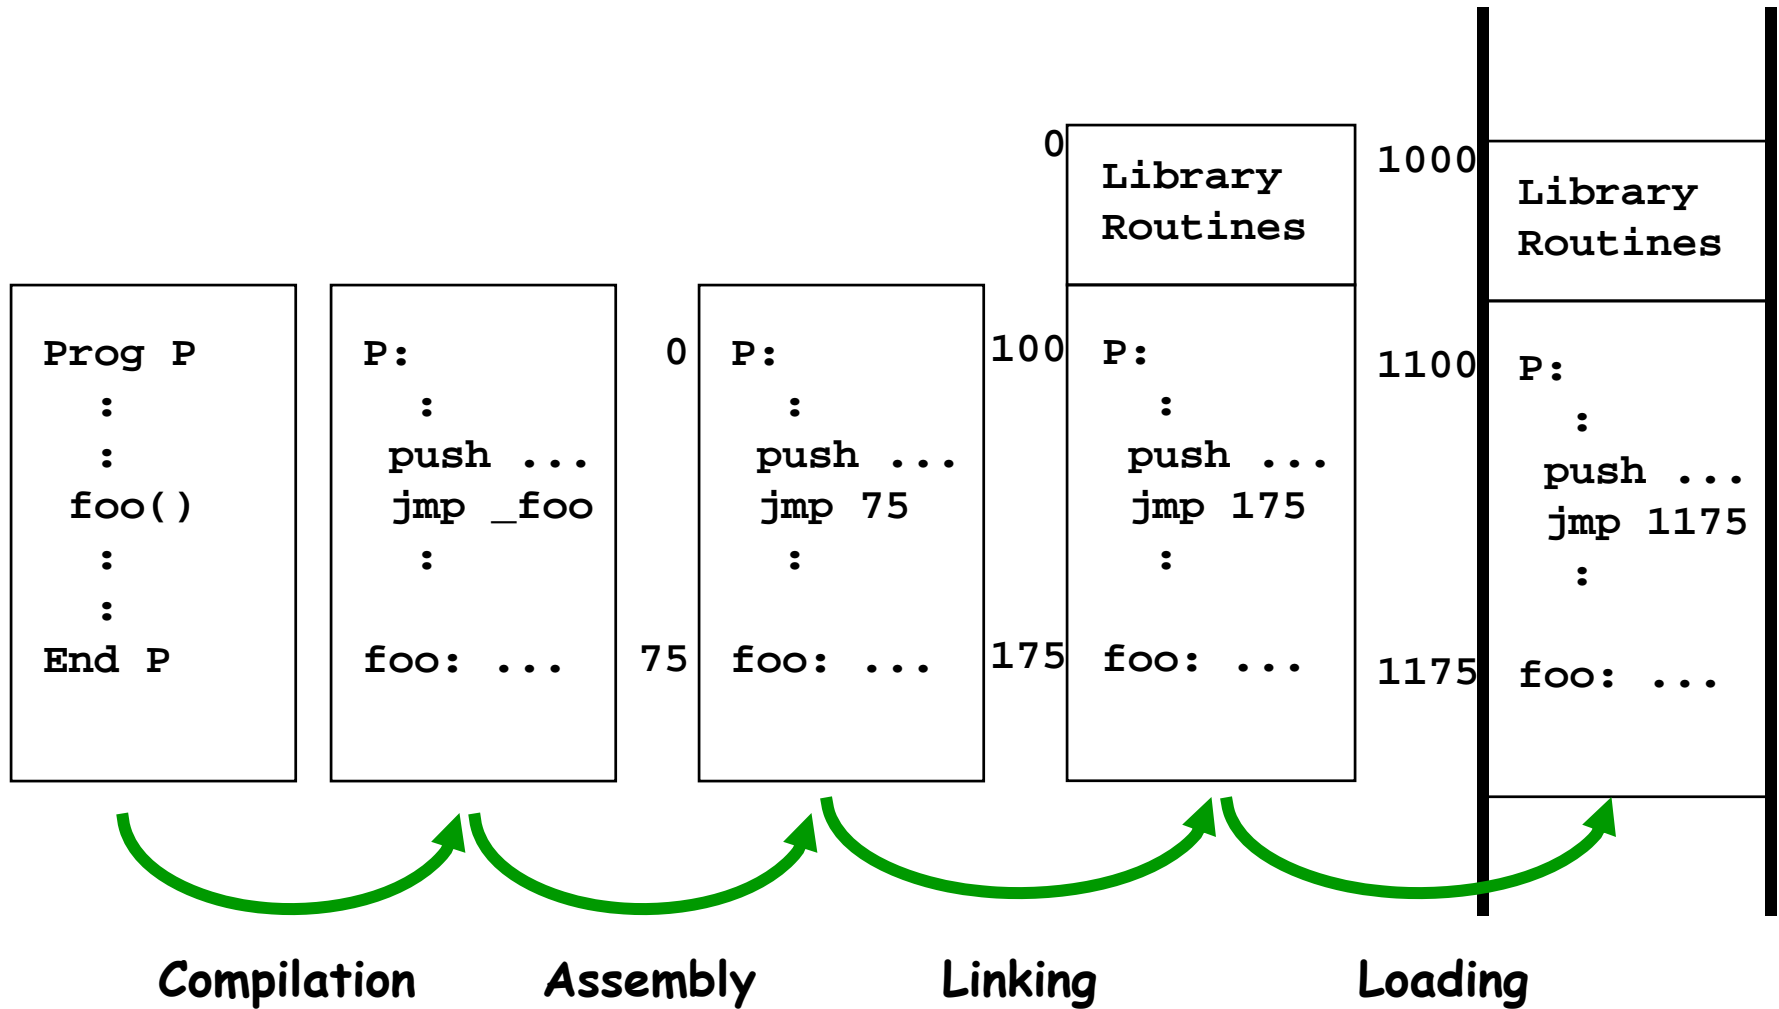
Memory management

- **Memory** - a linear array of bytes
 - ❖ Holds O.S. and programs (processes)
 - ❖ Each cell (byte) is named by a unique memory address
- **Recall**, processes are defined by an *address space*, consisting of text, data, and stack regions
- **Process execution**
 - ❖ CPU fetches instructions from the text region according to the value of the program counter (PC)
 - ❖ Each instruction may request additional operands from the data or stack region

Addressing memory

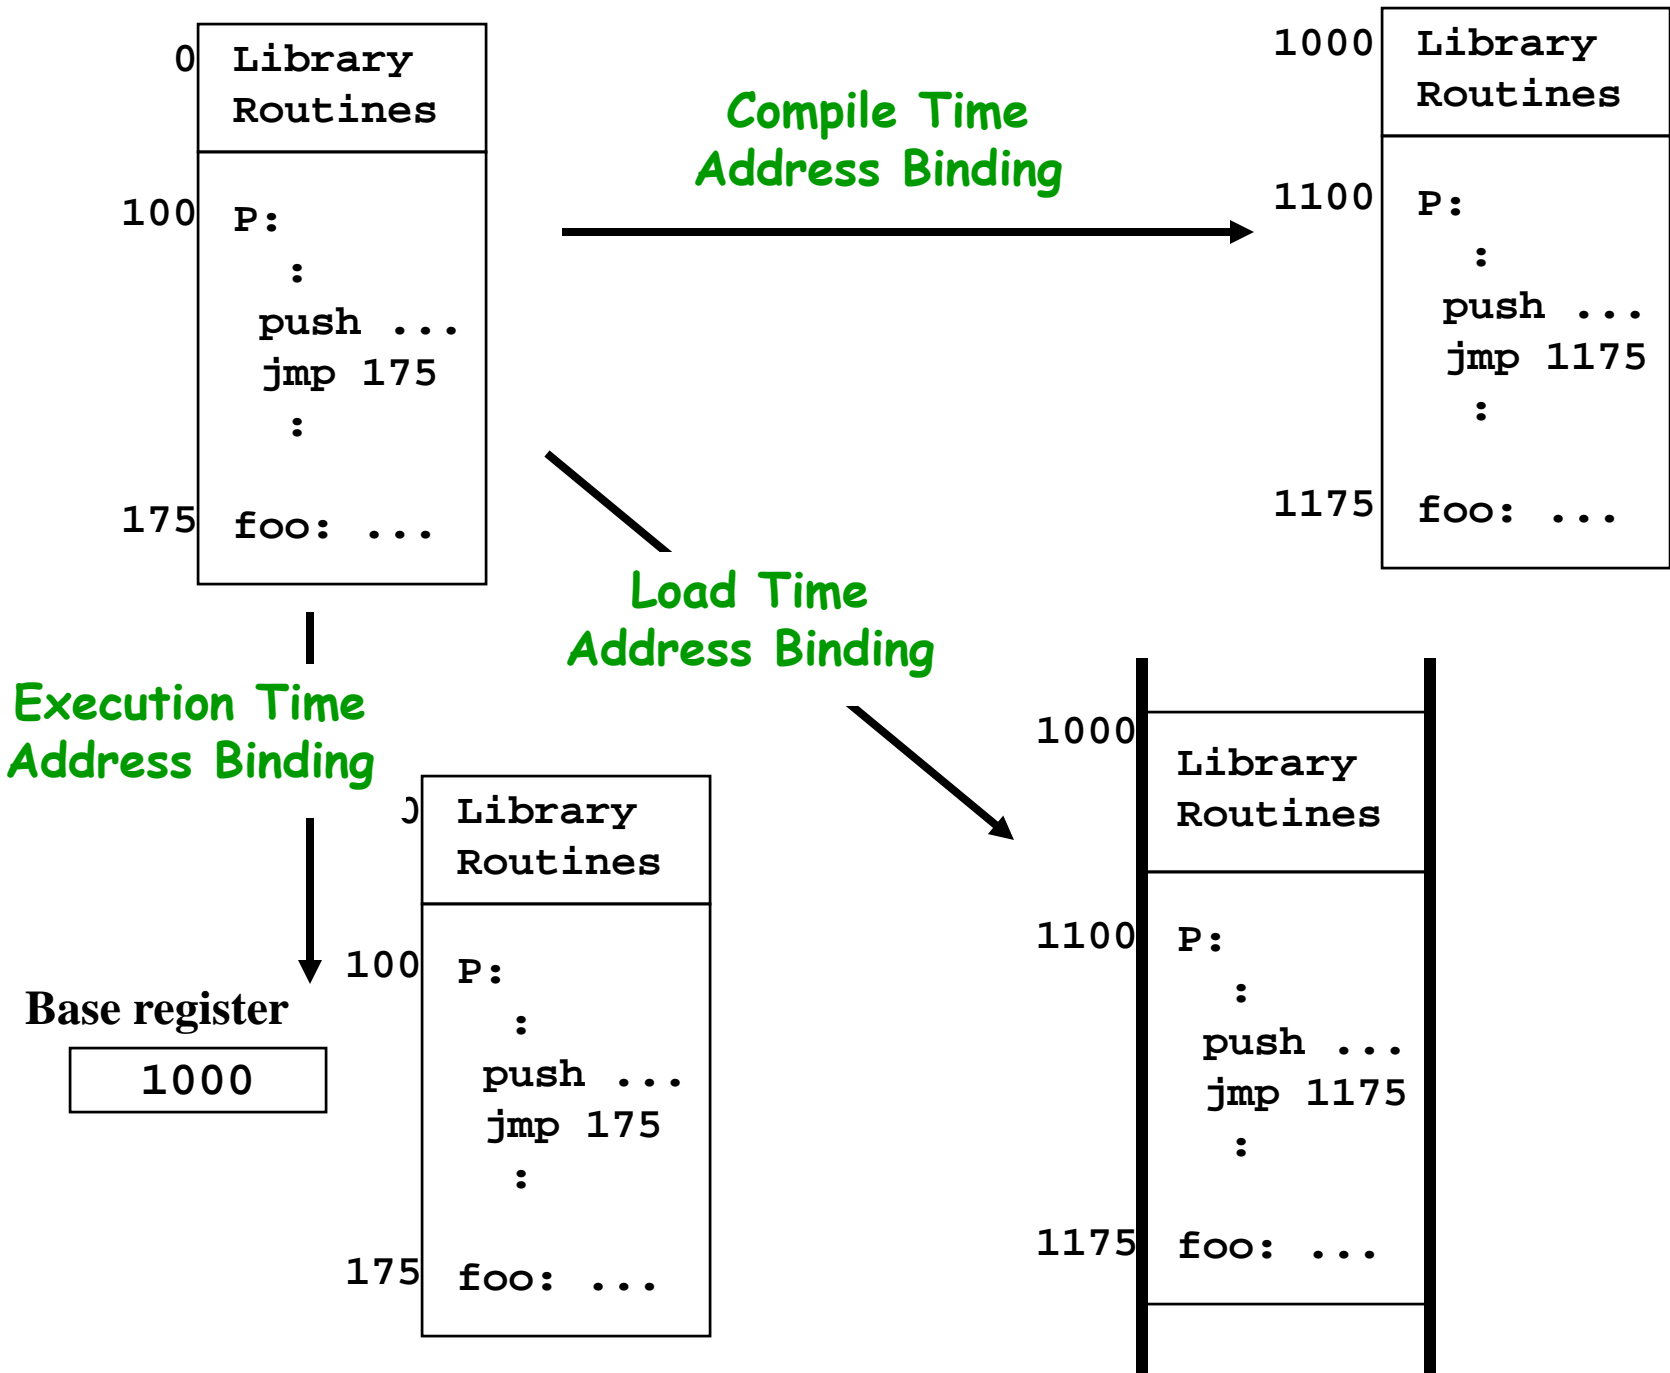
- ❑ Cannot know ahead of time where in memory a program will be loaded!
- ❑ Compiler produces code containing embedded addresses
 - ❖ these addresses can't be absolute (physical addresses)
- ❑ Linker combines pieces of the program
 - ❖ Assumes the program will be loaded at address 0
- ❑ We need to **bind** the compiler/linker generated addresses to the actual memory locations

Relocatable address generation



Address binding

- **Address binding**
 - ❖ fixing a physical address to the logical address of a process' address space
- ***Compile time binding***
 - ❖ if program location is fixed and known ahead of time
- ***Load time binding***
 - ❖ if program location in memory is unknown until run-time AND location is fixed
- ***Execution time binding***
 - ❖ if processes can be moved in memory during execution
 - ❖ Requires hardware support!

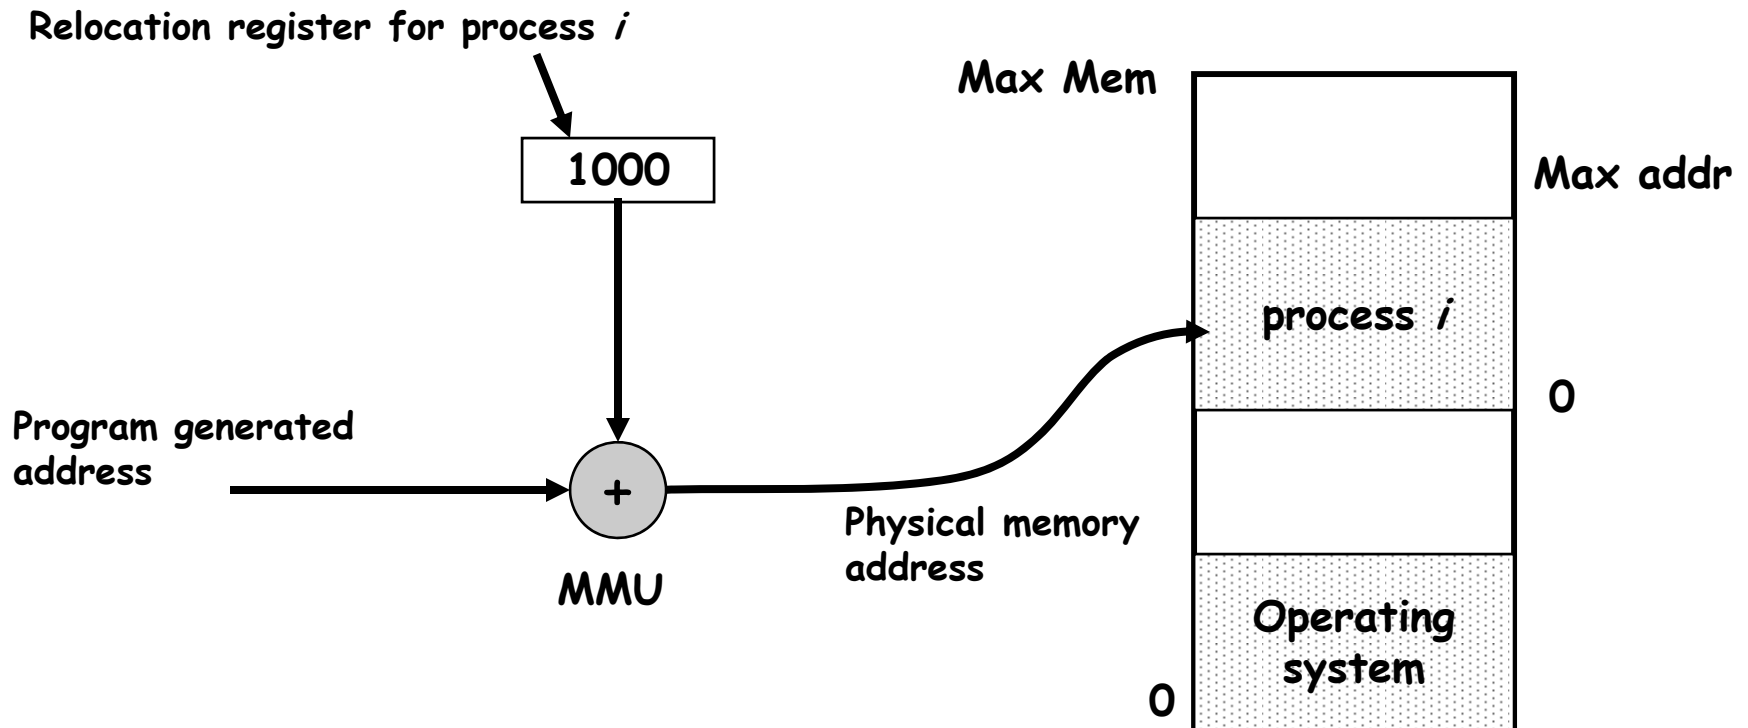


Runtime binding - base & limit registers

- Simple runtime relocation scheme
 - ❖ Use 2 registers to describe a partition
- For every address generated, at runtime...
 - ❖ Compare to the **limit** register (& abort if larger)
 - ❖ Add to the **base** register to give **physical** memory address

Dynamic relocation with a base register

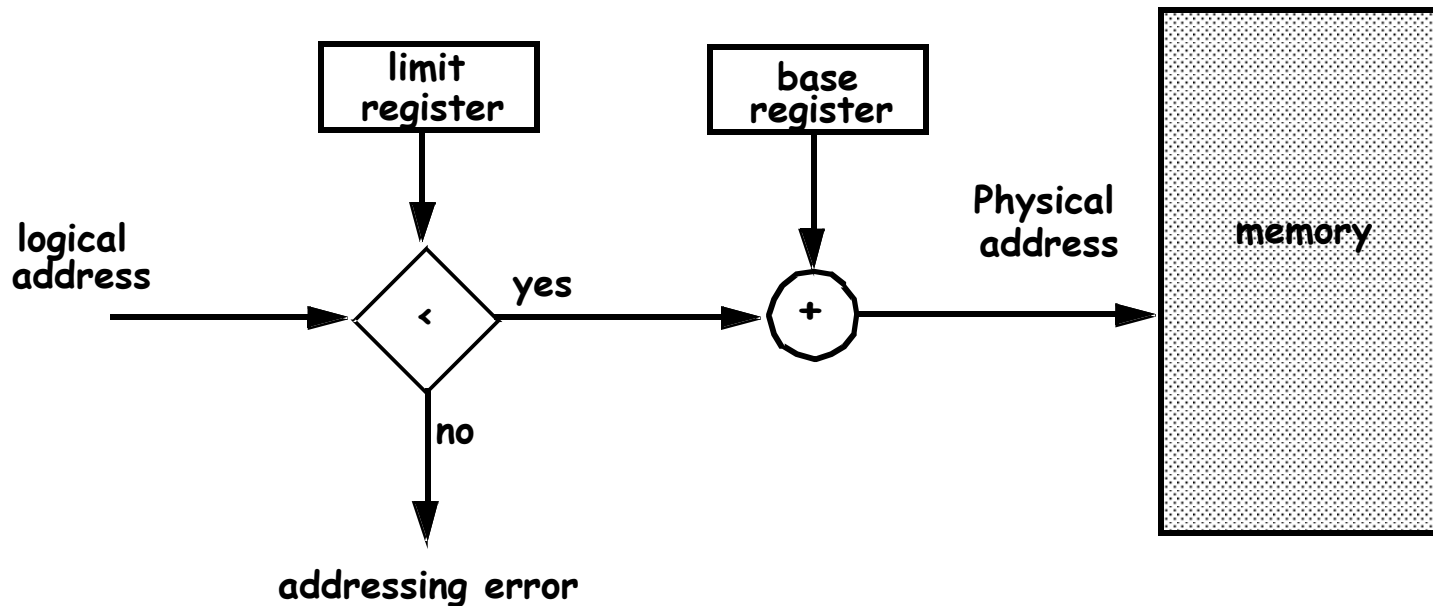
- ❑ **Memory Management Unit (MMU)** - dynamically converts logical addresses into physical address
- ❑ **MMU** contains base address register for running process



Protection using base & limit registers

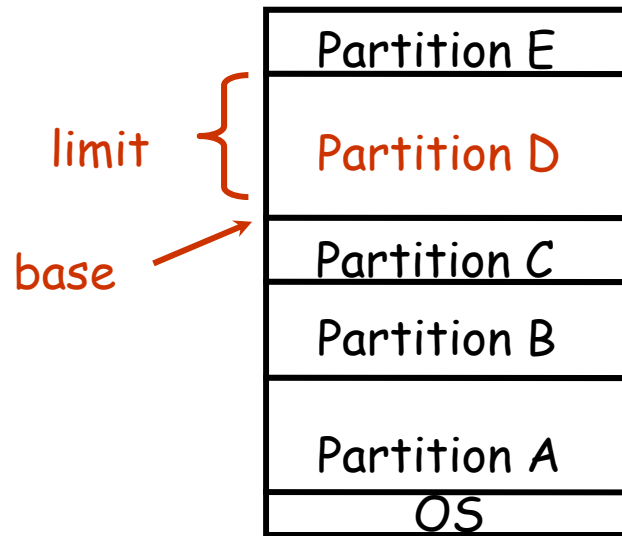
□ Memory protection

- ❖ **Base** register gives starting address for process
- ❖ **Limit** register limits the offset accessible from the relocation register



Multiprogramming with base and limit registers

- ❑ Multiprogramming: a separate partition per process
- ❑ What happens on a context switch?
 - ❖ Store process A's **base** and **limit** register values
 - ❖ Load new values into **base** and **limit** registers for process B



Swapping

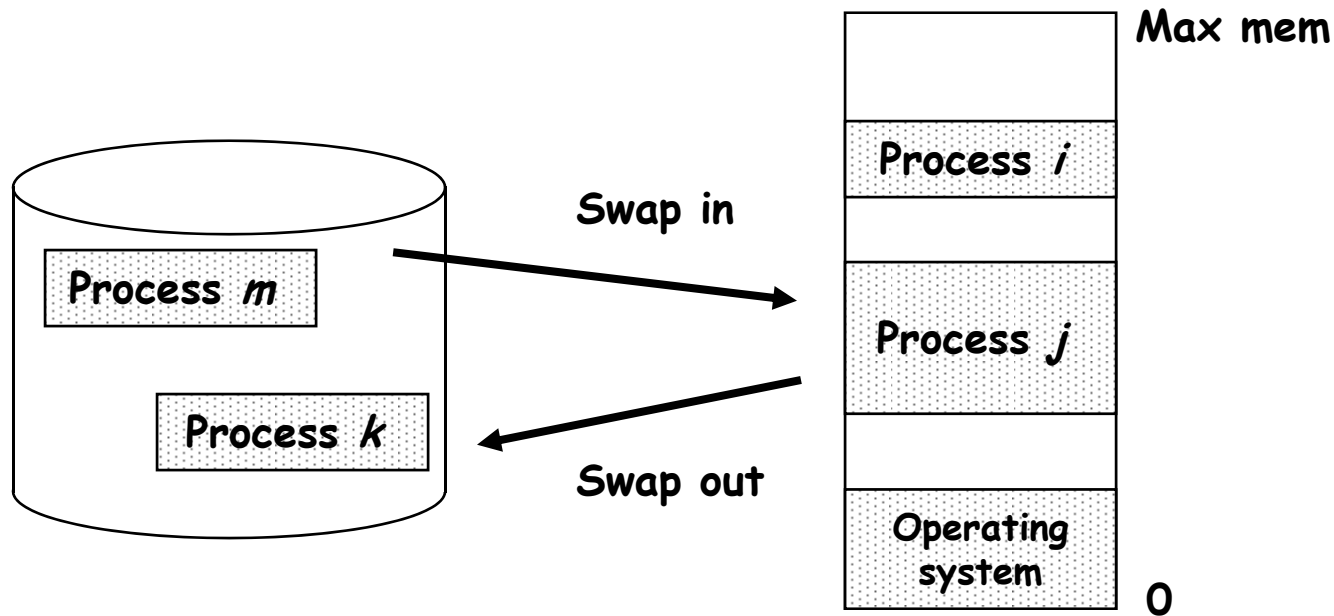
- **When a program is running...**
 - ❖ The entire program must be in memory
 - ❖ Each program is put into a single **partition**

- **When the program is not running...**
 - ❖ May remain resident in memory
 - ❖ May get "*swapped*" out to disk


- **Over time...**
 - ❖ Programs come into memory when they get swapped in
 - ❖ Programs leave memory when they get swapped out


Basics - swapping


- **Benefits of swapping:**
 - ❖ Allows multiple programs to be run concurrently
 - ❖ ... more than will fit in memory at once





Swapping can lead to fragmentation



























Dealing with fragmentation


- *Compaction* - from time to time shift processes around to collect all free space into one contiguous block
 - ❖ Memory to memory copying overhead
 - memory to disk to memory for compaction via swapping



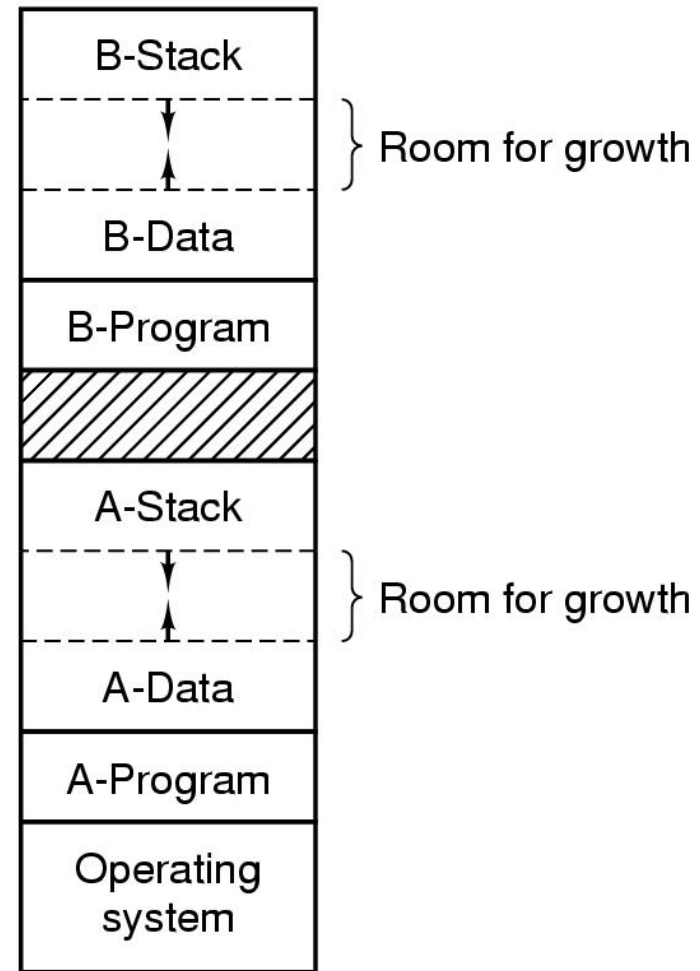
How big should partitions be?

- **Programs may want to grow during execution**
 - ❖ More room for stack, heap allocation, etc
- **Problem:**
 - ❖ If the partition is too small programs must be moved
 - ❖ Requires copying overhead
 - ❖ Why not make the partitions a little larger than necessary to accommodate "some" cheap growth?

Allocating extra space within partitions



(a)



(b)


Managing memory

- Each chunk of memory is either
 - ❖ Used by some process or unused ("free")
- Operations
 - ❖ **Allocate** a chunk of unused memory big enough to hold a new process
 - ❖ **Free** a chunk of memory by returning it to the **free pool** after a process terminates or is swapped out

Managing memory with bit maps

- ❑ **Problem - how to keep track of used and unused memory?**
- ❑ **Technique 1 - Bit Maps**
 - ❖ A long bit string
 - ❖ One bit for every chunk of memory
 - 1 = in use
 - 0 = free
 - ❖ Size of allocation unit influences space required
 - **Example: unit size = 32 bits**
 - overhead for bit map: $1/33 = 3\%$
 - **Example: unit size = 4Kbytes**
 - overhead for bit map: $1/32,769$

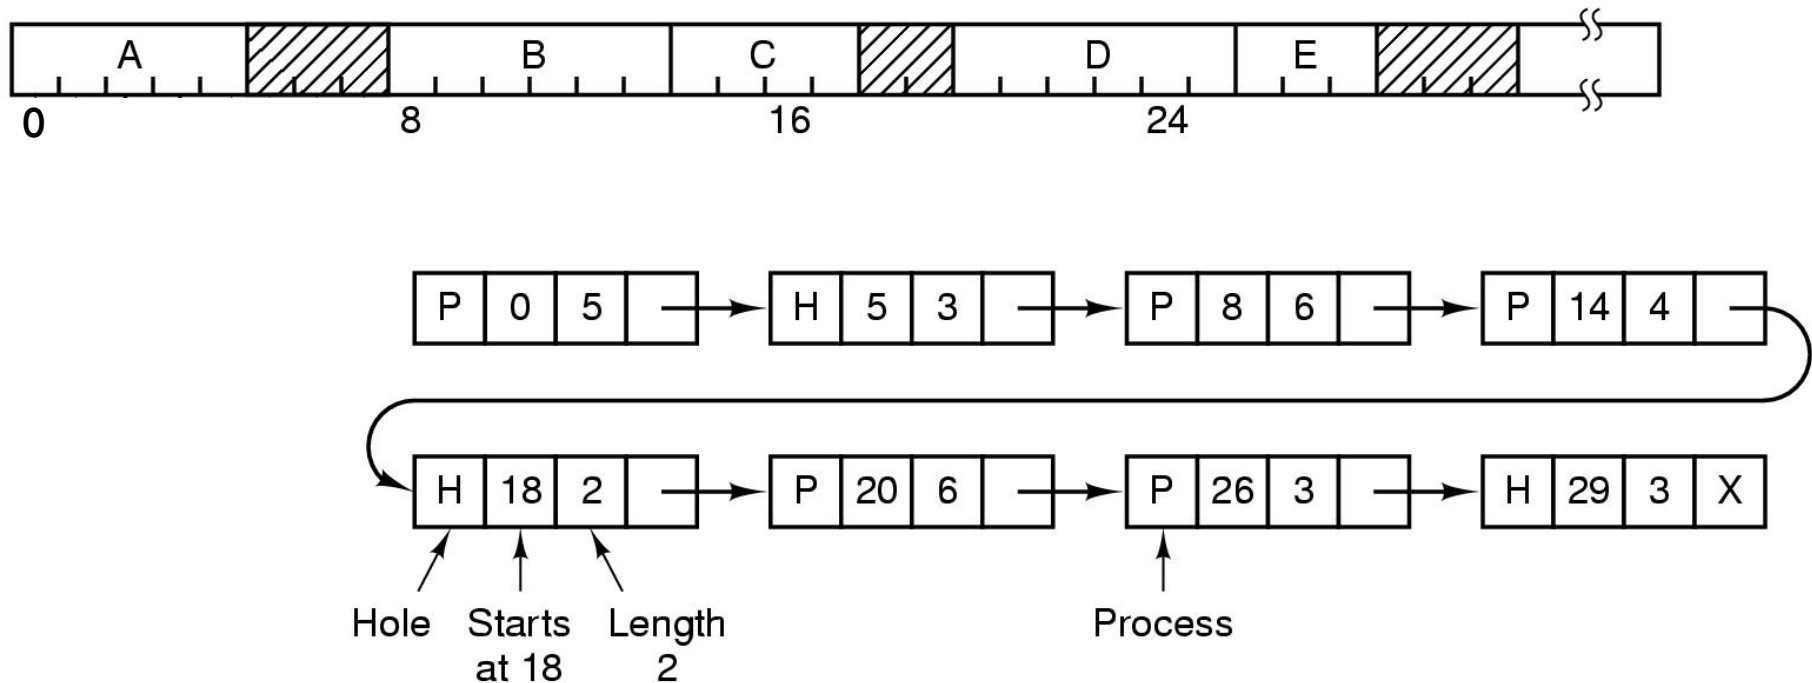
Managing memory with bit maps



Managing memory with linked lists

- ❑ **Technique 2 - Linked List**
- ❑ **Keep a list of elements**
- ❑ **Each element describes one unit of memory**
 - ❖ Free / in-use Bit ("P=process, H=hole")
 - ❖ Starting address
 - ❖ Length
 - ❖ Pointer to next element

Managing memory with linked lists




Merging holes

- Whenever a unit of memory is freed we want to merge adjacent holes!


Merging holes

Before X terminates




becomes

After X terminates





Merging holes

Before X terminates




becomes

After X terminates






becomes




Merging holes

Before X terminates




becomes


After X terminates



becomes







becomes




Merging holes

Before X terminates




becomes


After X terminates




becomes



becomes



becomes



Managing memory with linked lists

- Searching the list for space for a new process
 - ❖ First Fit
 - ❖ Next Fit
 - Start from current location in the list
 - ❖ Best Fit
 - Find the smallest hole that will work
 - Tends to create lots of really small holes
 - ❖ Worst Fit
 - Find the largest hole
 - Remainder will be big
 - ❖ Quick Fit
 - Keep separate lists for common sizes

Fragmentation

- ❑ Memory is divided into partitions
- ❑ Each partition has a different size
- ❑ Processes are allocated space and later freed
- ❑ After a while memory will be full of small holes!
 - ❖ No free space large enough for a new process even though there is enough free memory in total
- ❑ If we allow free space within a partition we have internal fragmentation
- ❑ Fragmentation:
 - ❖ External fragmentation = unused space between partitions
 - ❖ Internal fragmentation = unused space within partitions

Solution to fragmentation?

- ❑ **Compaction requires high copying overhead**
- ❑ **Why not allocate memory in non-contiguous equal fixed size units?**
 - ❖ no external fragmentation!
 - ❖ internal fragmentation < 1 unit per process
- ❑ **How big should the units be?**
 - ❖ The smaller the better for internal fragmentation
 - ❖ The larger the better for management overhead
- ❑ **The key challenge for this approach:**


"How can we do secure dynamic address translation?"

Using pages for non-contiguous allocation

- **Memory divided into fixed size page frames**
 - ❖ Page frame size = 2^n bytes
 - ❖ Lowest n bits of an address specify byte offset in a page
- **But how do we associate page frames with processes?**
 - ❖ And how do we map memory addresses within a process to the correct memory byte in a page frame?
- **Solution - address translation**
 - ❖ Processes use **virtual addresses**
 - ❖ CPU uses **physical addresses**
 - ❖ hardware support for virtual to physical **address translation**

Virtual addresses

- Virtual memory addresses (what the process uses)
 - ❖ Page number plus byte offset in page
 - ❖ Low order n bits are the byte offset
 - ❖ Remaining high order bits are the page number




Example: 32 bit virtual address

Page size = $2^{12} = 4\text{KB}$

Address space size = 2^{32} bytes = 4GB

Physical addresses

- Physical memory addresses (what the CPU uses)
 - ❖ Page "frame" number plus byte offset in page
 - ❖ Low order n bits are the byte offset
 - ❖ Remaining high order bits are the frame number



Example: 24 bit physical address


Frame size = 2^{12} = 4KB

Max physical memory size = 2^{24} bytes = 16MB

Address translation

- Hardware maps **page** numbers to **frame** numbers
- Memory management unit (MMU) has multiple registers for multiple pages
 - ❖ Like a base register except its value is substituted for the page number rather than added to it
 - ❖ Why don't we need a limit register for each page?

Memory Management Unit (MMU)




Virtual address spaces

- Here is the virtual address space
 - ❖ (as seen by the process)

Lowest address


Highest address



Virtual Addr Space


Virtual address spaces

- The address space is divided into “pages”
 - ❖ In BLITZ, the page size is 8K




Virtual address spaces

- In reality, only some of the pages are used




Physical memory

- Physical memory is divided into "*page frames*"
 - (Page size = frame size)




Virtual and physical address spaces

- Some frames are used to hold the pages of this process




Virtual and physical address spaces

- Some frames are used for other processes




Virtual address spaces

- Address **mappings** say which frame has which page



Page tables

- ❑ Address mappings are stored in a *page table* in memory
- ❑ One page table entry per page...
 - ❖ Is this page in memory? If so, which frame is it in?



Address mappings and translation

- **Address mappings** are stored in a **page table** in memory
 - ❖ Typically one page table for each process
- **Address translation** is done by **hardware** (ie the MMU)
- How does the MMU get the address mappings?
 - ❖ Either the MMU holds the entire page table (too expensive)
 - or it knows where it is in physical memory and goes there for every translation (too slow)
 - ❖ Or the MMU holds a portion of the page table
 - MMU **caches** page table entries
 - Cache is called a translation look-aside buffer (**TLB**)
 - ... and knows how to deal with TLB misses

Address mappings and translation


- What if the TLB needs a mapping it doesn't have?
- **Software managed TLB**
 - ❖ it generates a **TLB-miss fault** which is handled by the operating system (like interrupt or trap handling)
 - ❖ The operating system looks in the page tables, gets the mapping from the right entry, and puts it in the TLB
- **Hardware managed TLB**
 - ❖ it looks in a pre-specified physical memory location for the appropriate entry in the page table
 - ❖ The hardware architecture defines where page tables must be stored in physical memory
 - **OS must load current process page table there on context switch!**

The BLITZ architecture

- **Page size**
 - ❖ 8 Kbytes
- **Virtual addresses ("logical addresses")**
 - ❖ 24 bits --> 16 Mbyte virtual address space
 - ❖ 2^{11} Pages --> 11 bits for page number

The BLITZ architecture

- **Page size**
 - ❖ 8 Kbytes
- **Virtual addresses ("logical addresses")**
 - ❖ 24 bits --> 16 Mbyte virtual address space
 - ❖ 2^{11} Pages --> 11 bits for page number
- **An address:**




The BLITZ architecture

- Physical addresses
 - ❖ 32 bits --> 4 Gbyte installed memory (max)
 - ❖ 2^{19} Frames --> 19 bits for frame number


The BLITZ architecture

- Physical addresses
 - ❖ 32 bits --> 4 Gbyte installed memory (max)
 - ❖ 2^{19} Frames --> 19 bits for frame number



The BLITZ architecture

- The page table mapping:
 - ❖ Page --> Frame
- Virtual Address:



- Physical Address:



The BLITZ page table

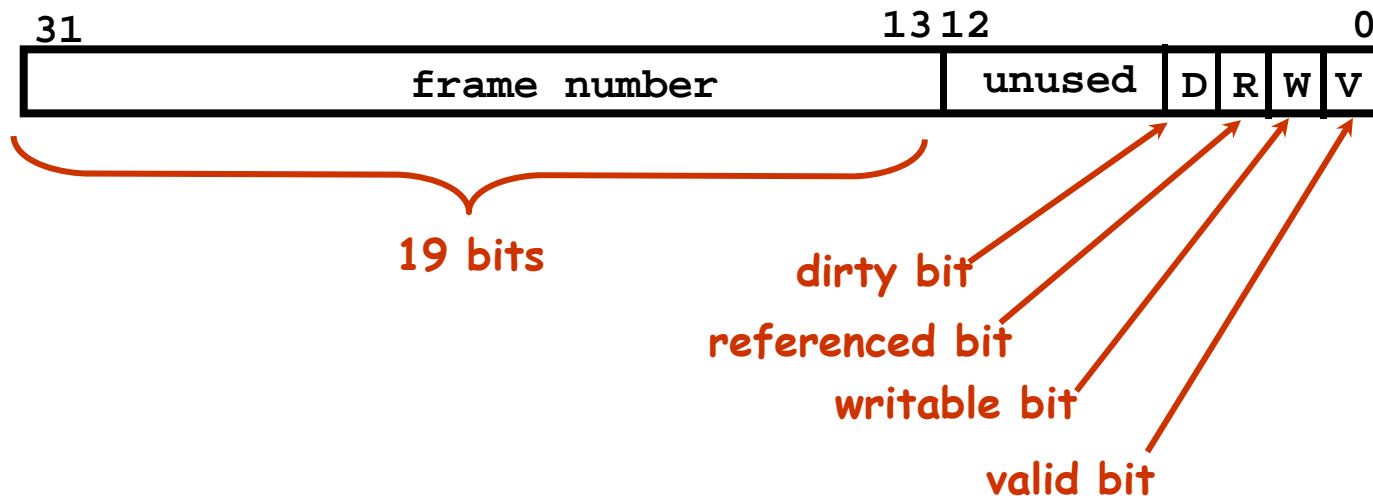
- An array of “*page table entries*”
 - ❖ Kept in memory
- 2^{11} pages in a virtual address space?
 - ❖ ---> 2K entries in the table
- Each entry is 4 bytes long
 - ❖ 19 bits The Frame Number
 - ❖ 1 bit Valid Bit
 - ❖ 1 bit Writable Bit
 - ❖ 1 bit Dirty Bit
 - ❖ 1 bit Referenced Bit
 - ❖ 9 bits Unused (and available for OS algorithms)

The BLITZ page table

- Two page table related registers in the CPU
 - ❖ Page Table Base Register
 - ❖ Page Table Length Register
- These define the “current” page table
 - ❖ This is how the CPU knows which page table to use
 - ❖ Must be saved and restored on context switch
 - ❖ They are essentially the Blitz MMU
- Bits in the CPU “status register”
 - “System Mode”
 - “Interrupts Enabled”
 - “Paging Enabled”
 - 1 = Perform page table translation for every memory access
 - 0 = Do not do translation

The BLITZ page table

□



The BLITZ page table

□


page table base register

	31	13	12				0
0	frame number		unused	D	R	W	V
1	frame number		unused	D	R	W	V
2	frame number		unused	D	R	W	V
	frame number		unused	D	R	W	V
2K	frame number		unused	D	R	W	V

Indexed by the page number

The BLITZ page table

□




virtual address

page table base register

	31		13	12		0		
0	frame number			unused	D	R	W	V
1	frame number			unused	D	R	W	V
2	frame number			unused	D	R	W	V
	frame number			unused	D	R	W	V
2K	frame number			unused	D	R	W	V

The BLITZ page table


□



virtual address

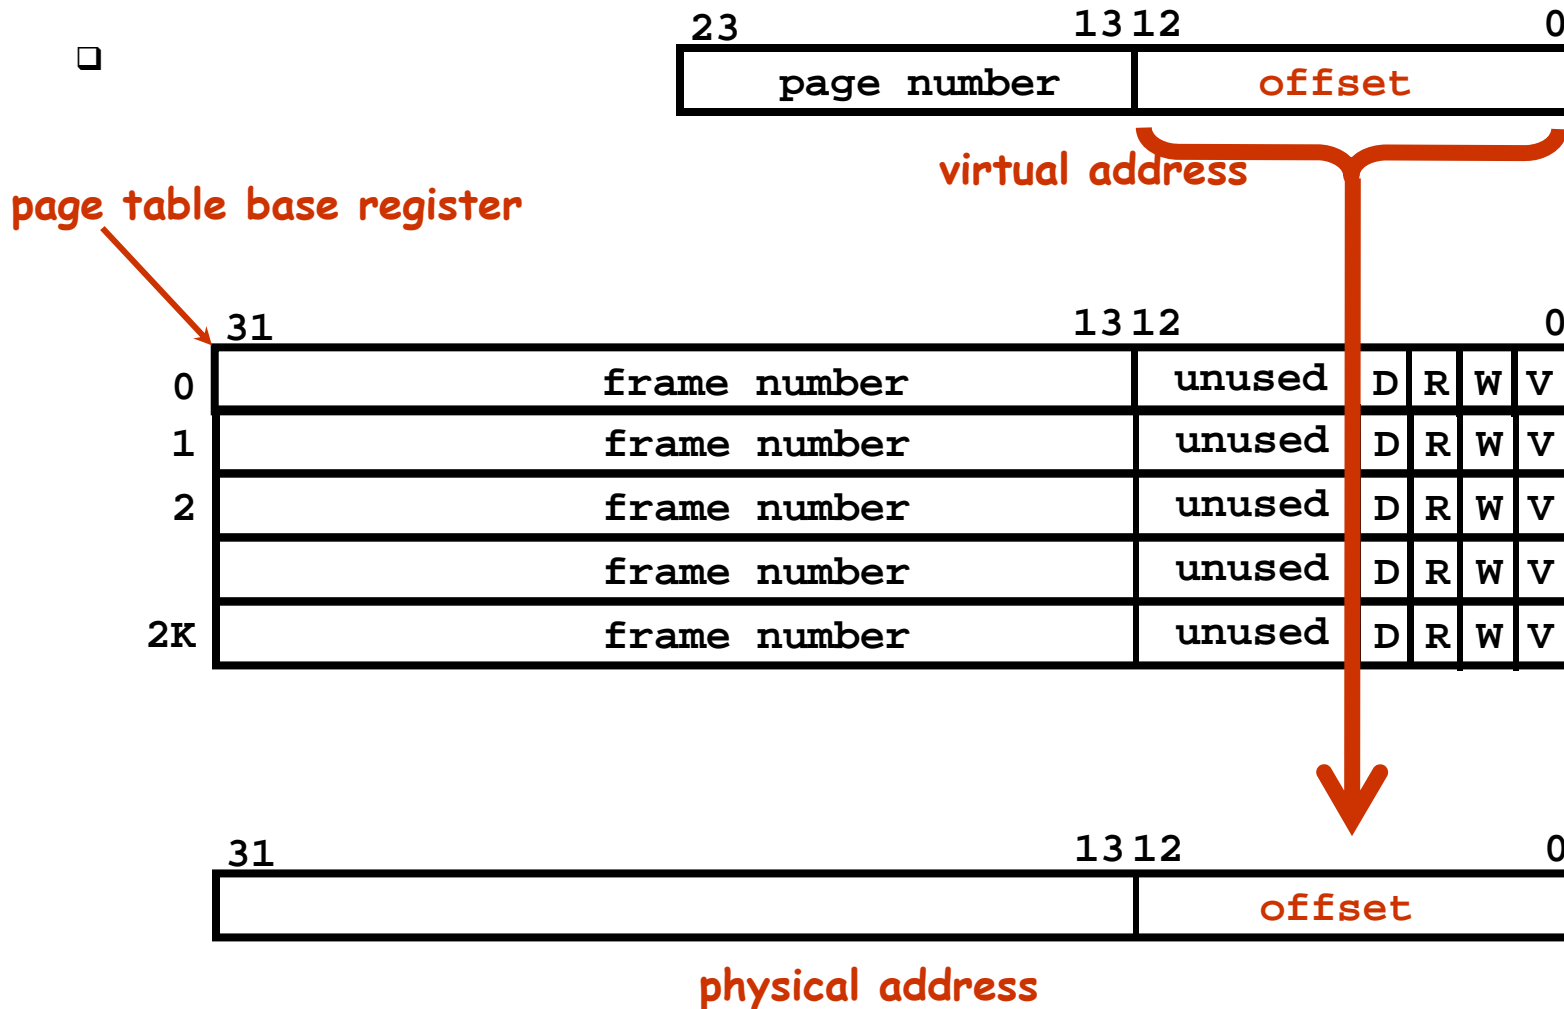
page table base register

	31		13	12		0		
0	frame number			unused	D	R	W	V
1	frame number			unused	D	R	W	V
2	frame number			unused	D	R	W	V
	frame number			unused	D	R	W	V
2K	frame number			unused	D	R	W	V

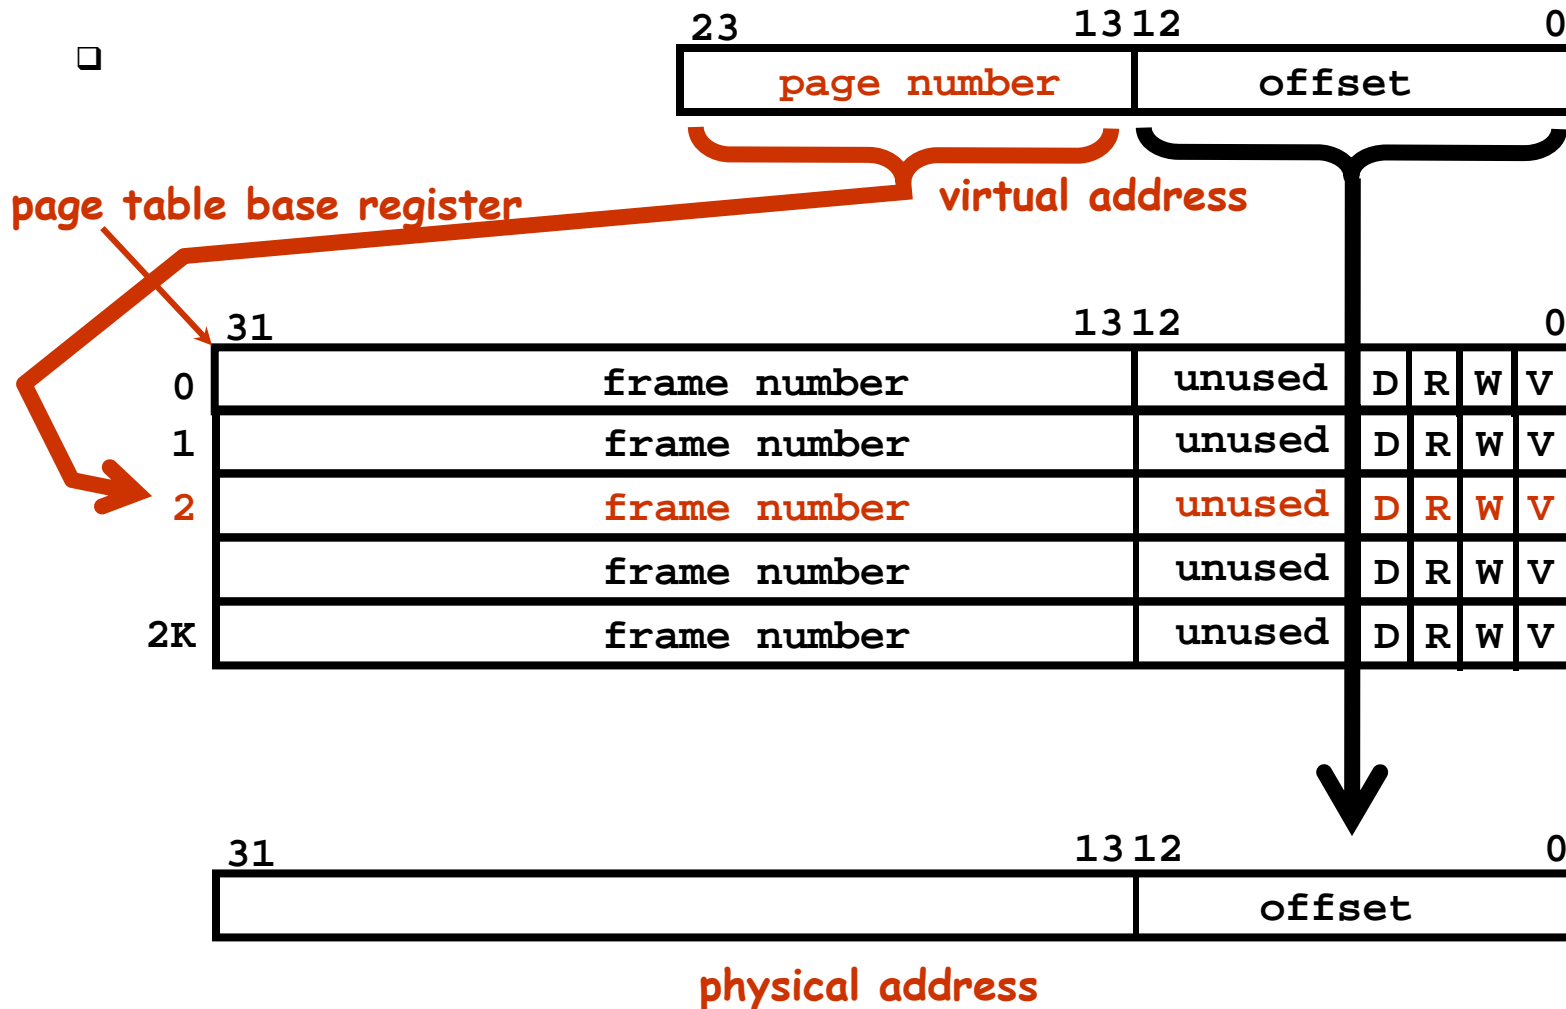


physical address

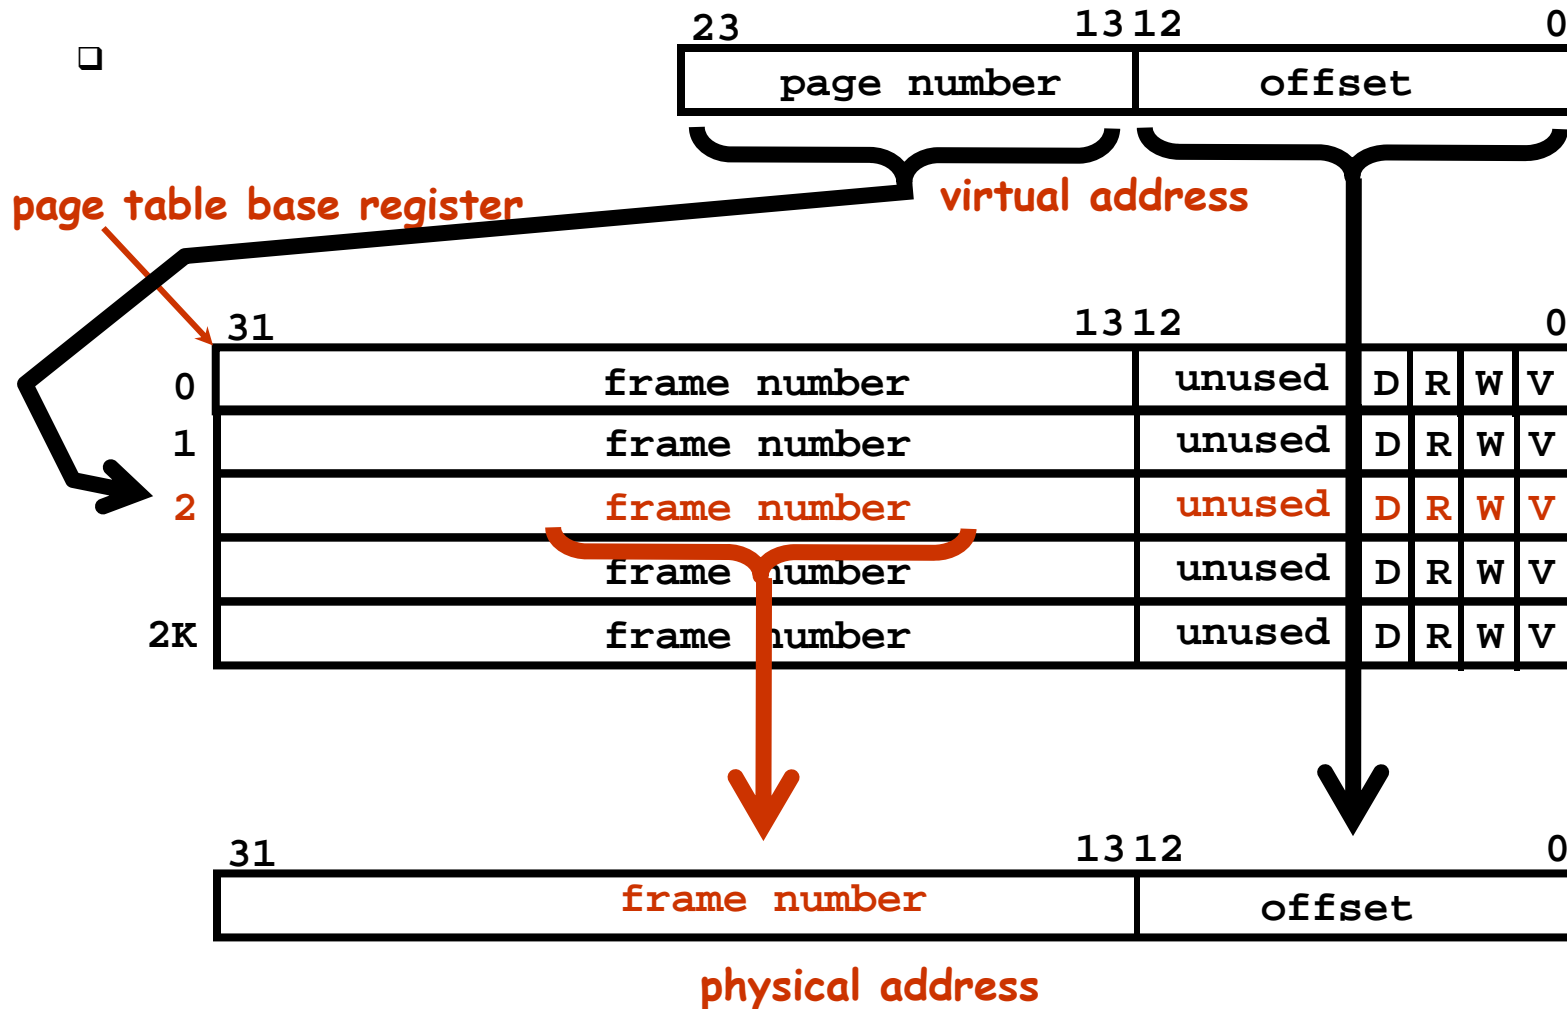
The BLITZ page table



The BLITZ page table



The BLITZ page table



Quiz

- ❑ What is the difference between a virtual and a physical address?
- ❑ What is address binding?
- ❑ Why are programs not usually written using physical addresses?
- ❑ Why is hardware support required for dynamic address translation?
- ❑ What is a page table used for?
- ❑ What is a TLB used for?
- ❑ How many address bits are used for the page offset in a system with 2KB page size?