

CS 333


Introduction to Operating Systems

Class 11 - Virtual Memory (1)

Jonathan Walpole
Computer Science
Portland State University

Virtual addresses

- Virtual memory addresses (what the process uses)
 - ❖ Page number plus byte offset in page
 - ❖ Low order n bits are the byte offset
 - ❖ Remaining high order bits are the page number




Example: 32 bit virtual address

Page size = $2^{12} = 4\text{KB}$

Address space size = 2^{32} bytes = 4GB

Physical addresses

- Physical memory addresses (what memory uses)
 - ❖ **Frame** number plus **byte offset** in frame
 - ❖ Low order n bits are the byte offset
 - ❖ Remaining high order bits are the **frame** number



Example: 24 bit physical address

Frame size = 2^{12} = 4KB


Max physical memory size = 2^{24} bytes = 16MB

Address translation

- Complete set of address mappings for a process are stored in a page table in memory
 - ❖ But accessing the table for every address translation is too expensive
 - ❖ So hardware support is used to map **page** numbers to **frame** numbers at full CPU speed
 - *Memory management unit (MMU) has multiple registers for multiple pages and knows how to access page tables*
 - *Also called a translation look aside buffer (TLB)*
 - *Essentially a cache of page table entries*

The BLITZ architecture

- The page table mapping:
 - ❖ Page --> Frame
- Virtual Address (24 bit in Blitz):



- Physical Address (32 bit in Blitz):



The BLITZ page table

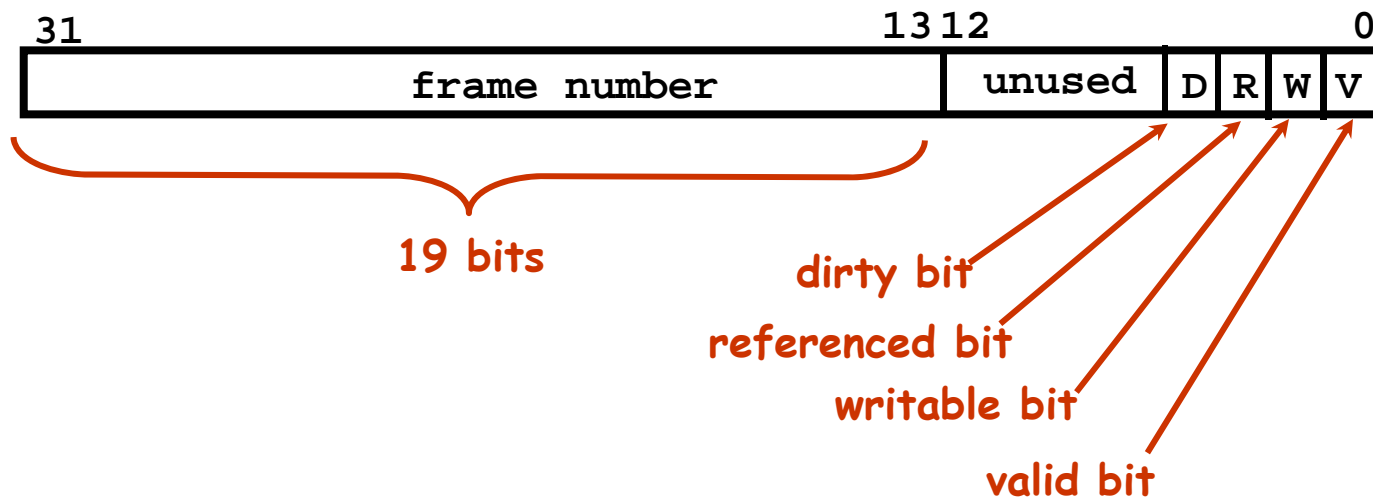
- An array of “*page table entries*”
 - ❖ Kept in memory
- 2^{11} pages in a virtual address space
 - ❖ ---> 2K entries in the table
- Each entry is 4 bytes long
 - ❖ 19 bits The Frame Number
 - ❖ 1 bit Valid Bit
 - ❖ 1 bit Writable Bit
 - ❖ 1 bit Dirty Bit
 - ❖ 1 bit Referenced Bit
 - ❖ 9 bits Unused (and available for OS algorithms)

The BLITZ page table

- Two page table related registers in the CPU
 - ❖ Page Table Base Register
 - ❖ Page Table Length Register
- These define the page table for the “current” process
 - ❖ Must be saved and restored on process context switch
- Bits in the CPU “status register”
 - “System Mode”
 - “Interrupts Enabled”
 - “Paging Enabled”
 - 1 = Perform page table translation for every memory access
 - 0 = Do not do translation

The BLITZ page table

- A page table entry



The BLITZ page table

- The full page table


page table base register

	31	13	12	0			
0	frame number	unused	D	R	W	V	
1	frame number	unused	D	R	W	V	
2	frame number	unused	D	R	W	V	
	frame number	unused	D	R	W	V	
2K	frame number	unused	D	R	W	V	

Indexed by the page number

The BLITZ page table

□




virtual address

page table base register

	31		13	12		0		
0	frame number			unused	D	R	W	V
1	frame number			unused	D	R	W	V
2	frame number			unused	D	R	W	V
	frame number			unused	D	R	W	V
2K	frame number			unused	D	R	W	V

The BLITZ page table


□



virtual address

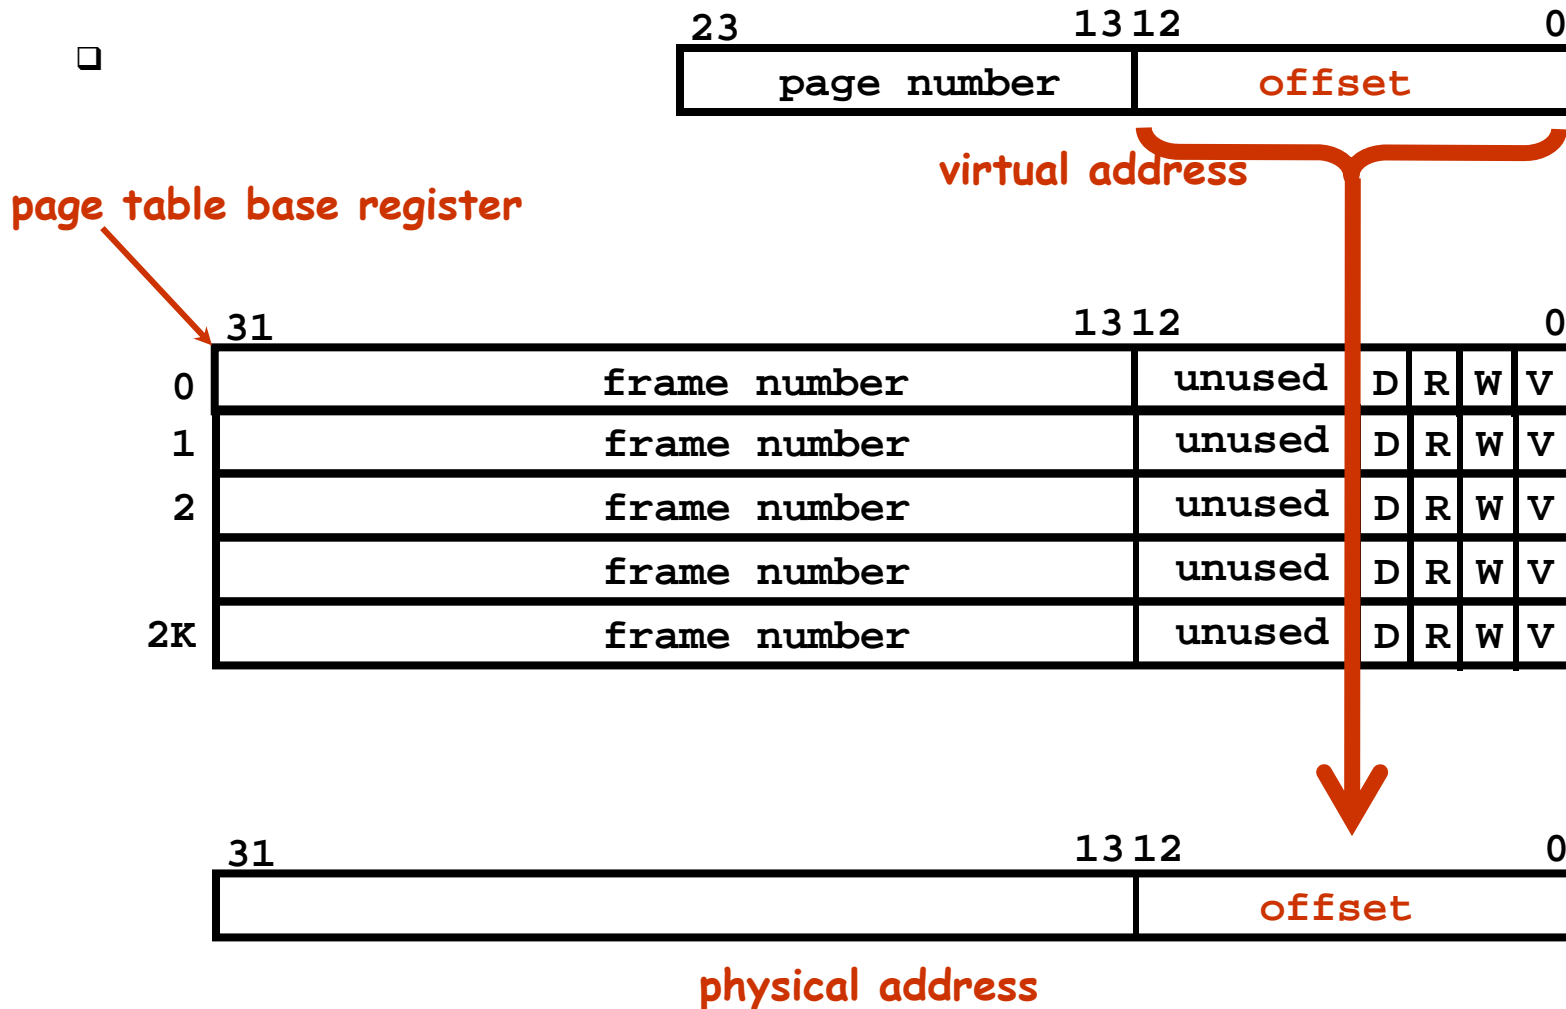
page table base register

	31		13	12		0		
0	frame number			unused	D	R	W	V
1	frame number			unused	D	R	W	V
2	frame number			unused	D	R	W	V
	frame number			unused	D	R	W	V
2K	frame number			unused	D	R	W	V

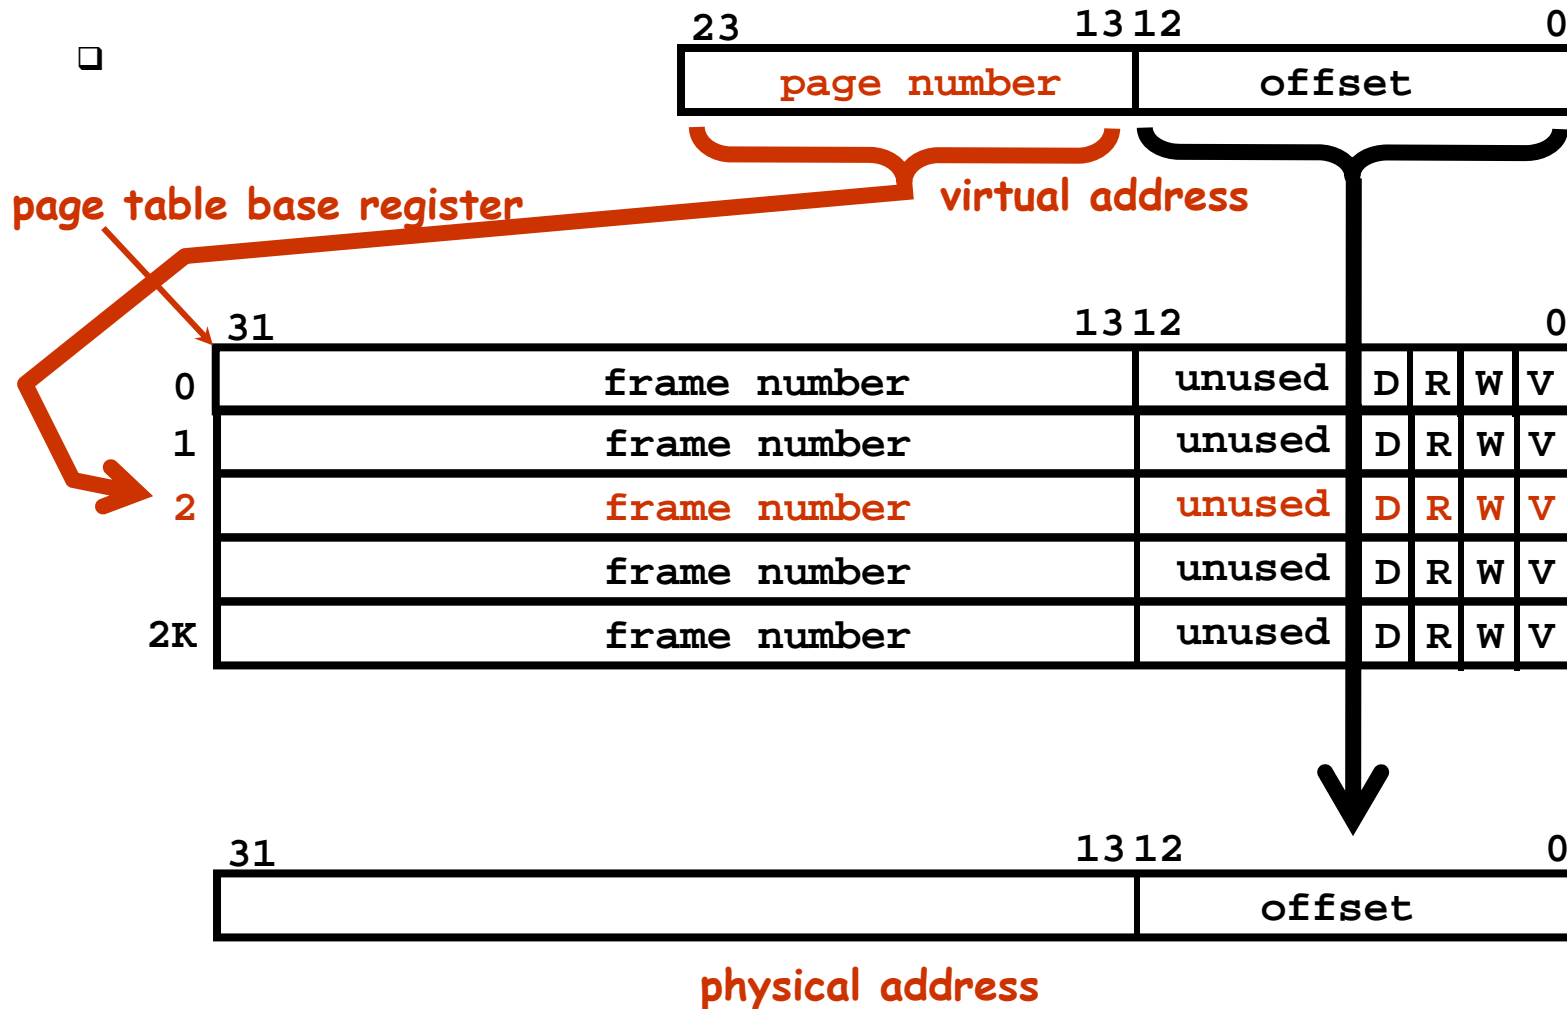


physical address


The BLITZ page table



The BLITZ page table



The BLITZ page table



Page tables

- **When and why do we access a page table?**
 - ❖ On every instruction to translate virtual to physical addresses?

Page tables

- **When and why do we access a page table?**
 - ❖ On every instruction to translate virtual to physical addresses?
 - ❖ In Blitz, YES, but in real machines NO!
 - ❖ In real machines it is only accessed
 - *On TLB miss faults to refill the TLB*
 - *During process creation and destruction*
 - *When a process allocates or frees memory?*

Translation Lookaside Buffer (TLB)

- **Problem:**

- ❖ MMU can't keep up with the CPU if it goes to the page table on every memory access!

Translation Lookaside Buffer (TLB)


❑ Problem:

- ❖ MMU can't keep up with the CPU if it goes to the page table on every memory access!

❑ Solution:

- ❖ Cache the page table entries in a hardware cache
- ❖ Small number of entries (e.g., 64)
- ❖ Each entry contains
 - Page number
 - Other stuff from page table entry
- ❖ Associatively indexed on page number
 - ie. You can do a lookup in a single cycle

Translation lookaside buffer



Hardware operation of TLB


□

Key					
Page Number	Frame Number	Other			
23	37	unused	D	R	W V
17	50	unused	D	R	W V
92	24	unused	D	R	W V
5	19	unused	D	R	W V
12	6	unused	D	R	W V

Hardware operation of TLB

□

virtual address



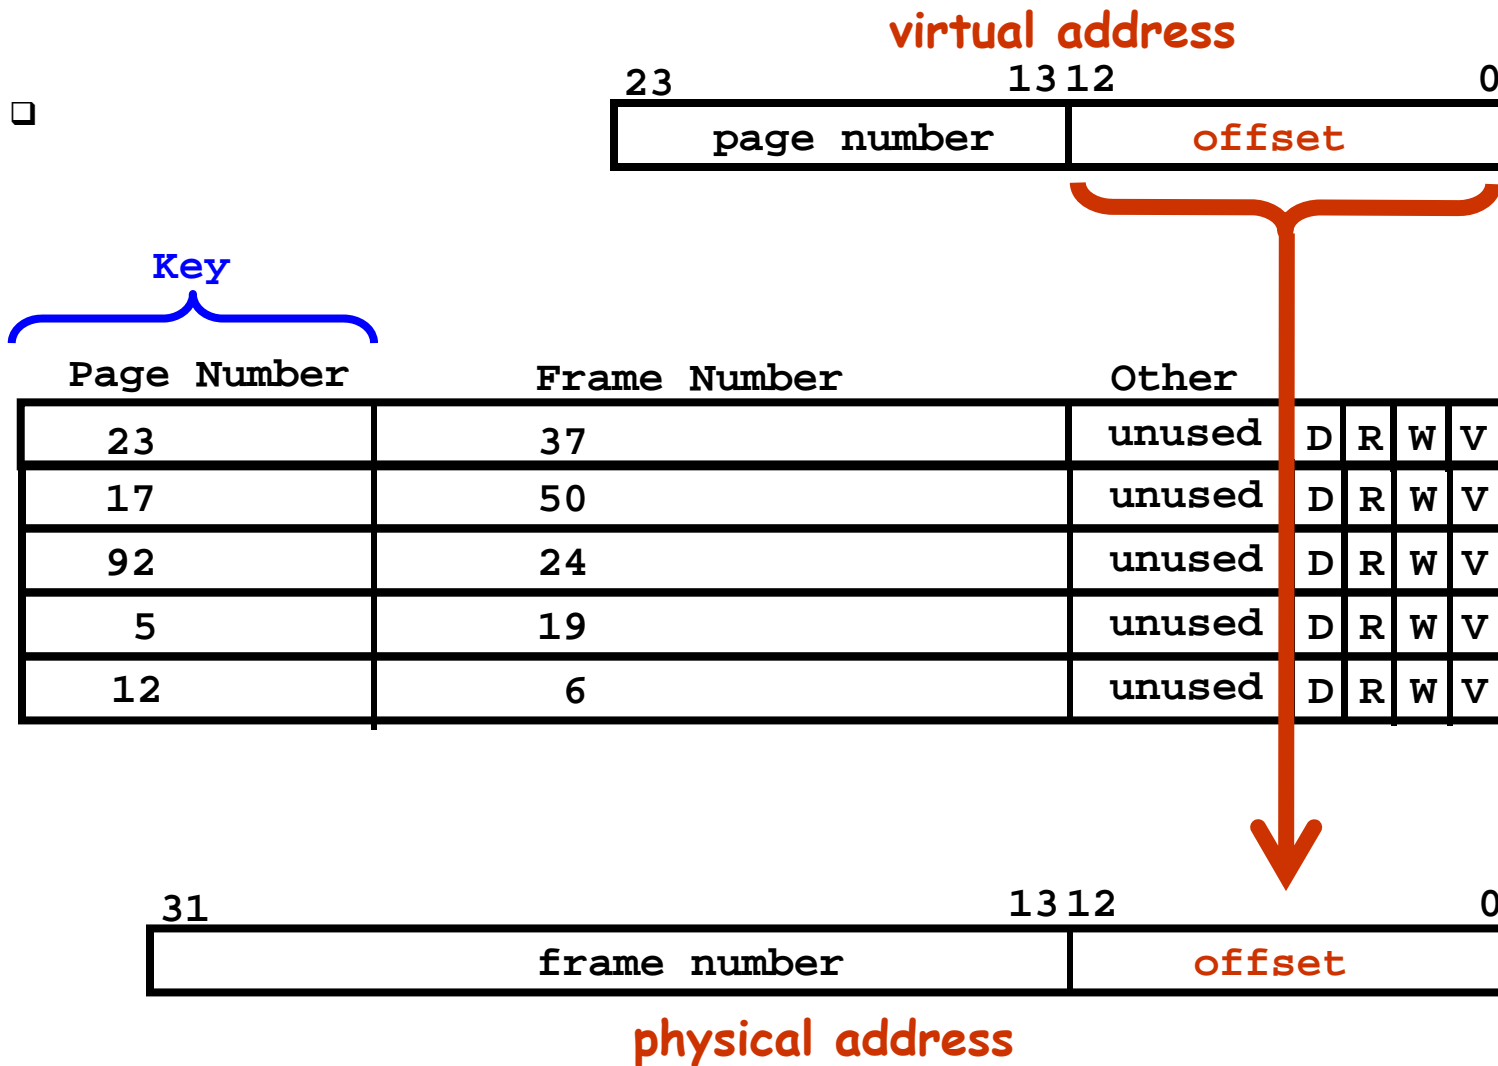
Key

Page Number		Frame Number	Other			
23	37	unused	D	R	W	V
17	50	unused	D	R	W	V
92	24	unused	D	R	W	V
5	19	unused	D	R	W	V
12	6	unused	D	R	W	V

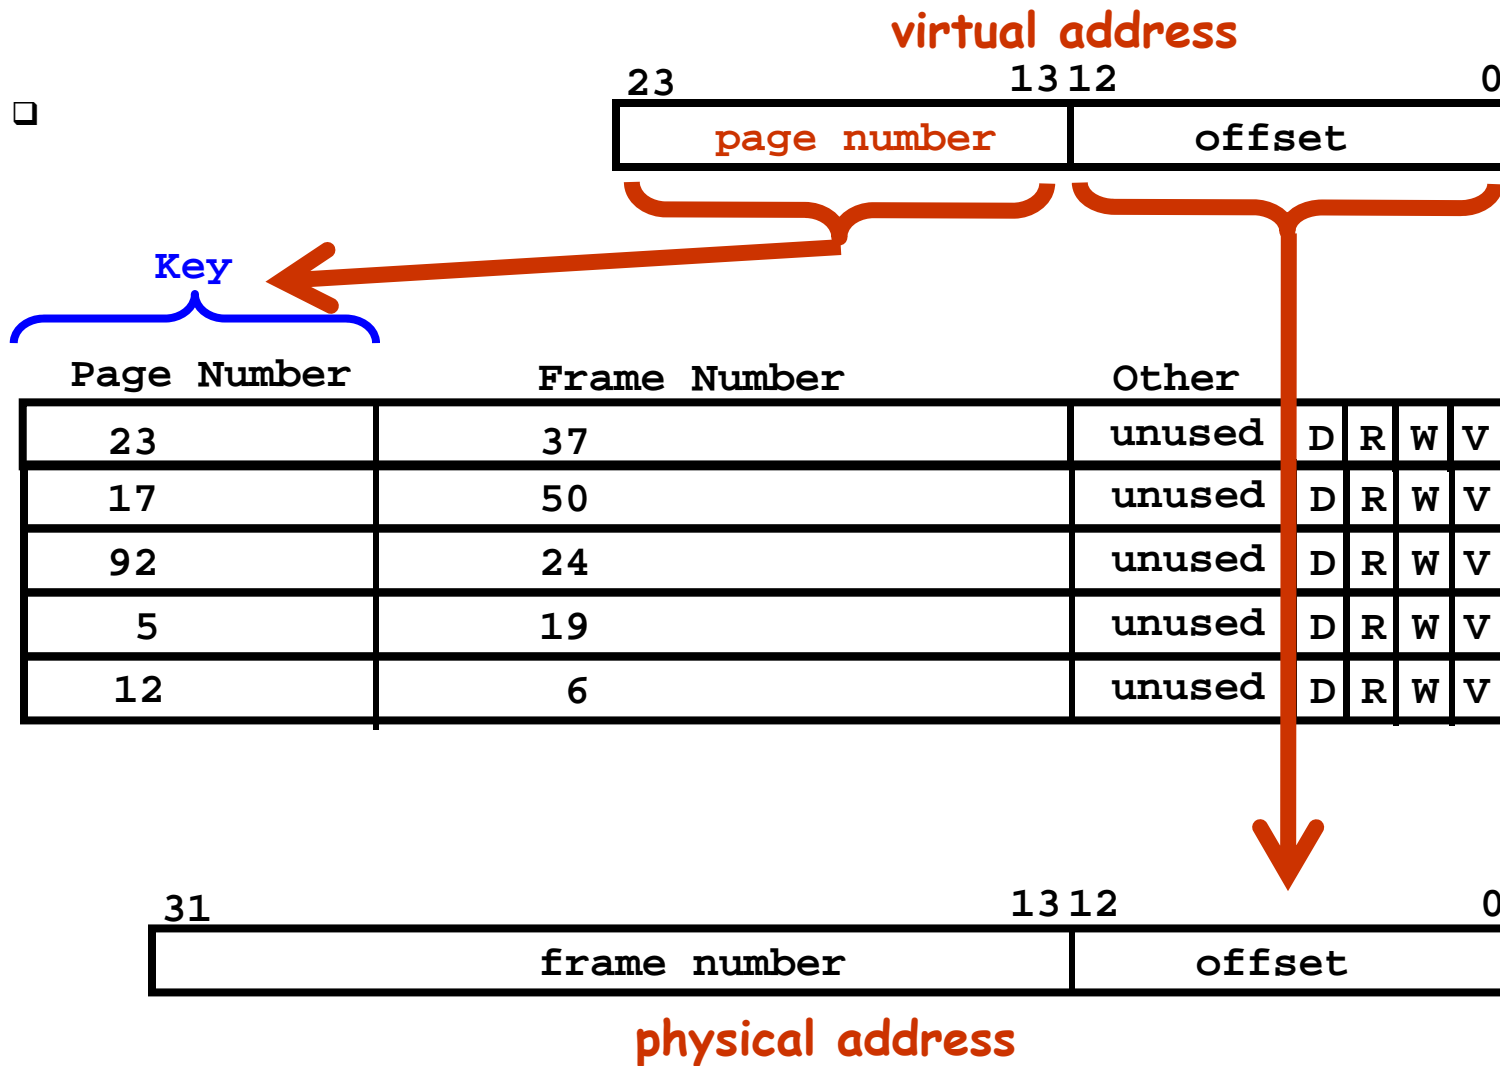


physical address

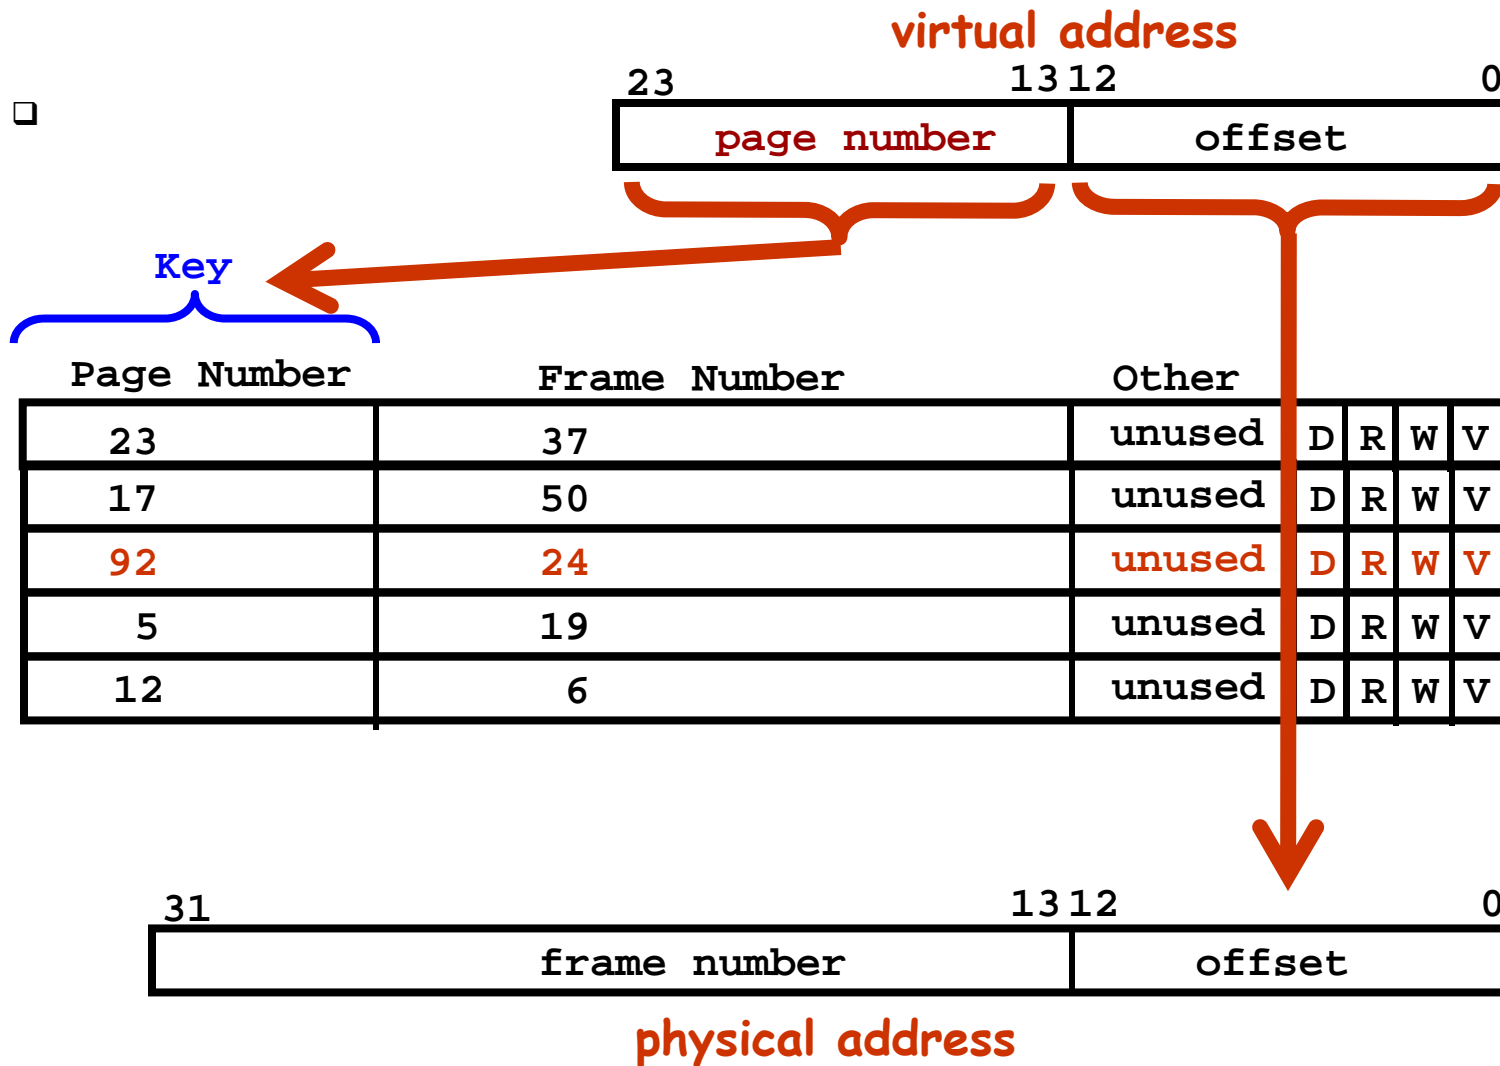
Hardware operation of TLB



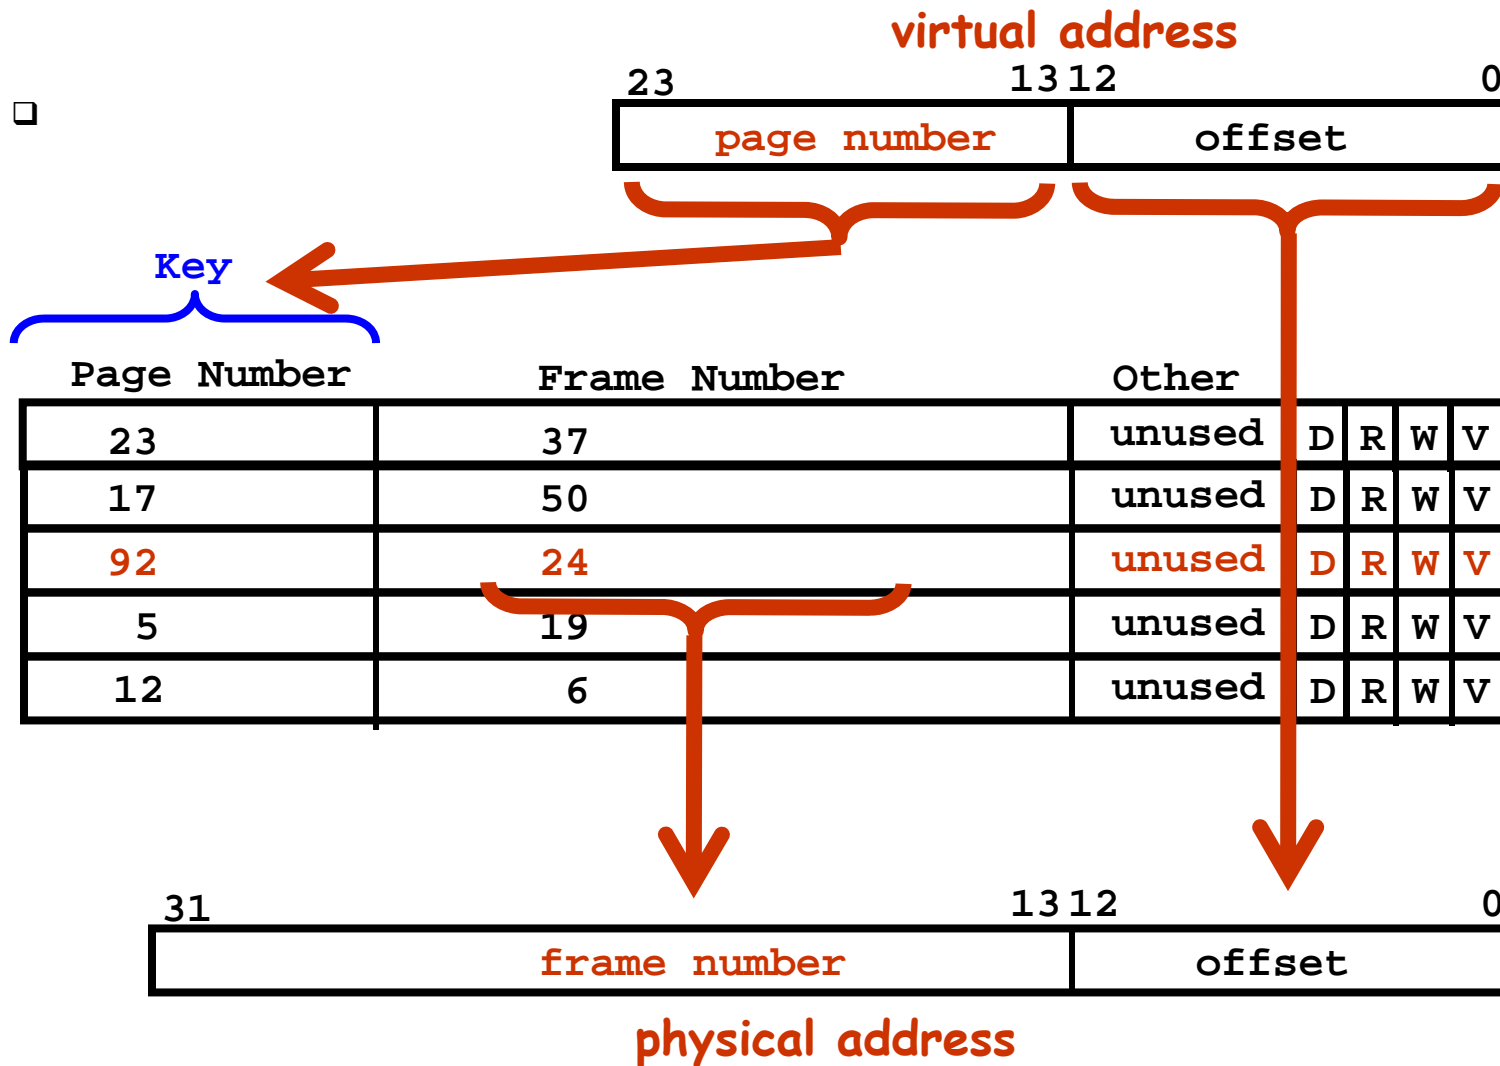
Hardware operation of TLB



Hardware operation of TLB



Hardware operation of TLB



Software operation of TLB

- What if the entry is not in the TLB?
 - ❖ Go look in the page table in memory
 - ❖ Find the right entry
 - ❖ Move it into the TLB
 - ❖ But which TLB entry should be replaced?

Software operation of TLB

- ❑ **Hardware TLB refill**
 - ❖ Page tables in specific location and format
 - ❖ TLB hardware handles its own misses
 - ❖ Replacement policy fixed by hardware
- ❑ **Software refill**
 - ❖ Hardware generates trap (**TLB miss fault**)
 - ❖ Lets the OS deal with the problem
 - ❖ Page tables become entirely a OS data structure!
 - ❖ Replacement policy managed in software

Software operation of TLB

- What should we do with the TLB on a context switch?
- How can we prevent the next process from using the last process's address mappings?
 - ❖ Option 1: empty the TLB
 - New process will generate faults until its pulls enough of its own entries into the TLB
 - ❖ Option 2: just clear the "Valid Bit"
 - New process will generate faults until its pulls enough of its own entries into the TLB
 - ❖ Option 3: the hardware maintains a process id **tag** on each TLB entry
 - Hardware compares this to a process id held in a specific register ... on every translation

Page tables

- Do we access a page table when a process allocates or frees memory?

Page tables


- Do we access a page table when a process allocates or frees memory?
 - ❖ Not necessarily
 - ❖ Library routines (malloc) can service small requests from a pool of free memory already allocated within a process address space
 - ❖ When these routines run out of space a new page must be allocated and its entry inserted into the page table
 - This allocation is requested using a system call

Page table design issues

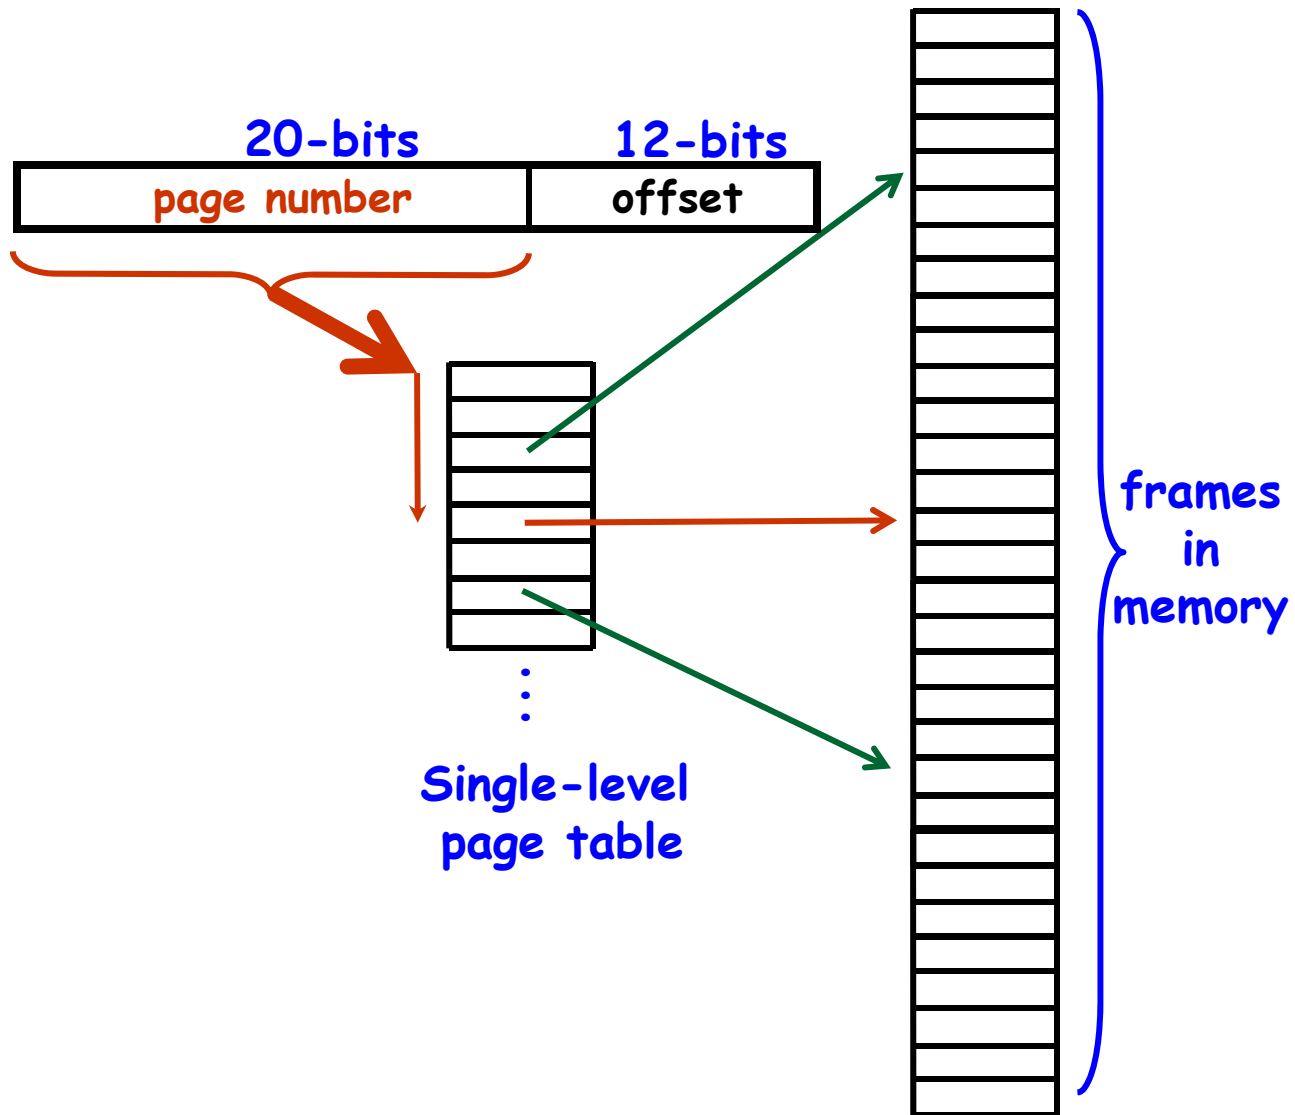
- ❑ **Page table size depends on**
 - ❖ Page size
 - ❖ Virtual address length
- ❑ **Memory used for page tables is overhead!**
 - ❖ How can we save space?
 - ❖ ... and still find entries quickly?
- ❑ **Three options**
 - ❖ Single-level page tables
 - ❖ Multi-level page tables
 - ❖ Inverted page tables

Single-level page tables

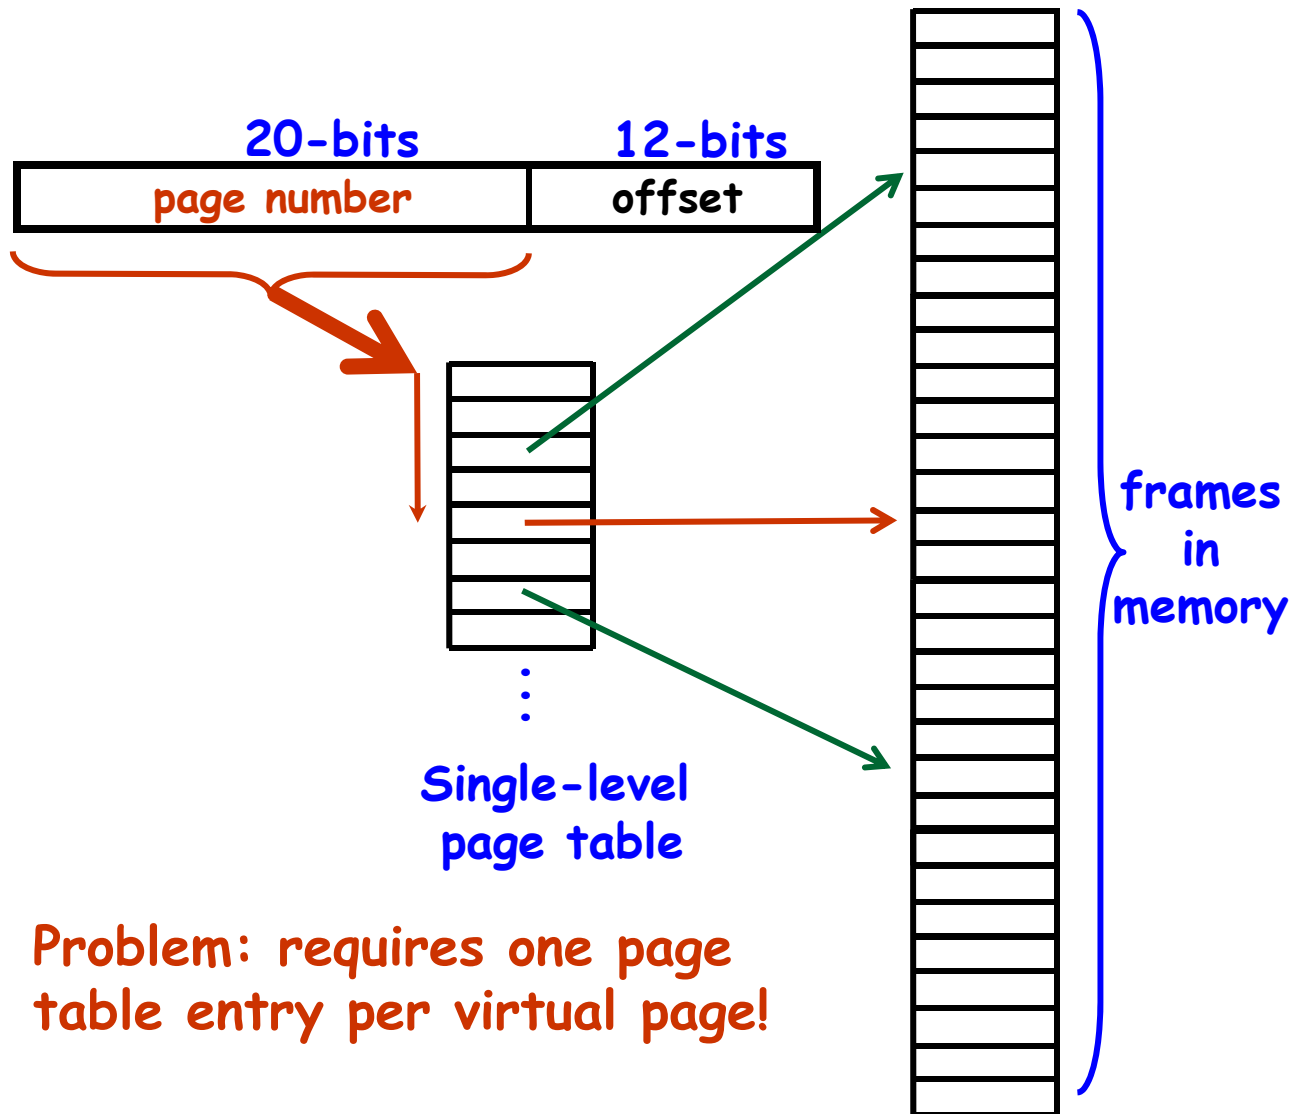
A Virtual Address (32 bit):



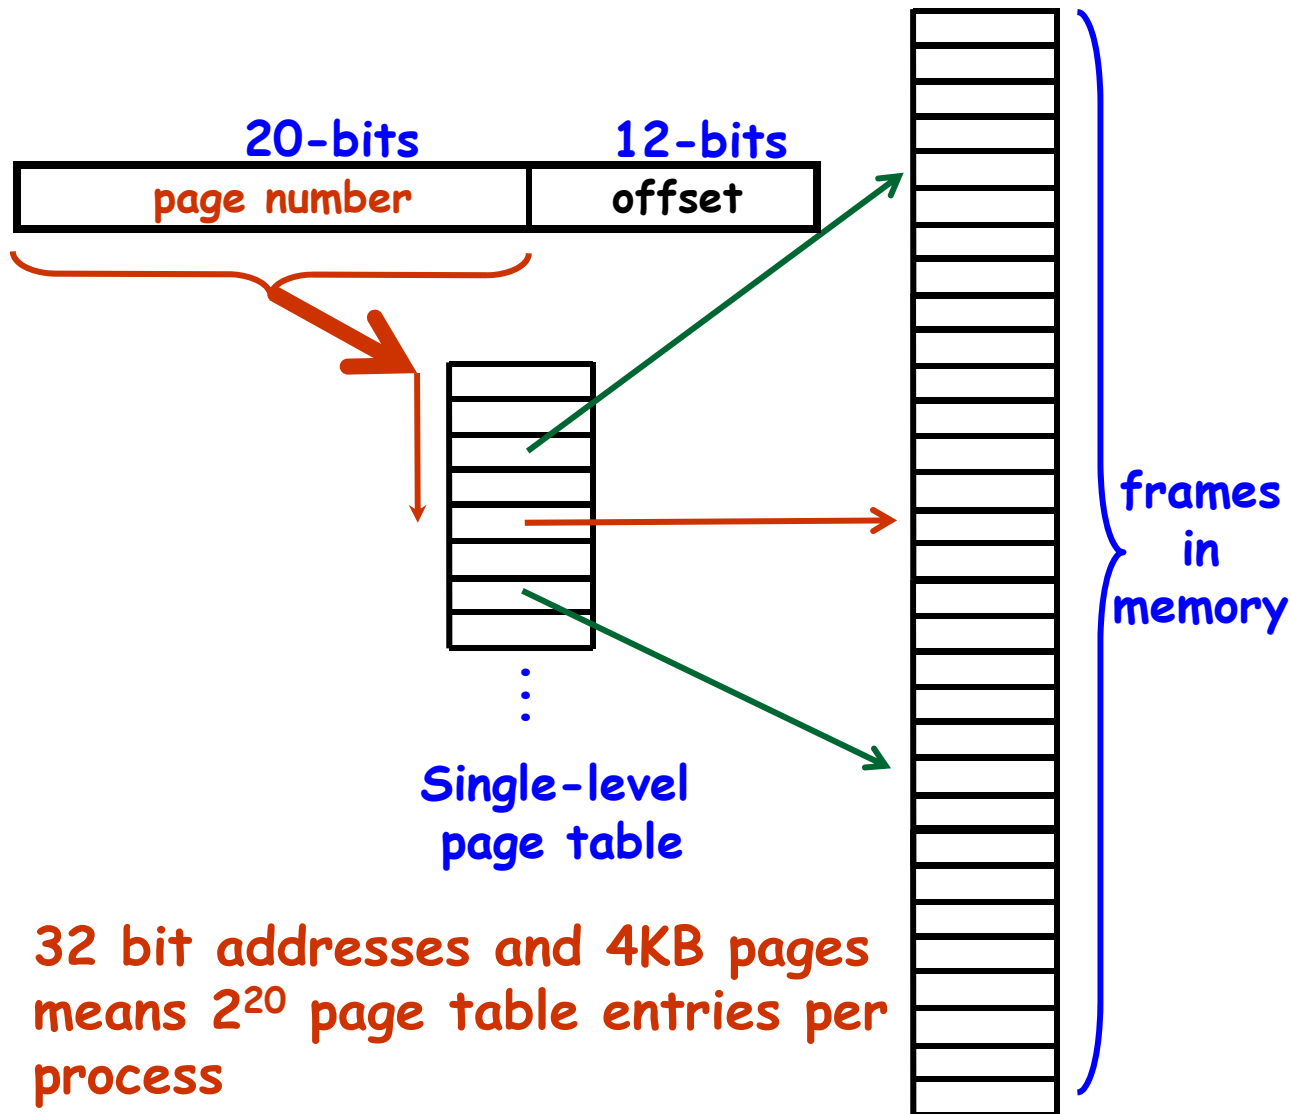
Single-level page tables



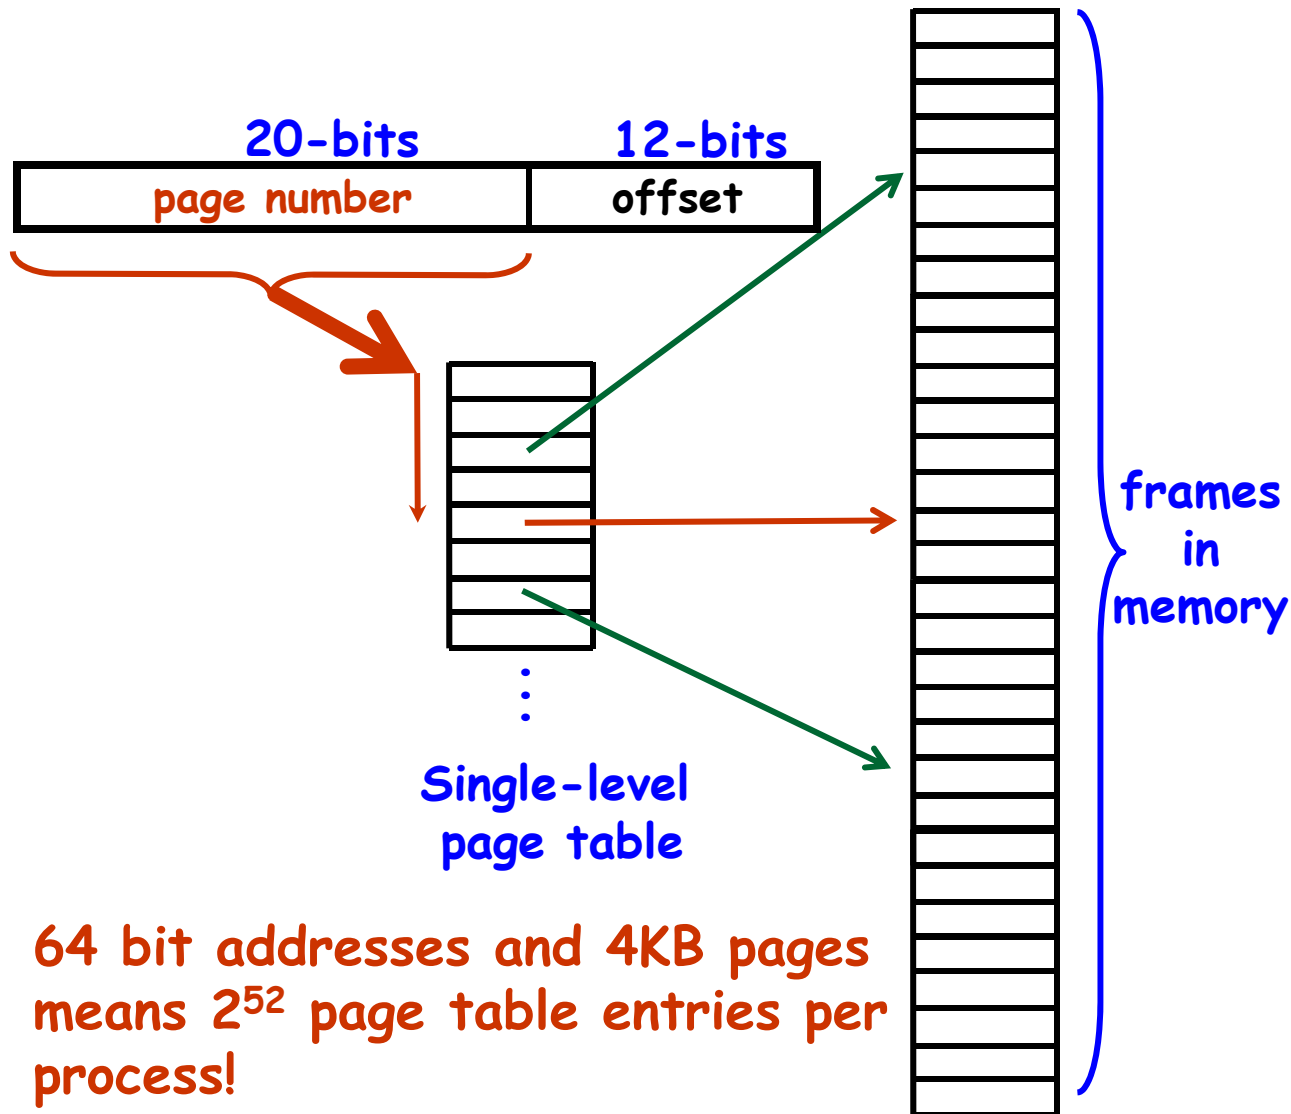
Single-level page tables




Single-level page tables



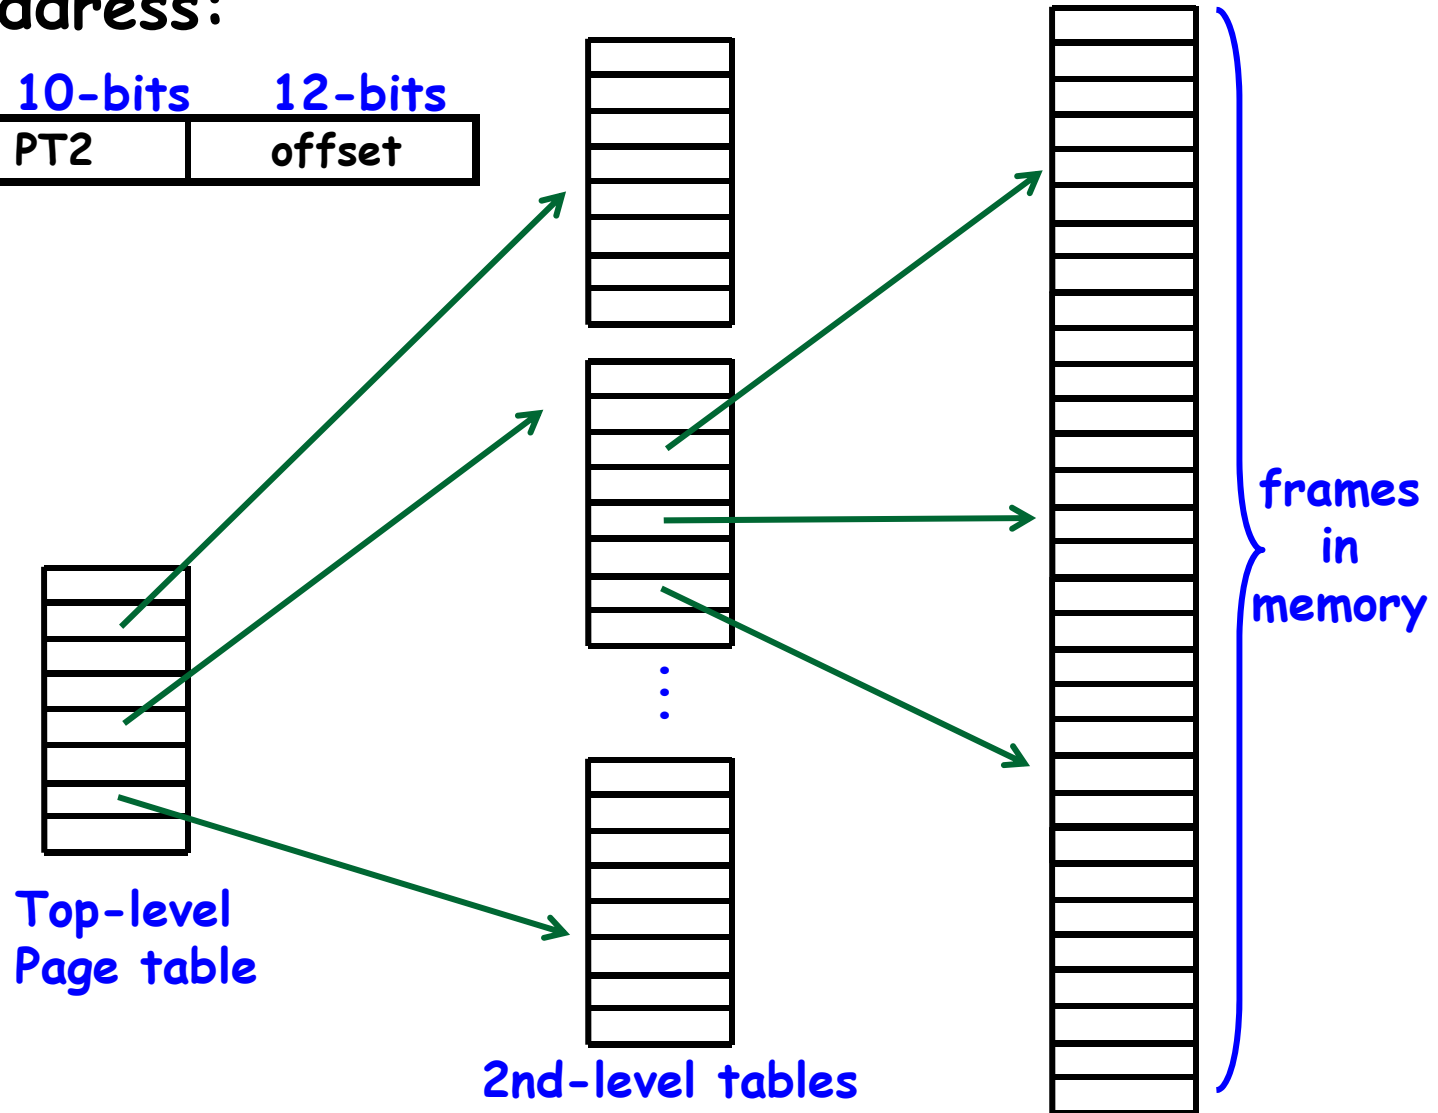

Single-level page tables



Multi-level page tables





- ❑ **A Virtual Address:**





Multi-level page tables

□ A Virtual Address:





Multi-level page tables

□ A Virtual Address:





Multi-level page tables

□ A Virtual Address:



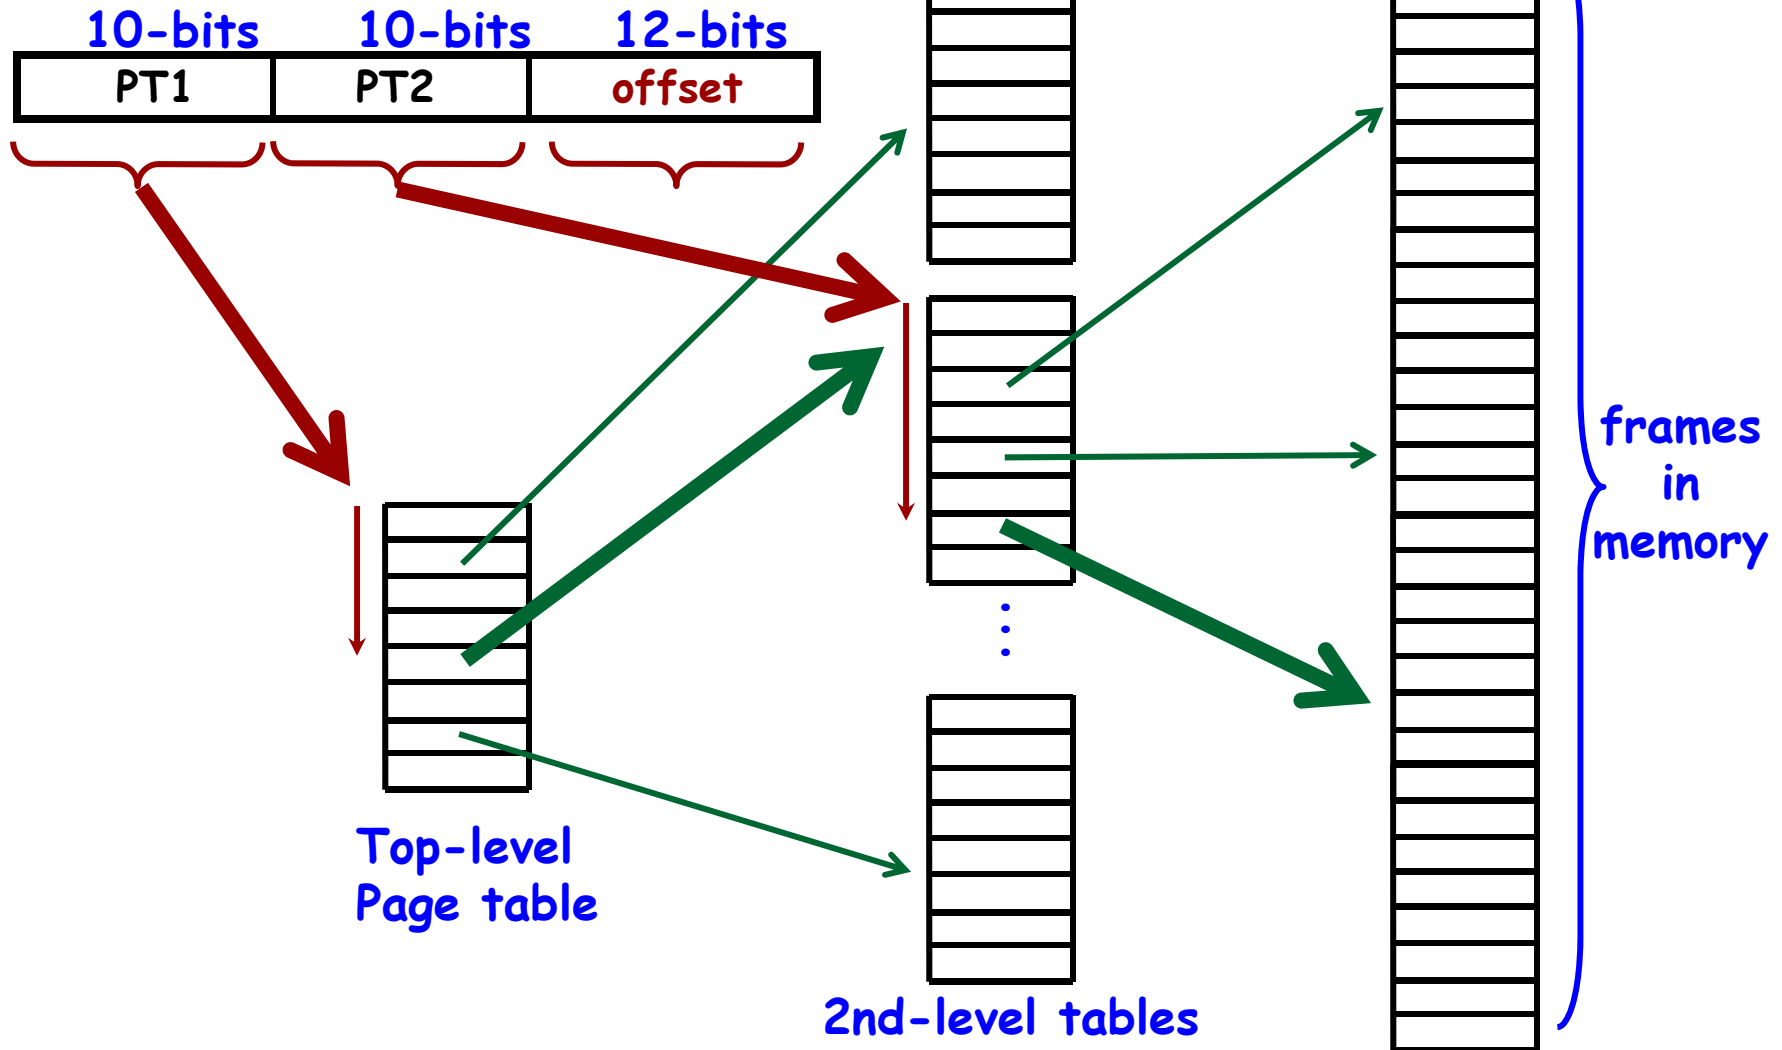
Multi-level page tables

□ A Virtual Address:



Multi-level page tables

□ A Virtual Address:



Multi-level page tables

- Ok, but how exactly does this save space?

Multi-level page tables

- ❑ Ok, but how exactly does this save space?
- ❑ Not all pages within a virtual address space are **allocated**
 - ❖ Not only do they not have a page frame, but that range of virtual addresses is not being used
 - ❖ So no need to maintain complete information about it
 - ❖ Some intermediate page tables are empty and not needed
- ❑ We could also page the page table
 - ❖ This saves space but slows access ... a lot!

VM puzzle

