

CS 333

Introduction to Operating Systems

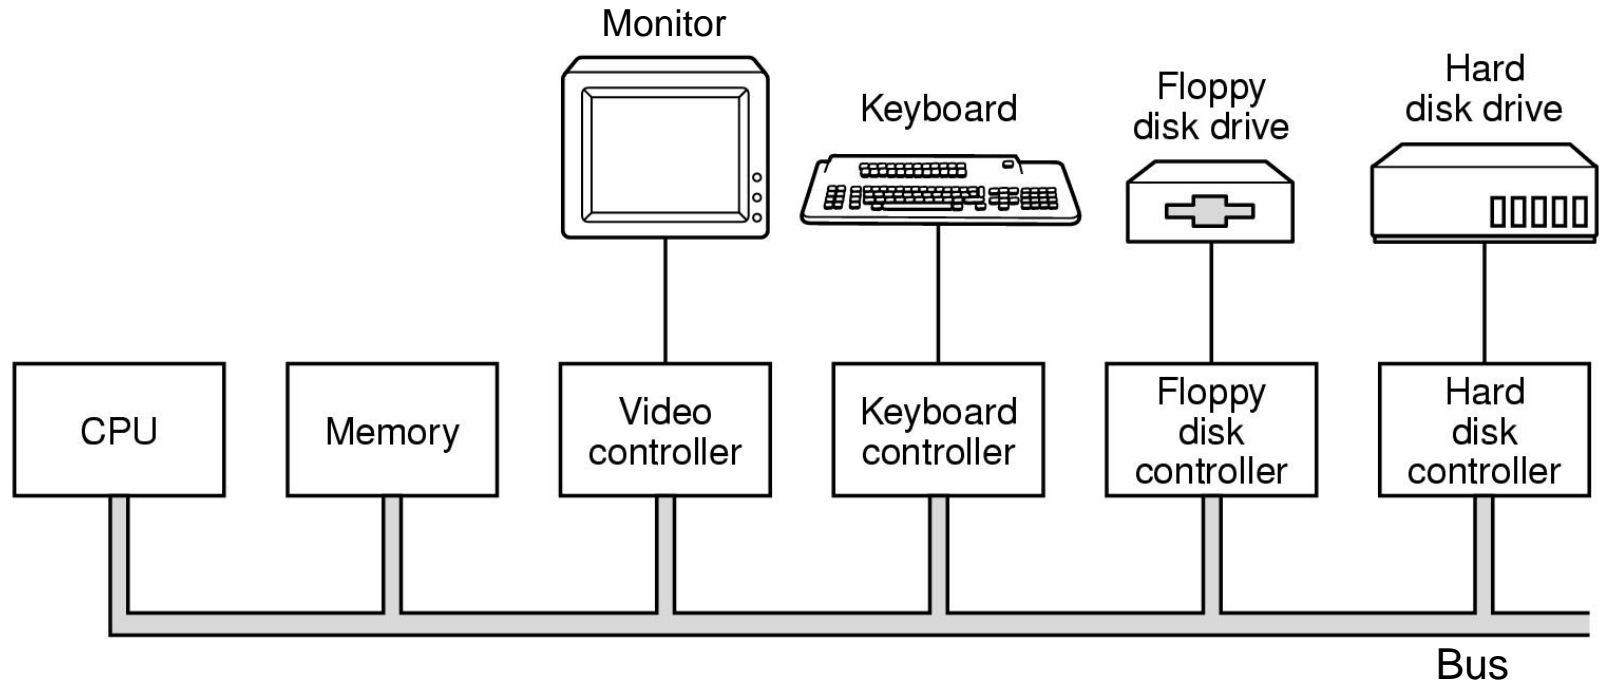
Class 15 - Input/Output

Jonathan Walpole
Computer Science
Portland State University

I/O devices - terminology

- ❑ Device (mechanical hardware)
- ❑ Device controller (electrical hardware)
- ❑ Device driver (software)

Example devices and their controllers




- **Components of a simple personal computer**

Device controllers

- ❑ **The Device vs. its Controller**
- ❑ **Some duties of a device controller:**
 - ❖ Interface between CPU and the Device
 - ❖ Start/Stop device activity
 - ❖ Convert serial bit stream to a block of bytes
 - ❖ Deal with errors
 - **Detection / Correction**
 - ❖ Move data to/from main memory
- ❑ **Some controllers may handle several (similar) devices**

How to communicate with a device?

- Hardware supports I/O ports or memory mapped I/O for accessing device controller registers and buffers



I/O ports

- Each port has a separate number.

- CPU has special I/O instructions

- ❖ `in r4,3`

- ❖ `out 3,r4`

The I/O Port Number

- Port numbers form an “address space”... separate from main memory


- Contrast with

- ❖ `load r4,3`

- ❖ `store 3,r4`

Memory-mapped I/O

- **One address space for**
 - ❖ main memory
 - ❖ I/O devices
- **CPU has no special instructions**
 - ❖ `load r4, addr`
 - ❖ `store addr, r4`
- **I/O devices are “mapped” into**
 - ❖ very high addresses



Wide range of I/O device speeds





Device	Data rate
Keyboard	10 bytes/sec
Mouse	100 bytes/sec
56K modem	7 KB/sec
Telephone channel	8 KB/sec
Dual ISDN lines	16 KB/sec
Laser printer	100 KB/sec
Scanner	400 KB/sec
Classic Ethernet	1.25 MB/sec
USB (Universal Serial Bus)	1.5 MB/sec
Digital camcorder	4 MB/sec
IDE disk	5 MB/sec
40x CD-ROM	6 MB/sec
Fast Ethernet	12.5 MB/sec
ISA bus	16.7 MB/sec
EIDE (ATA-2) disk	16.7 MB/sec
FireWire (IEEE 1394)	50 MB/sec
XGA Monitor	60 MB/sec
SONET OC-12 network	78 MB/sec
SCSI Ultra 2 disk	80 MB/sec
Gigabit Ethernet	125 MB/sec
Ultrium tape	320 MB/sec
PCI bus	528 MB/sec
Sun Gigaplane XB backplane	20 GB/sec

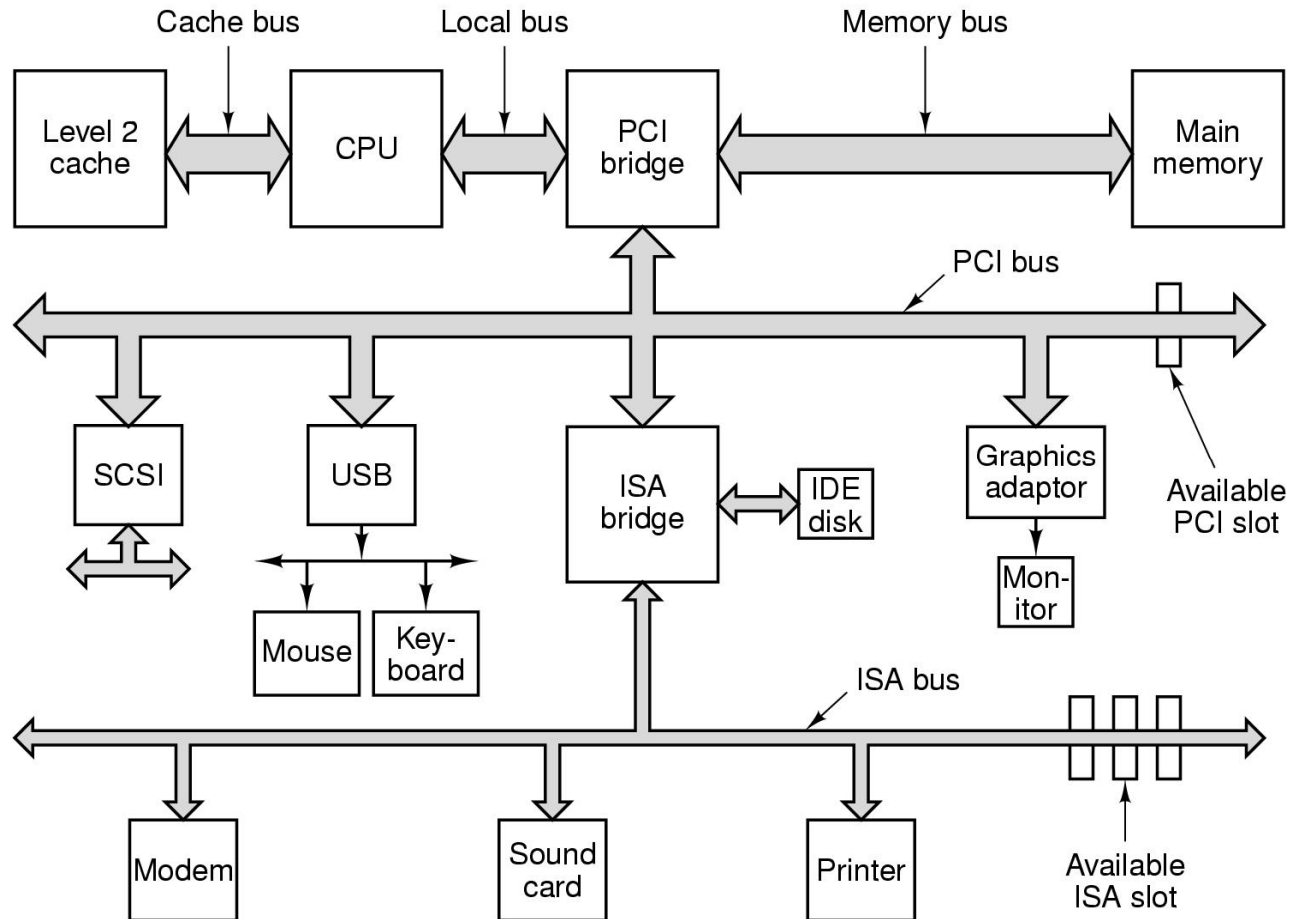
Performance challenges: I/O hardware

- How to prevent slow devices from slowing down memory due to bus contention
 - ❖ What is bus contention?
- How to access I/O addresses without interfering with memory performance

Single vs. dual bus architecture



Hardware view of Pentium




Structure of a large Pentium system

Performance challenges: I/O software

- How to prevent CPU throughput from being limited by I/O device speed (for *slow devices*)
 - ❖ Why would slow devices affect the CPU?
- How to prevent I/O throughput from being limited by CPU speed (for *fast devices*)
 - ❖ Why would device throughput be limited by the CPU?
- How to achieve good utilization of CPU and I/O devices
- How to meet the real-time requirements of devices

Programmed I/O



Steps in printing a string

Programmed I/O

- **Example:**
 - ❖ Writing a string to a serial output
 - Printing a string on the printer

```
CopyFromUser(virtAddr, kernelBuffer, byteCount)
for i = 0 to byteCount-1
    while *serialStatusReg != READY
    endwhile
    *serialDataReg = kernelBuffer[i]
endFor
return
```

- Called “Busy Waiting” or “Polling”
- Problem: CPU is continually busy working on I/O!

Interrupt-Driven I/O

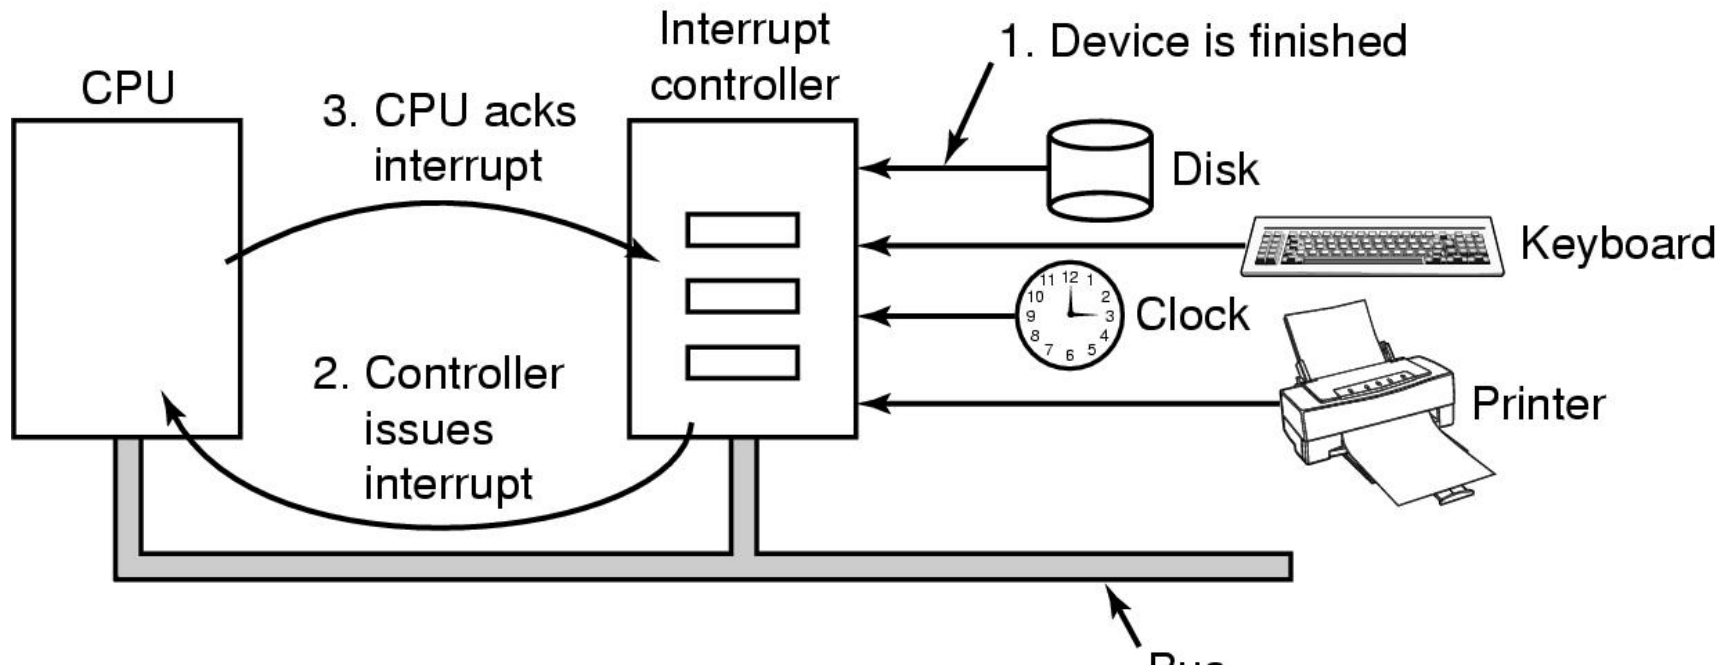
- **Getting the I/O started:**

```
CopyFromUser(virtAddr, kernelBuffer, byteCount)
EnableInterrupts()
while *serialStatusReg != READY
endWhile
*serialDataReg = kernelBuffer[0]
Sleep ()
```

- **The Interrupt Handler:**

```
if i == byteCount
    Wake up the user process
else
    *serialDataReg = kernelBuffer[i]
    i = i + 1
endIf
Return from interrupt
```

Hardware support for interrupts



How interrupts happen. Connections between devices and interrupt controller actually use interrupt lines on the bus rather than dedicated wires

Problem with Interrupt driven I/O

□ Problem:

- ❖ CPU is still involved in every data transfer
- ❖ Interrupt handling overhead is high
- ❖ Overhead cost is not amortized over much data
- ❖ Overhead is too high for fast devices
 - Gbps networks
 - Disk drives

Direct Memory Access (DMA)

- ❑ Data transferred from device straight to/from memory
- ❑ CPU not involved
- ❑ *The DMA controller:*
 - ❖ Does the work of moving the data
 - ❖ CPU sets up the DMA controller ("programs it")
 - ❖ CPU continues
 - ❖ The DMA controller moves the bytes

Sending data to a device using DMA

- **Getting the I/O started:**

`CopyFromUser(virtAddr, kernelBuffer, byteCount)`

`Set up DMA controller`

`Sleep ()`

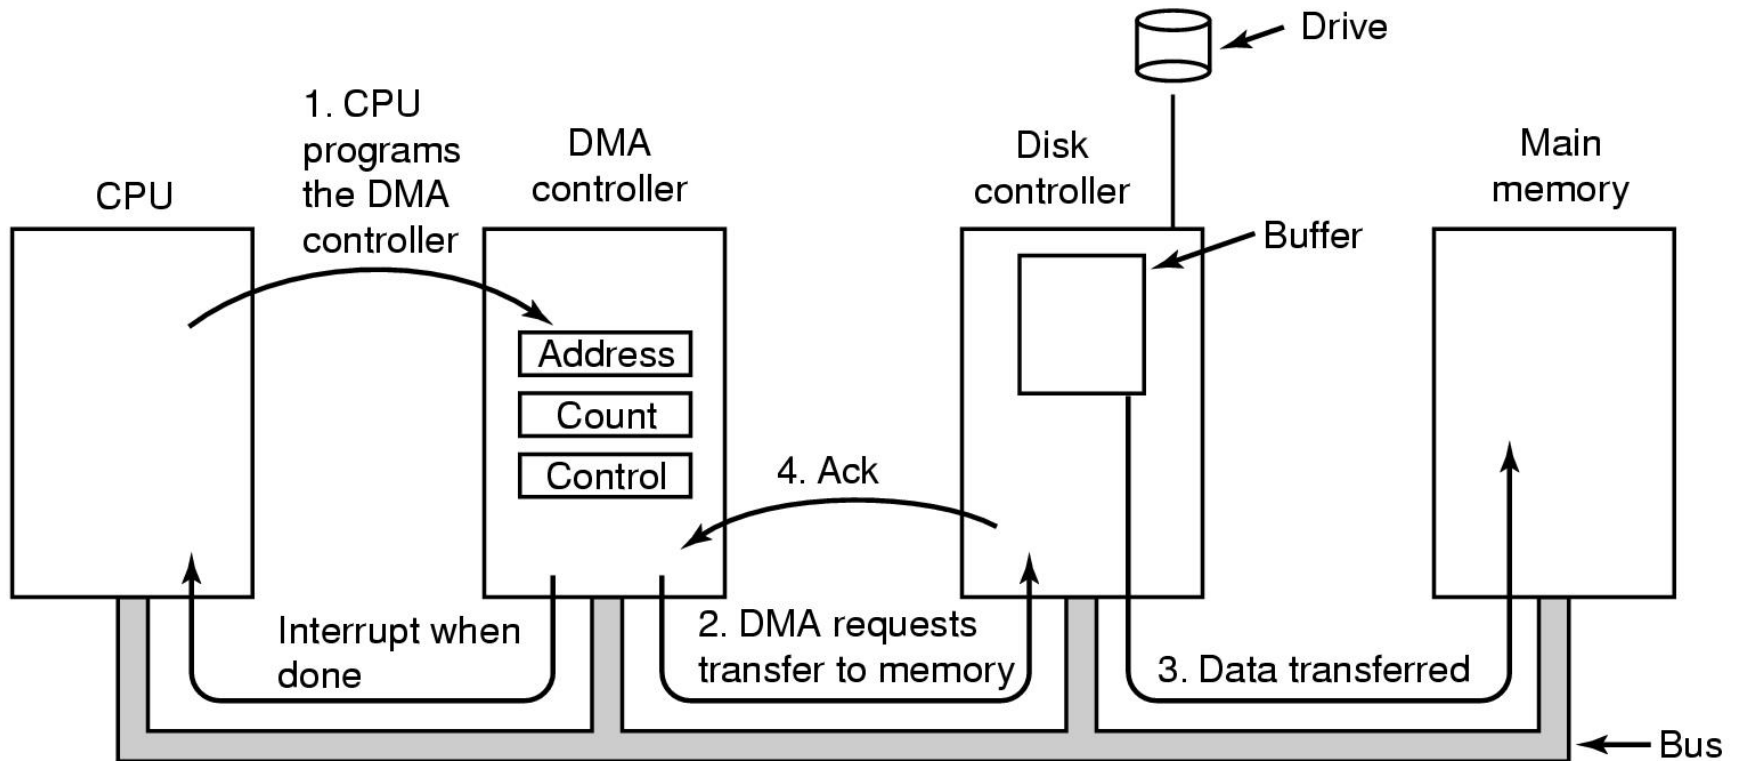
- **The Interrupt Handler:**

`Acknowledge interrupt`

`Wake up the user process`

`Return from interrupt`

Direct Memory Access (DMA)



Direct Memory Access (DMA)

- Cycle Stealing
 - ❖ DMA Controller acquires control of bus
 - ❖ Transfers a single byte (or word)
 - ❖ Releases the bus
 - ❖ The CPU is slowed down due to bus contention

- Burst Mode
 - ❖ DMA Controller acquires control of bus
 - ❖ Transfers all the data
 - ❖ Releases the bus
 - ❖ The CPU operation is temporarily suspended

Direct Memory Access (DMA)

□ Cycle Stealing

- ❖ DMA controller acquires control of bus
- ❖ Transfers a single byte (or word)
- ❖ Releases the bus
- ❖ The CPU is slowed down due to bus contention
- ❖ *Responsive but not very efficient*

□ Burst Mode

- ❖ DMA Controller acquires control of bus
- ❖ Transfers all the data
- ❖ Releases the bus
- ❖ The CPU operation is suspended
- ❖ *Efficient but interrupts may not be serviced in a timely way*

Principles of I/O software

- **Device Independence**
 - ❖ Programs can access any I/O device
 - Hard Drive, CD-ROM, Floppy,...
 - ... without specifying the device in advance
- **Uniform Naming**
 - ❖ Devices / Files are named with simple strings
 - ❖ Names should not depend on the device
- **Error Handling**
 - ❖ ...should be as close to the hardware as possible
 - ❖ ... because its often device-specific

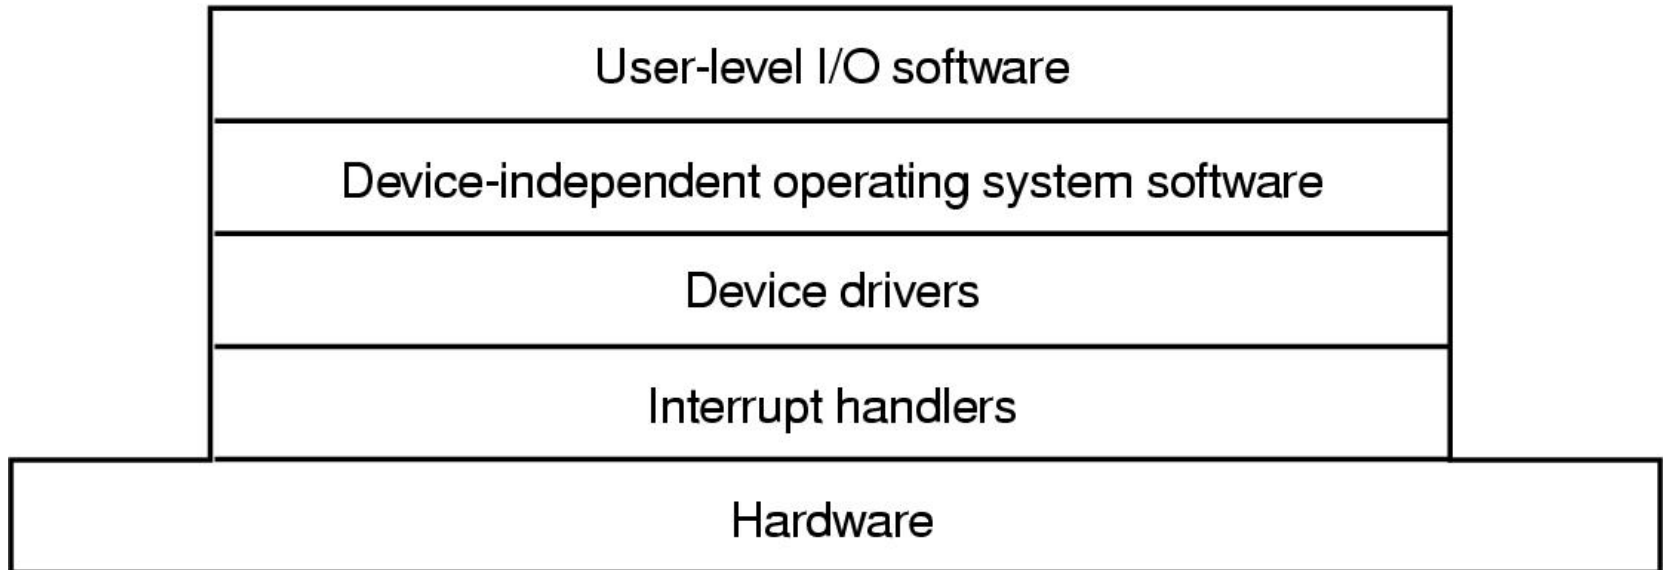
Principles of I/O software

- ❑ **Synchronous vs. Asynchronous Transfers**
 - ❖ Process is blocked vs. interrupt-driven or polling approaches
- ❑ **Buffering**
 - ❖ Data comes off a device
 - ❖ May not know the final destination of the data
 - e.g., a network packet... Where to put it???
- ❑ **Sharable vs. Dedicated Devices**
 - ❖ Disk should be sharable
 - ❖ Keyboard, Screen dedicated to one process

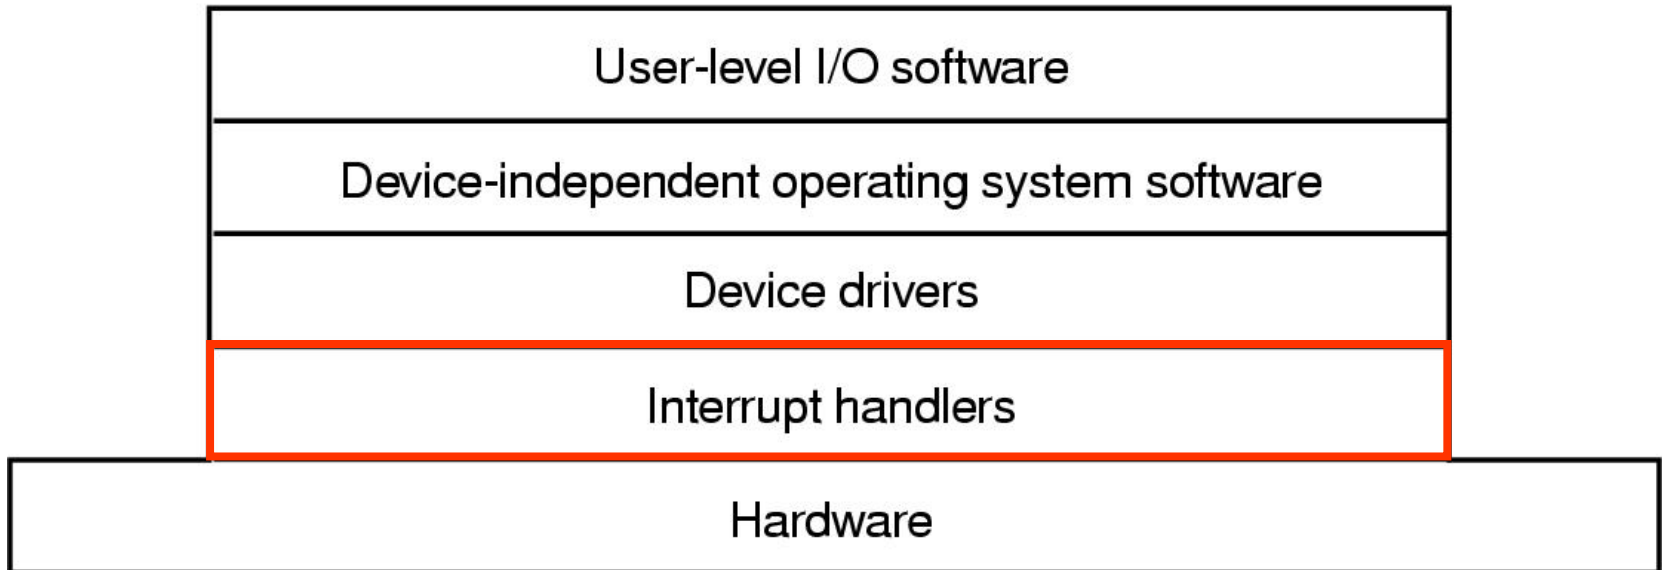
Software engineering-related challenges

- How to remove the complexities of I/O handling from application programs
 - ❖ Solution
 - standard I/O APIs (libraries and system calls)
- How to support a wide range of device types on a wide range of operating systems
 - ❖ Solution
 - standard interfaces for device drivers (DDI)
 - standard/published interfaces for access to kernel facilities (DKI)

I/O software layers



I/O software layers



Interrupt handling

- ❑ **I/O Device Driver starts the operation**
 - ❖ Then blocks until an interrupt occurs
 - ❖ Then it wakes up, finishes, & returns

- ❑ **The Interrupt Handler**
 - ❖ Does whatever is immediately necessary
 - ❖ Then unblocks the driver

- ❑ **Example: The BLITZ “DiskDriver”**
 - ❖ Start I/O and block (waits on semaphore)
 - ❖ Interrupt routine signals the semaphore & returns

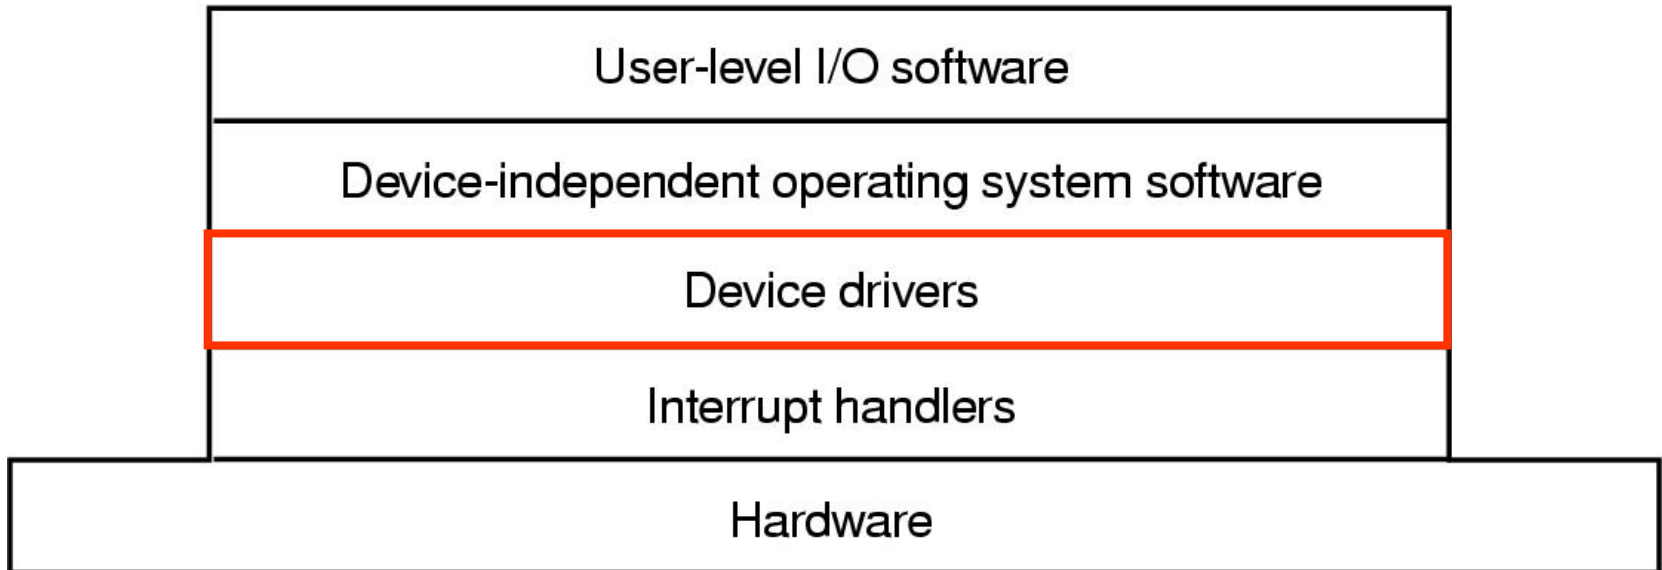
Interrupt handlers - top/bottom halves

- ❑ Interrupt handlers are divided into *scheduled* and *non scheduled* tasks
- ❑ Non-scheduled tasks execute immediately on interrupt and run in the context of the interrupted thread
 - ❖ Ie. There is no VM context switch
 - ❖ They should do a minimum amount of work so as not to disrupt progress of interrupted thread
 - ❖ They should minimize time during which interrupts are disabled
- ❑ Scheduled tasks are queued for processing by a designated thread
 - ❖ This thread will be scheduled to run later
 - ❖ May be scheduled preemptively or nonpreemptively

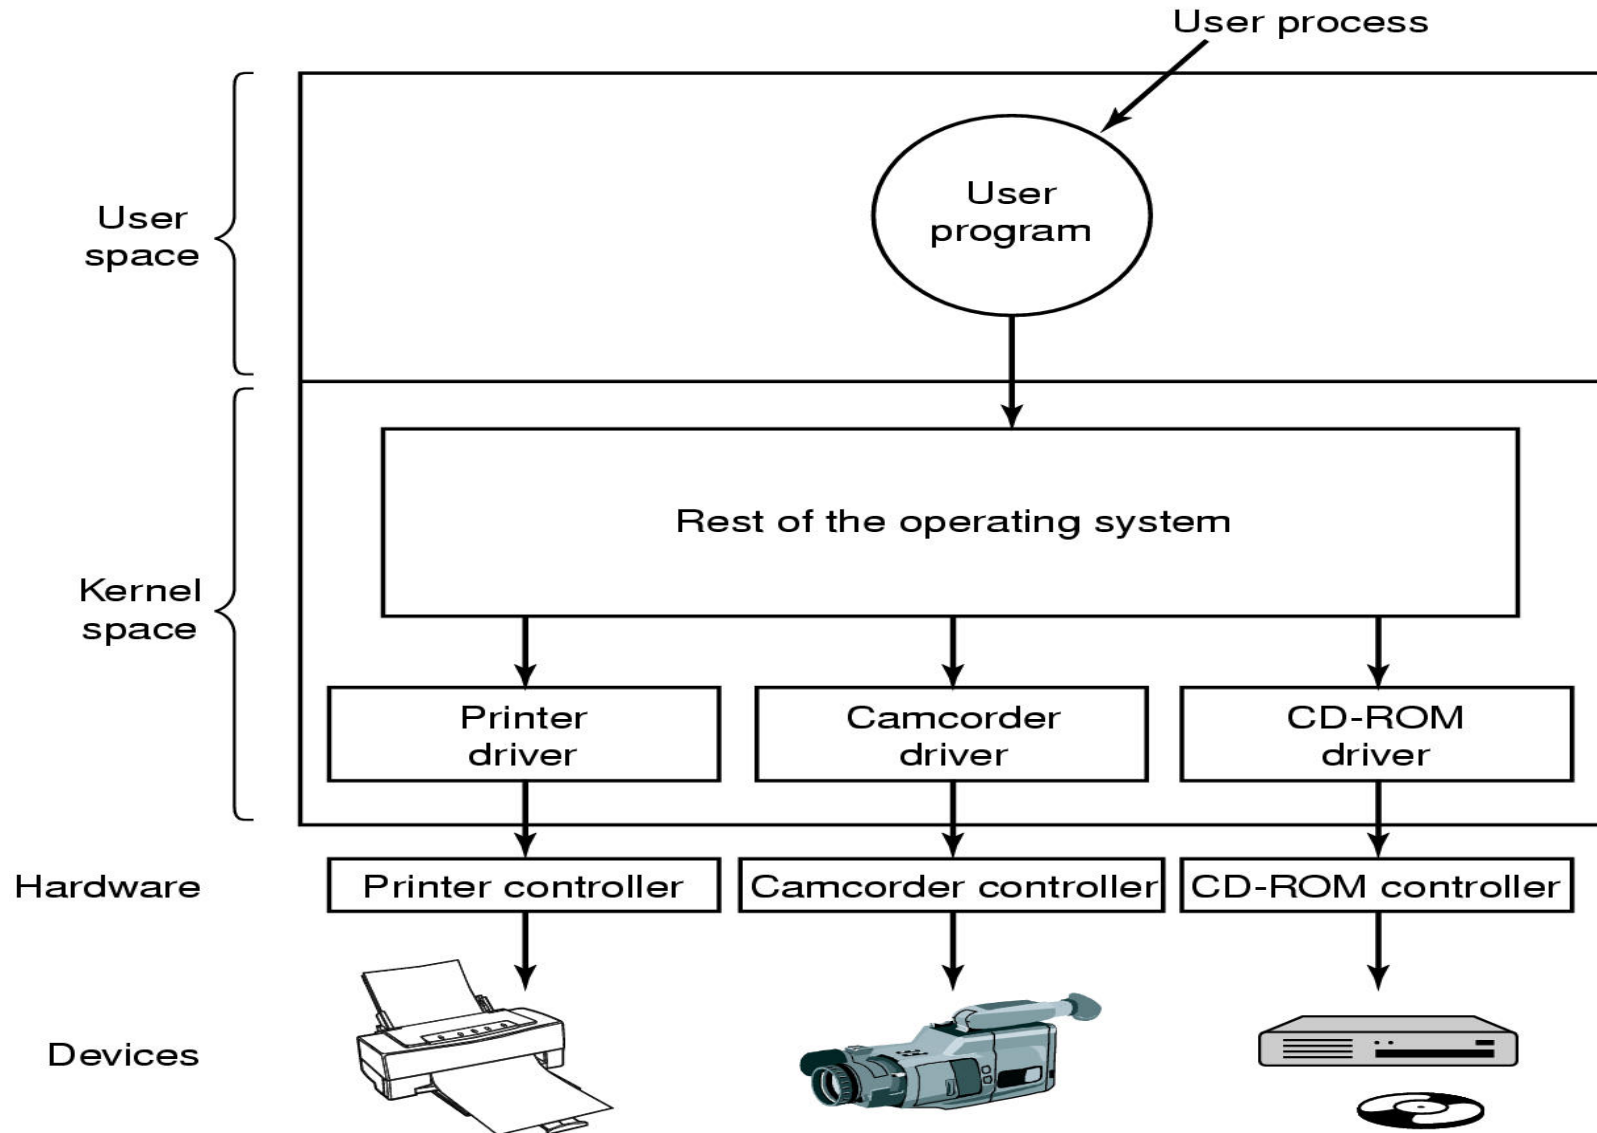
Basic activities of an interrupt handler

- ❑ Set up stack for interrupt service procedure
- ❑ Ack interrupt controller, reenale interrupts
- ❑ Copy registers from where saved
- ❑ Run service procedure

I/O software layers



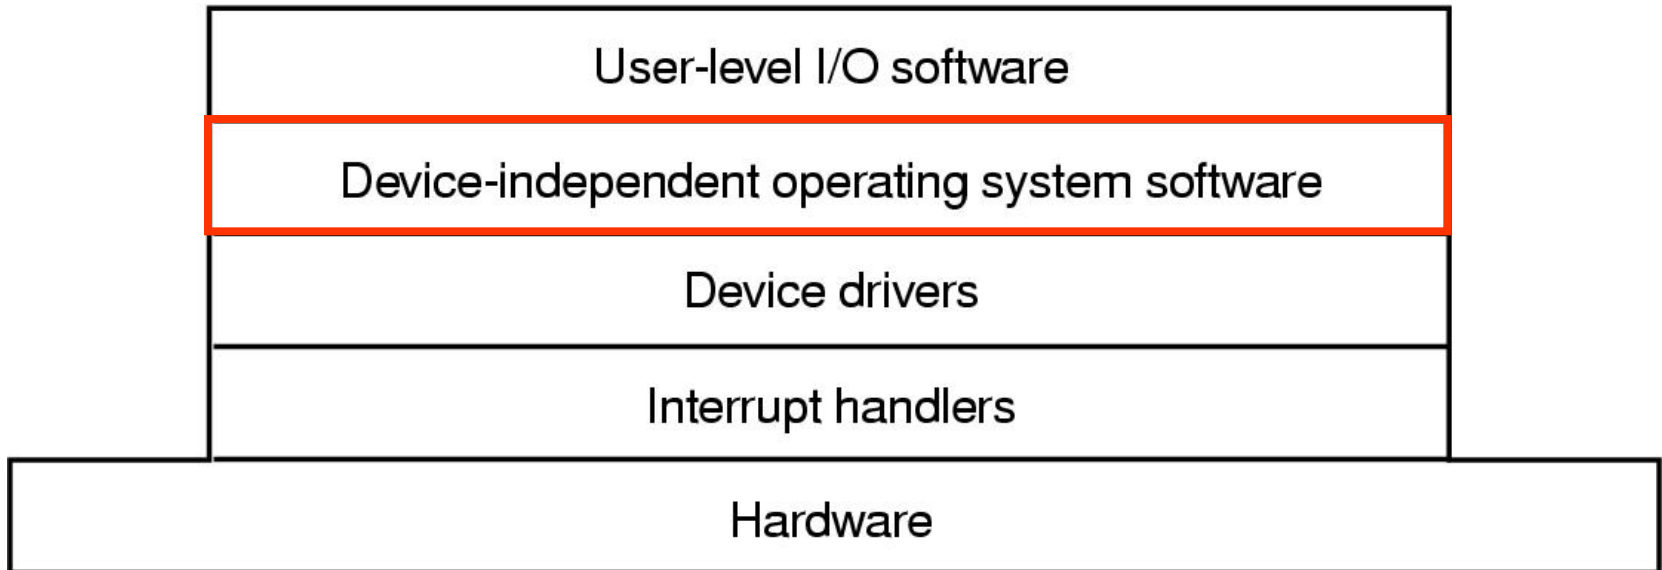
Device drivers in kernel space



Device drivers

- ❑ Device drivers are device-specific software that connects devices with the operating system
 - ❖ Typically a nasty assembly-level job
 - Must deal with hardware-specific details (and changes)
 - Must deal with O.S. specific details (and changes)
 - ❖ Goal: hide as many device-specific details as possible from higher level software
- ❑ Device drivers are typically given kernel privileges for efficiency
 - ❖ Bugs can bring down the O.S.!
 - ❖ Open challenge: how to provide efficiency and safety???

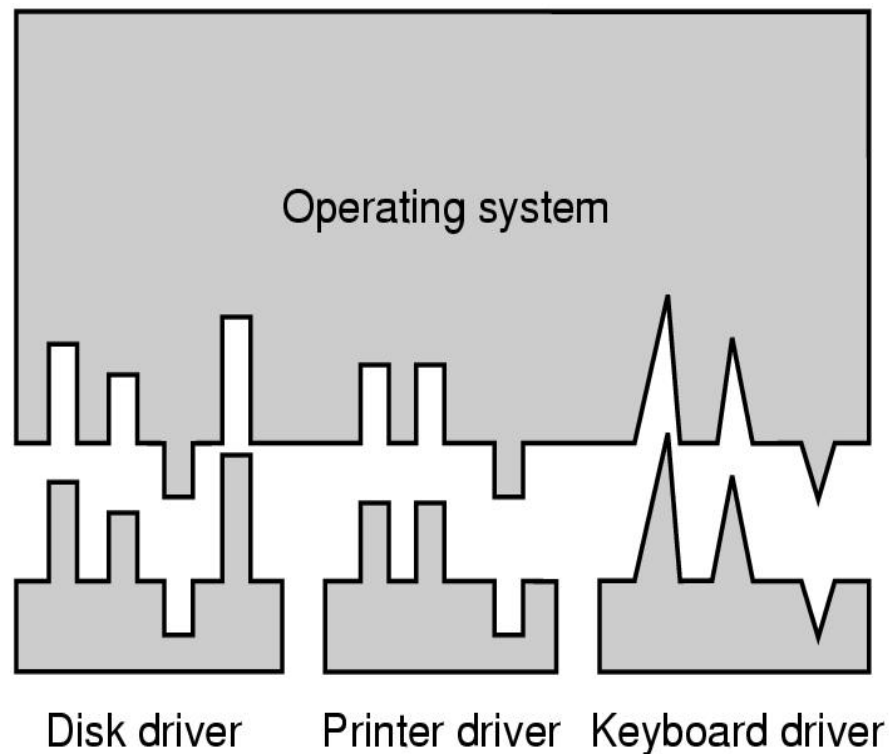
I/O software layers



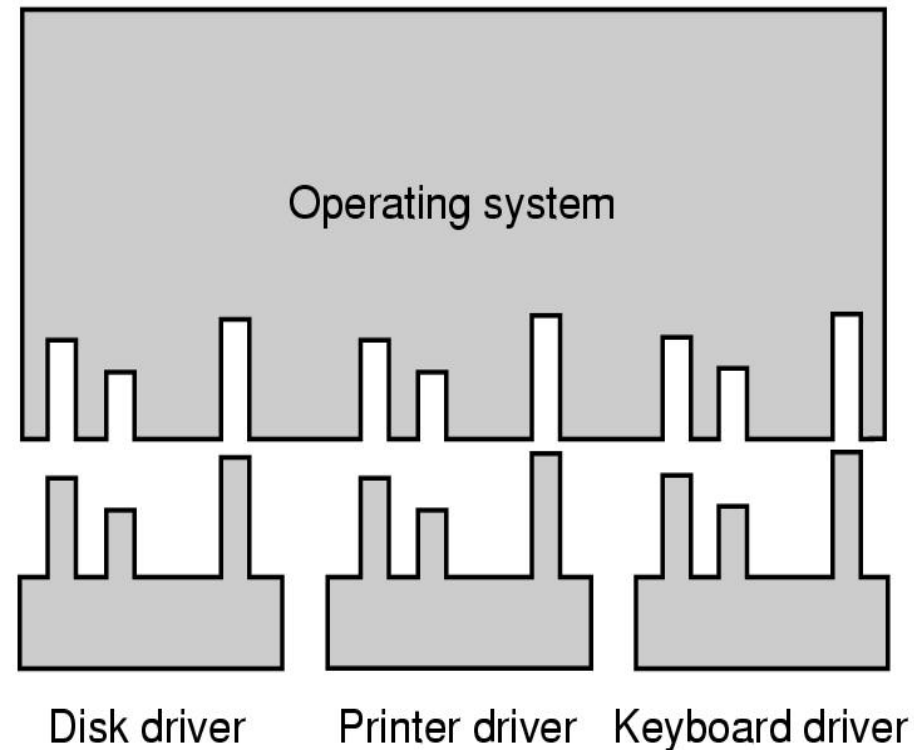
Device-independent I/O software

- **Functions and responsibilities**
 - ❖ Uniform interfacing for device drivers
 - ❖ Buffering
 - ❖ Error reporting
 - ❖ Allocating and releasing dedicated devices
 - ❖ Providing a device-independent block size

Device-independent I/O software



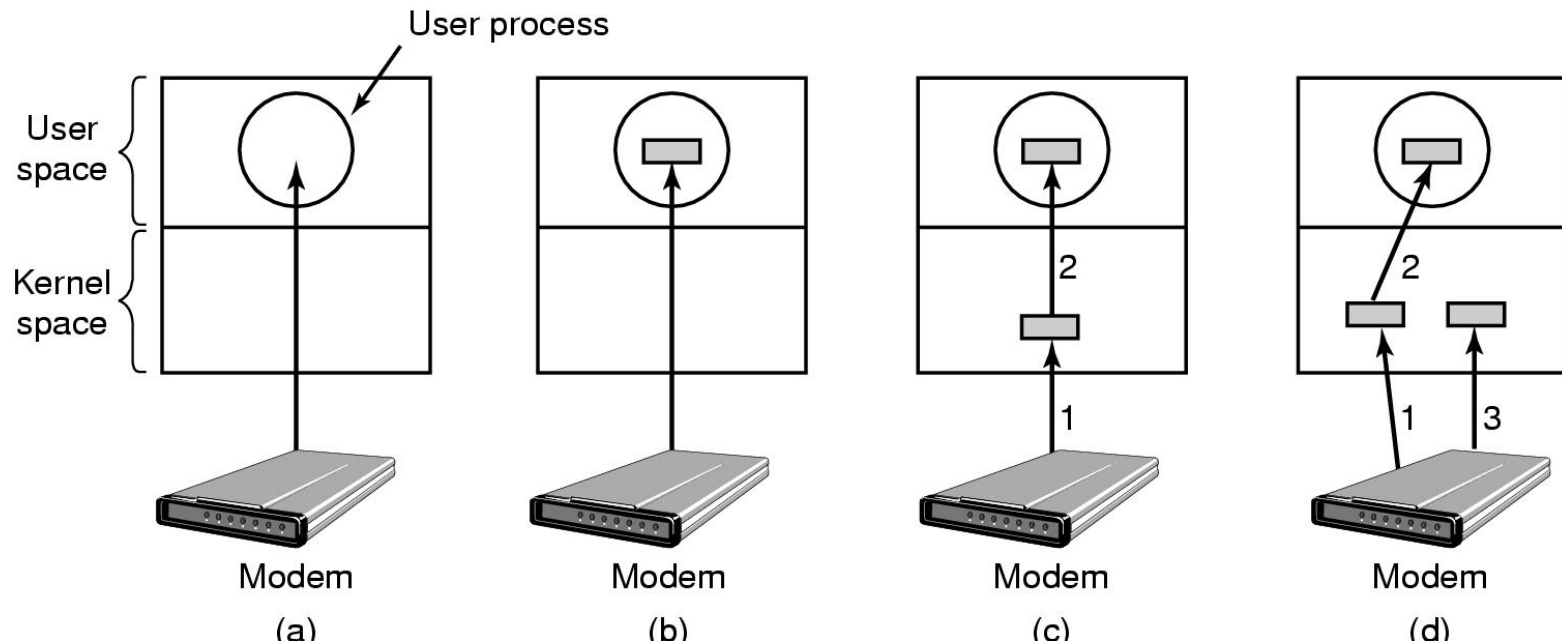
(a)



(b)

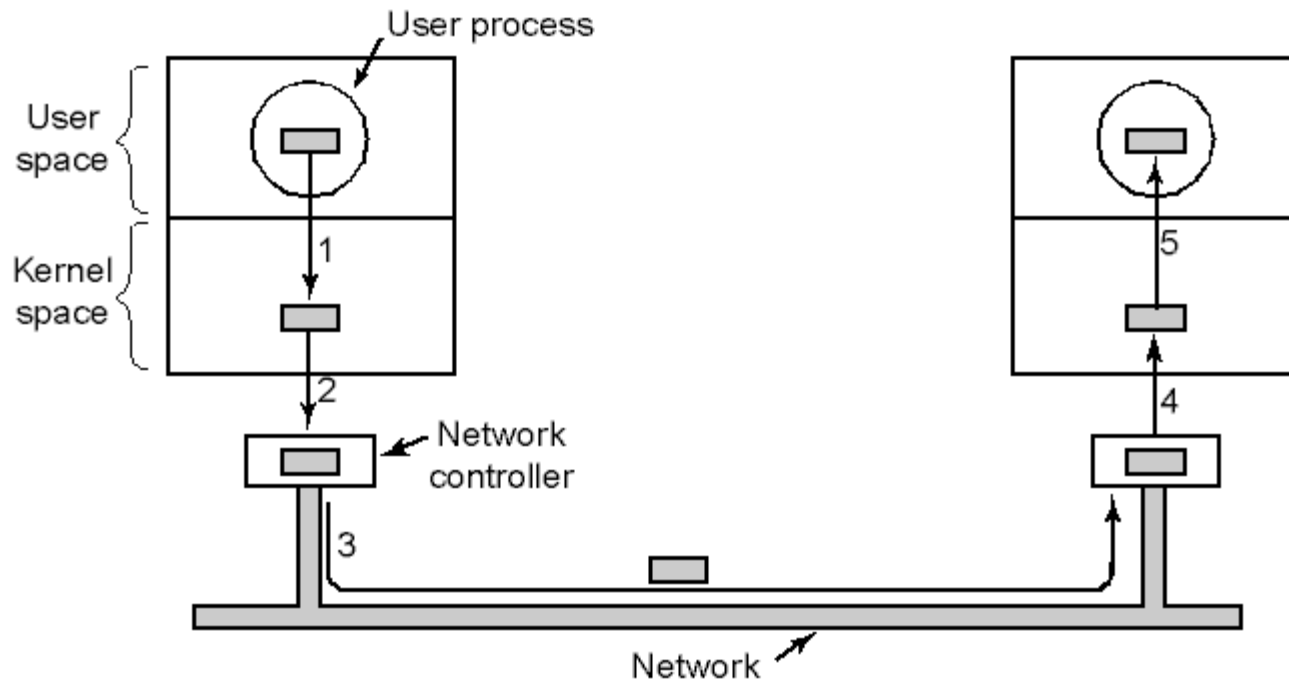
**Device Driver Interface (DDI) and Device Kernel Interface (DKI)
without/with standardization**

Device-independent I/O software buffering




- (a) Unbuffered input
- (b) Buffering in user space
- (c) Buffering in the kernel followed by copying to user space
- (d) Double buffering in the kernel

Copying overhead in network I/O





Networking may involve many copies

Devices as files



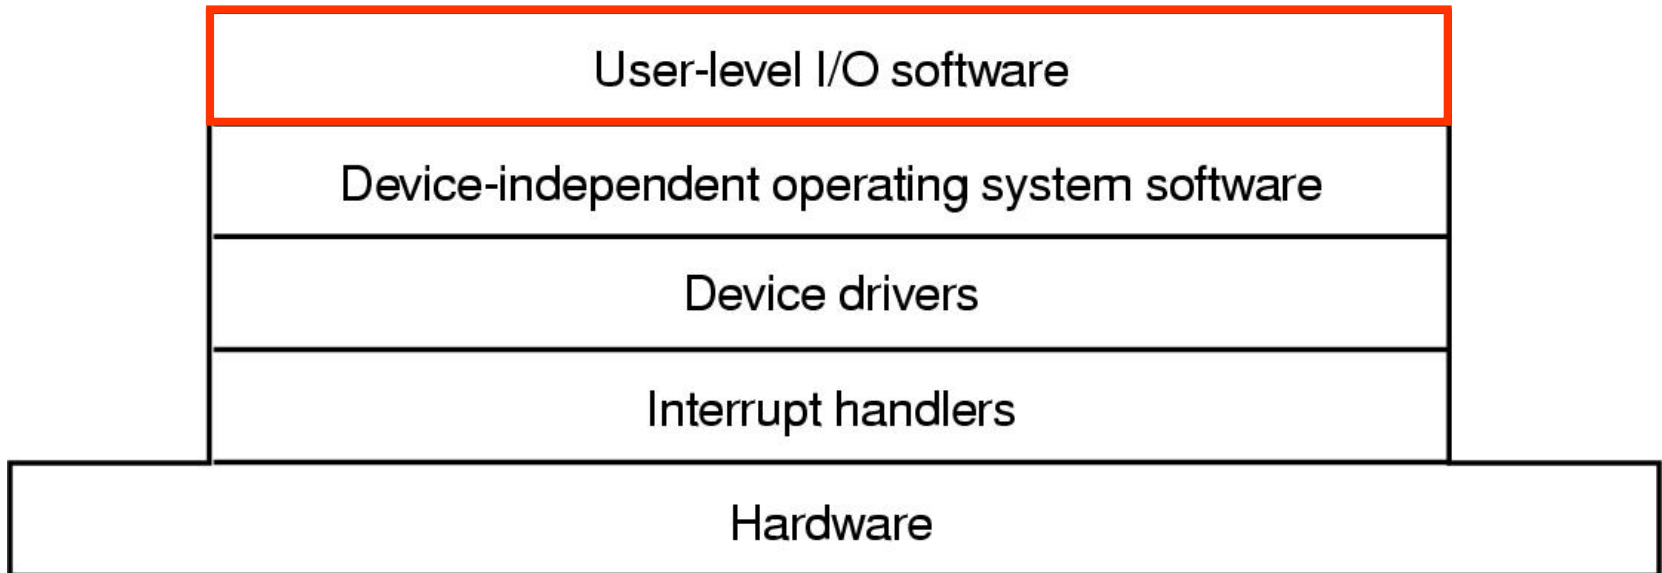
(a)



(b)

- ❑ **Before mounting,**
 - ❖ files on floppy are inaccessible
- ❑ **After mounting floppy on b,**
 - ❖ files on floppy are part of file hierarchy

I/O software layers



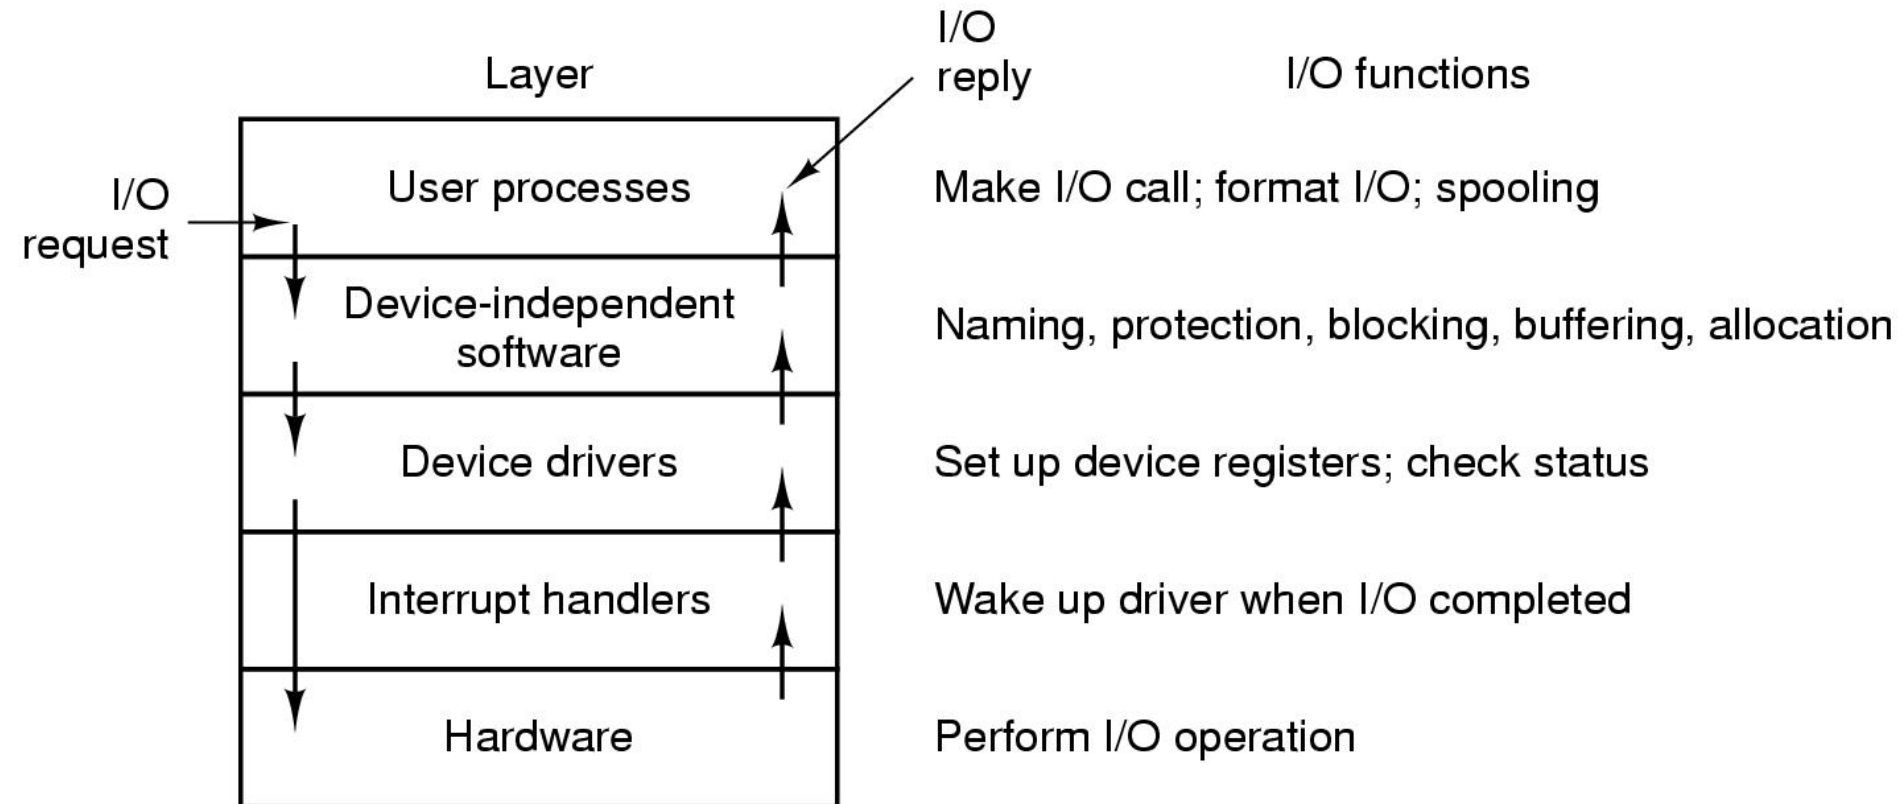
User-space I/O software

- In user's (C) program

```
count = write (fd, buffer, nbytes);  
printf ("The value of %s is %d\n", str, i);
```

- Linked with library routines.
- The library routines contain:
 - ❖ Lots of code
 - ❖ Buffering
 - ❖ The syscall to trap into the kernel


Communicating across the I/O layers



Some example I/O devices

- ❑ Timers
- ❑ Terminals
- ❑ Graphical user interfaces
- ❑ Network terminals

Programmable clocks



- One-shot mode:
 - ❖ Counter initialized then decremented until zero
 - ❖ At zero a single interrupt occurs
- Square wave mode:
 - ❖ At zero the counter is reinitialized with the same value
 - ❖ Periodic interrupts (called "clock ticks") occur

Time

- ❑ 500 MHz Crystal (oscillates every 2 nanoseconds)
- ❑ 32 bit register overflows in 8.6 seconds
 - ❖ So how can we remember what the time is?
- ❑ Backup clock
 - ❖ Similar to digital watch
 - ❖ Low-power circuitry, battery-powered
 - ❖ Periodically reset from the internet
 - ❖ UTC: Universal Coordinated Time
 - ❖ Unix: Seconds since Jan. 1, 1970
 - ❖ Windows: Seconds since Jan. 1, 1980


Goals of clock software

- **Maintain time of day**
 - ❖ Must update the time-of-day every tick
- **Prevent processes from running too long**
- **Account for CPU usage**
 - ❖ Separate timer for every process
 - ❖ Charge each tick to the current process
- **Handling the “Alarm” syscall**
 - ❖ User programs ask to be sent a signal at a given time
- **Providing watchdog timers for the OS itself**
 - ❖ E.g., when to spin down the disk
- **Doing profiling, monitoring, and statistics gathering**

Software timers

- A process can ask for notification (alarm) at time T
 - ❖ At time T , the OS will signal the process
- Processes can “go to sleep until time T ”
- Several processes can have active timers
- The CPU has only one clock
 - ❖ Must service the “alarms” in the right order
- Keep a sorted list of all timers
 - ❖ Each entry tells when the alarm goes off and what to do then

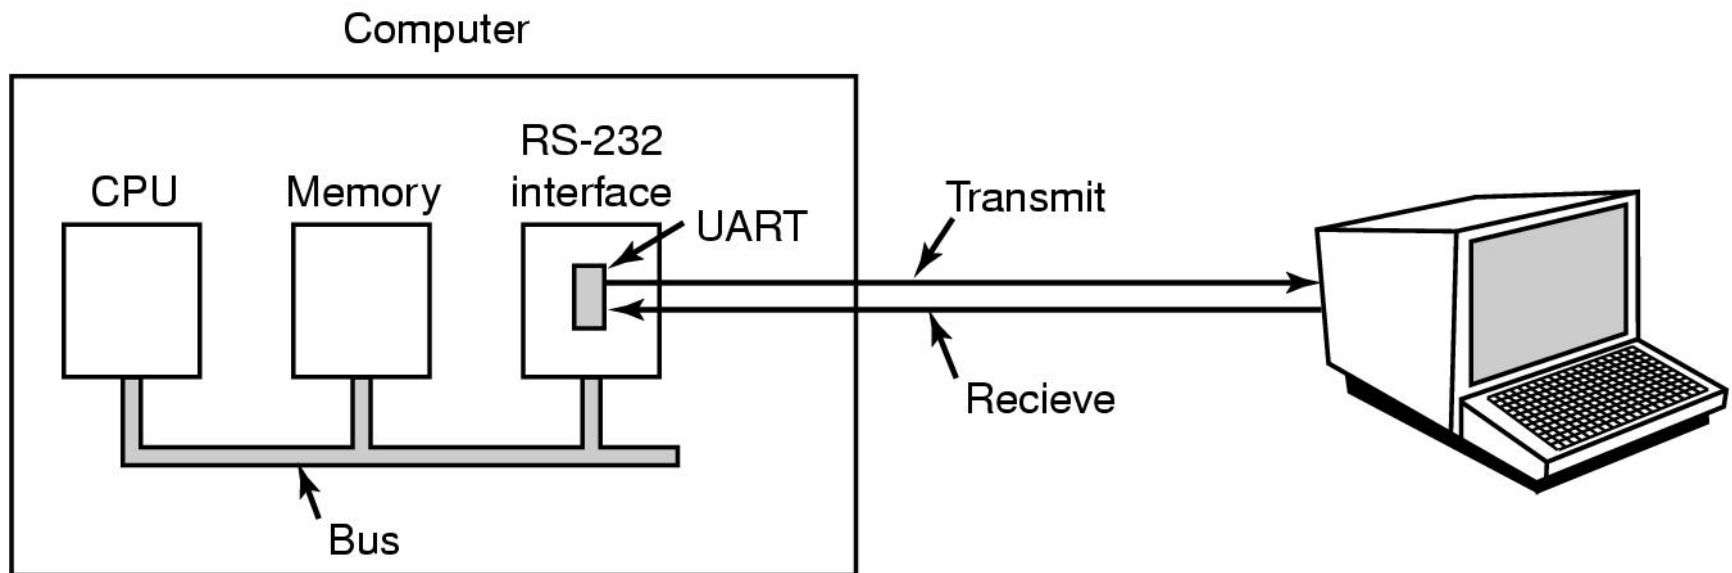
Software timers



- ❑ Alarms set for 4203, 4207, 4213, 4215 and 4216.
- ❑ Each entry tells how many ticks past the previous entry.
- ❑ On each tick, decrement the "NextSignal".
- ❑ When it gets to 0, then signal the process.

Character-oriented I/O

- ❑ RS-232 / Serial interface / Modem / Terminals / tty / COM
- ❑ Bit serial (9- or 25-pin connectors), only 3 wires used
- ❑ UART: Universal Asynchronous Receiver Transmitter
 - ❖ byte → serialize bits → wire → collect bits → byte



Terminals

- ❑ 56,000 baud = 56,000 bits per second = 7000 bytes / sec
 - ❖ Each is an ASCII character code
- ❑ Dumb terminals (CRTs / teletypes)
 - ❖ Very few control characters
 - newline, return, backspace
- ❑ Intelligent CRTs
 - ❖ Also accept "escape sequences"
 - ❖ Reposition the cursor, clear the screen, insert lines, etc.
 - ❖ The standard "terminal interface" for computers
 - Example programs: vi, emacs

Input software

- ❑ **Character processing**
 - ❖ User types "hella←o"
 - ❖ Computer echoes as: "hella←_←o"
 - ❖ Program will see "hello"
- ❑ **Raw mode**
 - ❖ The driver delivers all characters to application
 - ❖ No modifications, no echoes
 - ❖ vi, emacs, the BLITZ emulator, password entry
- ❑ **Cooked mode**
 - ❖ The driver does echoing and processing of special chars.
 - ❖ "Canonical mode"



Special control characters (in “cooked mode”)

Character	POSIX name	Comment
CTRL-H	ERASE	Backspace one character
CTRL-U	KILL	Erase entire line being typed
CTRL-V	LNEXT	Interpret next character literally
CTRL-S	STOP	Stop output
CTRL-Q	START	Start output
DEL	INTR	Interrupt process (SIGINT)
CTRL-\	QUIT	Force core dump (SIGQUIT)
CTRL-D	EOF	End of file
CTRL-M	CR	Carriage return (unchangeable)
CTRL-J	NL	Linefeed (unchangeable)

Cooked mode

- The terminal driver must...
 - ❖ Buffer an entire line before returning to application
 - ❖ Process special control characters
 - **Control-C, Backspace, line-erase, tabs**
 - ❖ Echo the character just typed
 - ❖ Accommodate type-ahead
 - **Ie., it needs an internal buffer**
- [Approach 1](#) (for computers with many terminals)
 - ❖ Have a pool of buffers to use as necessary
- [Approach 2](#) (for single-user computer)
 - ❖ Have one buffer (e.g., 500 bytes) per terminal

Central buffer pool vs. dedicated buffers



The end-of-line problem

- ❑ NL “newline” (ASCII 0x0A, \n)
 - ❖ Move cursor down one line (no horizontal movement)
- ❑ CR “return” (ASCII 0x0D, \r)
 - ❖ Move cursor to column 1 (no vertical movement)
- ❑ “ENTER key”
 - ❖ Behavior depends on the terminal specs
 - May send CR, may send NL, may send both
 - Software must be device independent
- ❑ Unix, Macintosh:
 - ❖ Each line (in a file) ends with a NL
- ❑ Windows:
 - ❖ Each line (in a file) ends with CR & NL

Control-D: EOF

- Typing Control-D ("End of file") causes the read request to be satisfied immediately
 - ❖ Do not wait for "enter key"
 - ❖ Do not wait for any characters at all
 - ❖ May return 0 characters
- Within the user program

```
count = Read (fd, buffer, buffSize)
if count == 0
    -- Assume end-of-file reached...
```


Outputting to a terminal

- ❑ The terminal accepts an “escape sequence”
- ❑ Tells it to do something special

*ESCAPE:
0x1B*

Example:

esc [3 ; 1 H esc [0 K esc [1 M

Move to position (3,1) on screen Erase the line Shift following lines up one

- ❑ Each terminal manufacturer had a slightly different specification
 - ❖ *Makes device independent software difficult*
 - ❖ Unix “termcap” file
 - Database of different terminals and their behaviors.

ANSI escape sequence standard


Escape sequence	Meaning
ESC [<i>n</i> A	Move up <i>n</i> lines
ESC [<i>n</i> B	Move down <i>n</i> lines
ESC [<i>n</i> C	Move right <i>n</i> spaces
ESC [<i>n</i> D	Move left <i>n</i> spaces
ESC [<i>m</i> ; <i>n</i> H	Move cursor to (<i>m</i> , <i>n</i>)
ESC [<i>s</i> J	Clear screen from cursor (0 to end, 1 from start, 2 all)
ESC [<i>s</i> K	Clear line from cursor (0 to end, 1 from start, 2 all)
ESC [<i>n</i> L	Insert <i>n</i> lines at cursor
ESC [<i>n</i> M	Delete <i>n</i> lines at cursor
ESC [<i>n</i> P	Delete <i>n</i> chars at cursor
ESC [<i>n</i> @	Insert <i>n</i> chars at cursor
ESC [<i>n</i> m	Enable rendition <i>n</i> (0=normal, 4=bold, 5=blinking, 7=reverse)
ESC M	Scroll the screen backward if the cursor is on the top line

Graphical user interfaces (GUIs)

- **Memory-mapped displays “bit-mapped graphics”**
- **Video driver moves bits into special memory region**
 - ❖ Changes appear on the screen
 - ❖ Video controller constantly scans video ram
- **Black and white displays**
 - ❖ 1 bit = 1 pixel
- **Color**
 - ❖ 24 bits = 3 bytes = 1 pixels
 - red (0-255)
 - green (0-255)
 - blue (0-255)


$$1280 * 854 * 3 \\ = 3 \text{ MB}$$

Graphical user interfaces (GUIs)

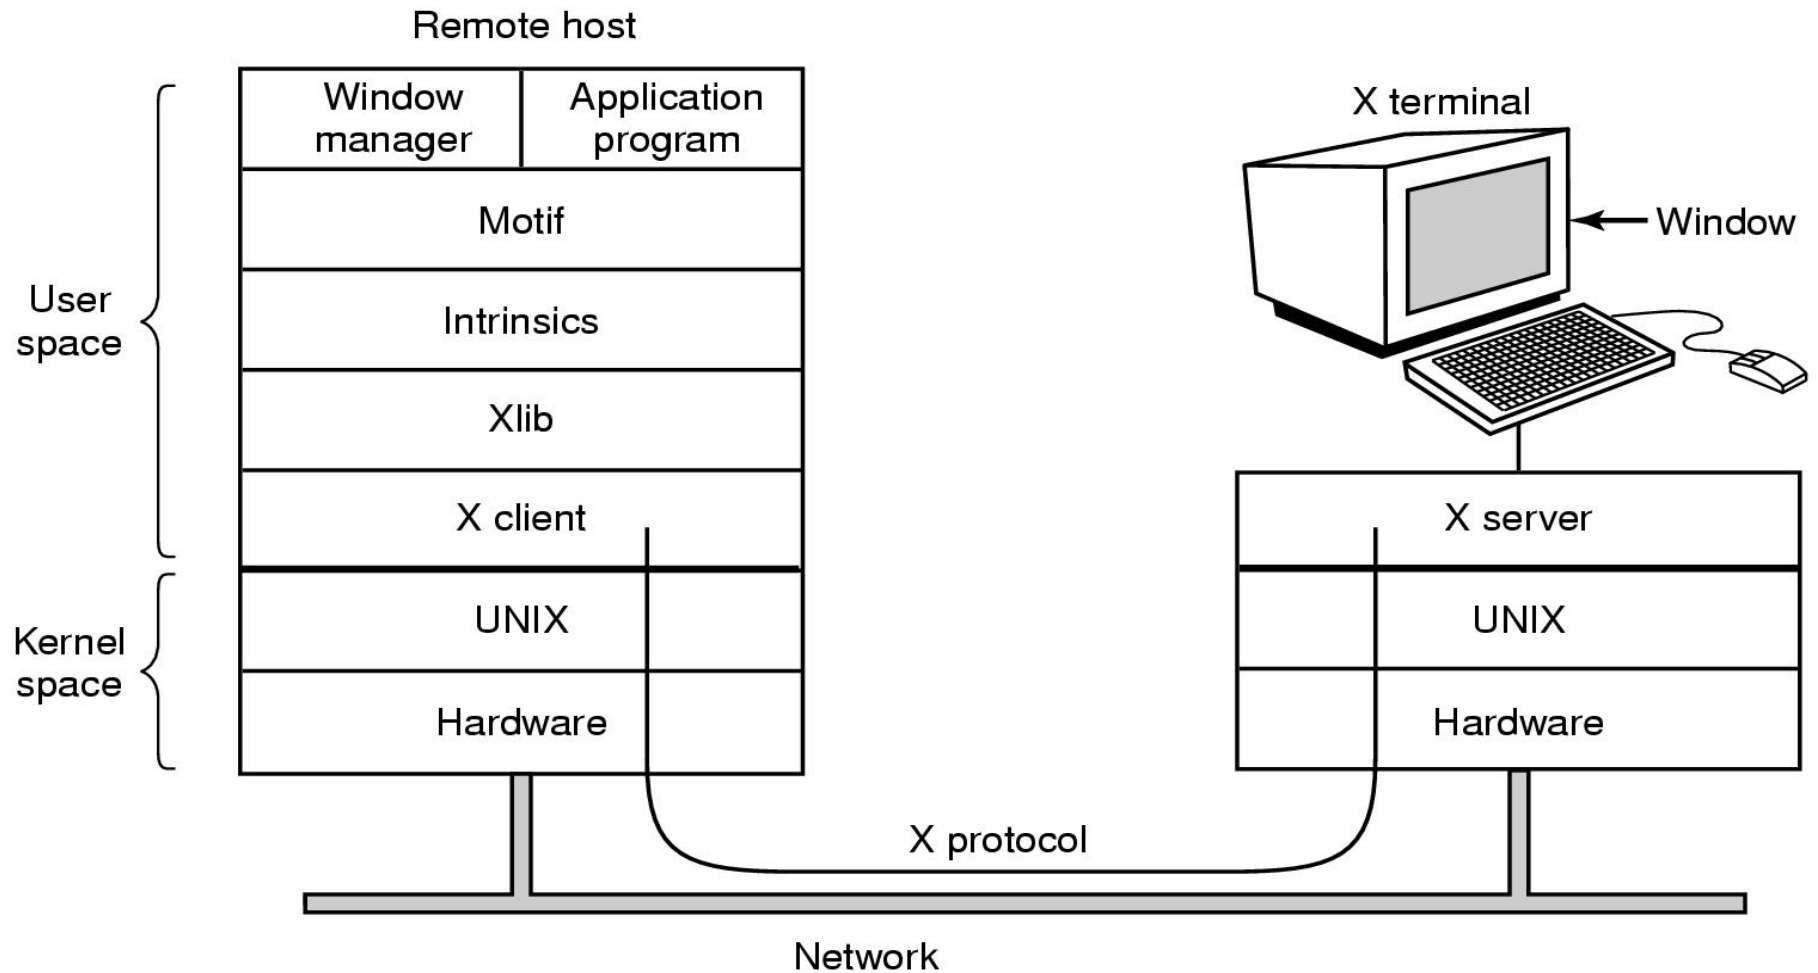


Spare slides

X Window System

- ❑ Client - Server
- ❑ Remote Procedure Calls (RPC)
- ❑ Client makes a call.
- ❑ Server is awakened; the procedure is executed.
- ❑ Intelligent terminals ("X terminals")
- ❑ The display side is the server.
- ❑ The application side is the client.
- ❑ The application (client) makes requests to the display server.
- ❑ Client and server are separate processes
 - ❖ (May be on the same or different machines)

X window system



X window system

- *X-Server*

- ❖ Display text and geometric shapes, move bits
- ❖ Collect mouse and keyboard status

- *X-Client*

- ❖ Xlib
 - library procedures; low-level access to X-Server
- ❖ Intrinsics
 - provide “*widgets*”
 - buttons, scroll bars, frames, menus, etc.
- ❖ Motif
 - provide a “look-and-feel” / style
- ❖ Window Manager
 - Application independent functionality
 - Create & move windows

The SLIM network terminal

- ❑ Stateless Low-level Interface Machine (SLIM)
 - ❖ Sun Microsystems
- ❑ Philosophy: *Keep the terminal-side very simple!*
- ❑ Back to “dumb” terminals”
- ❑ Interface to X-Server:
 - ❖ 100's of functions
- ❑ SLIM:
 - ❖ Just a few messages
 - ❖ The host tells which pixels to put where
 - ❖ The host contains all the intelligence

The SLIM network terminal

- The SLIM Protocol
 - ❖ from application-side (server)
 - ❖ to terminal (the "thin" client)

Message	Meaning
SET	Update a rectangle with new pixels
FILL	Fill a rectangle with one pixel value
BITMAP	Expand a bitmap to fill a rectangle
COPY	Copy a rectangle from one part of the frame buffer to another
CSCS	Convert a rectangle from television color (YUV) to RGB